

Deep Learning for Urban Transportation

A MIT Seminar 11.s938J/11.S196J Fall 2019

Peyman Noursalehi, Shenhao Wang and Jinhua Zhao
Massachusetts Institute of Technology

What's changing in
Transportation?

What's changing in transportation?

Technology

- Automation
- Electrification
- 5G/Connected
- Shared economy
- ...

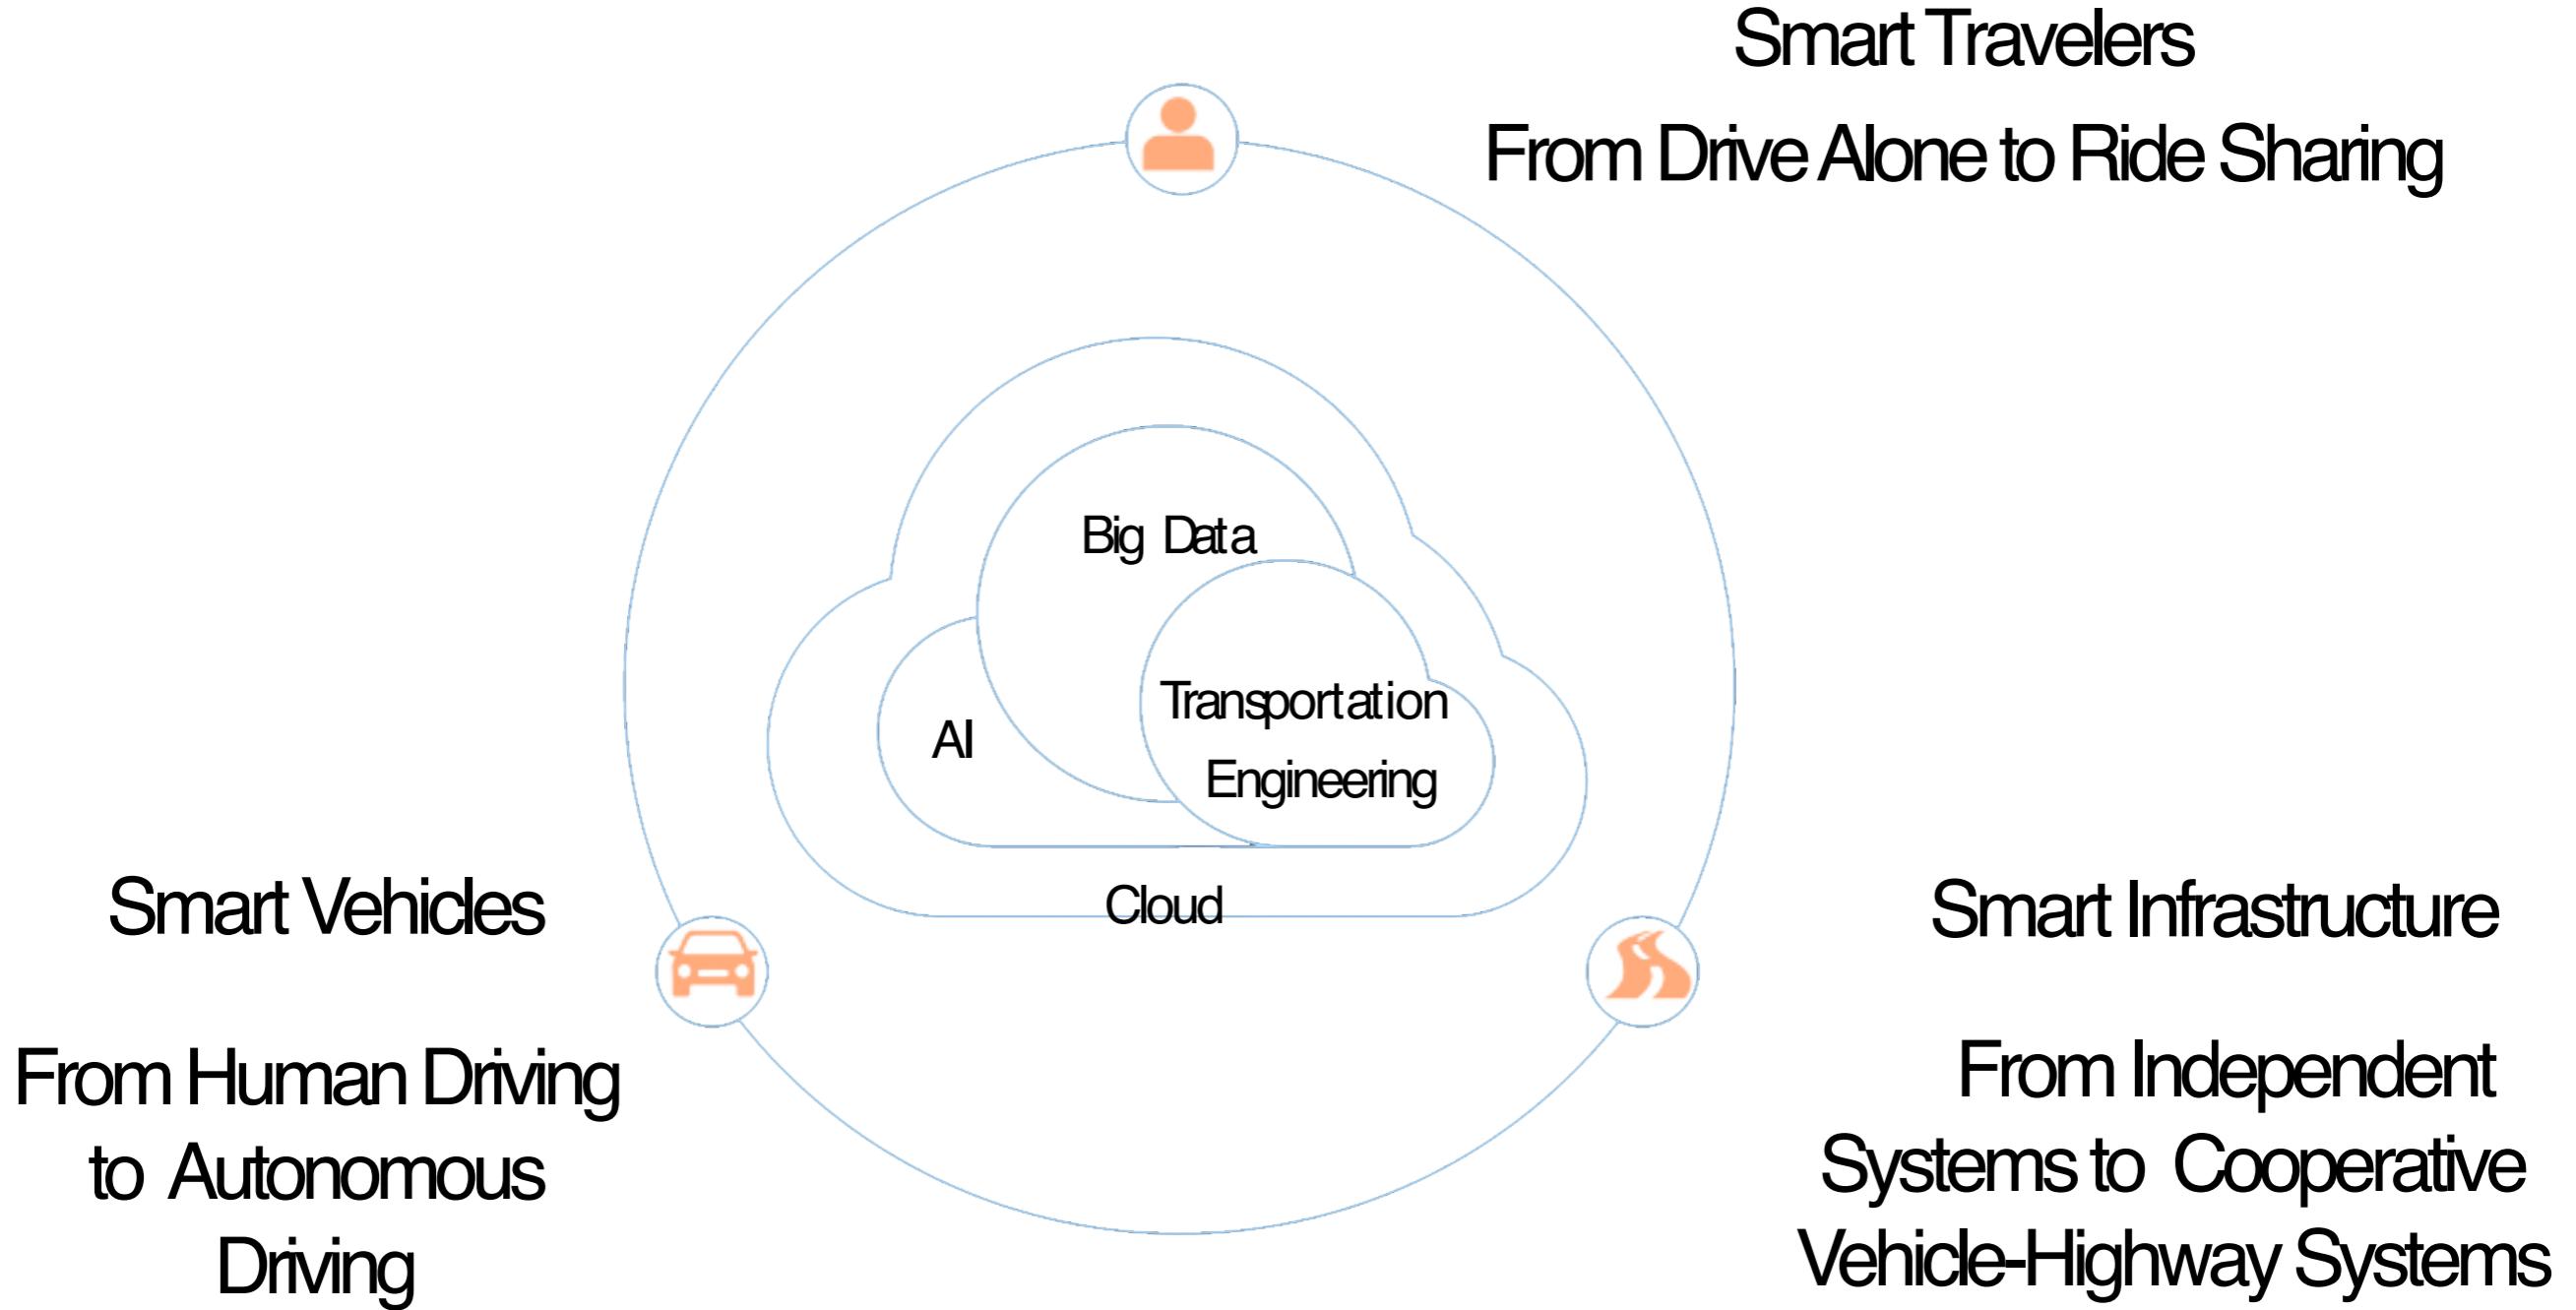
Data

- Ubiquitous sensing
- AI / computing
- Cybersecurity
- Clearing house
- ...

Value

- Climate change
- Future of work
- Public health
- Social justice
- Urban livability
- ...

People, Vehicles and Infrastructure



Tech Uncertainty → System Uncertainty

- Ownership model or access Model
- Single use or shared use
- Integrate with public transit or not
- Rational pricing structure or not
- Coordinated with land use or not
- Rational regulatory framework or not
- ...

Multimodal
Transportation



Data

The Economist

FEBRUARY 27TH-MARCH 3RD 2004

Economist.com

Obama the warrior

Misgoverning Argentina

The economic shift from West to East

Genetically modified crops blossom

The right to eat cats and dogs

The data deluge

AND HOW TO HANDLE IT: A 14-PAGE SPECIAL REPORT



Data Resource

- Location data, Trajectorydata



- Transaction data



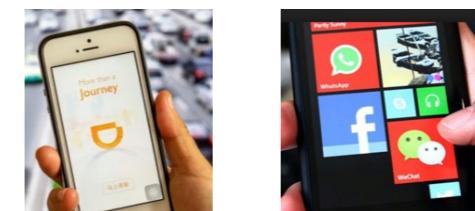
- Profile data



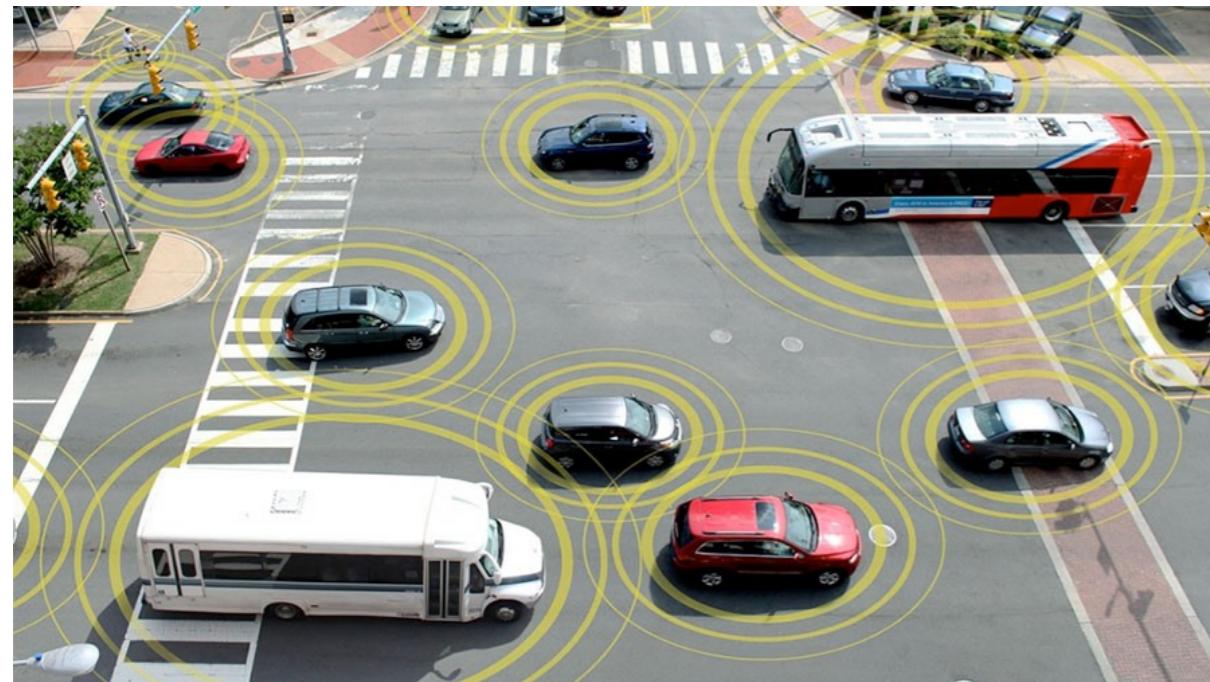
- Sensors: multimediacdata



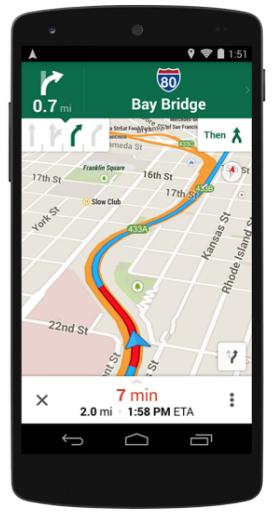
- Cross-platform identification



Location Data and Floating-Car Trajectory



Sensors: Loop detector, camera, microphone, mobile sensors . . .



Open Datasets



KDD2017
KDDCup 2017

Highway Tollgates Traffic
Flow Prediction



Uber Movement



Public Data



Federal Highway Administration
Next Generation Simulation
(NGSIM) Program



GAIA Open Dataset
Trajectory Data



- LTDS / CRM
- Oyster / Octopus
- iBus
- Mobile CDR / Wifi / RFID / Bluetooth
- Tracking Apps: MOVES / FMS / Strava
- Mobile App
- Text: Incidence Log, Customer Feedback, Twitter, Facebook, What's App...



Applying Machine Learning Techniques to Transportation Surveys

Jane Shepherd, Westat
Marcelo Simas, Westat
Anthony Fucci, Westat
Alexander Cates, Westat

Background



- Household Travel Surveys
 - Collect socio-economic and demographic data about households and individual members
 - Collect a travel diary for 1-2 days
 - Describe the how, why, when, and where of each place visited on the assigned travel day(s)
 - Recently deployed smartphone-based surveys
 - Geolocation – Auto-detects trip start/stops using geofences
 - Travel capture – GPS data informs arrival and departure times
 - Prompted recall
 - Past surveys
 - Asheville, Fairbanks, Albuquerque, South Jersey, Las Vegas, Michigan, Billings, NHTS
 - Present / future surveys
 - Chicago, Maryland, Laredo

How do we use machine learning?

1. Coding open-end responses using Natural Language Processing and Random Forest models.
2. Ascertaining Industry and Occupation in real time using Natural Language Processing and Vector Space models.
3. Determining place validity and predicting travel attributes using GPS and Accelerometer-derived features to train Random Forest models.

Coding Open-End Responses

- Problem
 - NHTS yielded over 180,000 open-ended responses
 - Around 52,000 of these belonged to the “Trip Purpose” question
- Traditional Solution
 - Analyst attempts to up-code each response by hand
 - Average 15 sec / response = ~ 750 hours
- ML Solution
 - Analyst up-codes a sample of responses.
 - Treat the sample as labeled training data to be modeled

Breadth of Coverage (User, Time, Space): Oyster + Mobile CDR



Statistical Data Fusion via
Machine Learning

Depth of Attributes (Semantics)
LTDS + CRM + Tracking



Statistical Data Fusion via
Machine Learning

Multi-modal Data Fusion and Analytics

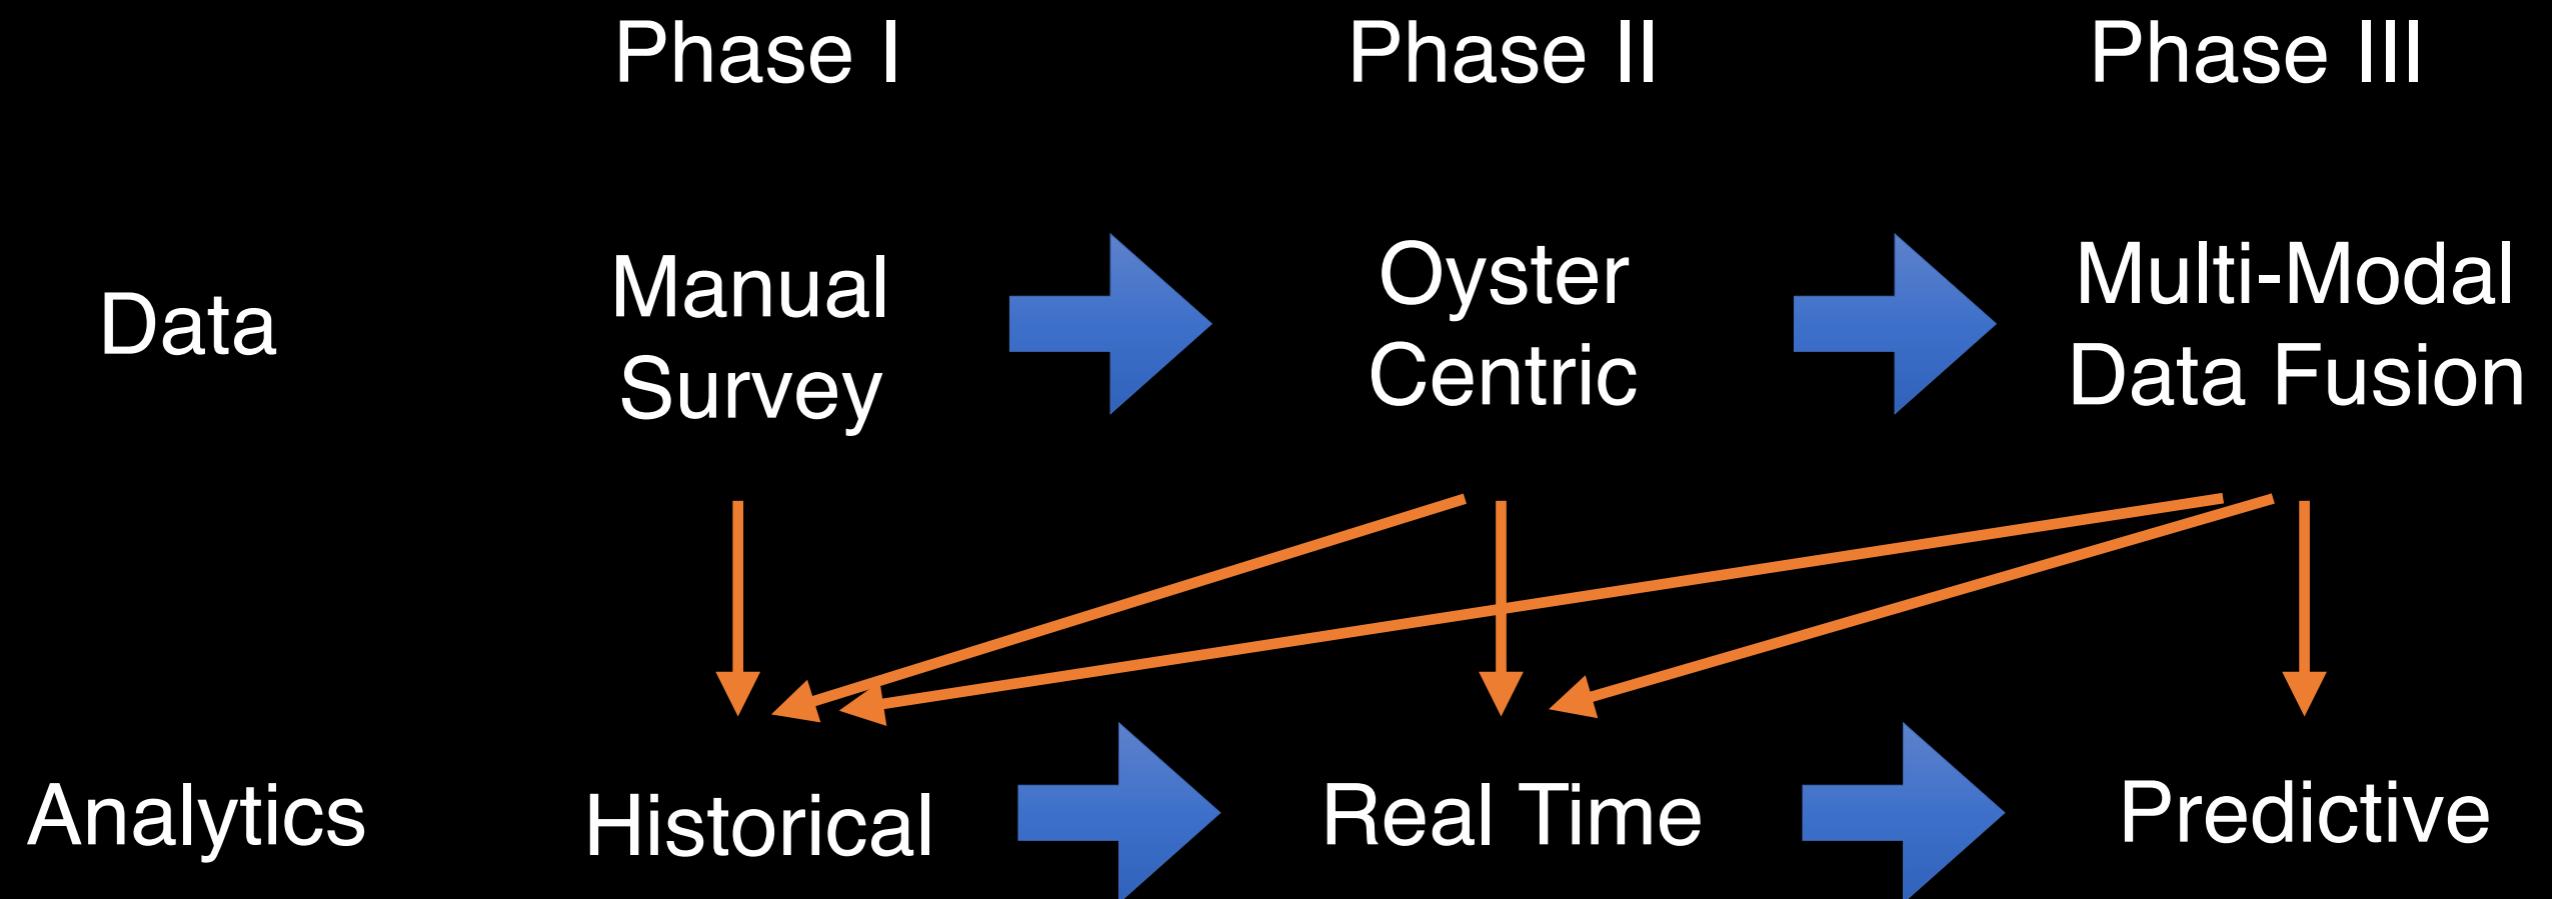
Multimodal Travel
Information Platform

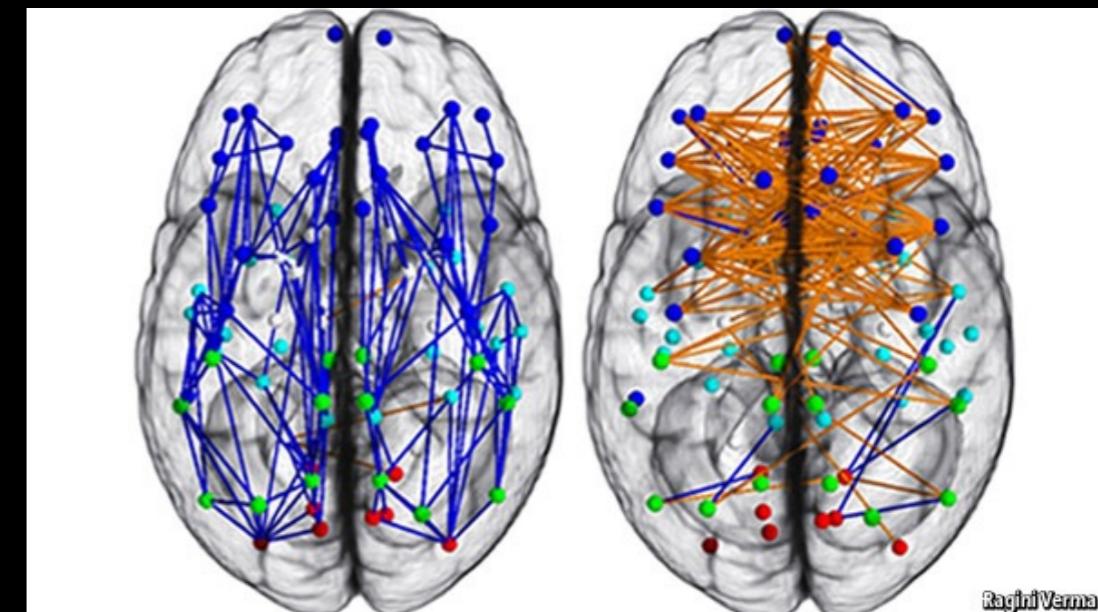
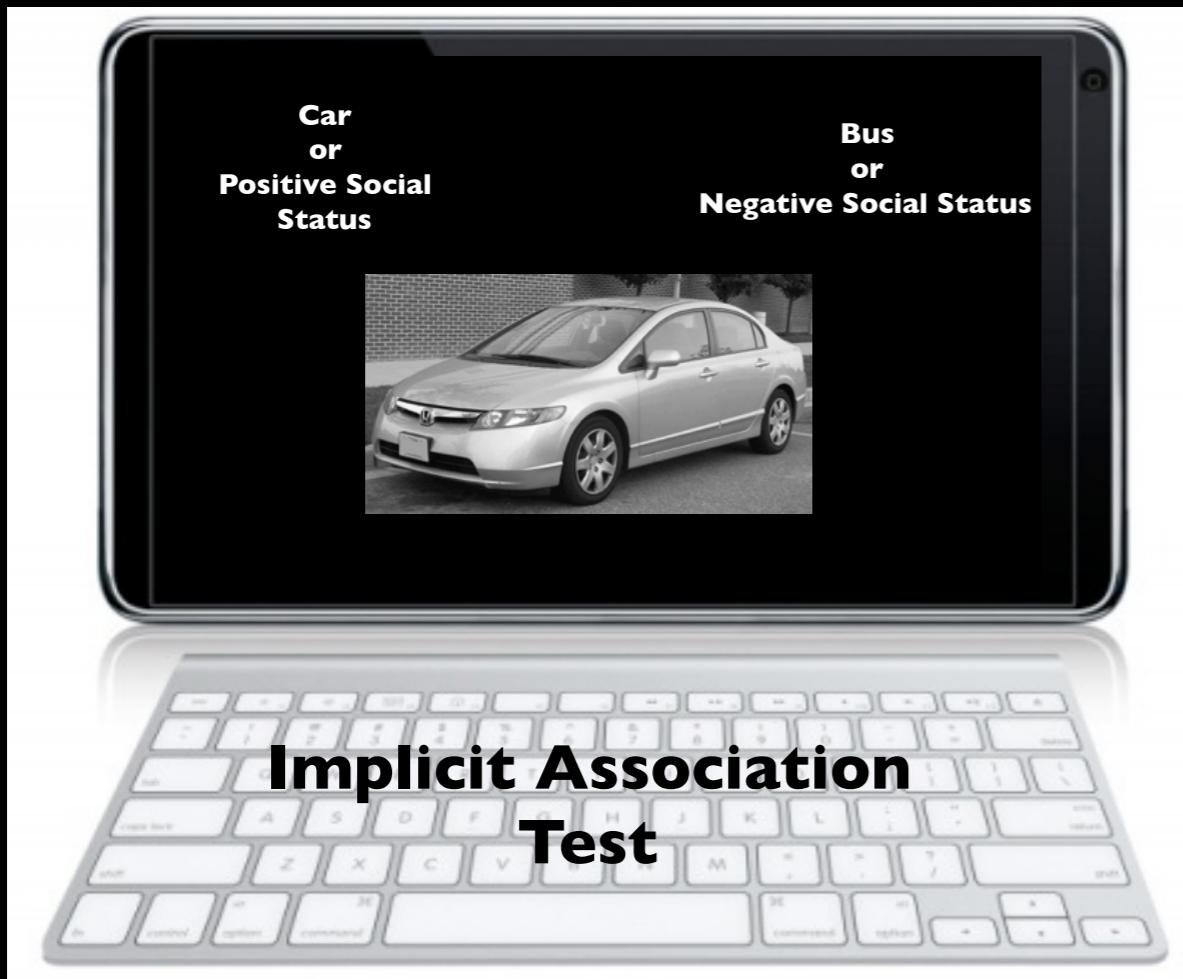


Statistical Data
Fusion

- Oyster
 - Mobile CDR / Wifi / RFID / Bluetooth
 - LTDS / CRM
 - Tracking Apps: MOVES / FMS
-
- Discrete Choice Modeling
 - Clustering: Unsupervised Learning
 - Machine Learning: Neural Network / Random Forest /Support Vector Machine ...

TfL Data Analytics





SP + RP

Multitask Learning Deep Neural Networks to Combine Revealed and Stated Preference Data

Shen Hao Wang

Qingyi Wang

Jinhua Zhao

Massachusetts Institute of Technology
ICMC 2019

Background

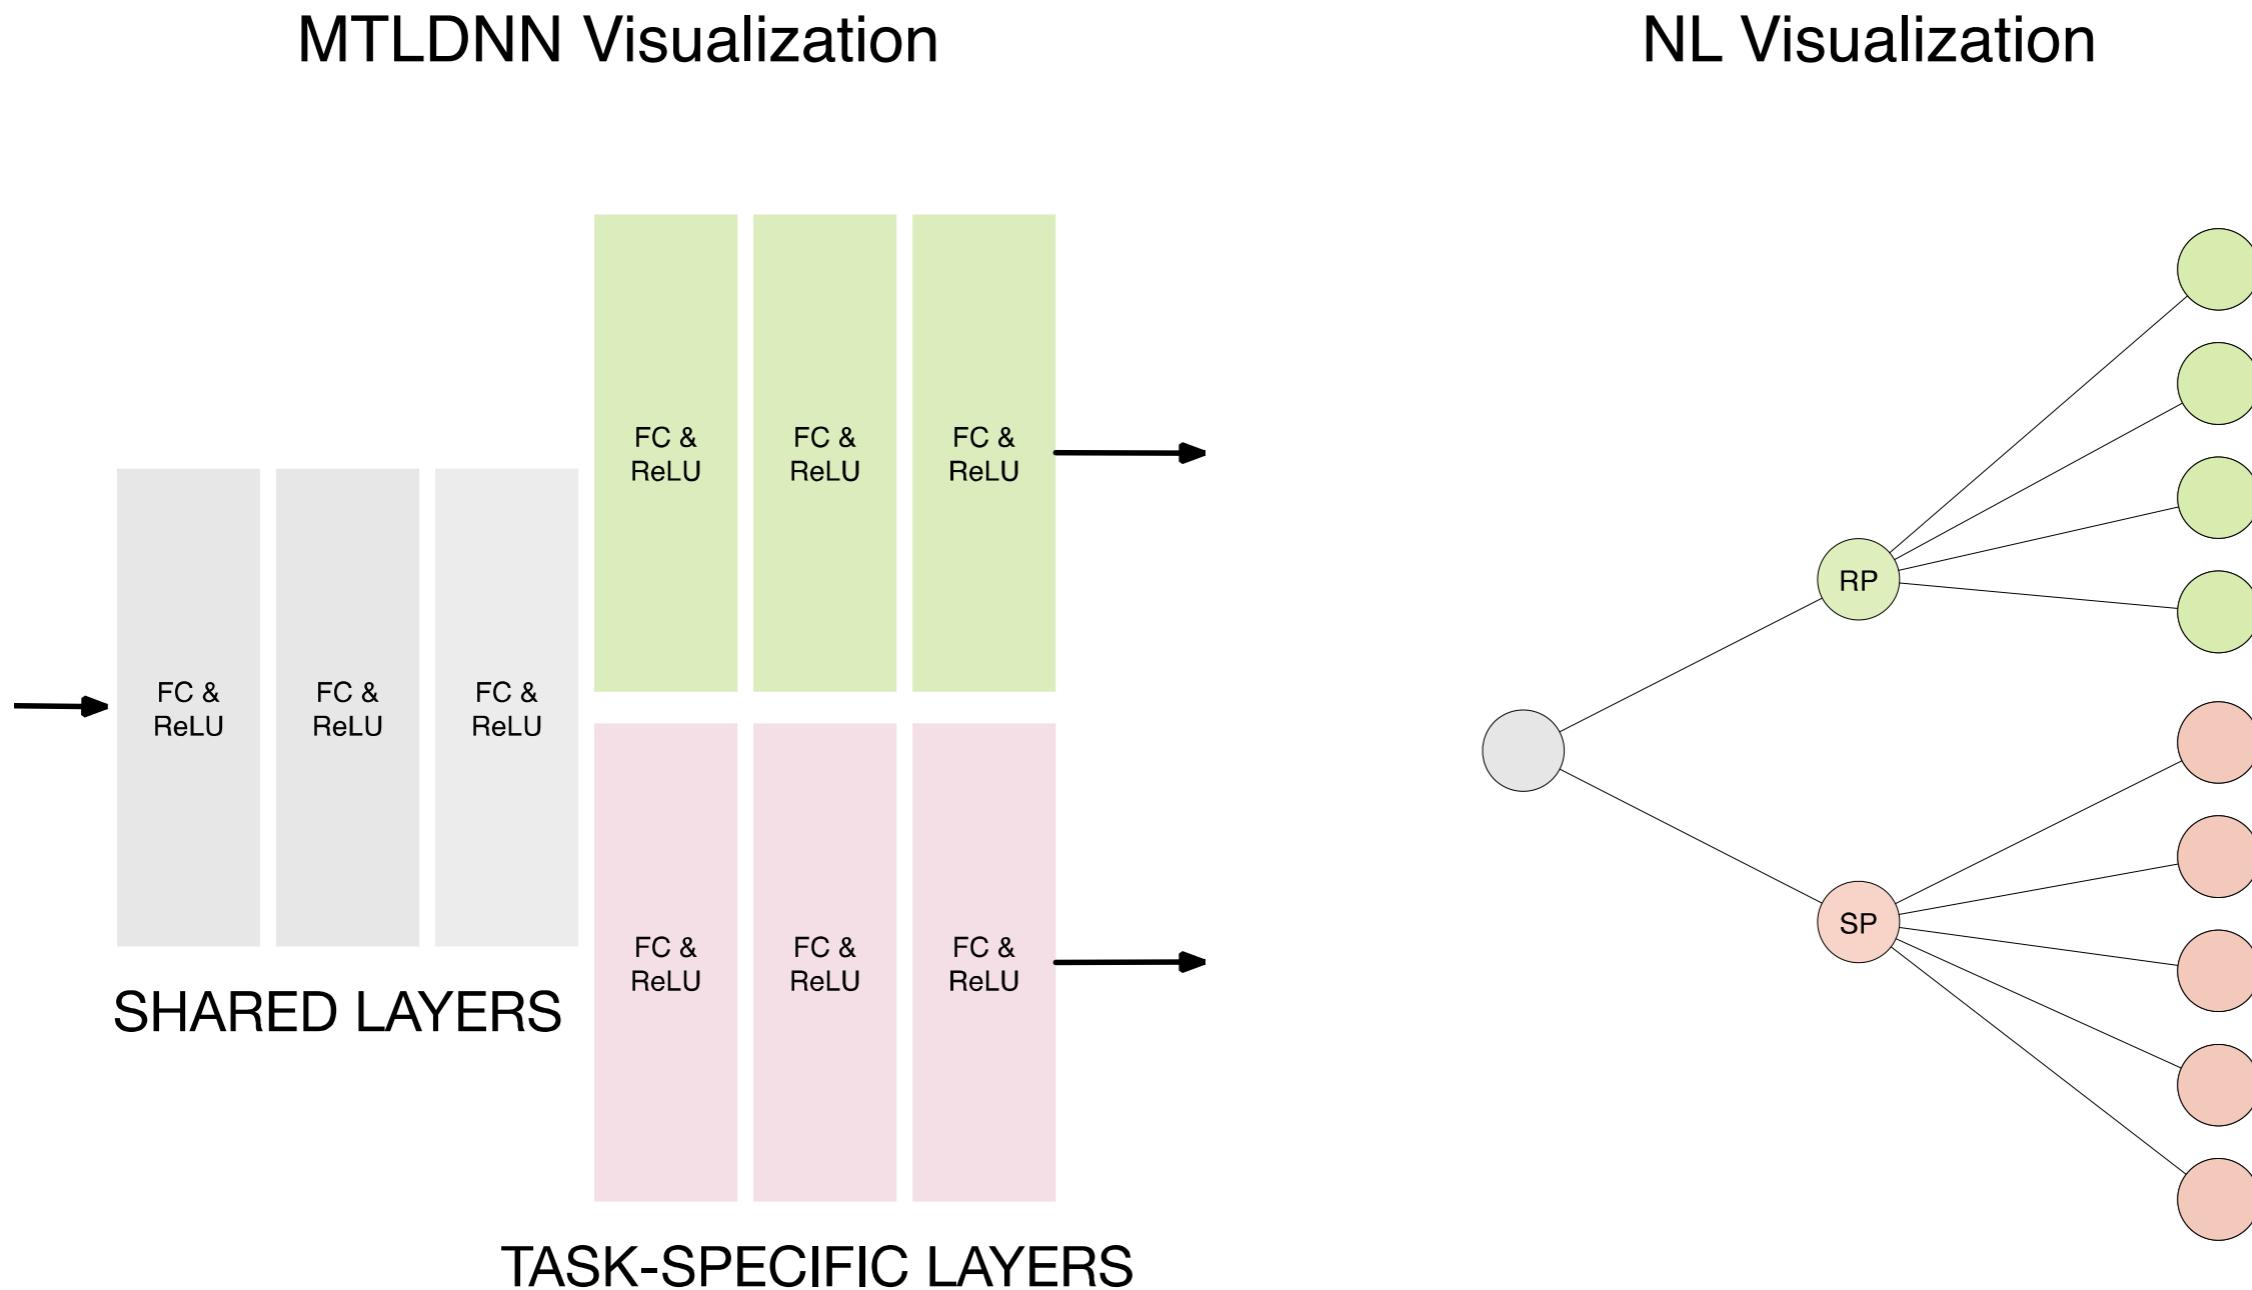
RP+SP as a Classical Question

- Pros and cons of RP and SP (Ben-Akiva et al., 1994; Hausman et al., 1998)
- Joint RP+SP (Ben-Akiva et al., 1994; Hensher and Bradley, 1993; Polydoropoulou and Ben-Akiva, 1994)
- Nested logit model as one classical method (Hensher and Bradley, 1993; Louviere et al. 1999)

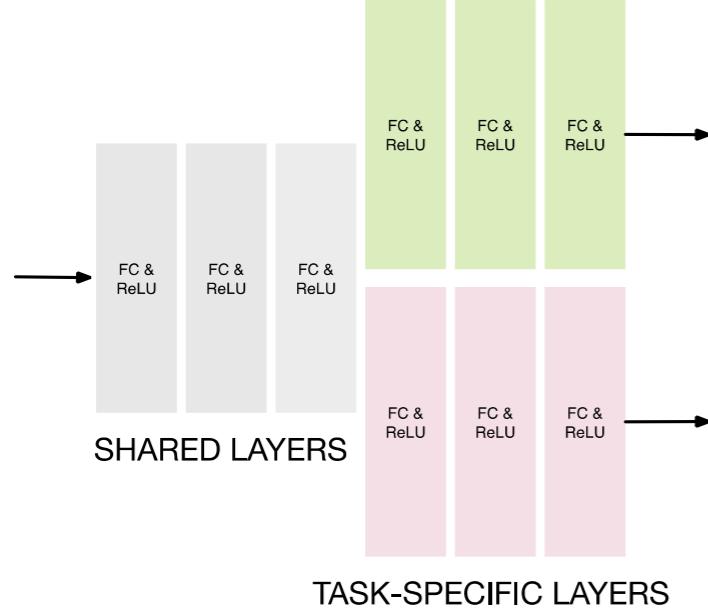
MTLDNNs as a New Method

- A multitask learning perspective
- "Multiple tasks arise naturally..." (Caruana, 1997)
- Wide applications: NLP (Collobert and Weston, 2008; Hashimoto et al. 2016); healthcare drug discovery (Ramsundar et al. 2015); etc.

An Intriguing Question: MTLDNNs and NLs



Formulation of MTLDNNs



Empirical Risk Minimization

$$\min_{w_r, w_s, w_0, T} R(X, Y; w_r, w_s, w_0, T; c_H) \equiv \text{Cross-Entropy Loss}$$

$$\begin{aligned} \min_{w_r, w_s, w_0, T} \Bigg\{ & -\frac{1}{N_r} \sum_{i=1}^{N_r} \sum_{k_r=1}^{K_r} y_{k_r} \log P(y_{k_r, i}; w_r, w_0; c_H) \\ & - \frac{\lambda_0}{N_s} \sum_{t=1}^{N_s} \sum_{k_s=1}^{K_s} y_{k_s} \log P(y_{k_s, t}; w_r, w_0, T; c_H) \\ & + [\lambda_1 \|w_0\|_2^2 + \lambda_2 \|w_s\|_2^2 + \lambda_3 \|w_s - w_r\|_2^2] \Bigg\} \end{aligned}$$

2 Regularizations: Scale Controls

MTLDNNs are More Generic than NLs.

MTLDNNs

1. Automatic feature learning
2. “Soft” constraints to describe the similarities between RP and SP
 - Architectural design (e.g. # of shared vs. task-specific layers)
 - Regularizations (e.g. λ_3)
3. Possible to specify the model with a massive number of tasks

NLs

1. Handcrafted feature learning
2. “Hard” constraints to describe the similarities between RP and SP
 - Shared vs. task-specific parameters (e.g. β_r vs. β_s)
3. Hard to specify the model with a massive number of tasks

Deep Learning for Urban Mobility

Part I: Introduction and DNN Basics

Sep 10 Introduction: Deep Learning Meets Transportation

Part II: Passenger and Traffic Flow Prediction

Sep 17 Demand prediction and Deep learning basics

Sep 24 Advanced modeling techniques: ConvLSTM, Attention, individualized predictions

Oct 1 Guest Lecture Prof. Justin Dauwels

Oct 8 Advanced Applications: Generative models (VAE), graph embeddings,

Part III: DNN and Demand Analysis

Oct 22 DNN and Discrete Choice 1

Oct 29 DNN and Discrete Choice 2

Nov 5 DNN and Prospect Theory

Part IV: Reinforcement Learning and Control in Transportation

Nov 12 RL and Control Part 1

Nov 19 RL and Control Part 2

Nov 26 Guest Lecture Prof. Cathy Wu

Part V: FAT and Summary

Dec 3 Fairness, Accountability and Transparency

Dec 10 Student Presentation

Style of the Class

- Balance: transportation ~ machine learning
- Seminar: a bit ad hoc, a survey
- Credit: 6, p/d/f
- Expectation of students
 - Class participation
 - Term project

Applications

ML for transit

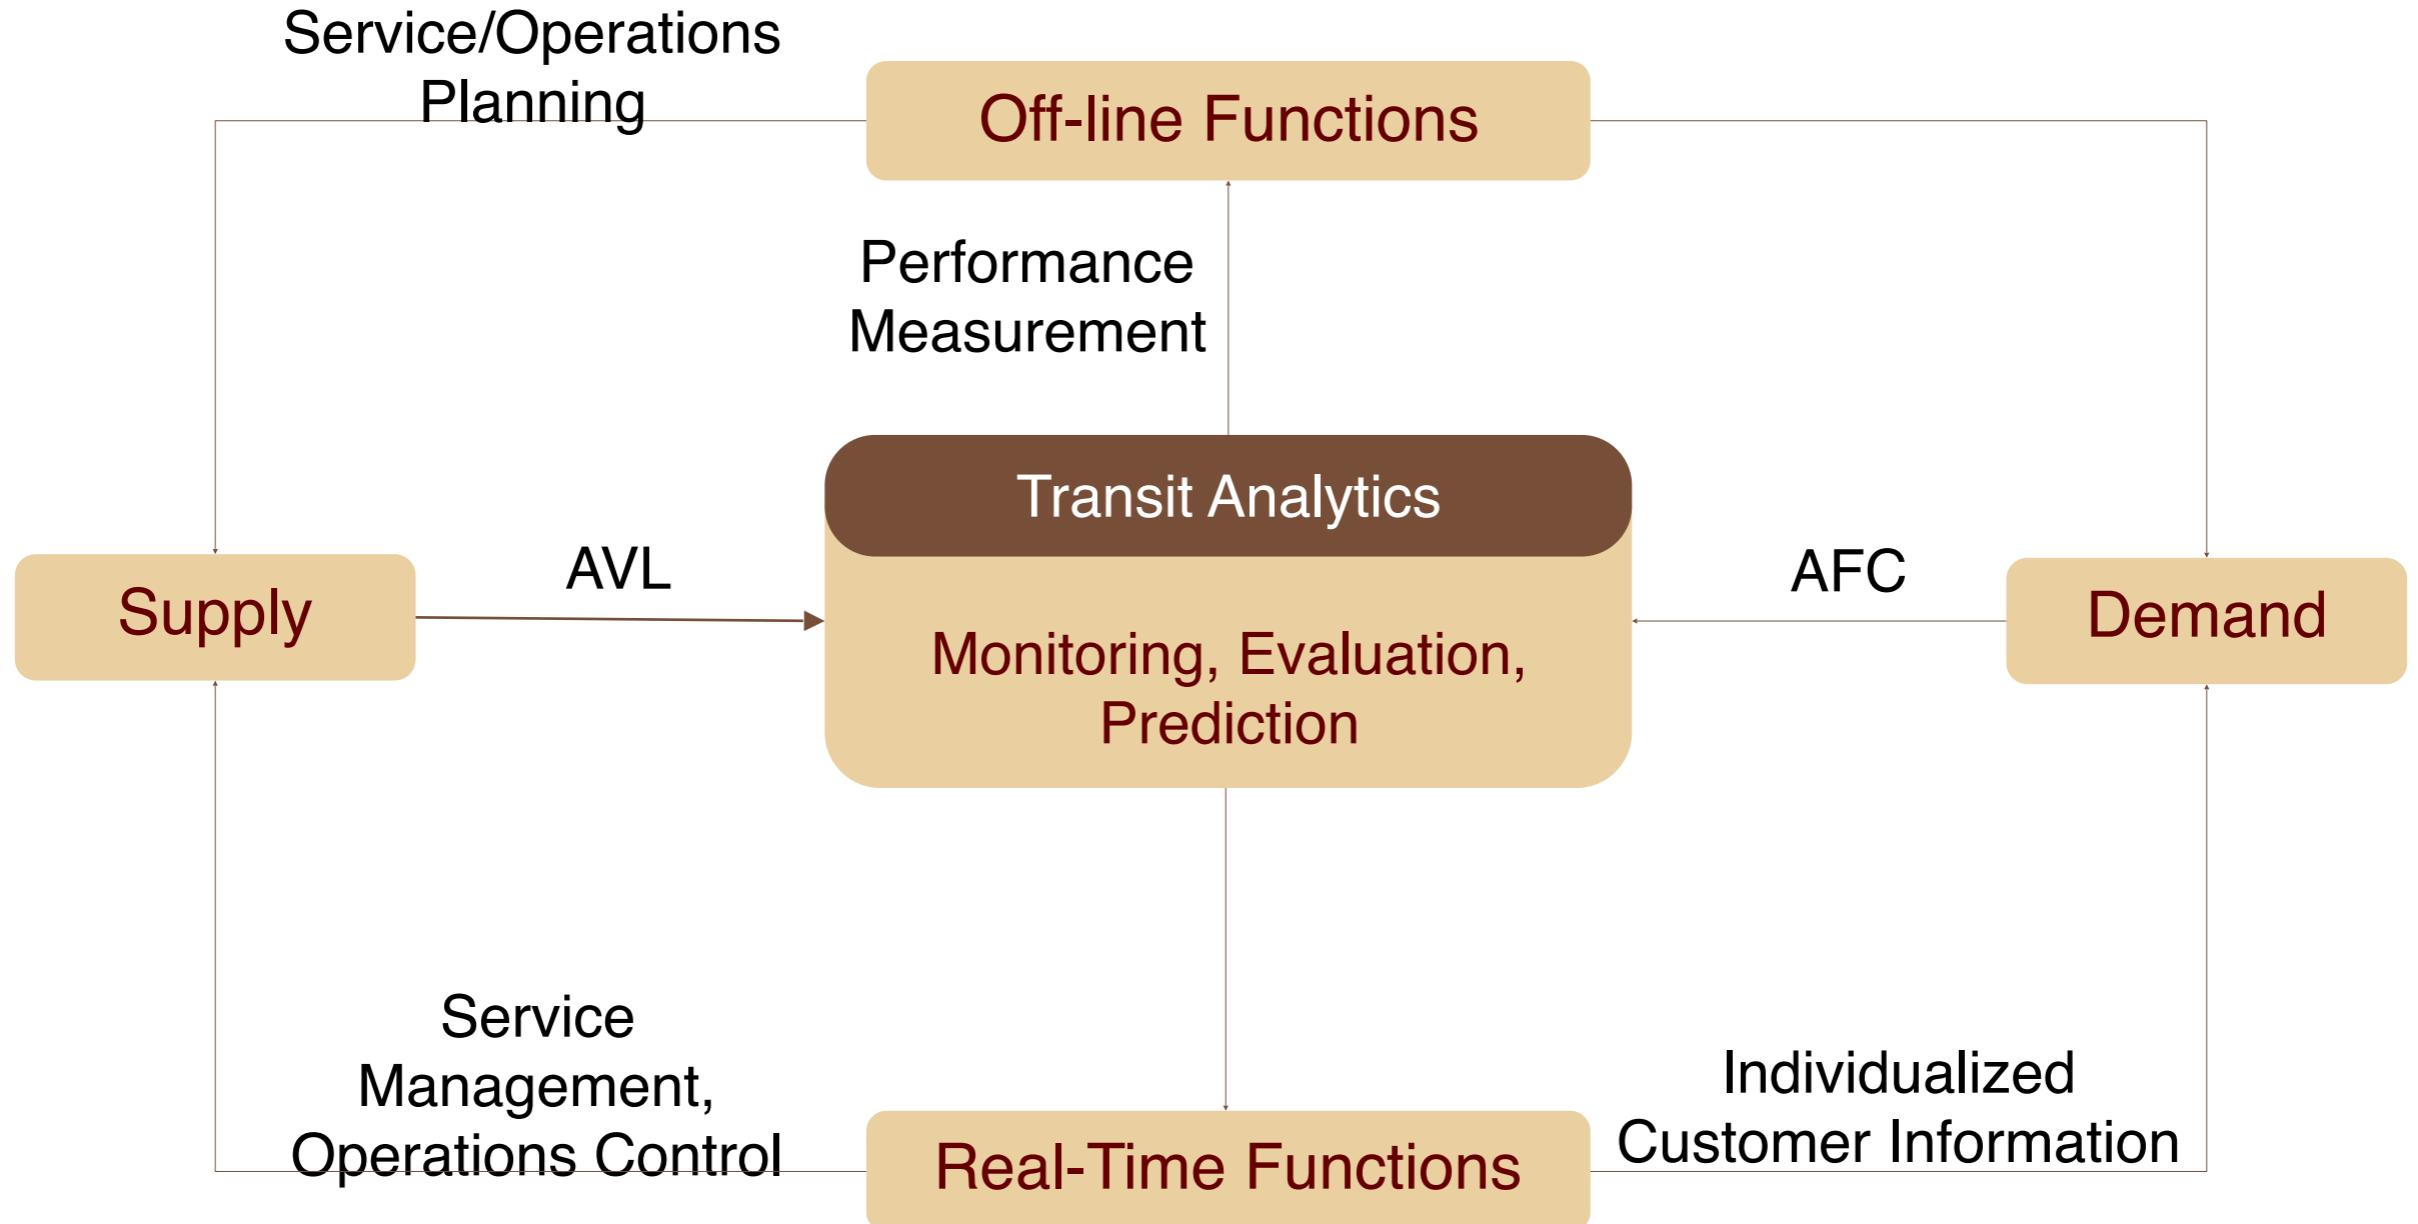


Challenges

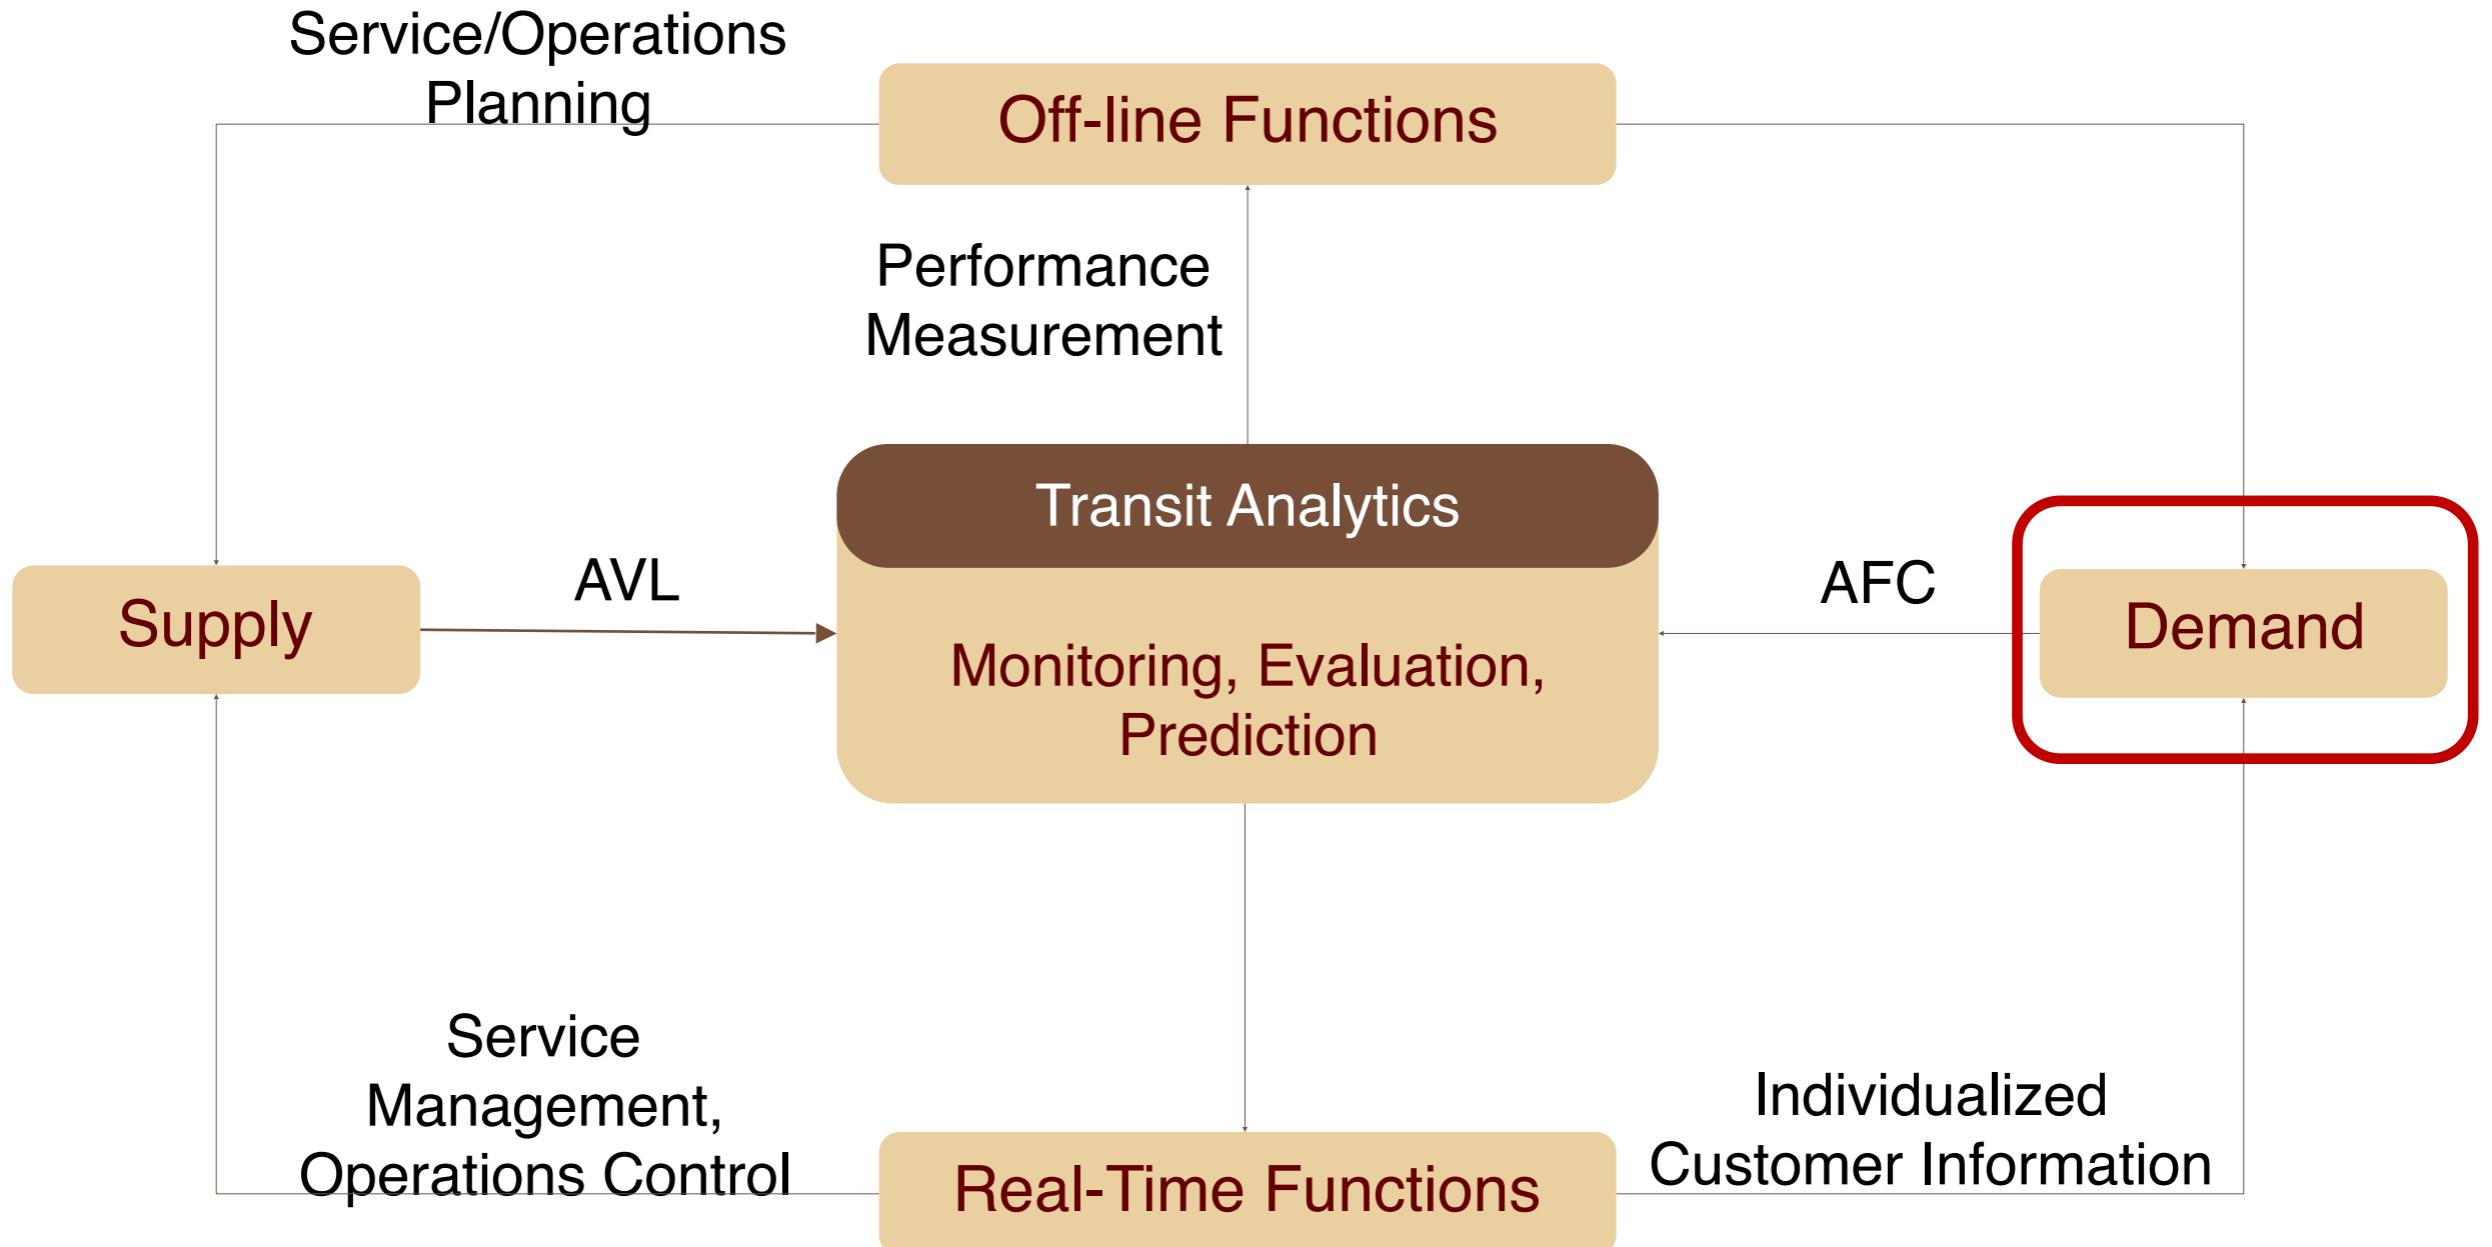
- Near capacity operations
- Crowding (train, platform)
- Passengers left behind
- Safety concerns
- Disruptions



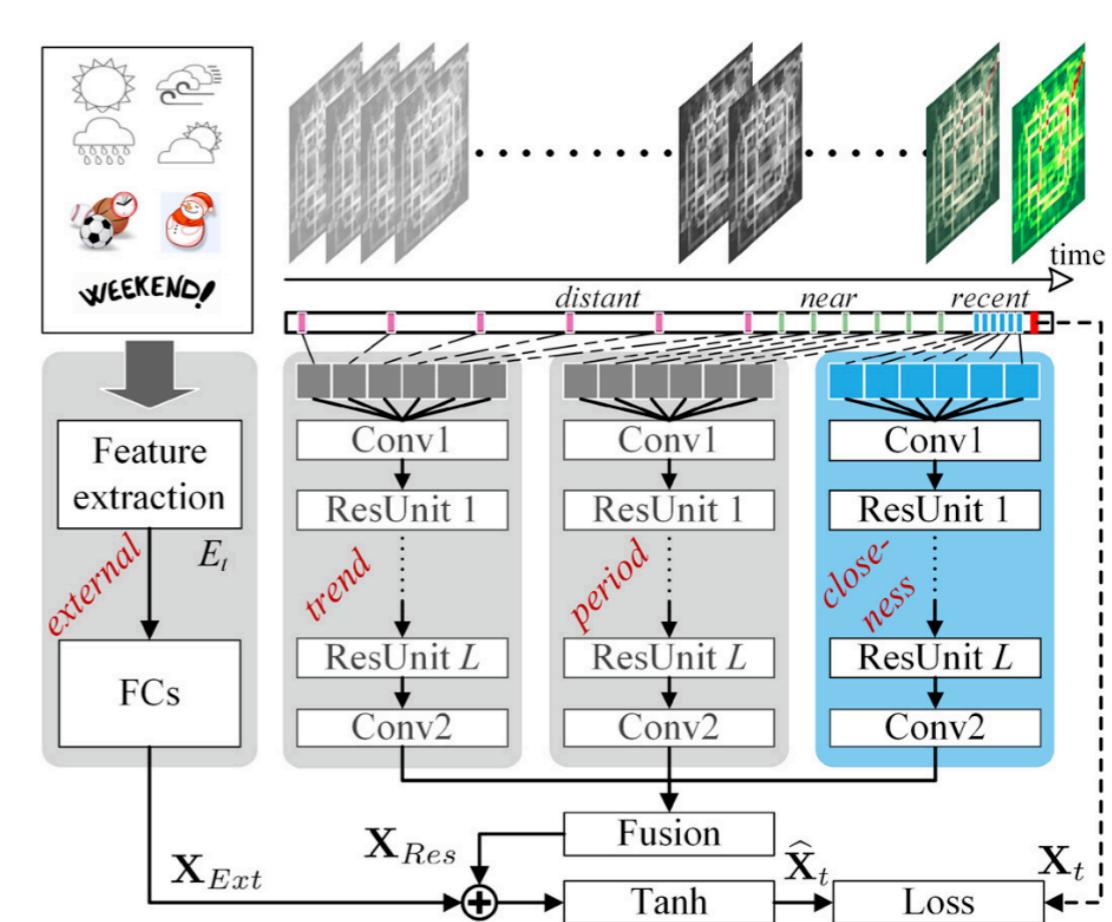
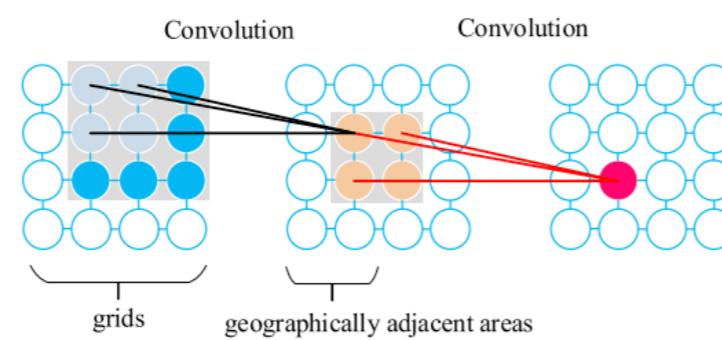
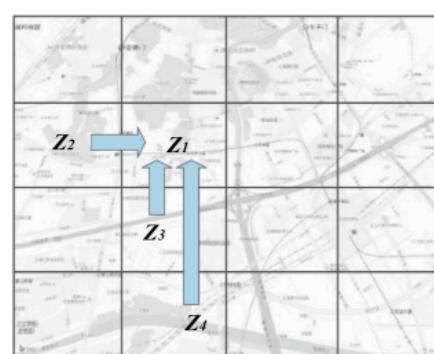
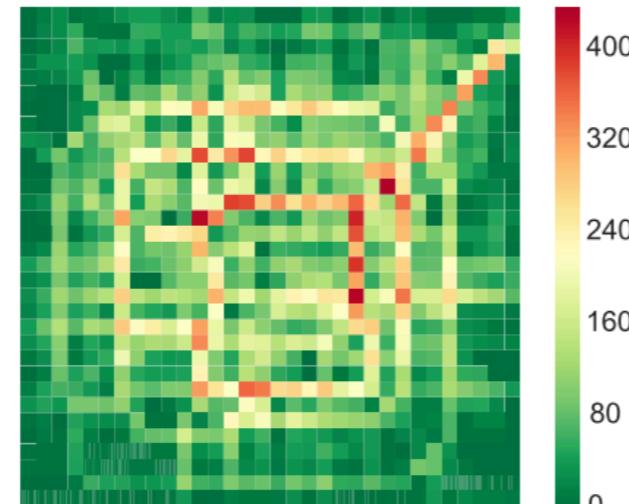
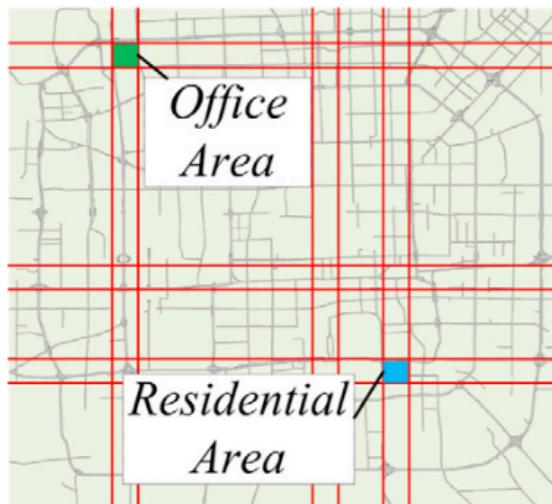
Overall framework



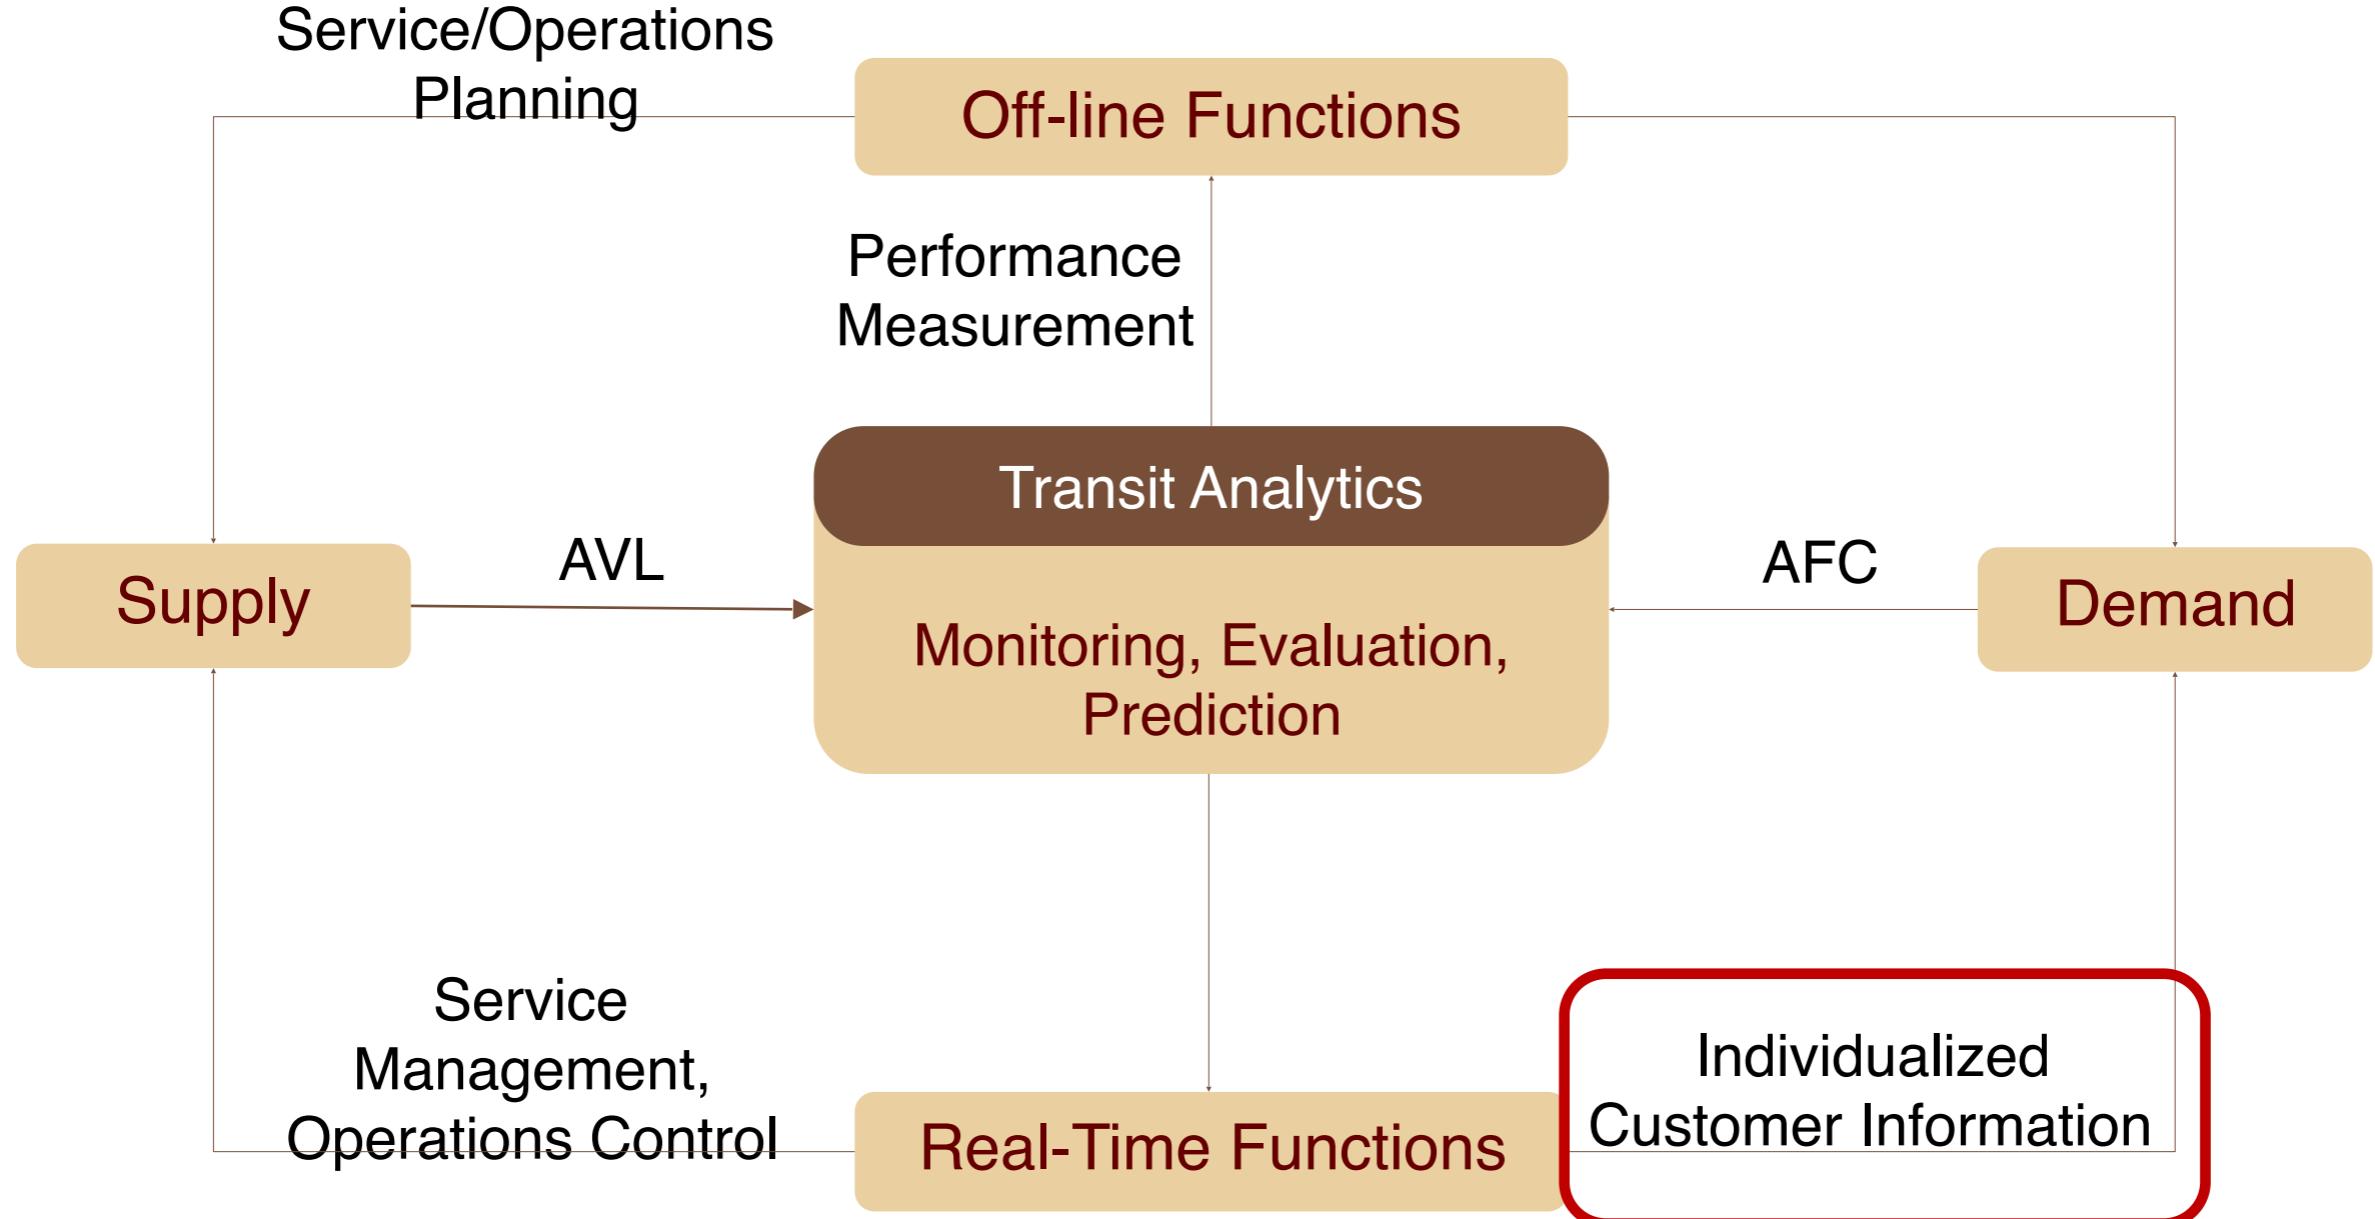
Overall framework



Aggregate demand prediction

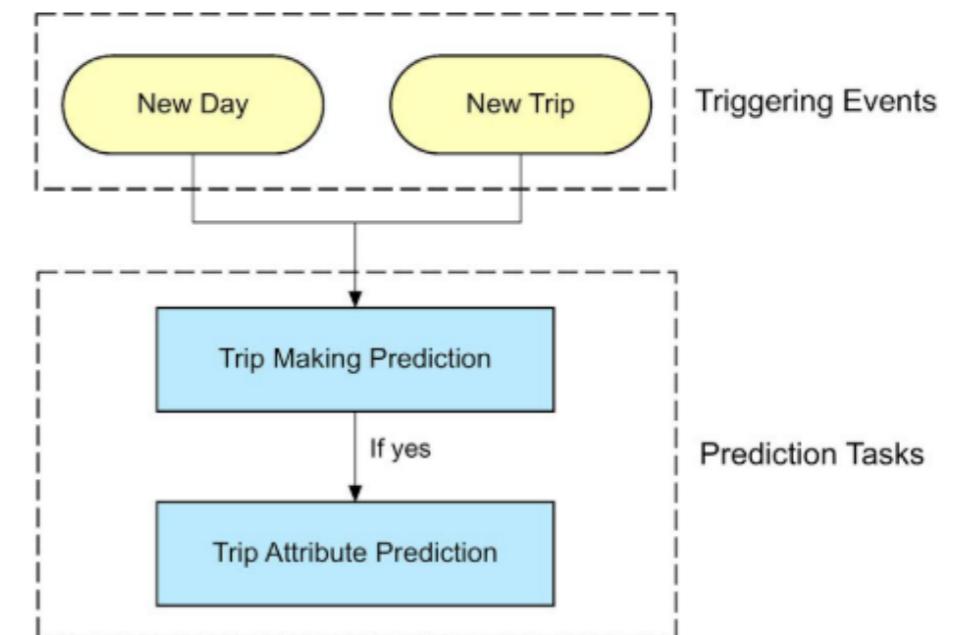
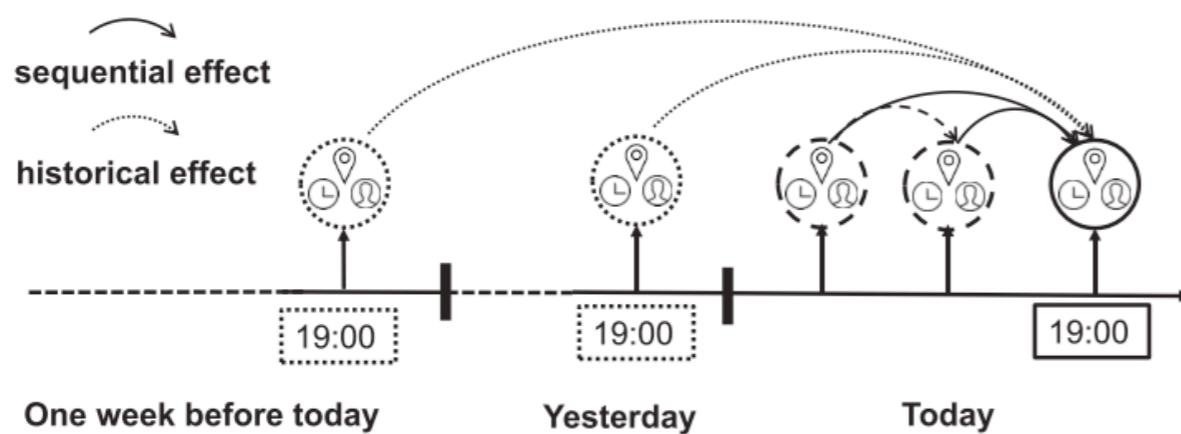


Overall framework



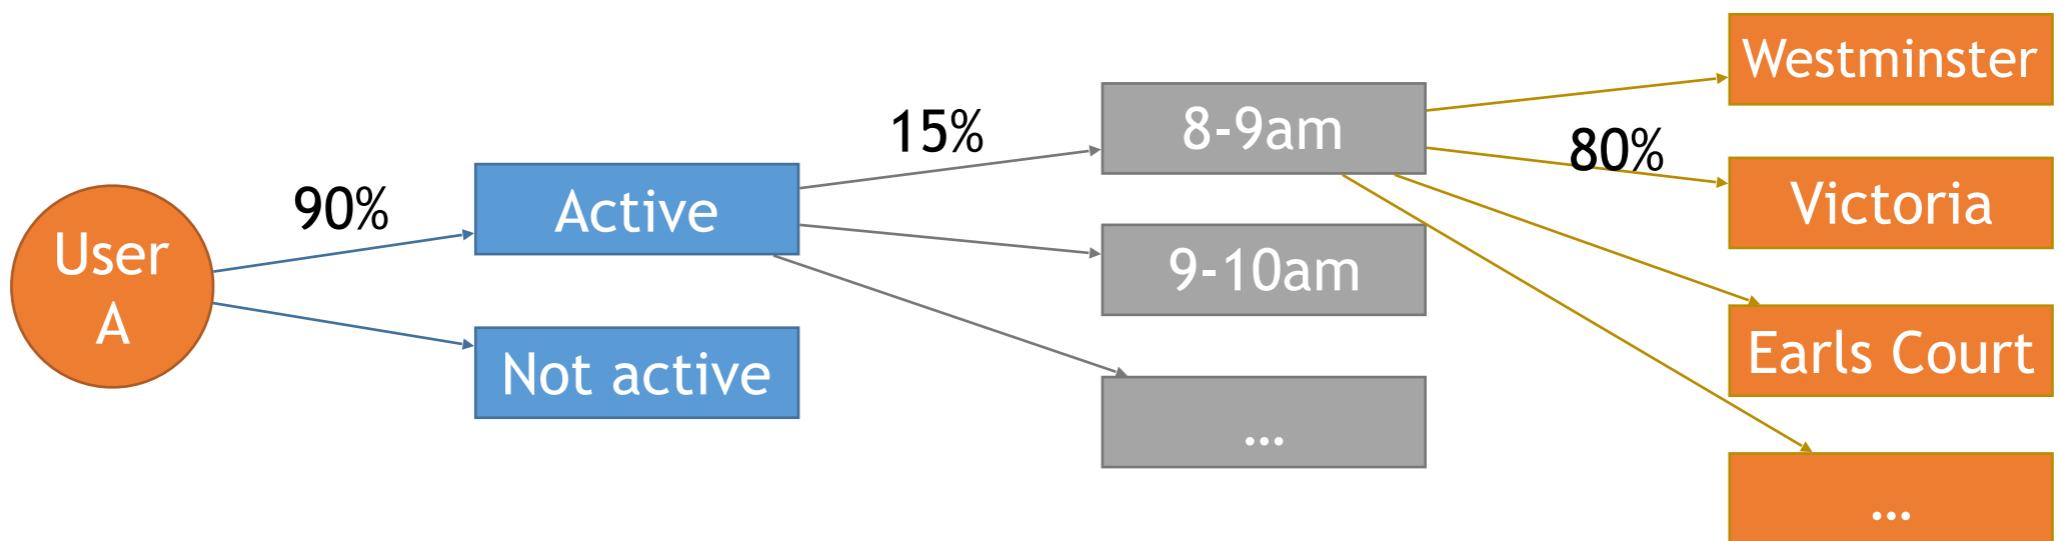
Individual mobility prediction

- When an incidence/delay occurs, send travel alerts to passengers who are likely to be affected



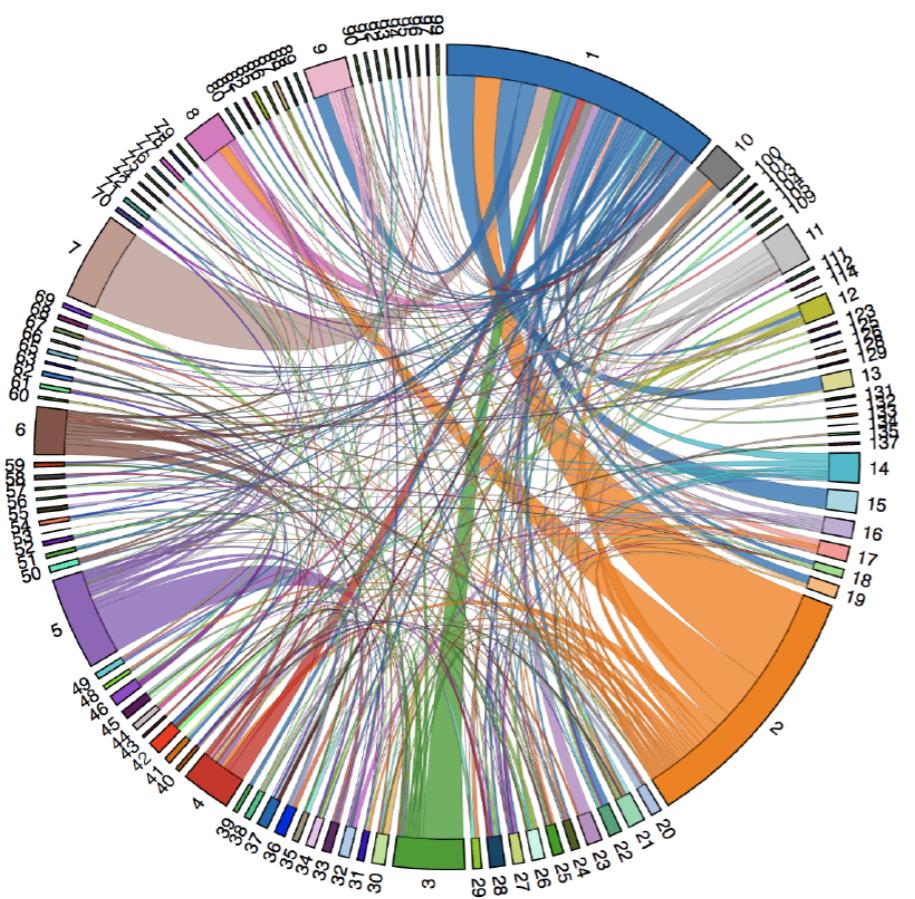
Individual mobility prediction

- Example: at 8 am, an incident occurs in Victoria Station, and will likely affect the service for the next hour

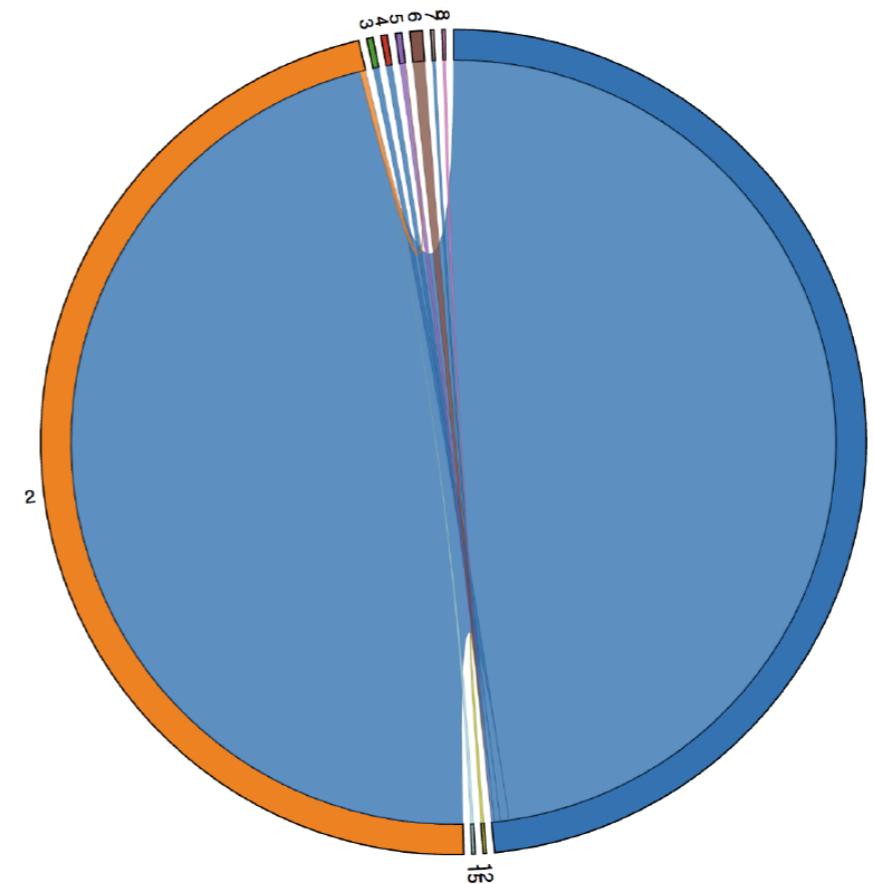


Individual mobility prediction

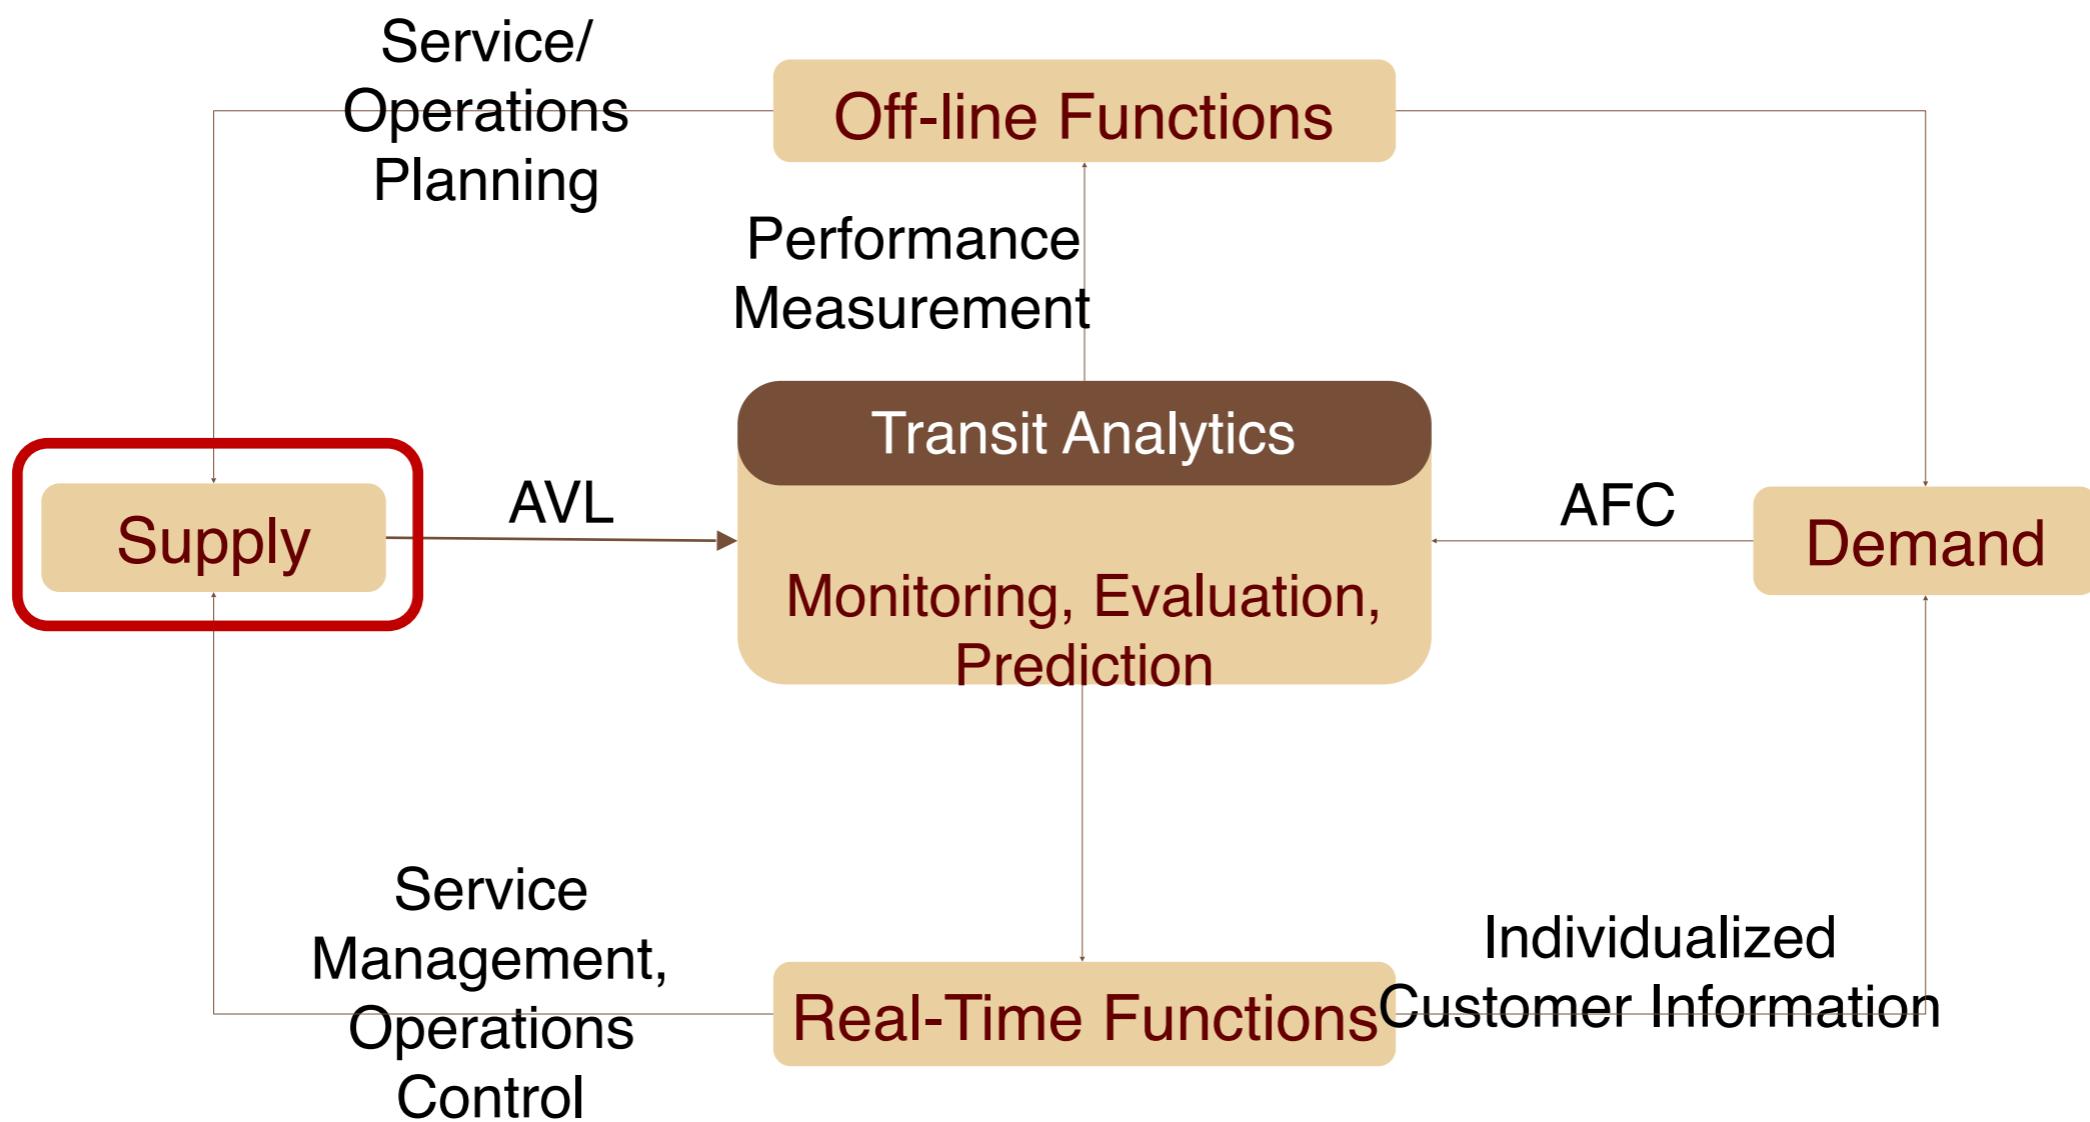
Individual 1



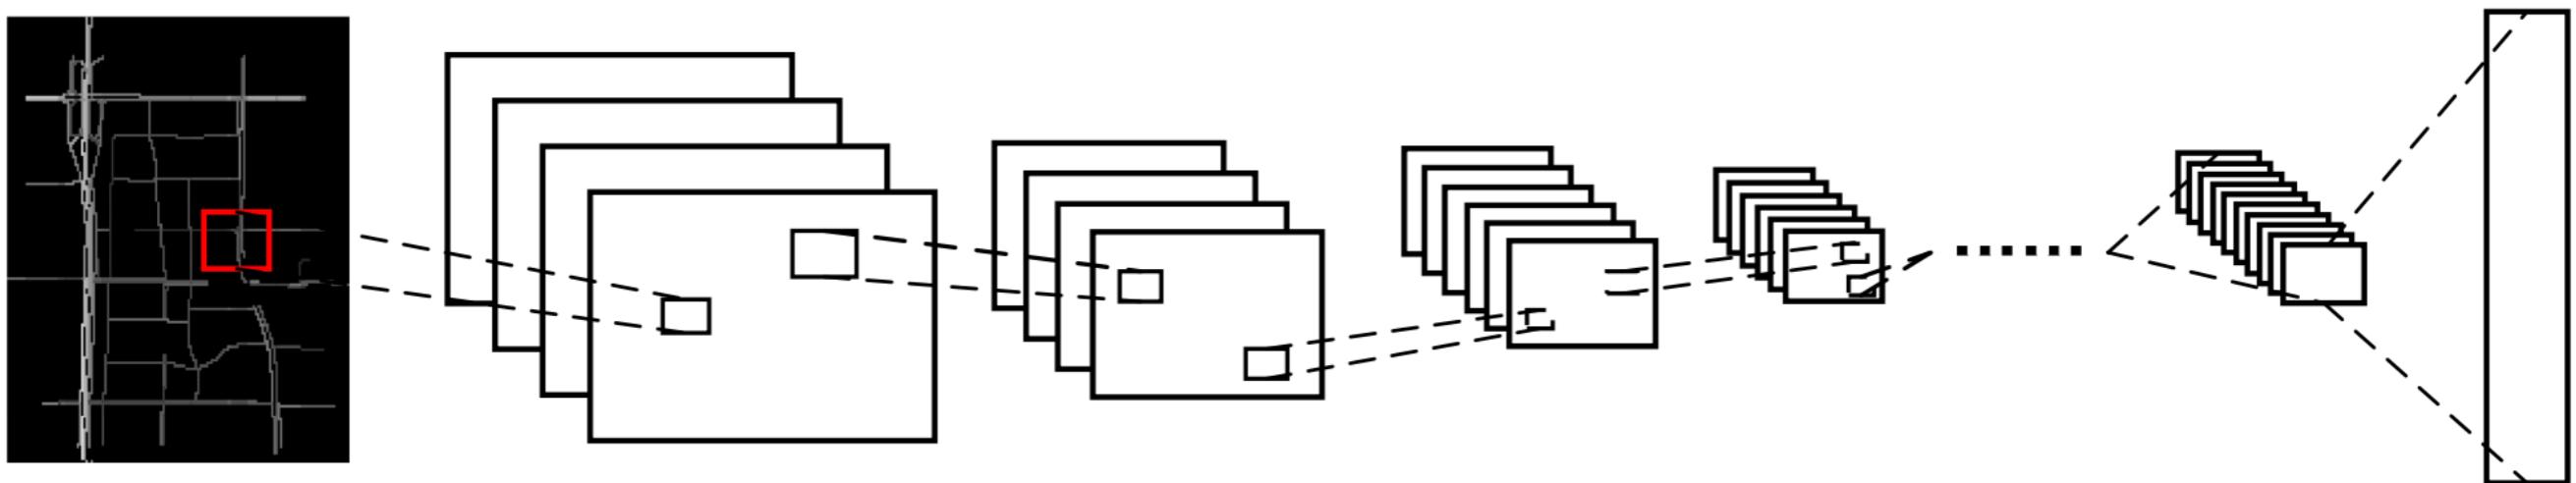
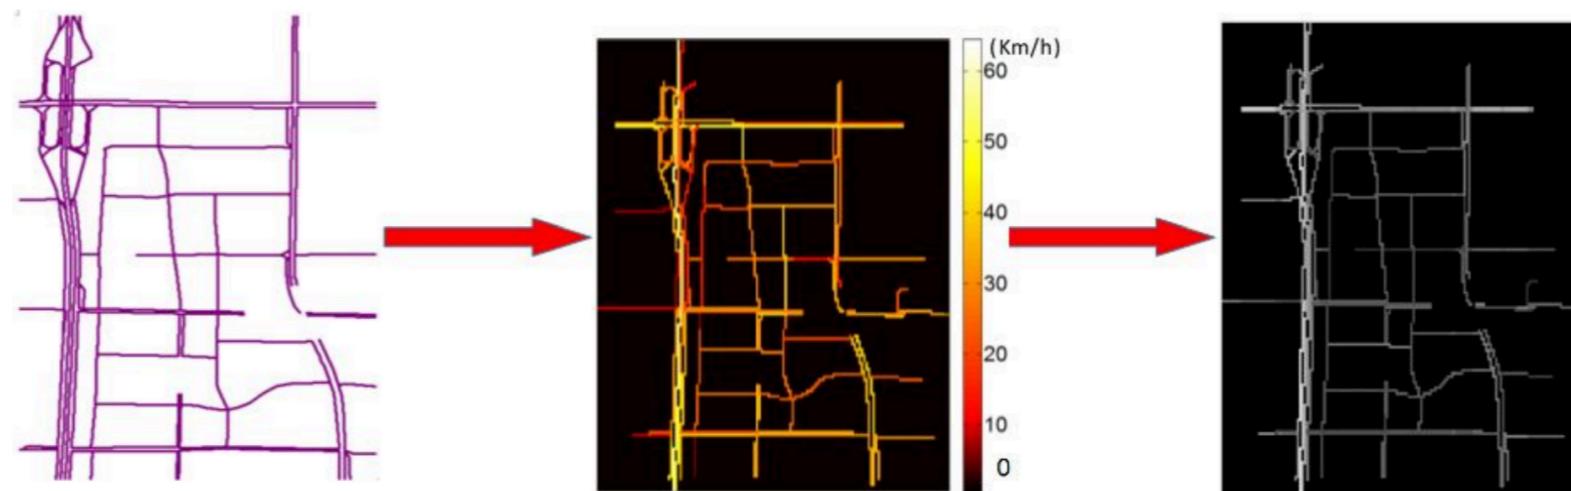
Individual 2



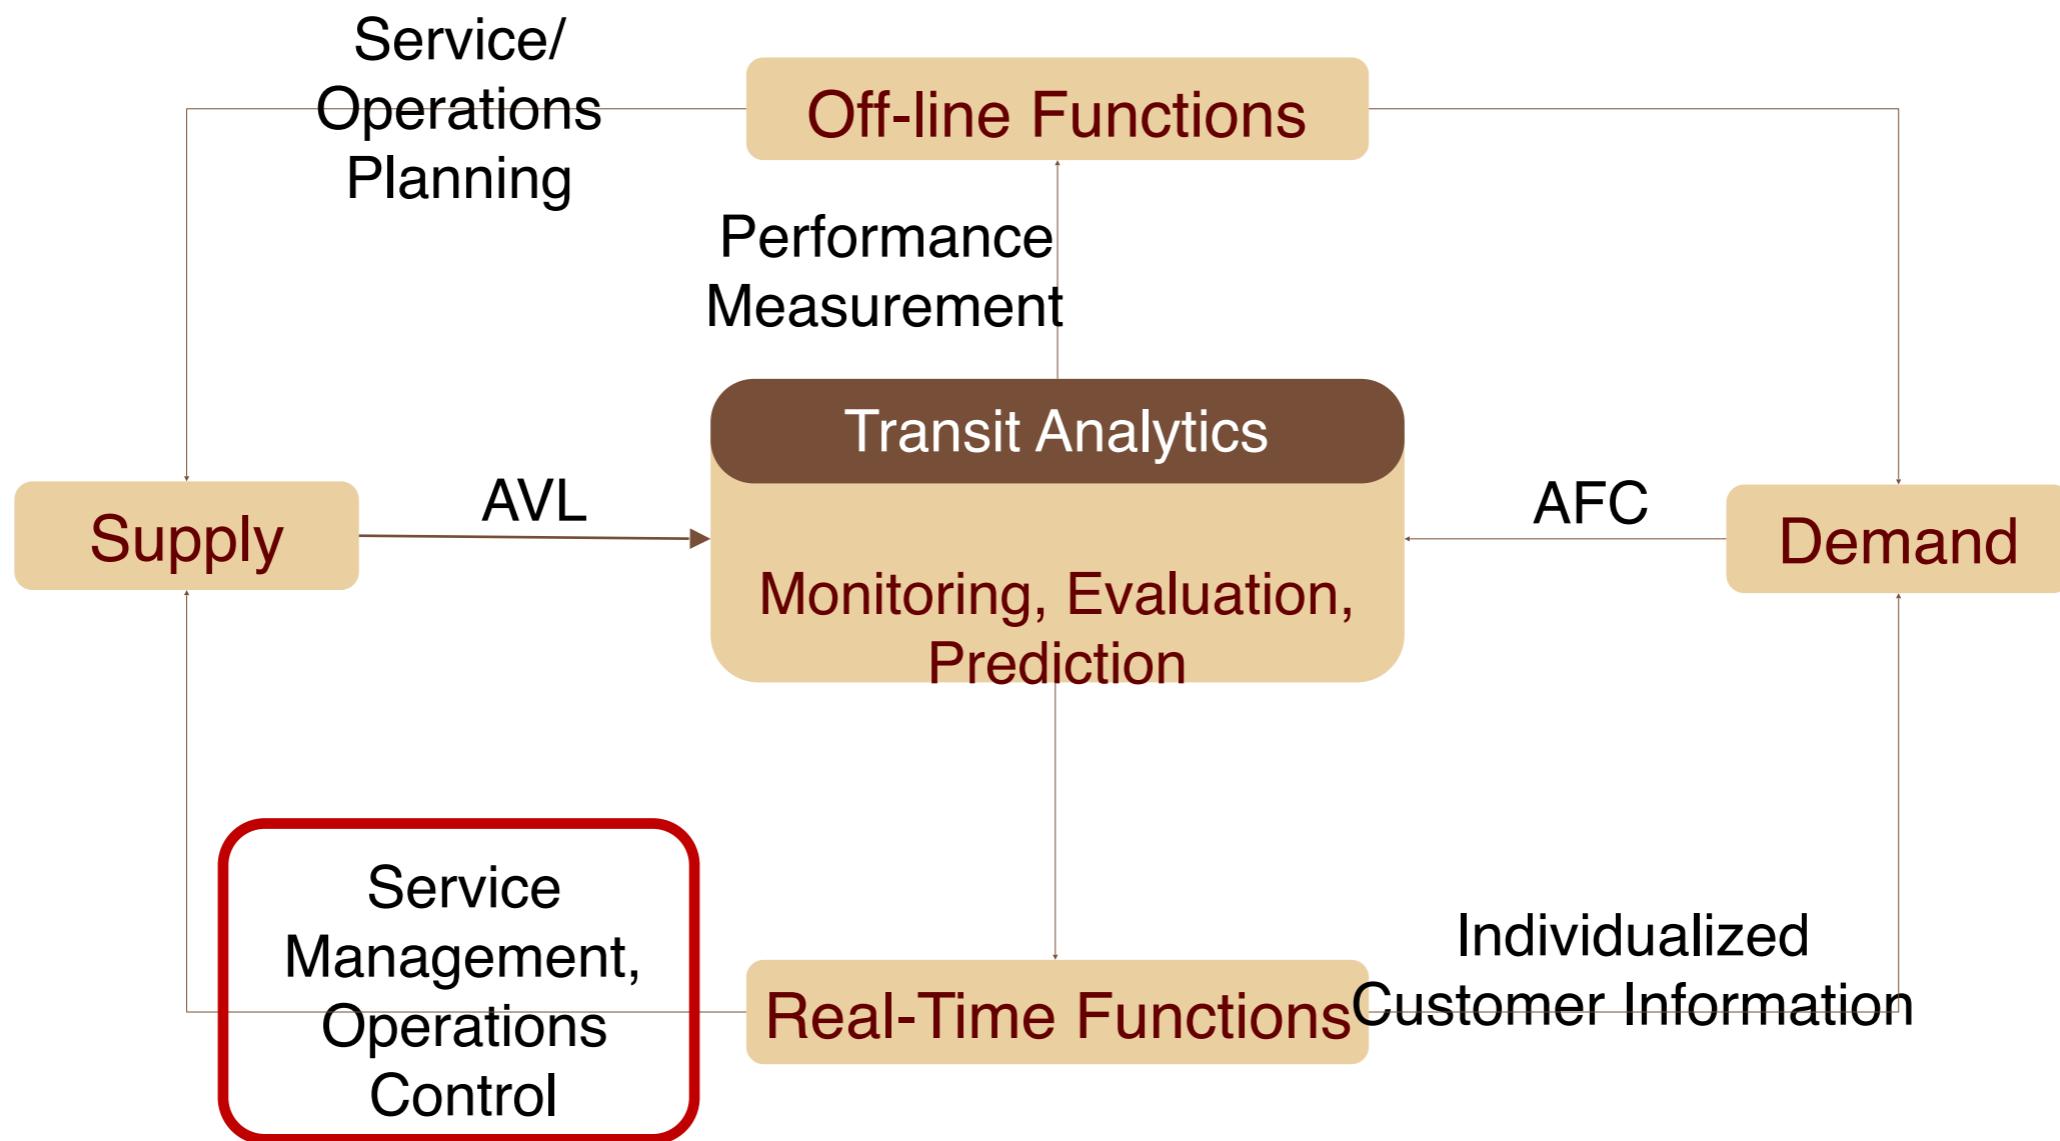
Overall framework



Traffic prediction



Overall framework



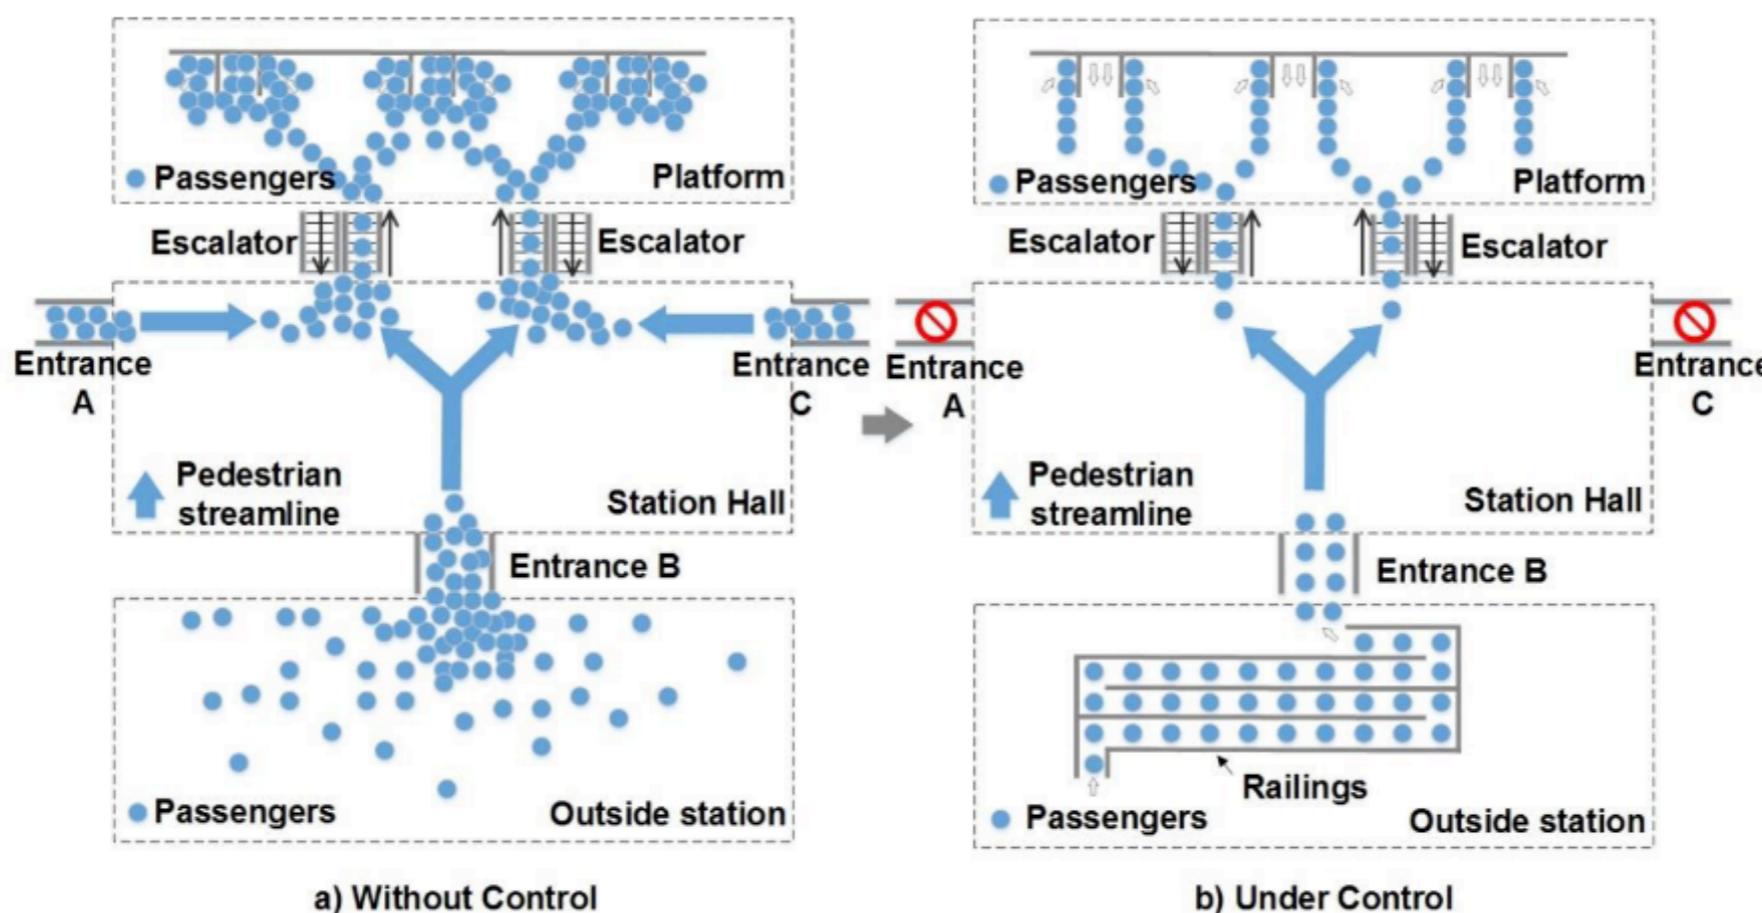
Opportunities

- Proactive crowd management

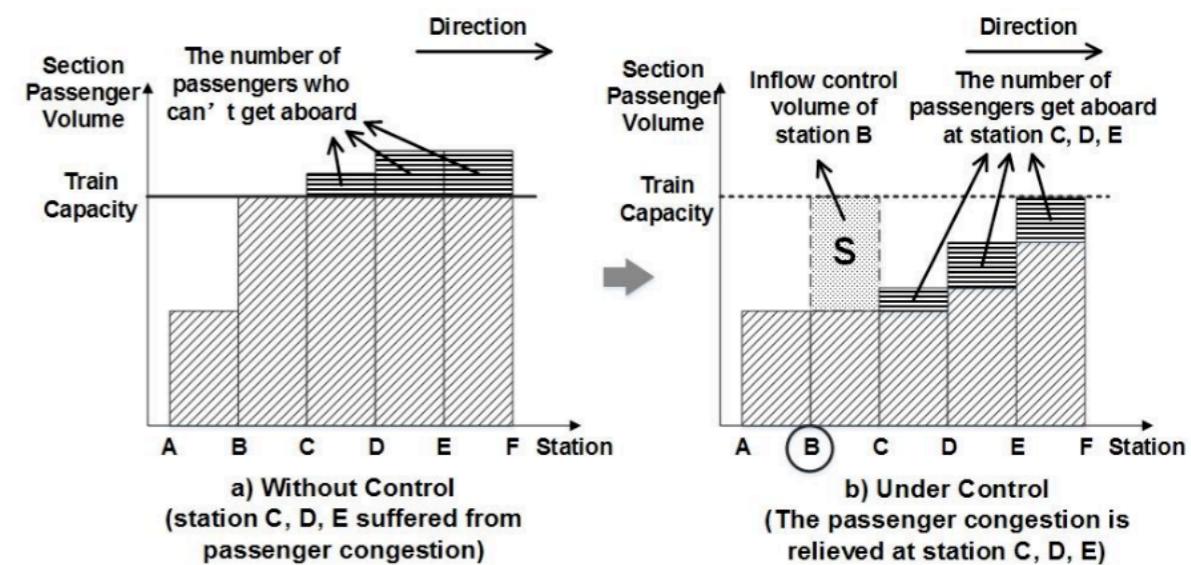
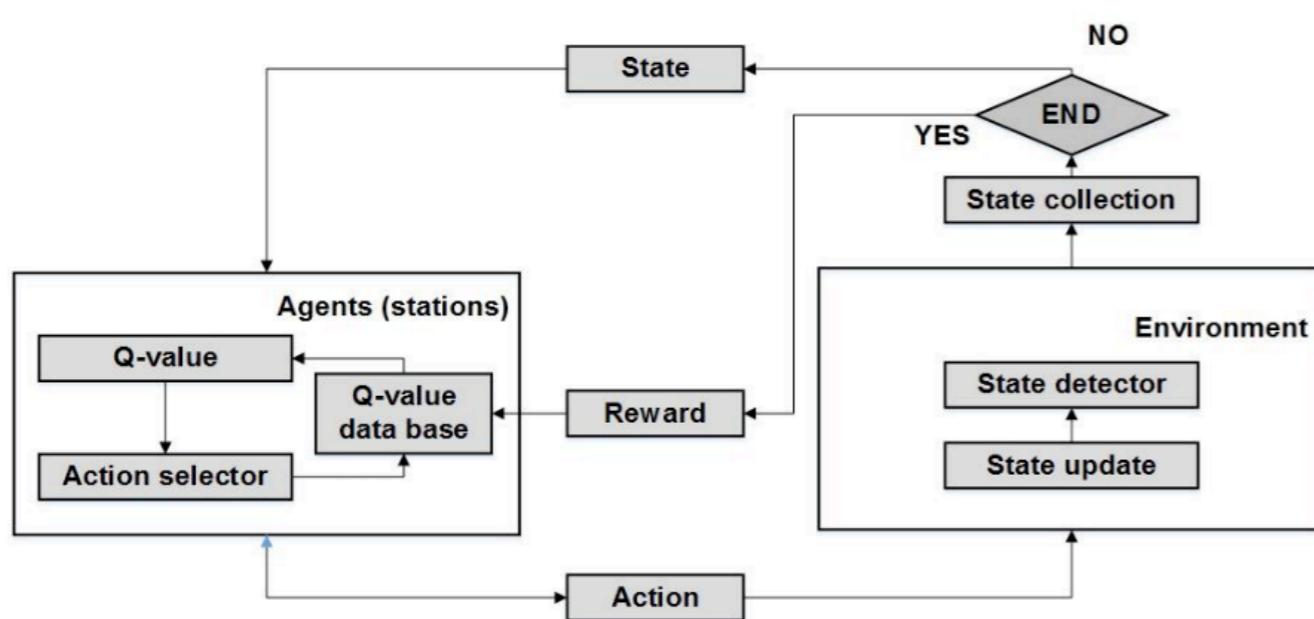


Proactive crowd management

- Example: Beijing Metro



Proactive crowd management



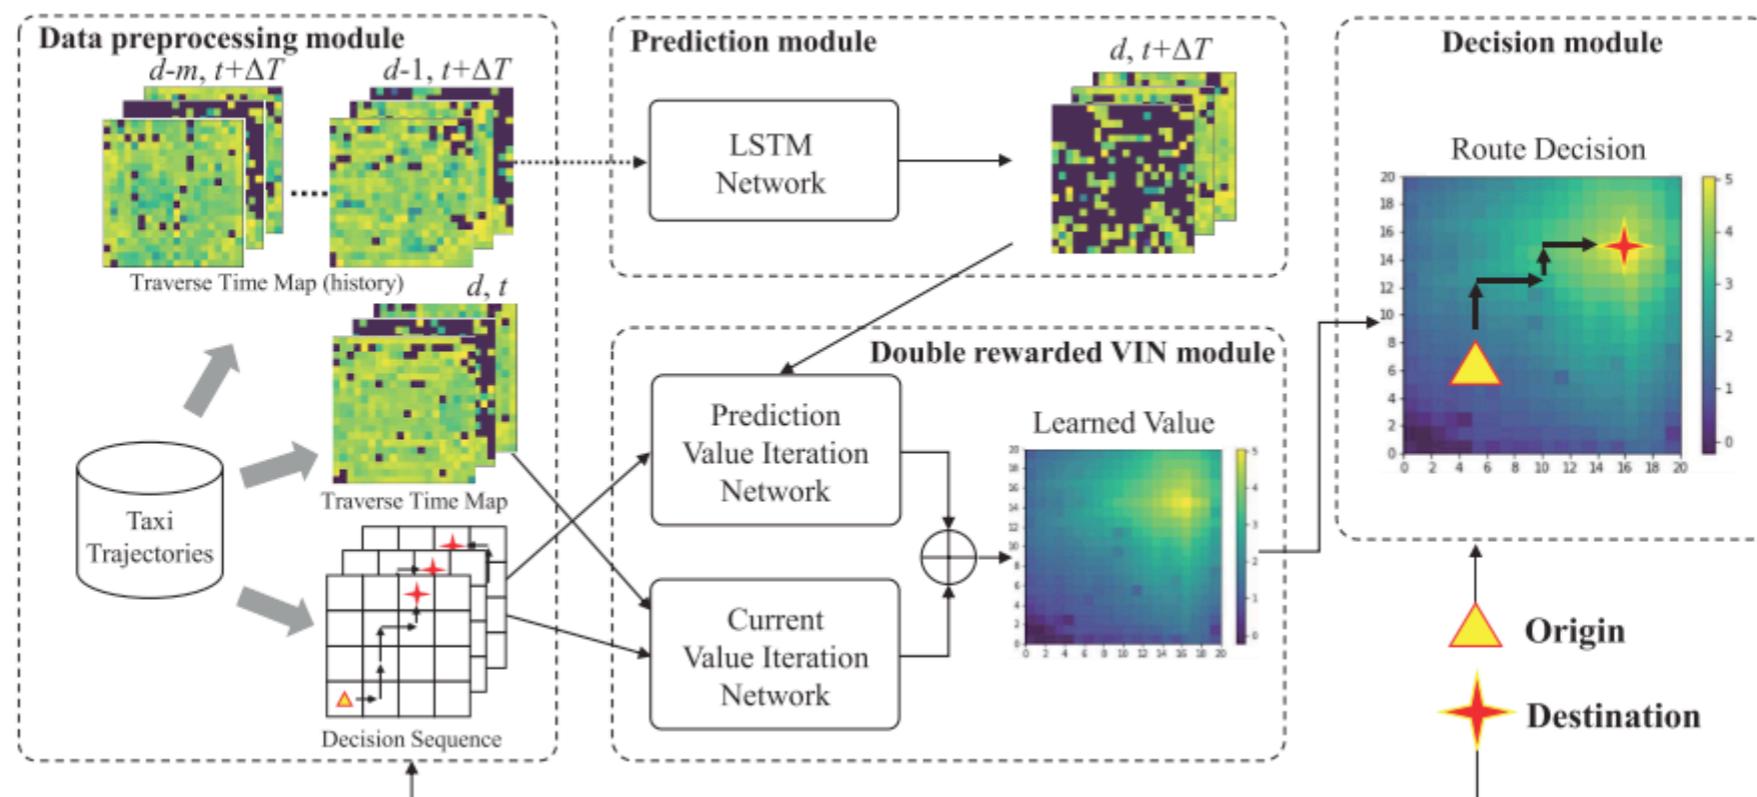
ML for ride hailing
services

Ride hailing systems

- Short-term demand forecasting is crucial for operation
- Examples: Uber, Didi, and Lyft
- Applications:
 - Dispatching efficiency
 - Driver to zone assignment
 - Adaptive pricing
 - Rebalancing
 - Passenger information

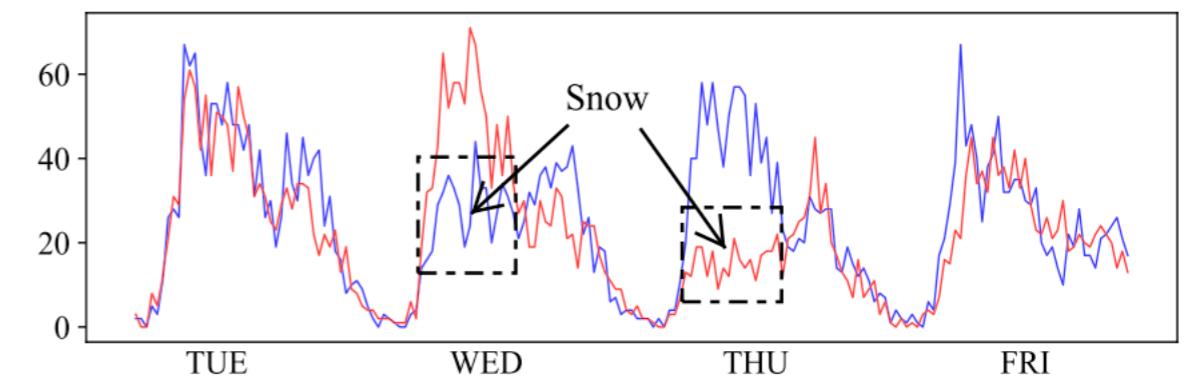
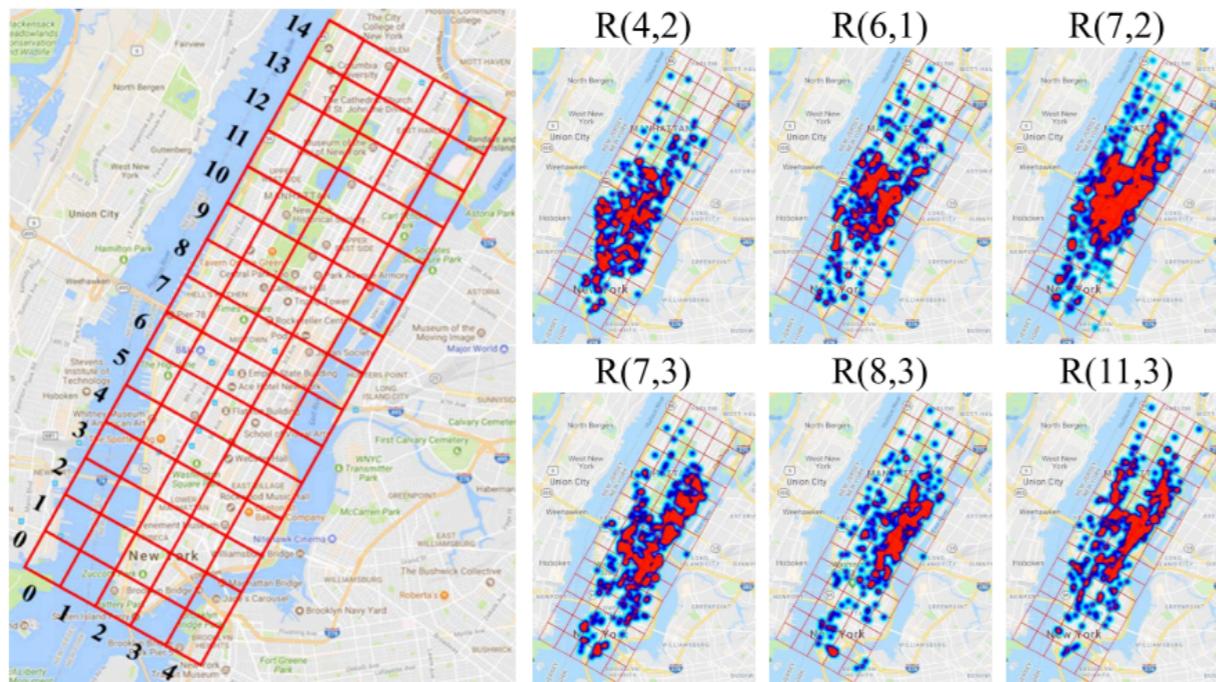
Ride hailing systems with AV

- RL for centrally controlled taxi operation



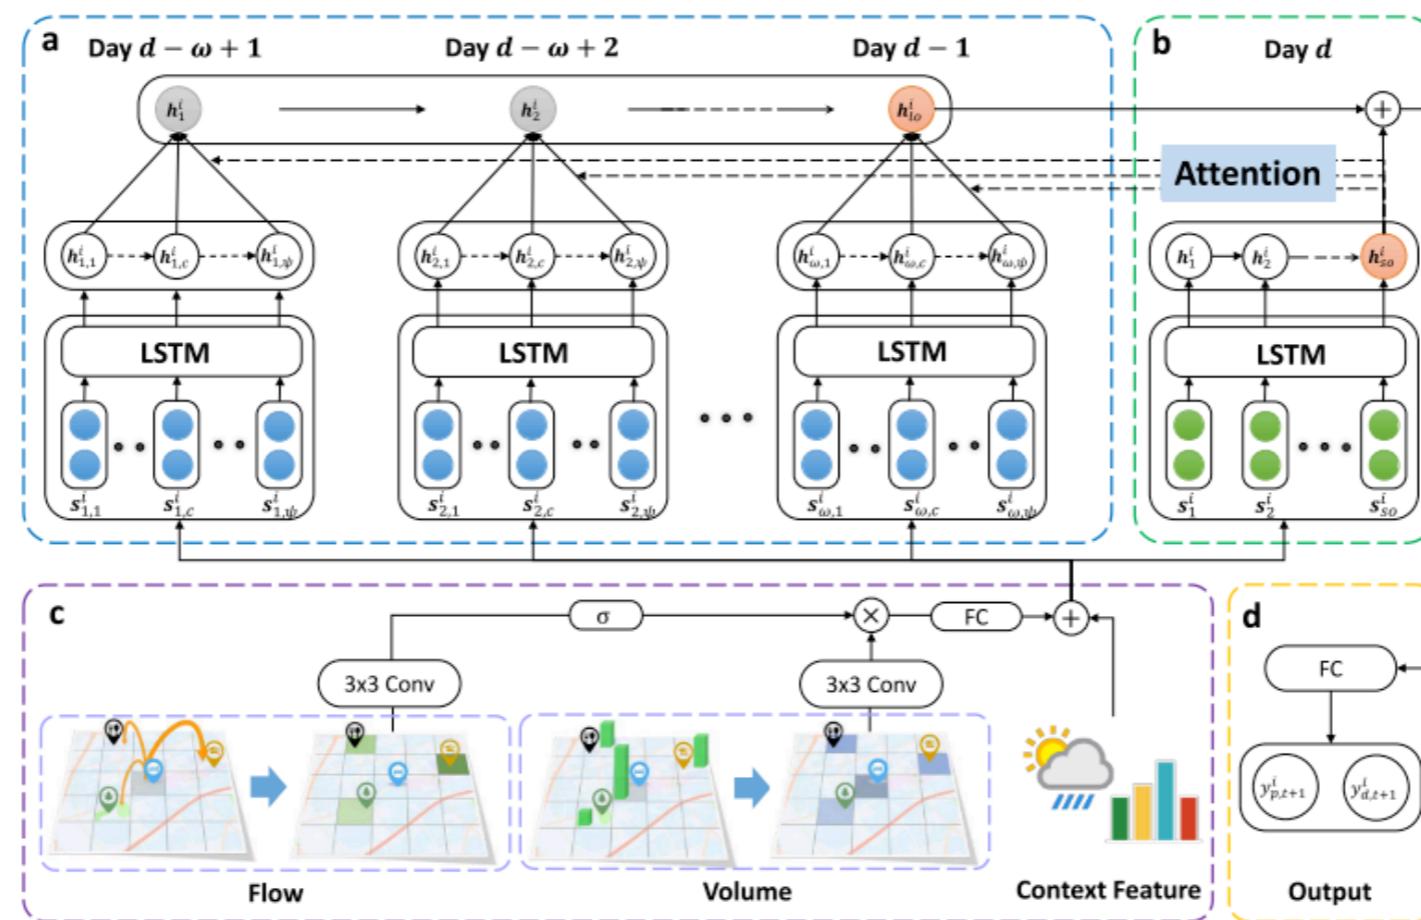
Large-scale Demand prediction

- Origin-destination demand for taxi in NYC



Large-scale Demand prediction

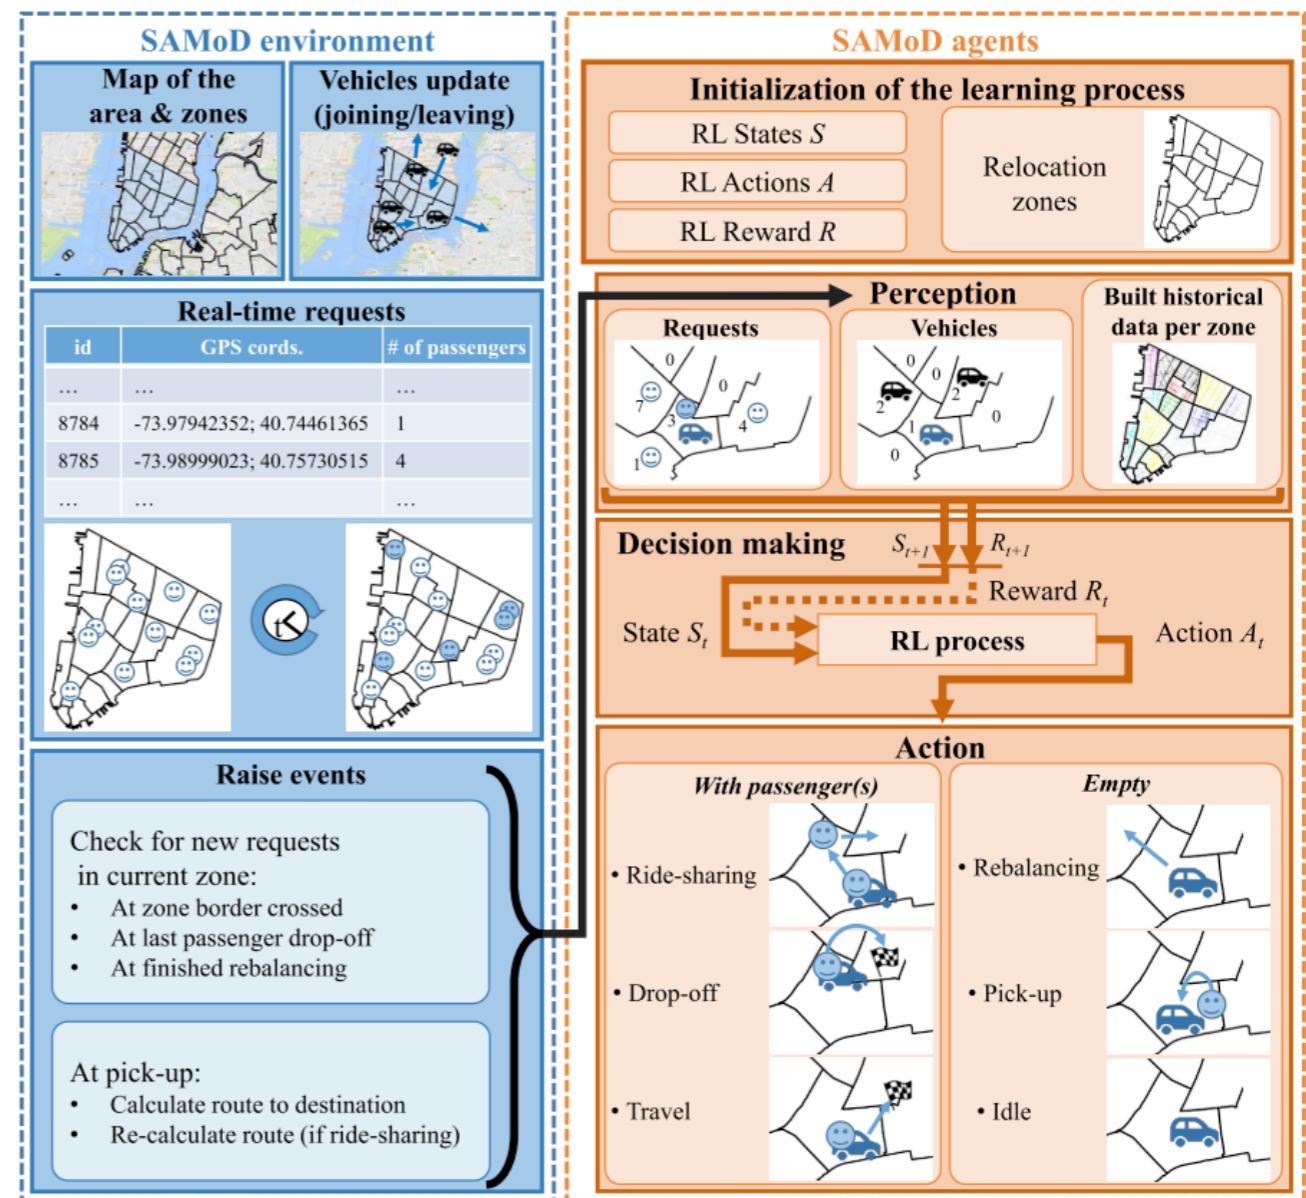
- Origin-destination demand for taxi in NYC



Ride hailing systems with AV

- RL for centrally controlled taxi operation

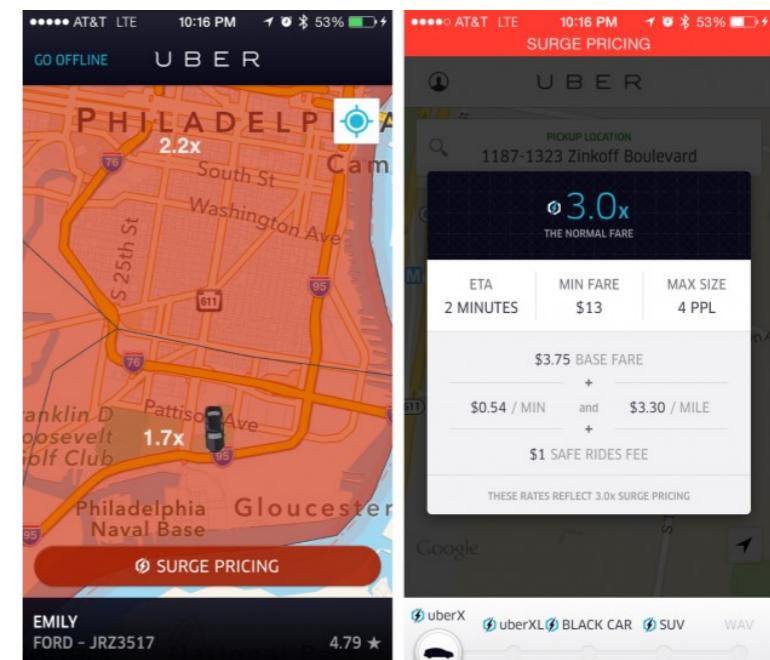
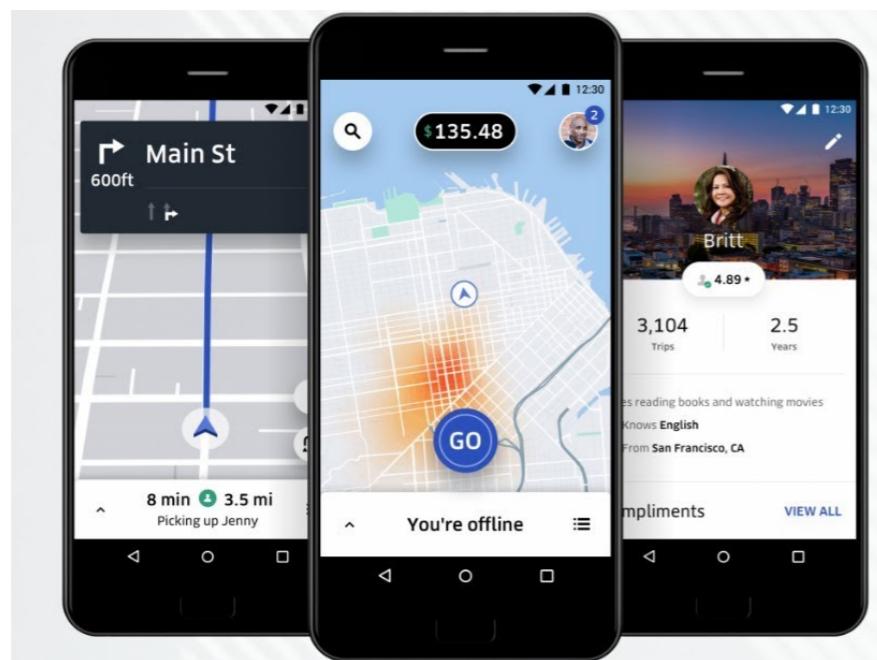
- Which zone to go to next?
- Should it pick up a passenger in the current zone, or move to another zone to maximize its future profits?



Driver-based Ride hailing systems

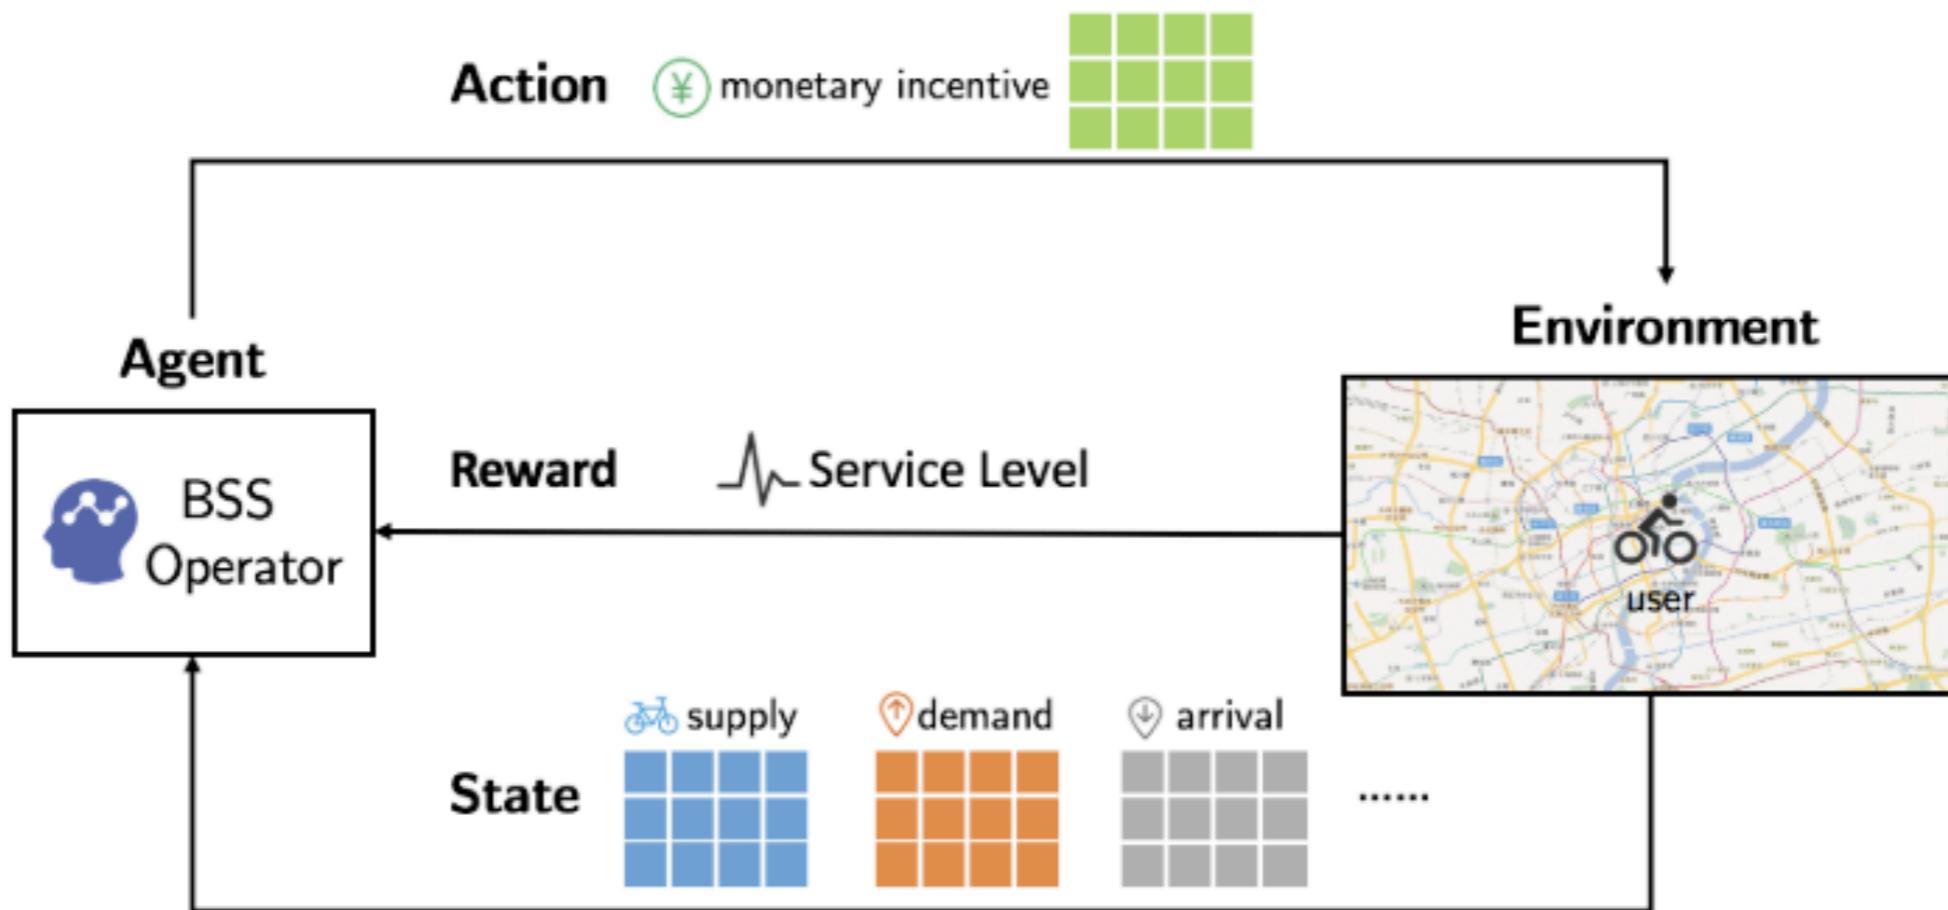
- “Drivers are independent agents. So they are our customers as much as riders are. We cannot tell the drivers where to move. We can suggest where we think they might want to move”

Phillips, R. 2017. Balancing supply and demand in a two-sided marketplace



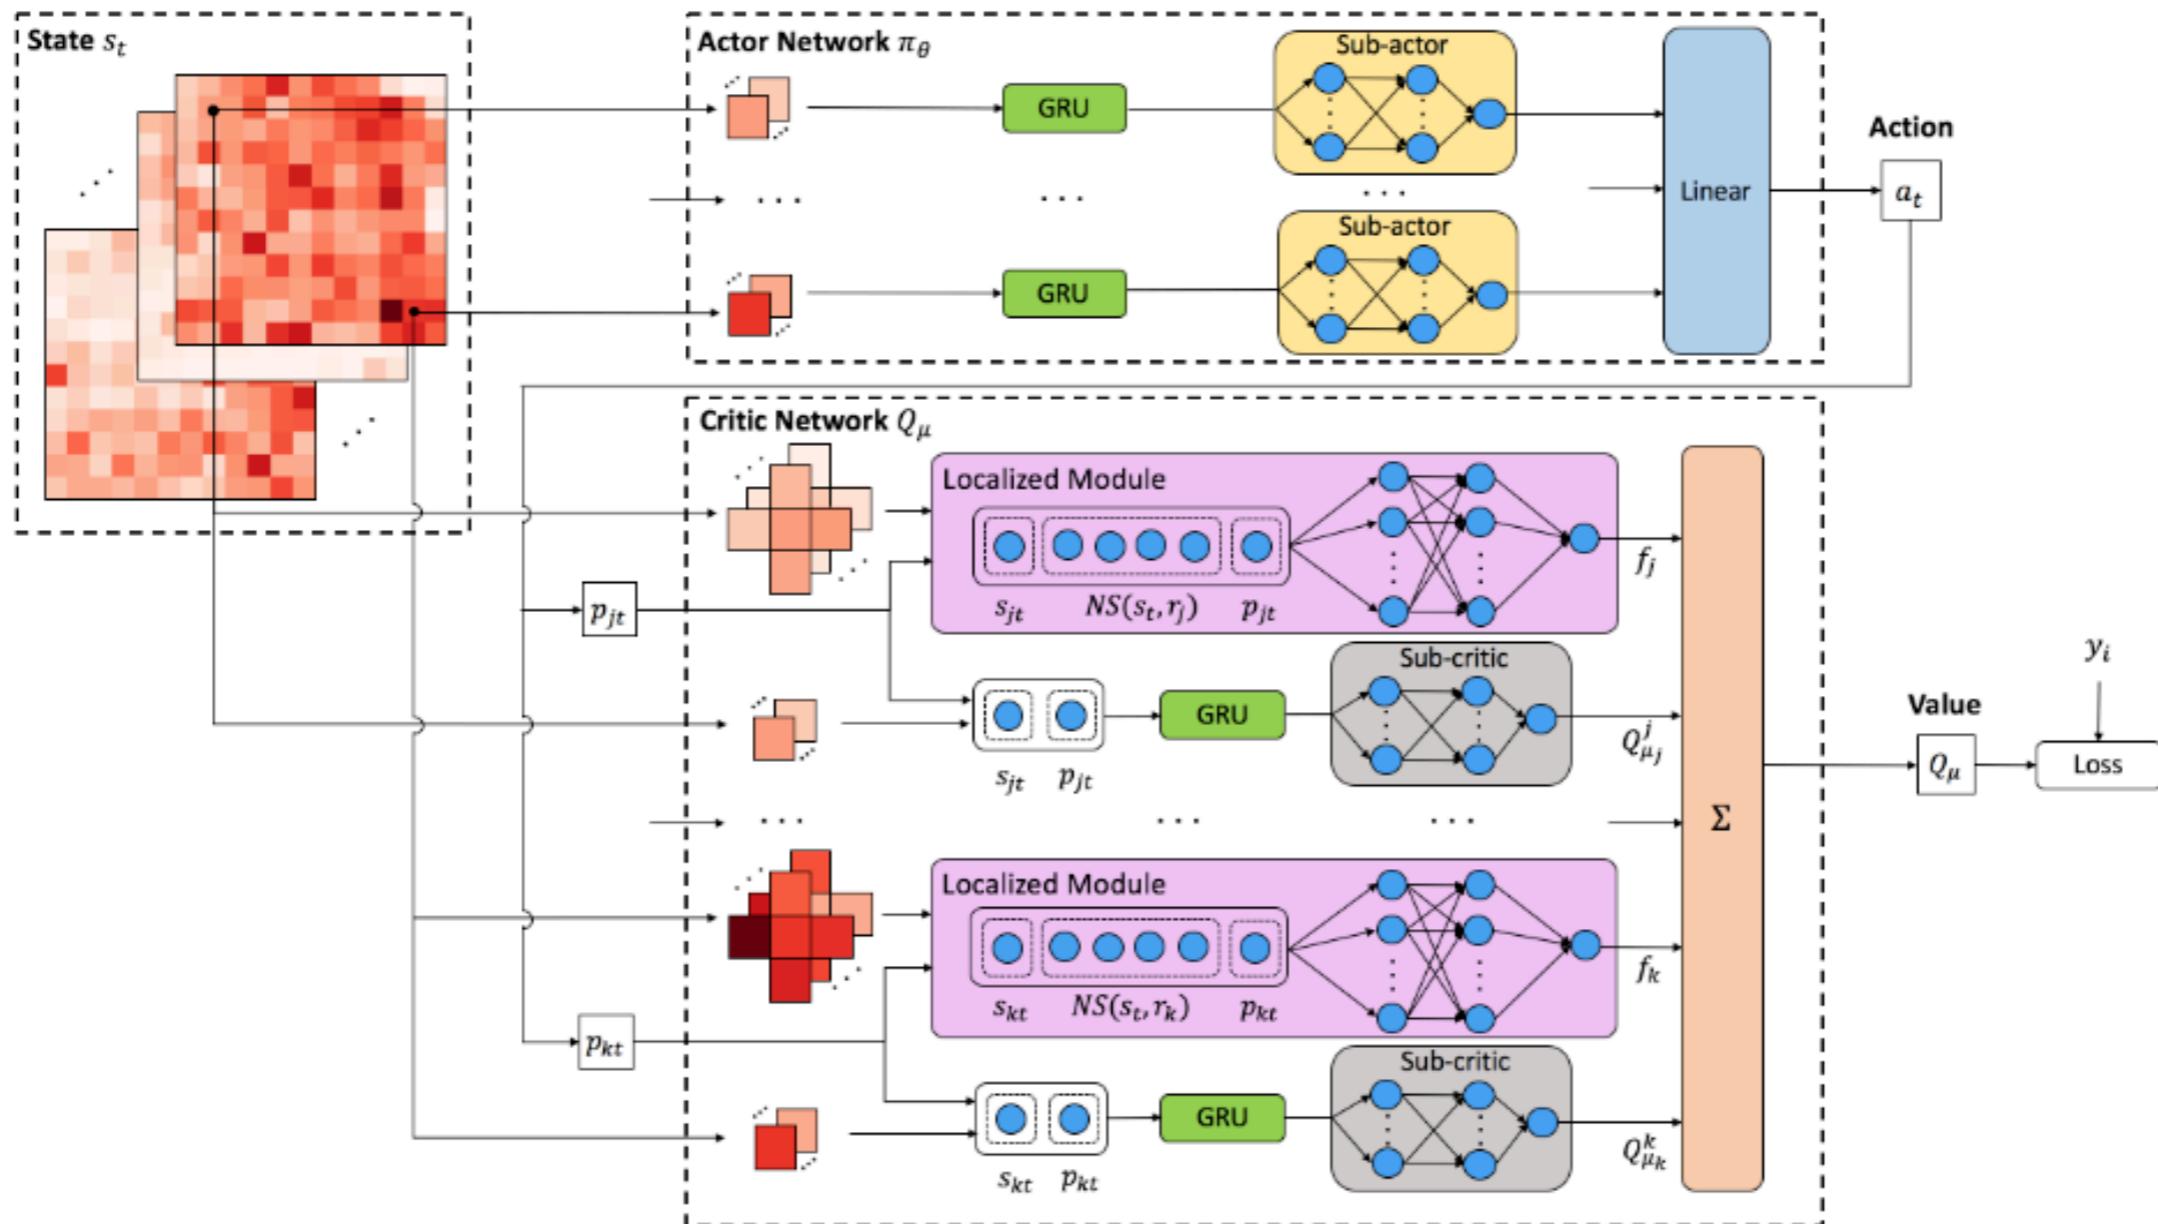
Incentive design

- Incentivize drivers to move to zones that are poorly served



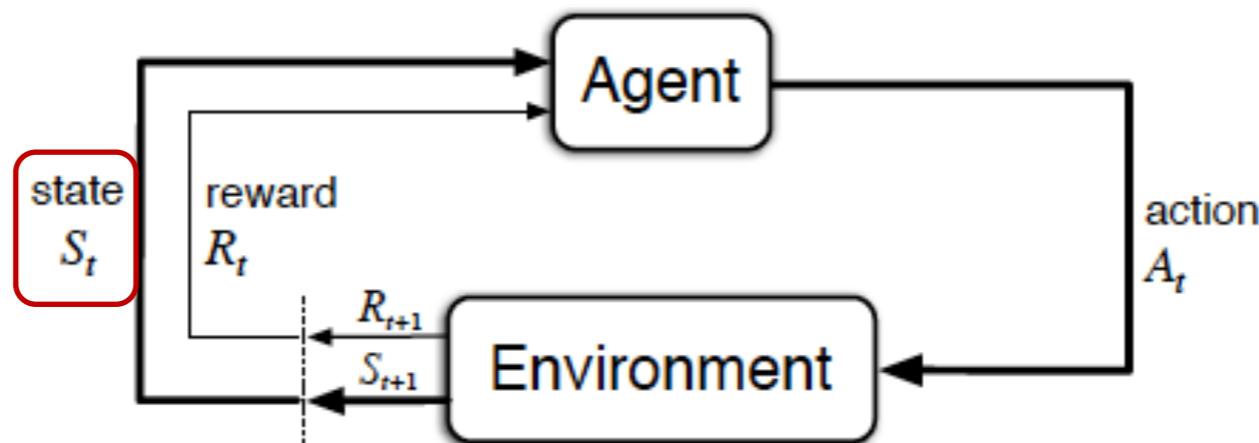
Incentive design

- RL for incentive design



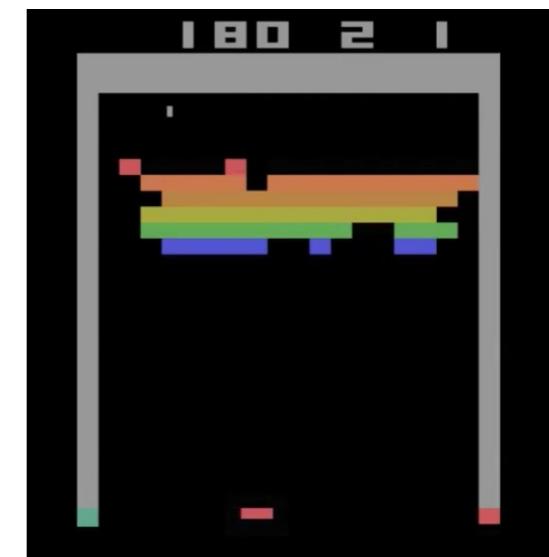
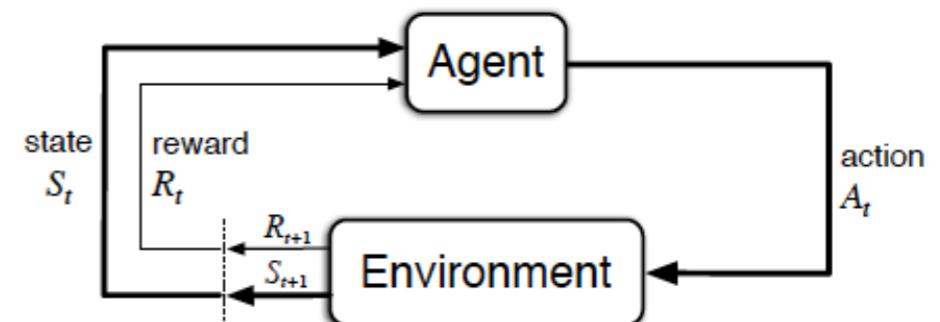
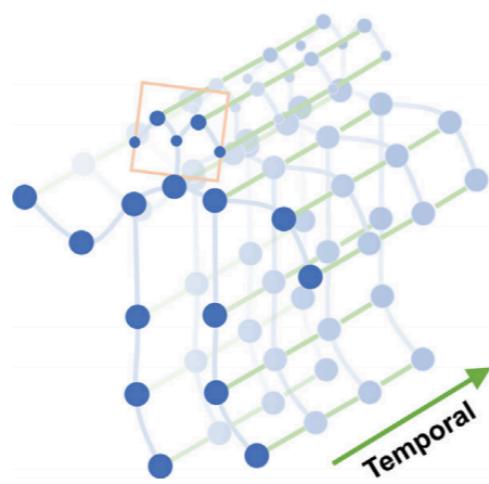
Representation

- Agent-environment interaction in a MDP
- How should we represent the state?



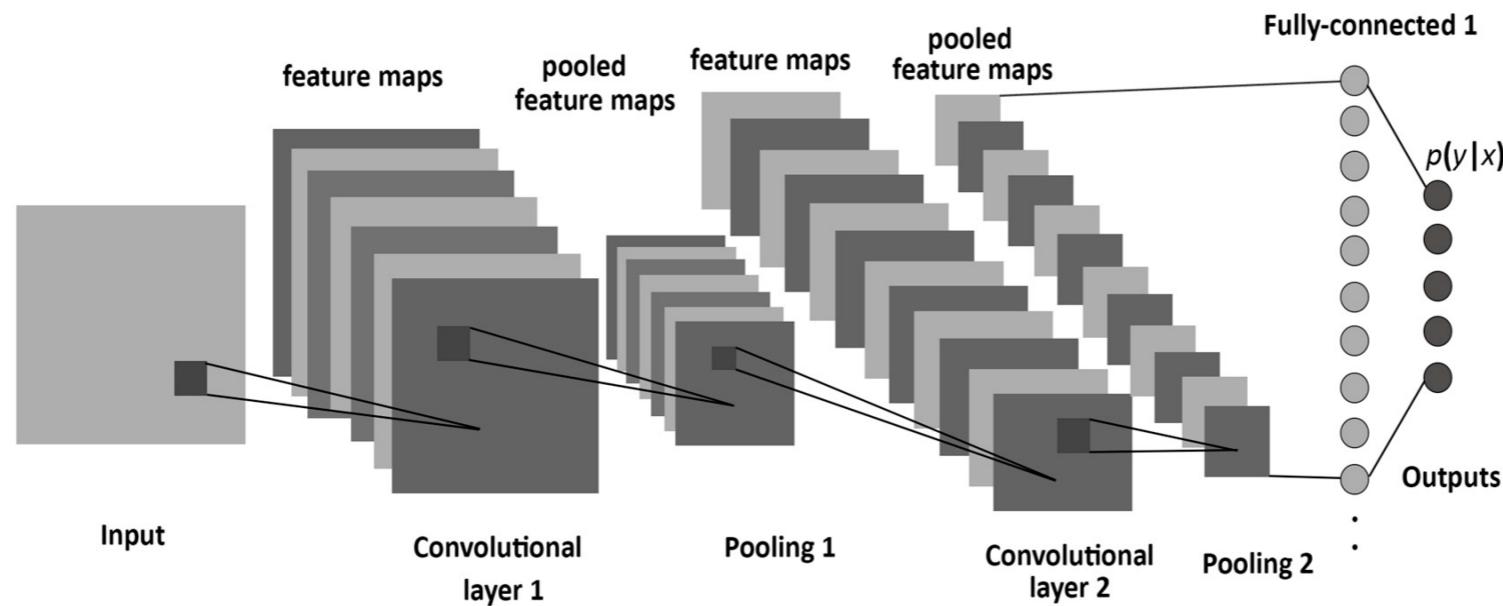
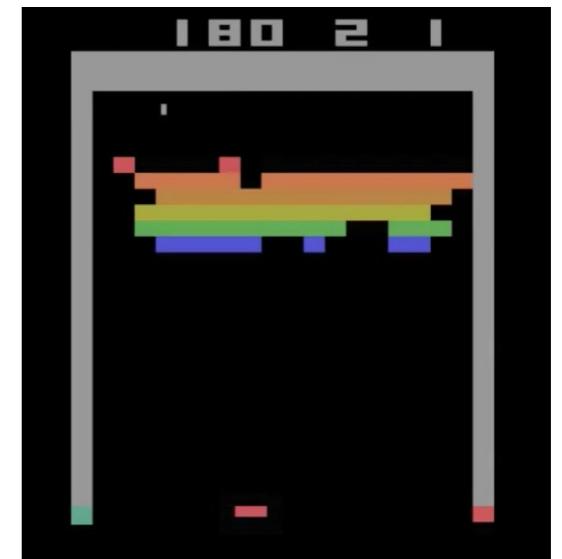
Representation

- State representation
 - Atari games
 - Pixels
 - Robotics
 - Joints



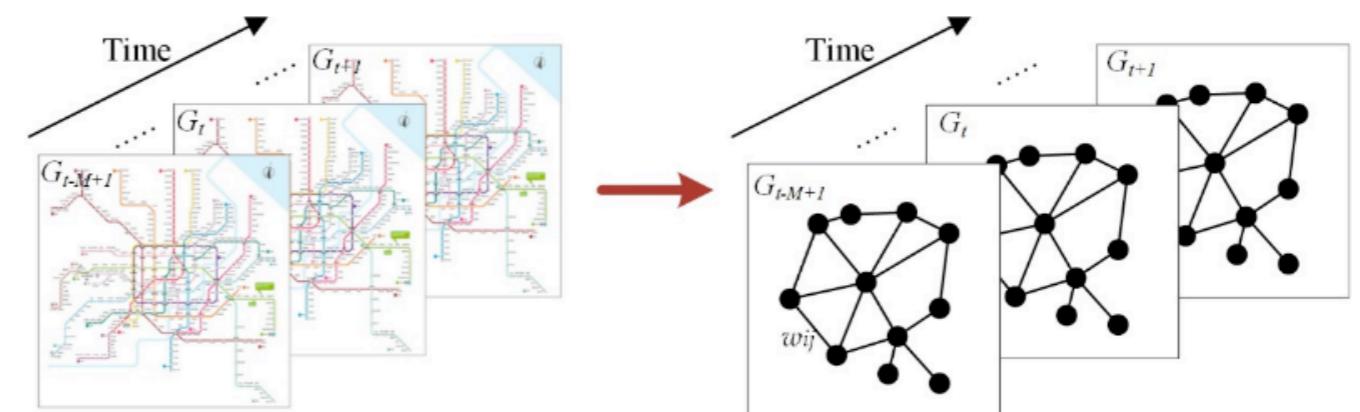
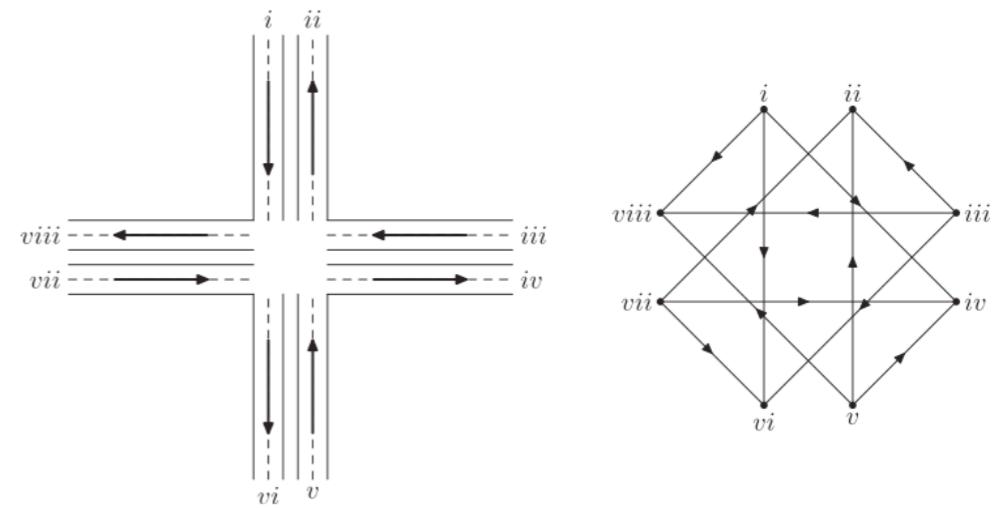
Representation

- State representation
 - Atari games
 - Pixels
 - CNN as feature extraction layer
 - Robotics
 - Joints



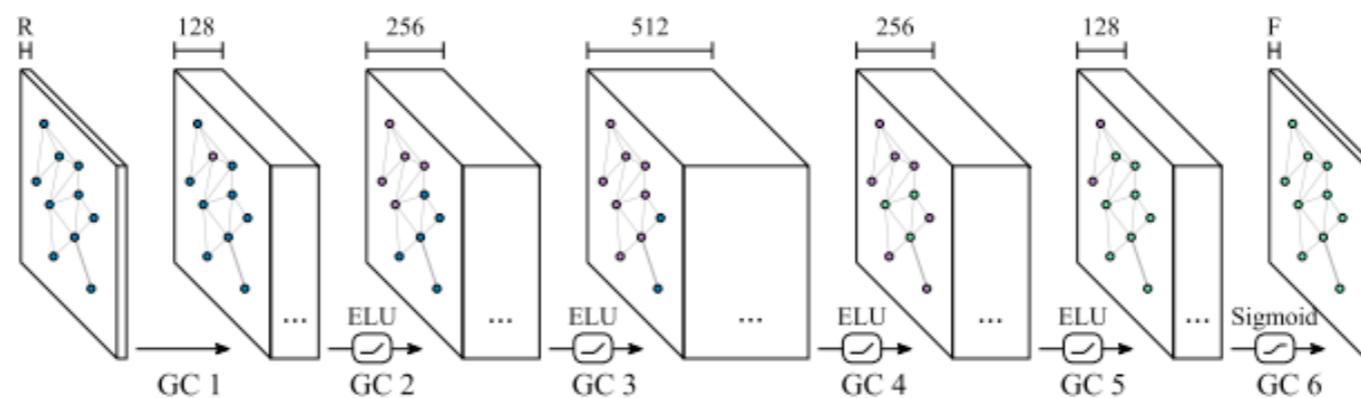
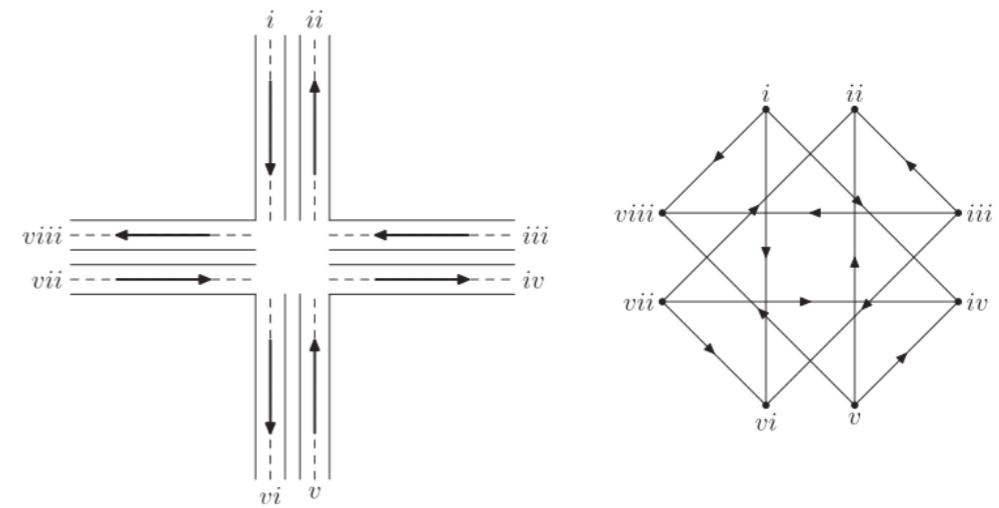
Representation

- What about transportation
- Typical approach:
 - Graphs!



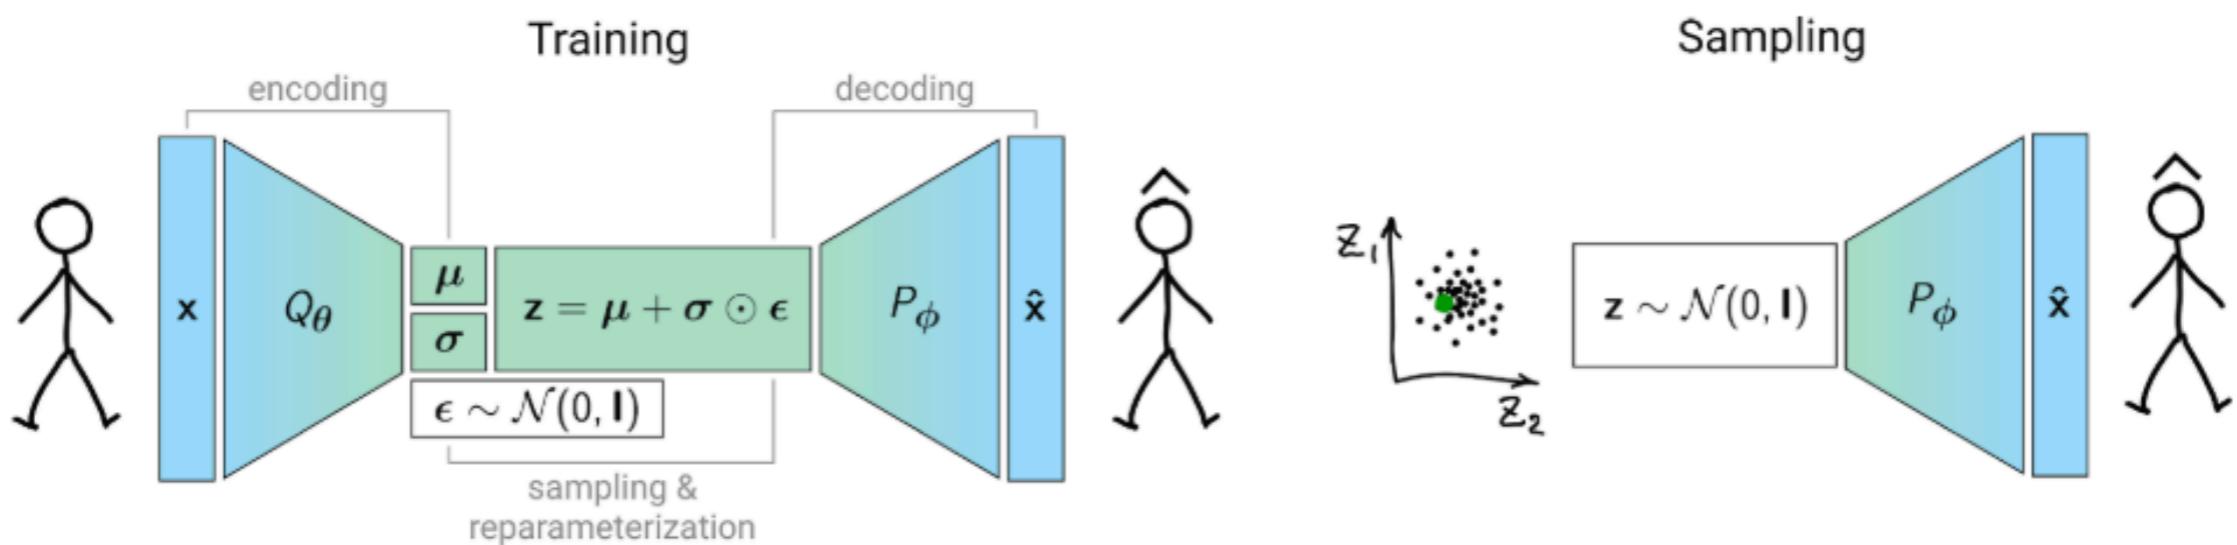
Representation

- What about transportation
- Typical approach:
 - Graphs!
- Could we use Graph Convolutional models?



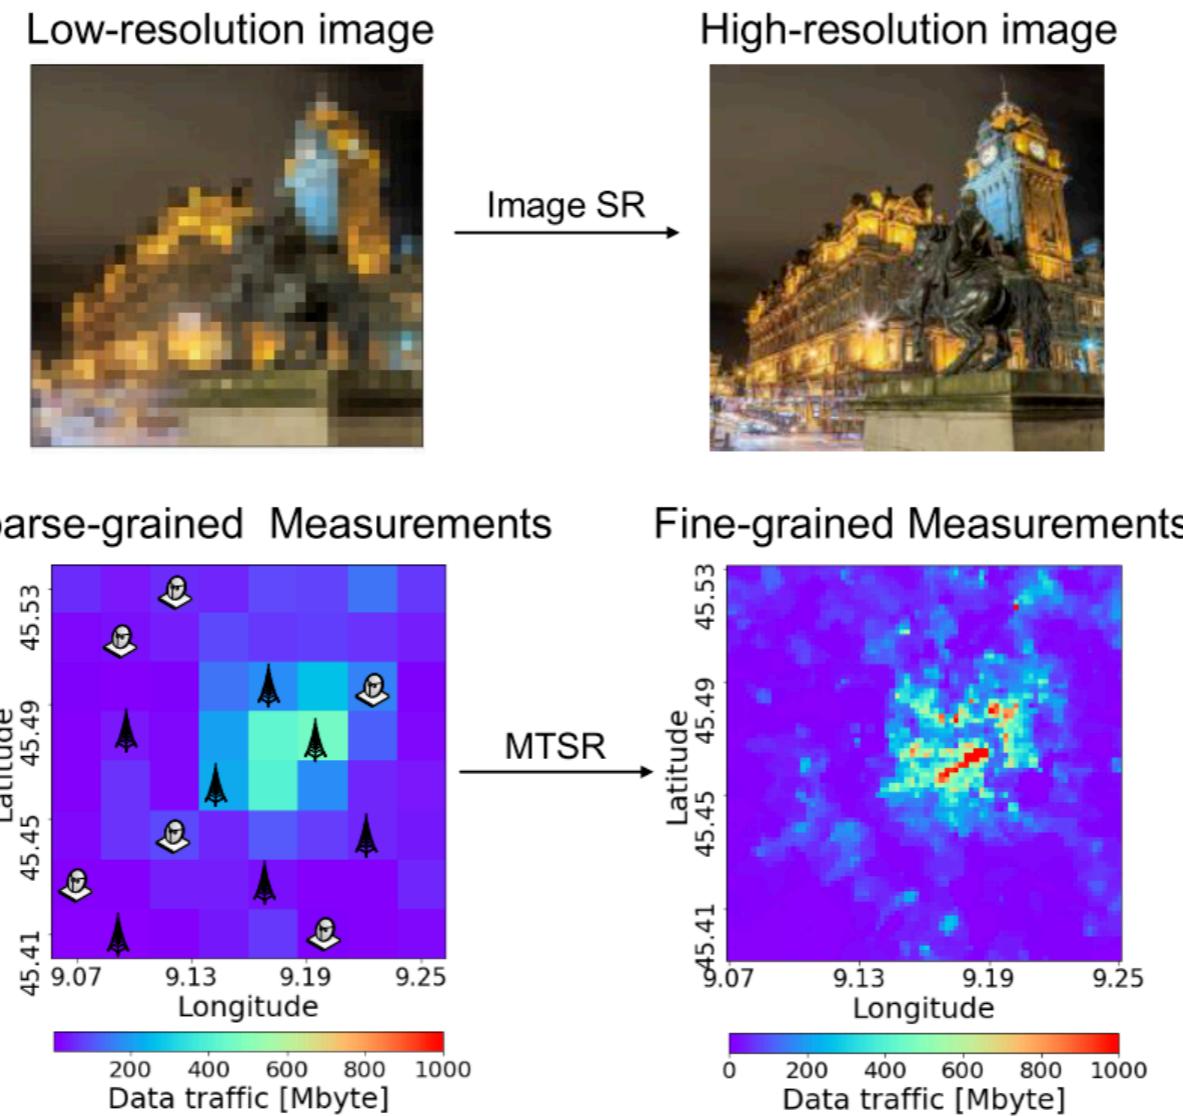
Population synthesis

- How can we learn a realistic representation of the population?



Population synthesis

- How can we learn a realistic representation of the population?



Planning Process

Conventional Travel Forecasting Approach

Data Inputs

Inventories and forecasts of population, land uses, travel behavior, etc.

Trip Generation

Predicts number of trips produced and attracted in a given zone

Trip Distribution

Produces trip production and attraction for each zone

Modal Split

Predicts mode share typically for auto and public transport (can include walk, bike)

Trip Assignment

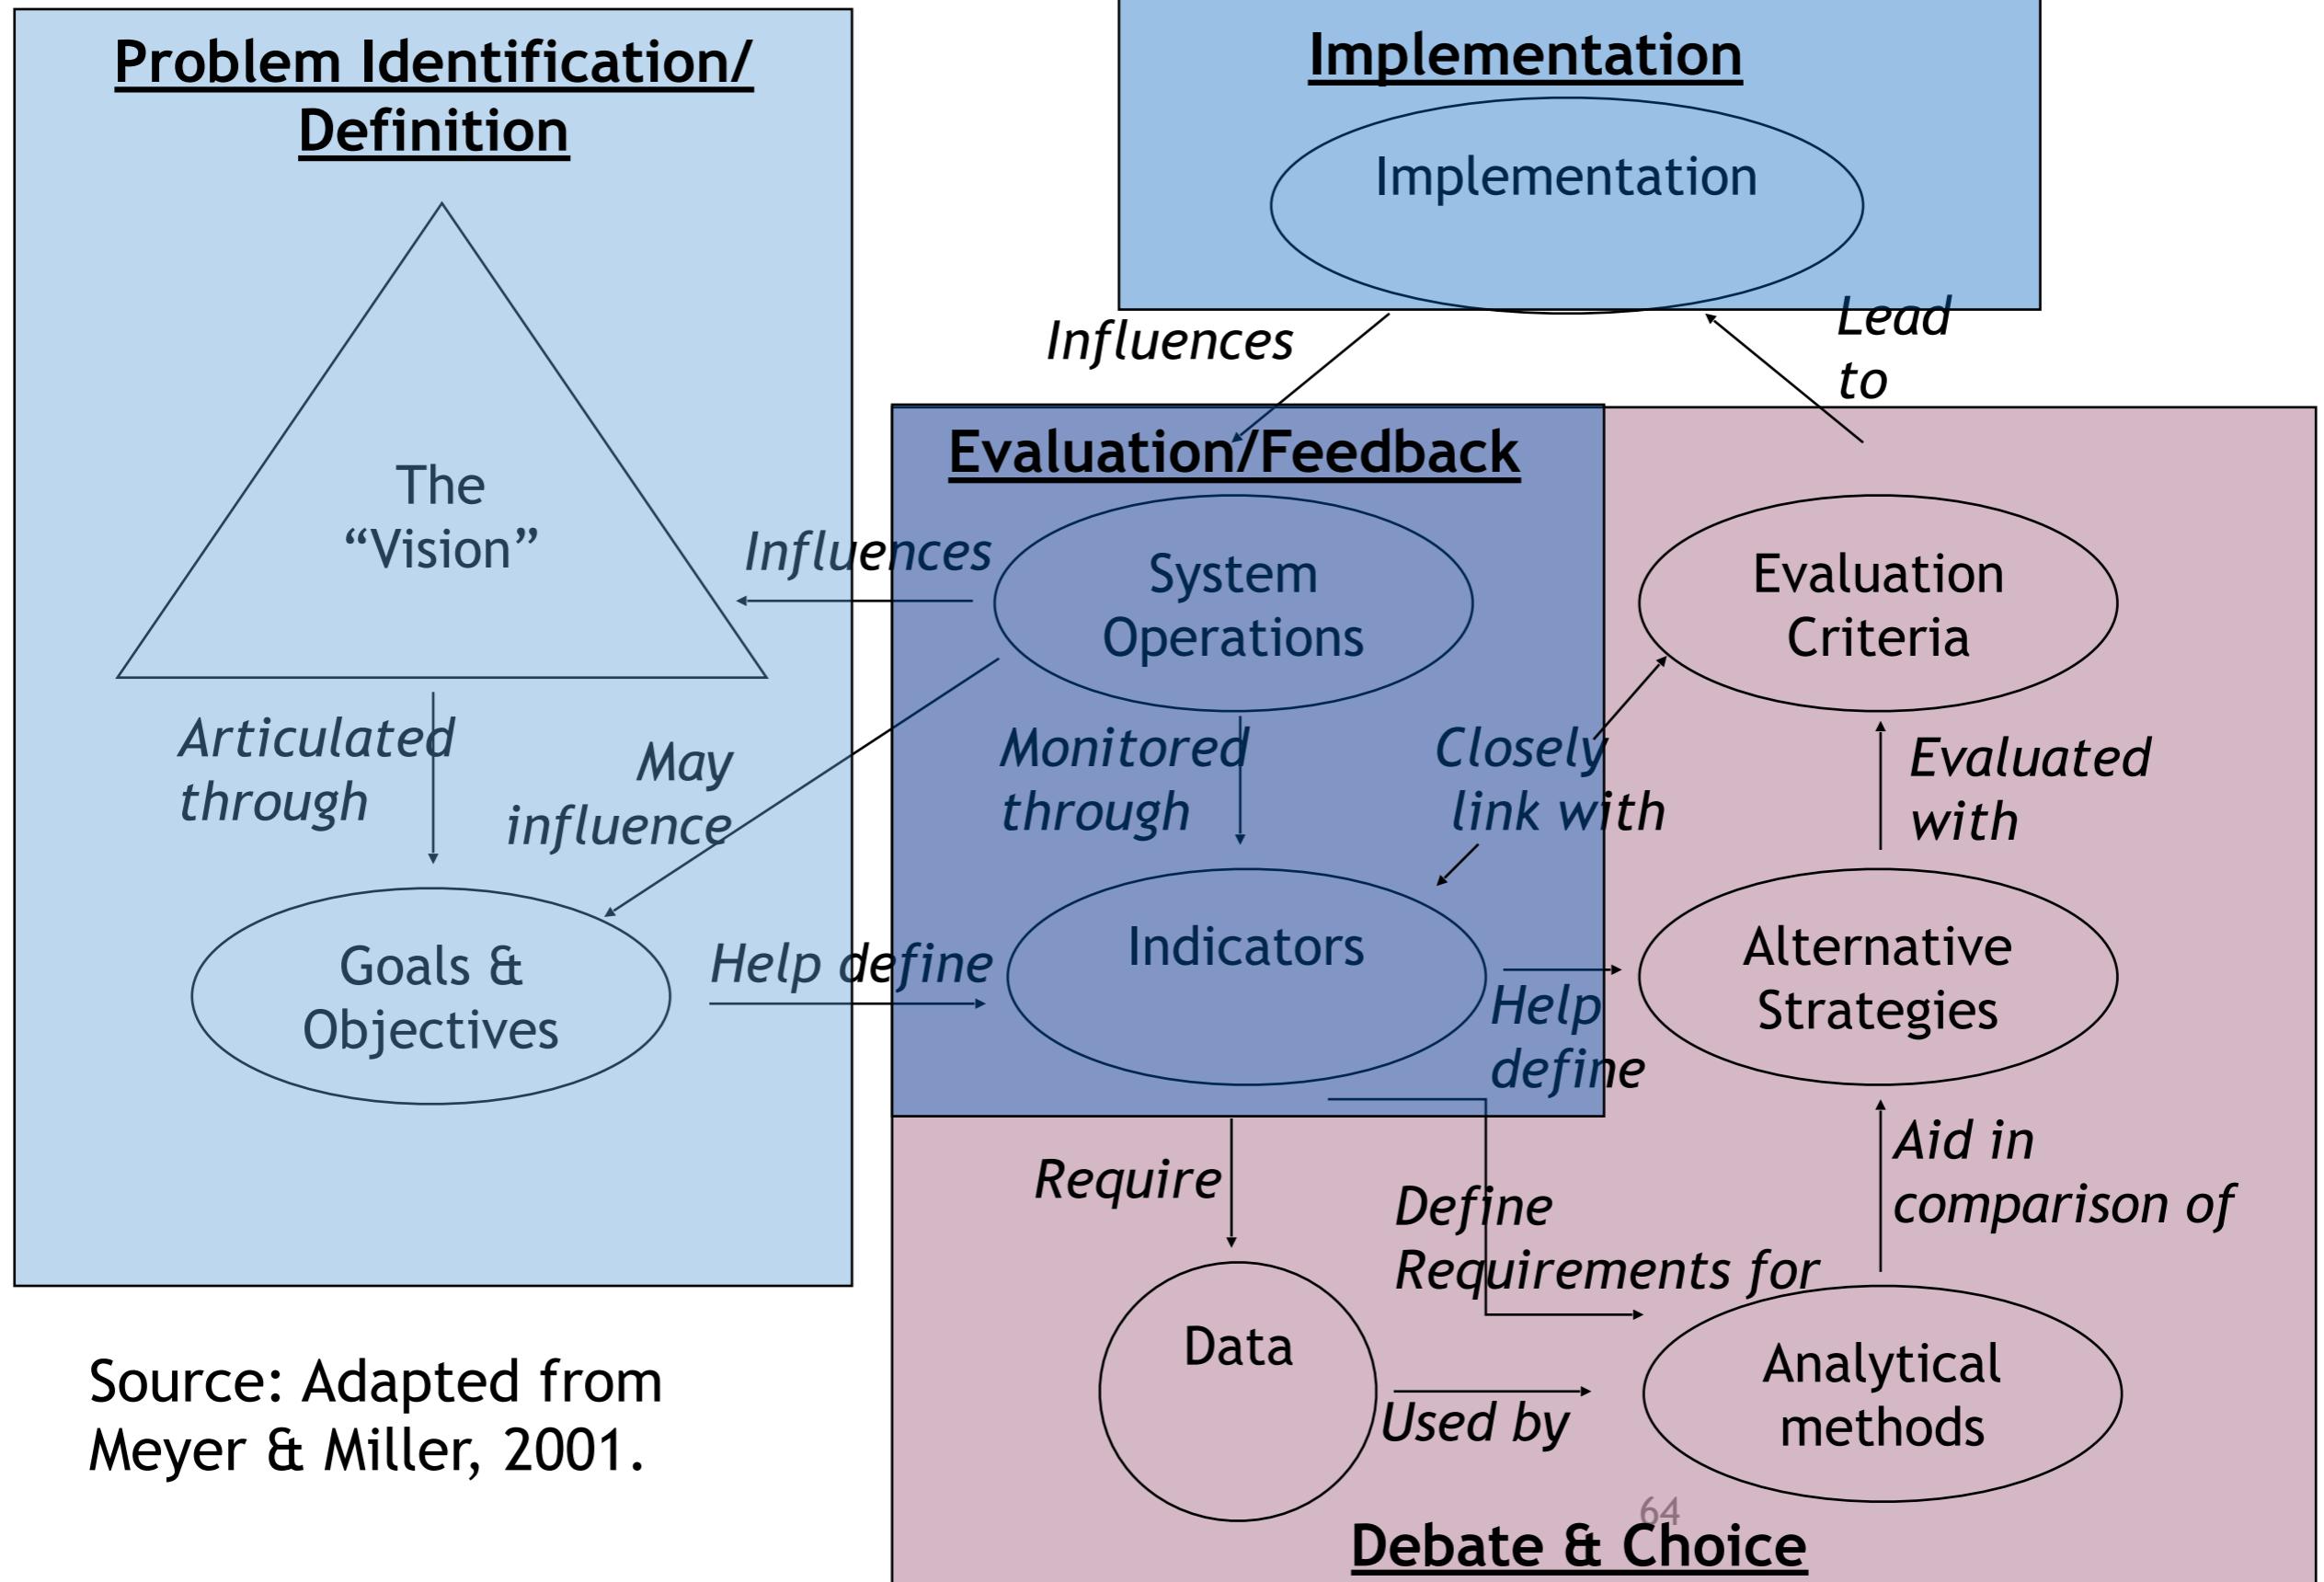
Assigns trips to their respective networks

System Outputs

Provides, for each link, data including traffic volumes, speeds, vehicle mix

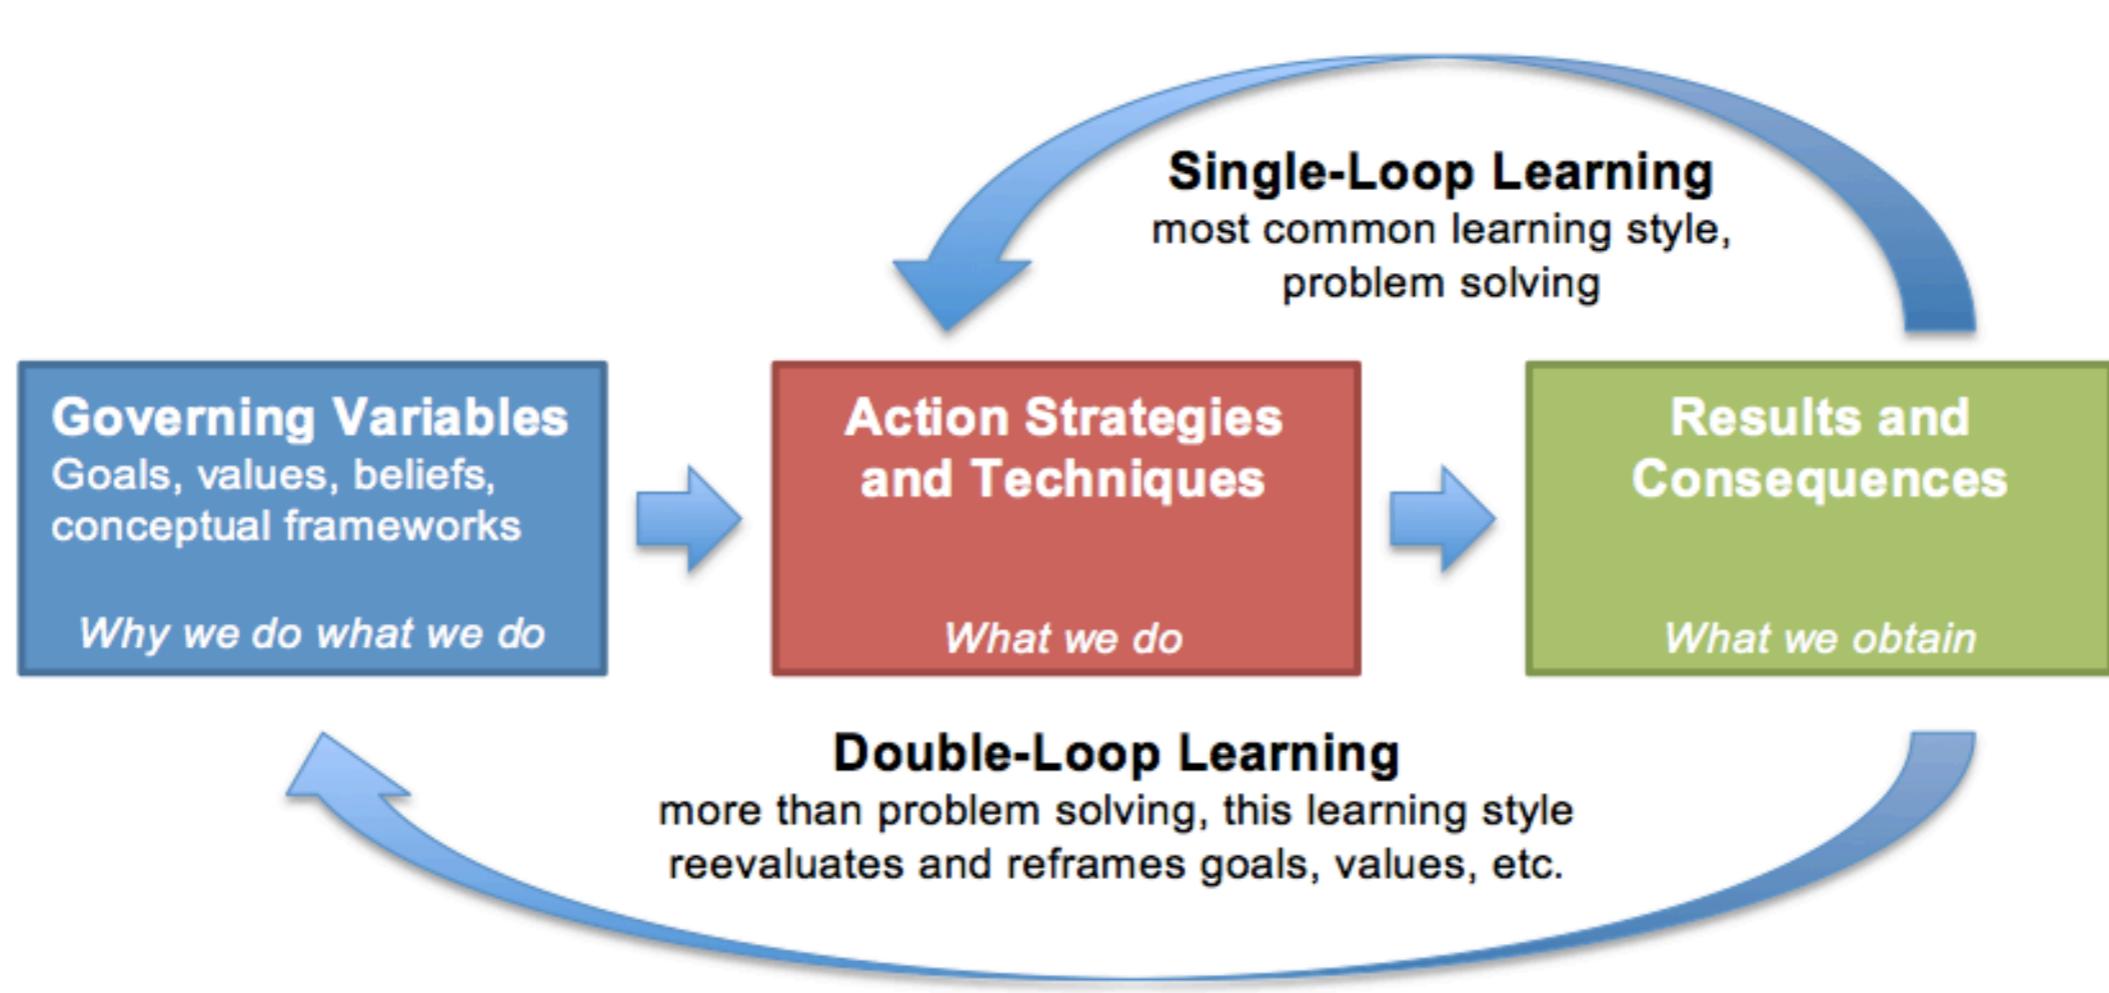
Evaluation

Rational Planning



Cracks in the rational model....

- Mutual learning: planners + people (Friedmann, 1973)
- Reflective Practitioner (Schon, 1983): reality is a social construction; planner must participate in construction of problem; technical tools fit in reflection in action



Tensions in DNN and Transportation

Tensions in DNN and Transportation

- Role of Model: explain or predict
- Role of Theory: Generic model vs domain specific
- Simple to complex vs complex to simple
- Causality vs correlation
- Understand vs action (control)
- Knowledge production: discover vs create

Criteria of a good model

- Prediction accuracy
- Interpretability (FAT)
- Robustness (security)
- Sparsity
- Practicality
- Difficulty

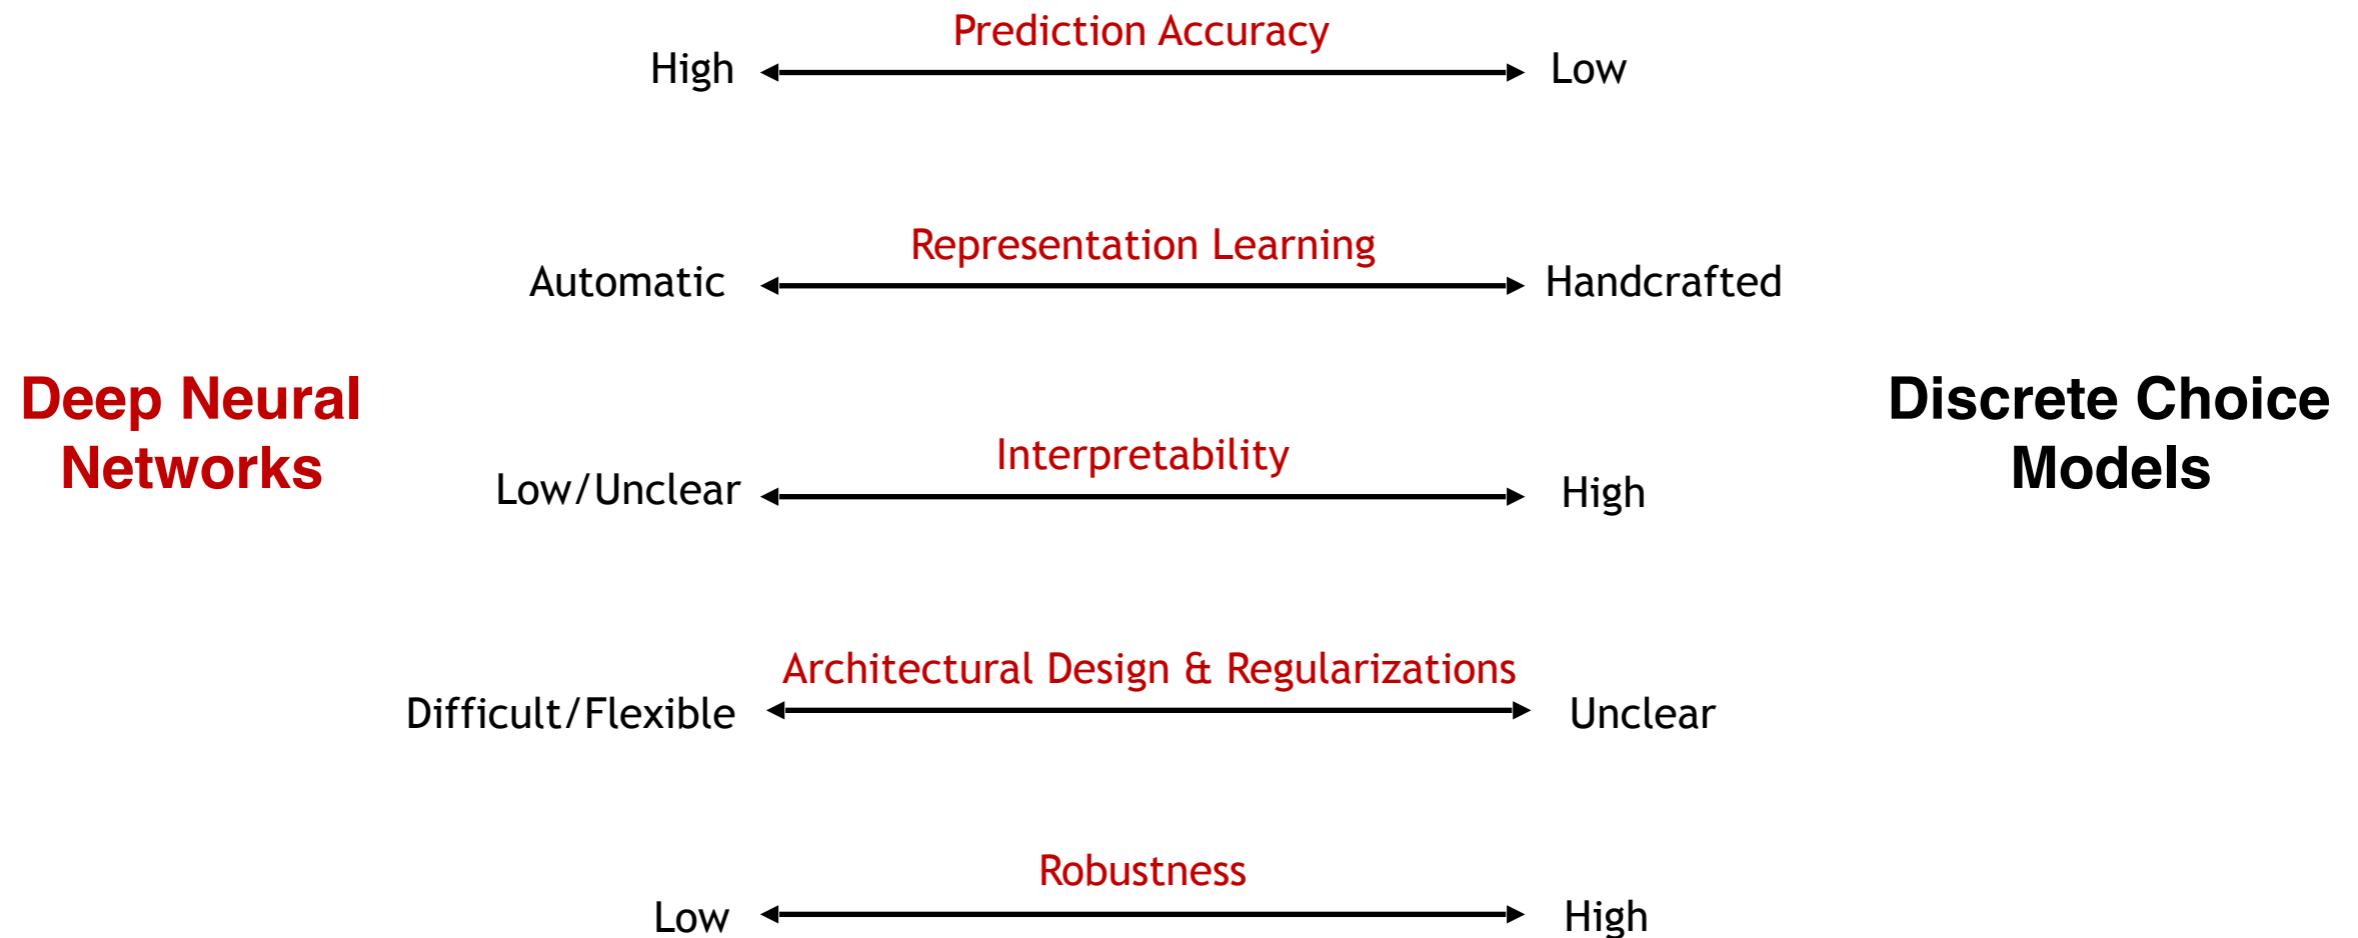
DNN and Behavioral Theory

Deep Neural Networks for Choice Analysis

Shenhao Wang

Dissertation defense for the degree of
Doctor of Philosophy in Computer and Urban Science
September, 2019

Two Modeling Paradigms for Choice Analysis



Prediction Accuracy: Nijkamp et al., (1996); Xie et al., (2003); Cantarella & de Luca, (2005); Celikoglu (2006); etc.

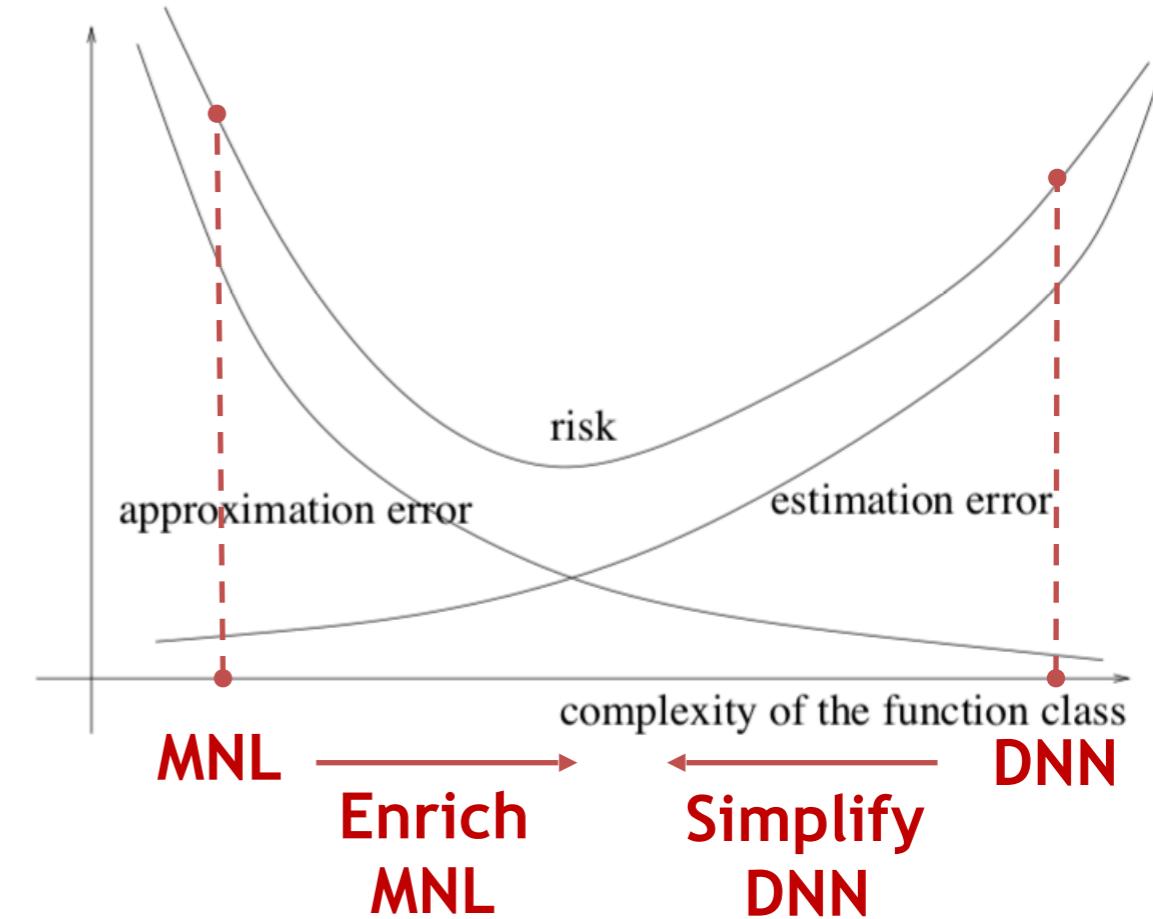
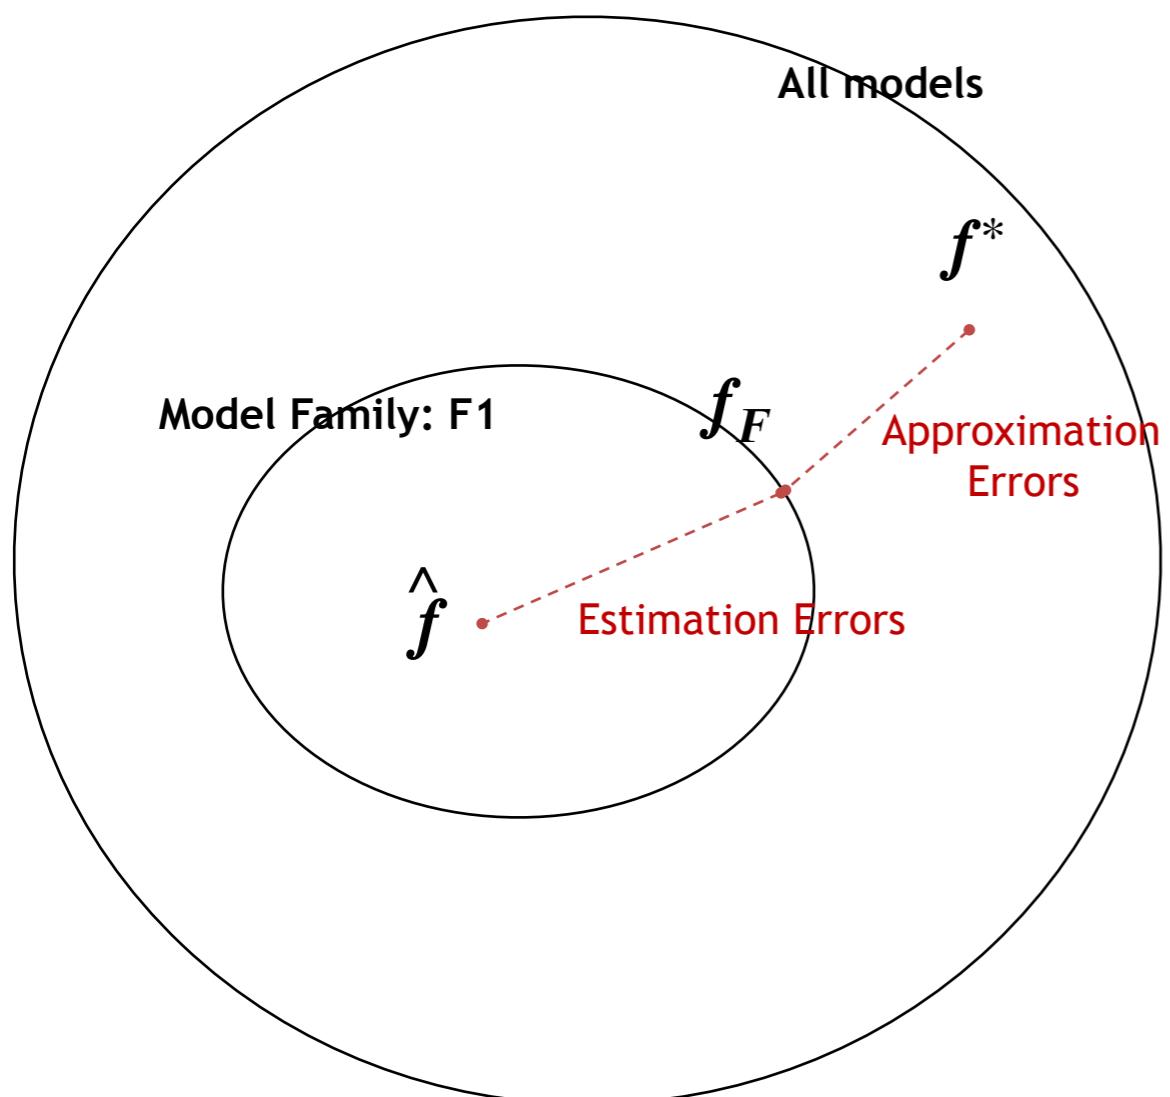
Representation Learning: LeCun et al. (2015); Bengio et al. (2013); etc.

Interpretability: Kim and Doshi-Velez (2017); Lipton (2016); Montavon et al. (2018); Ribeiro et al., (2016); etc.

Architectural Design & Regularizations: Krizhevsky, et al. 2012; Zoph et al., 2017; Martin and Bartlett (2009); etc.

Robustness: Szegedy (2014); Goodfellow (2015); Papernot (2016); Kurakin (2017); Matthew and Jegelka (2017), etc.

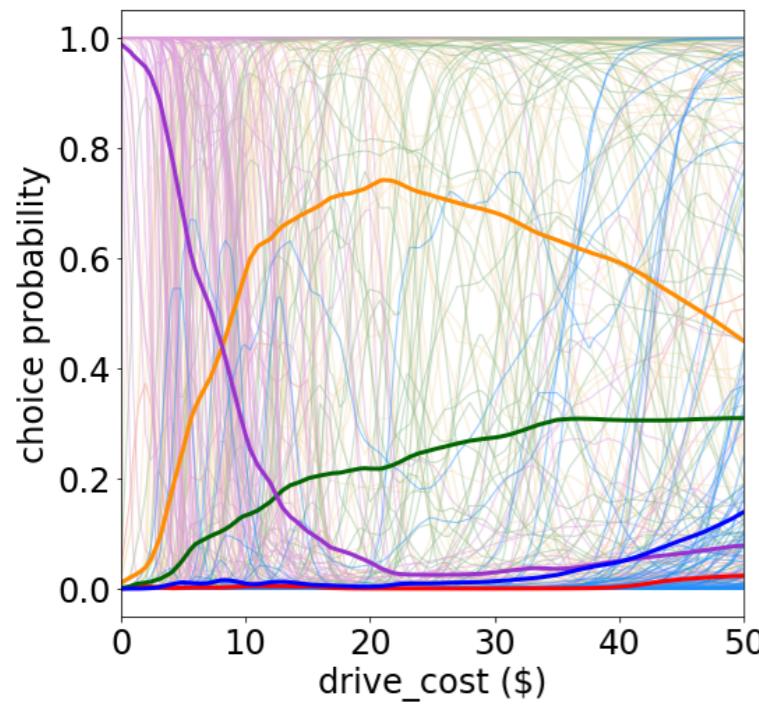
Prediction Errors = Approximation + Estimation Errors



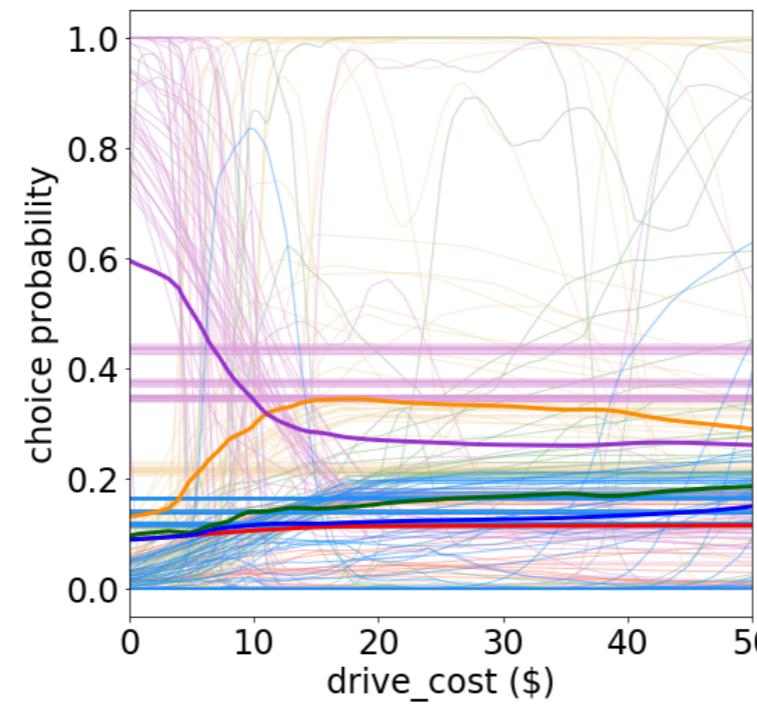
Substitution Patterns of Five Alternatives

$$s_{k_1}(x_j; x_{\setminus j})/s_{k_2}(x_j; x_{\setminus j})$$

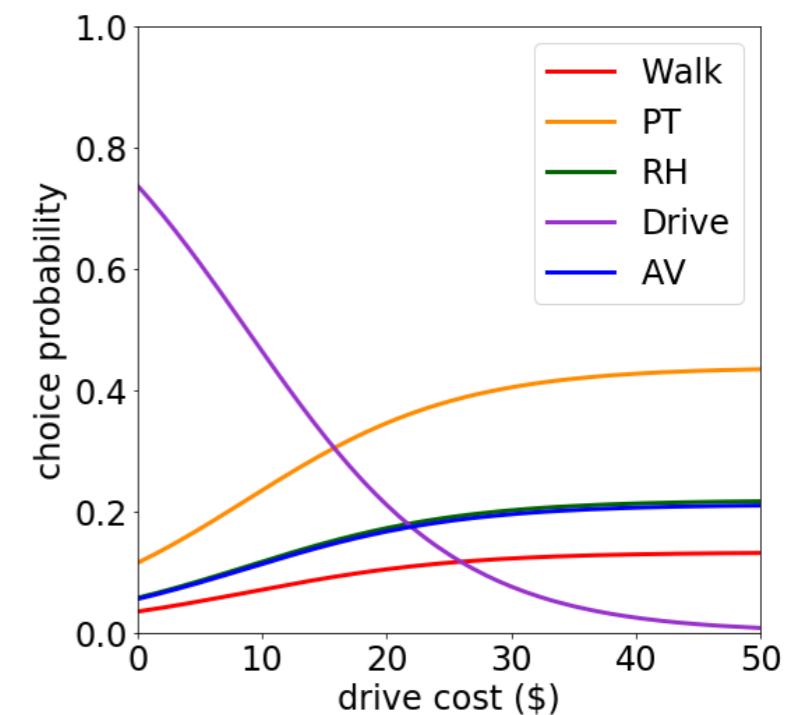
5L-DNNs



HP-DNNs



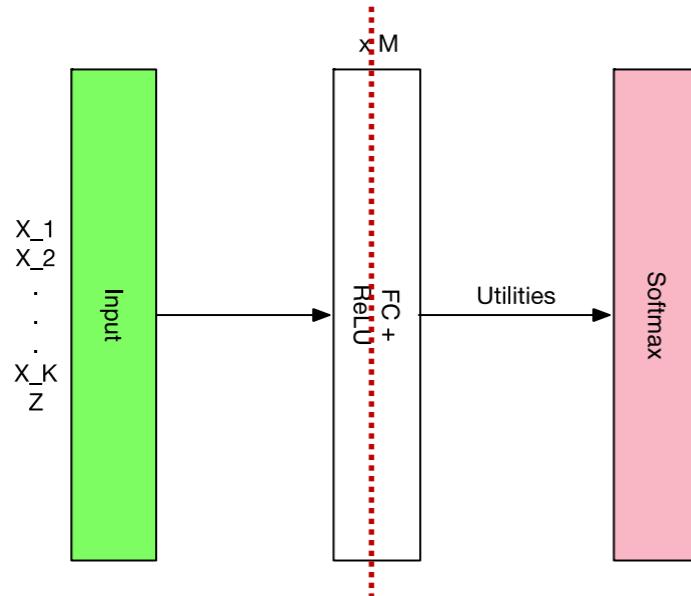
MNL



Synergy between DNN and DCM

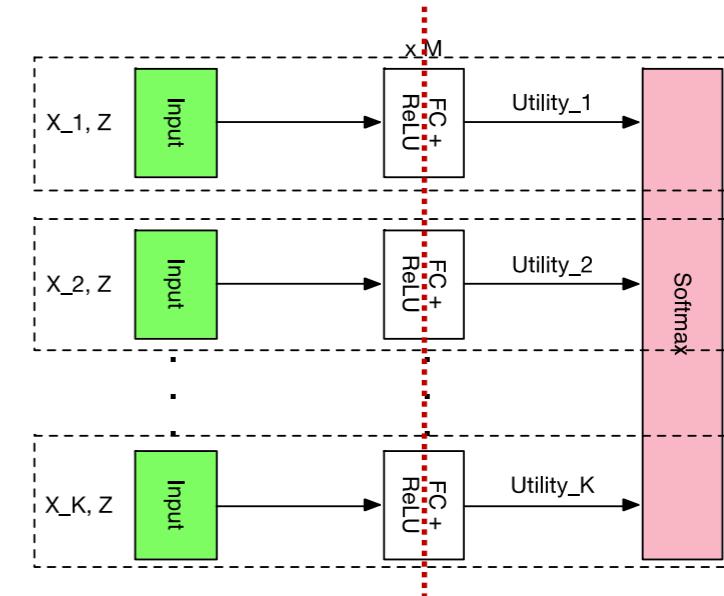
Using Utility Theory to Design DNN

Fully connected DNN (F-DNN)



$$V_k = f(x_1, x_2, \dots, x_K, z)$$

DNN with alternative-specific connection (ASU-DNN)



$$V_k = f(x_k, z)$$

Sparsity as Regularization Introduced by Utility Theory

$$WF-DNN = \begin{bmatrix} w_{1,1} & \dots & w_{1,T_1} \\ \vdots & \ddots & \vdots \\ w_{T_1,1} & \dots & w_{T_1,T_1} \\ w_{T_1+1,1} & \dots & w_{T_1+1,T_1+T_2} \\ \vdots & \ddots & \vdots \\ w_{T_1+T_2,1} & \dots & w_{T_1+T_2,T_1+T_2} \\ \vdots & & \vdots \\ w_{\sum_{i=1}^{k-1} T_i+1,1} & \dots & w_{\sum_{i=1}^{k-1} T_i+1,T_1} \\ \vdots & \ddots & \vdots \\ w_{\sum_{i=1}^K T_i,1} & \dots & w_{\sum_{i=1}^K T_i,T_1} \end{bmatrix}$$

$$W_{ASU-DNN} = \begin{bmatrix} w_{1,1} & \dots & w_{1,T_1} & 0 & \dots & 0 & 0 & \dots & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ w_{T_1,1} & \dots & w_{T_1,T_1} & 0 & \dots & 0 & 0 & \dots & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ 0 & \dots & 0 & w_{T_1+1,T_1+1} & \dots & w_{T_1+1,T_1+T_2} & 0 & \dots & 0 \\ 0 & \dots & 0 & w_{T_1+T_2,T_1+1} & \dots & w_{T_1+T_2,T_1+T_2} & \vdots & \ddots & \vdots \\ \vdots & & \vdots & \vdots & & \vdots & \vdots & & \vdots \\ 0 & \dots & 0 & w_{\sum_{i=1}^{k-1} T_i+1,\sum_{i=1}^{k-1} T_i+1} & \dots & w_{\sum_{i=1}^{k-1} T_i+1,\sum_{i=1}^K T_i} & w_{\sum_{i=1}^{k-1} T_i+1,\sum_{i=1}^{k-1} T_i+1} & \dots & w_{\sum_{i=1}^{k-1} T_i+1,\sum_{i=1}^K T_i} \\ \vdots & \ddots & \vdots & \vdots & & \vdots & \vdots & & \vdots \\ 0 & \dots & 0 & w_{\sum_{i=1}^K T_i,\sum_{i=1}^{k-1} T_i+1} & \dots & w_{\sum_{i=1}^K T_i,\sum_{i=1}^K T_i} & w_{\sum_{i=1}^K T_i,\sum_{i=1}^{k-1} T_i+1} & \dots & w_{\sum_{i=1}^K T_i,\sum_{i=1}^K T_i} \end{bmatrix}$$

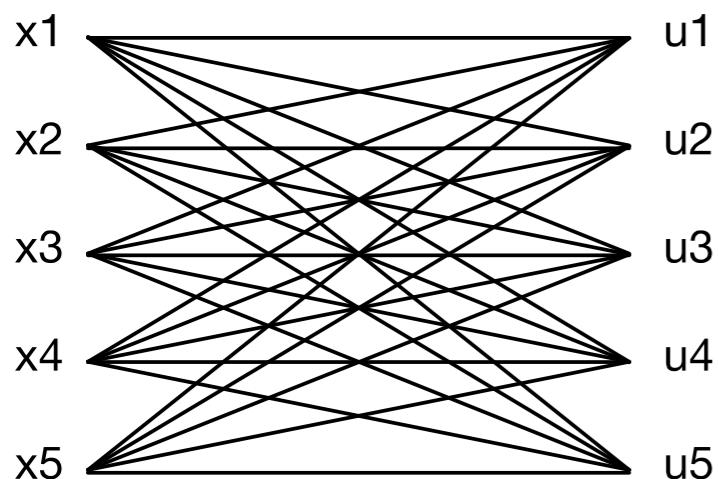
Reduce Estimation Error

$$\hat{R}_n(F_2|_S) \leq \frac{(\sqrt{2 \log(D)} + 1) \sqrt{\frac{1}{N} \sum_{i=1}^N \|x_i\|^2}}{\sqrt{N}} \times c^D T^D$$

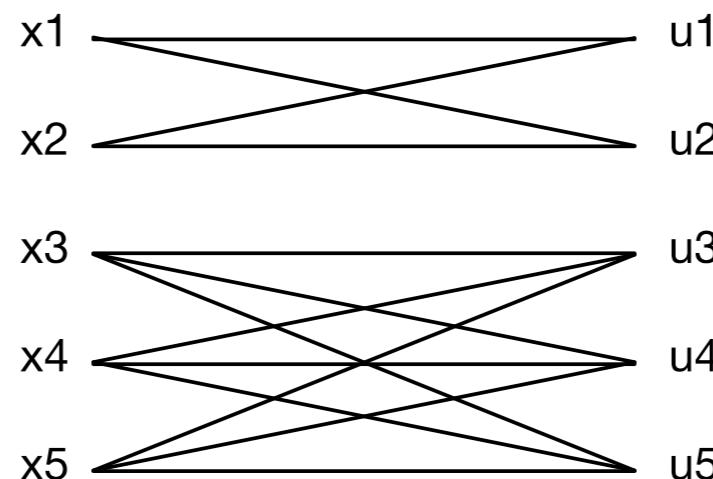
$$\hat{R}_n(F_1|_S) \leq \frac{(\sqrt{2 \log(D)} + 1) \sqrt{\frac{1}{N} \sum_{i=1}^N \|x_i\|^2}}{\sqrt{N}} \times \frac{c^D T^D}{K^{D/2}}$$

DNN Design with Utility Connectivity Graph

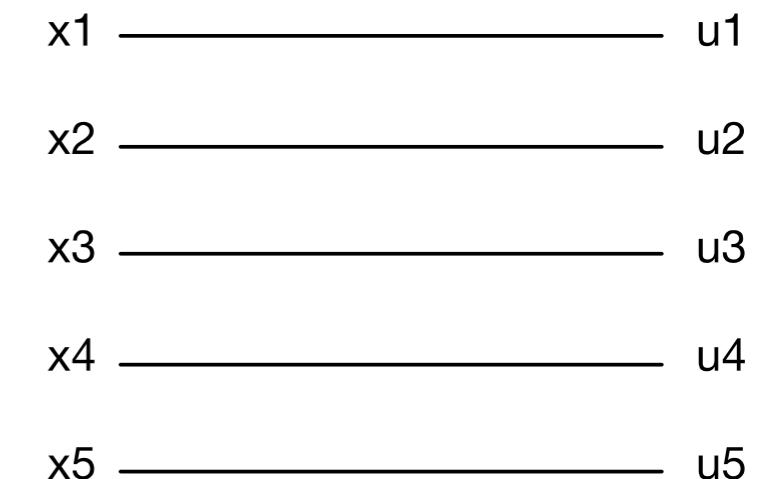
F-DNN



NSU-DNN



ASU-DNN



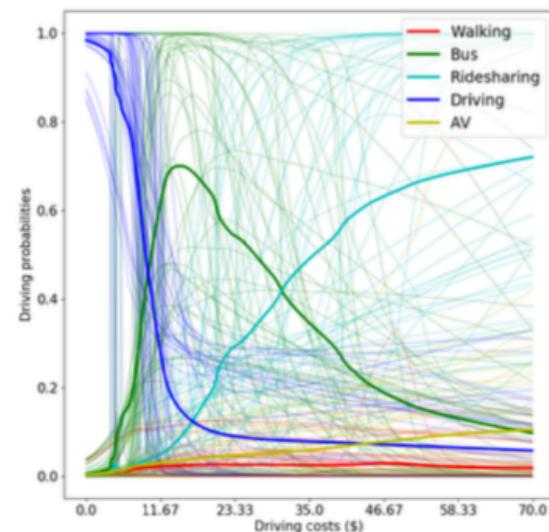
$$A_{F-DNN} = \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

$$A_{NSU-DNN} = \begin{bmatrix} 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \end{bmatrix}$$

$$A_{ASU-DNN} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

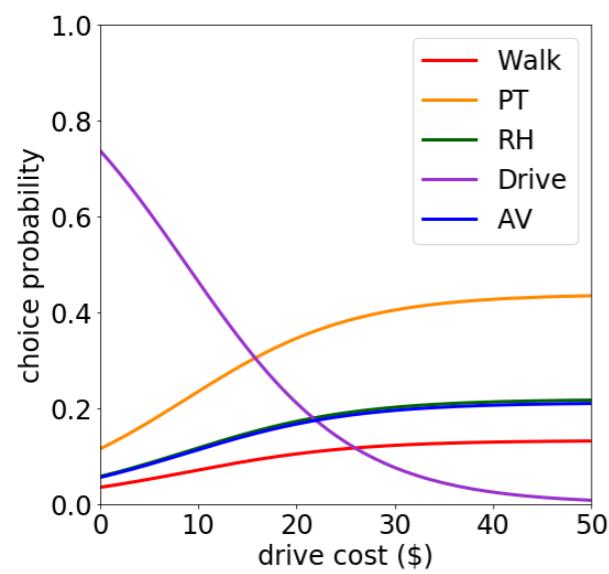
Empirical Result 3: Interpreting the Substitution Pattern

F-DNNs

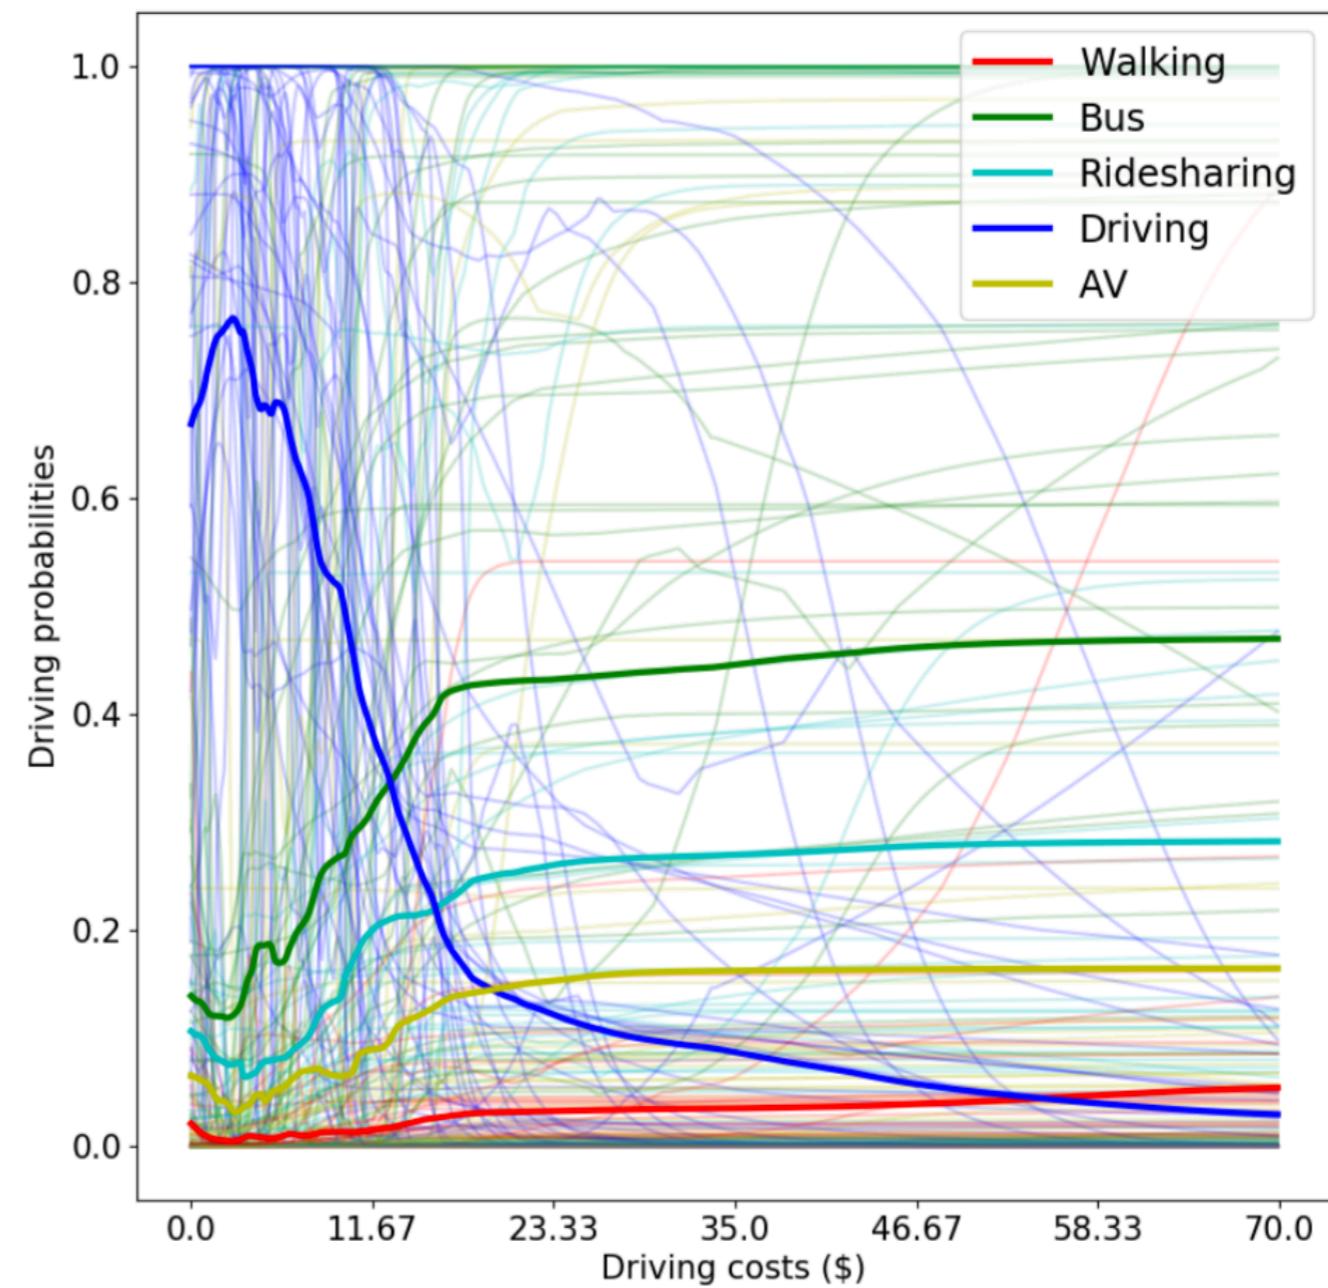


(c) F-DNN (Top 10 Models)

MNLs



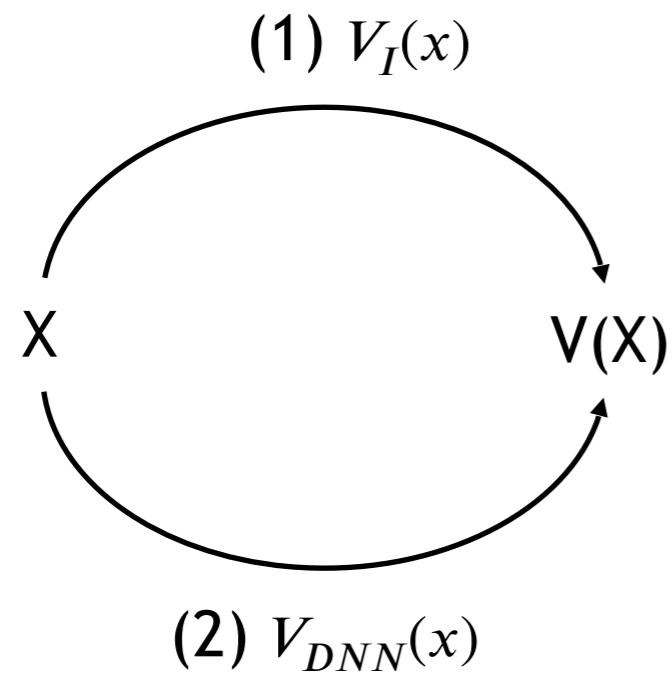
ASU-DNNs



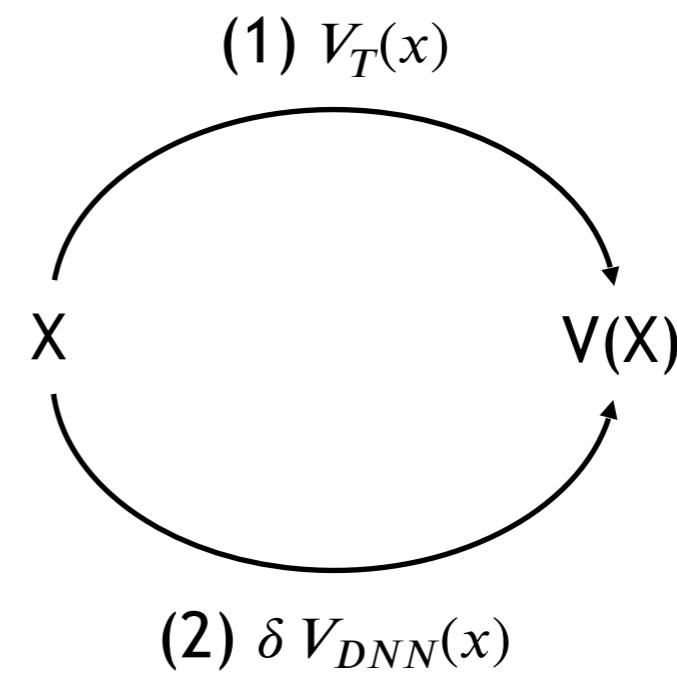
Theory-Based Residual Neural Network (TB-ResNet)

$$V(x) = V_T(x) + \delta V_{DNN}(x)$$

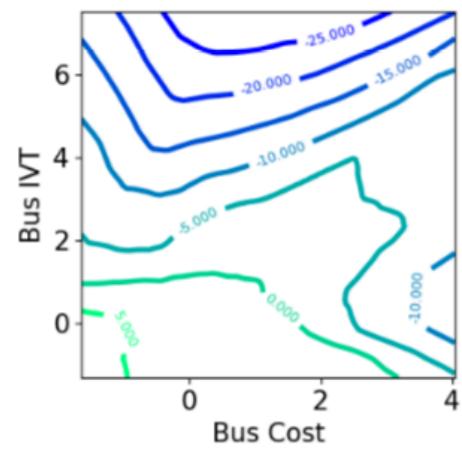
ResNet



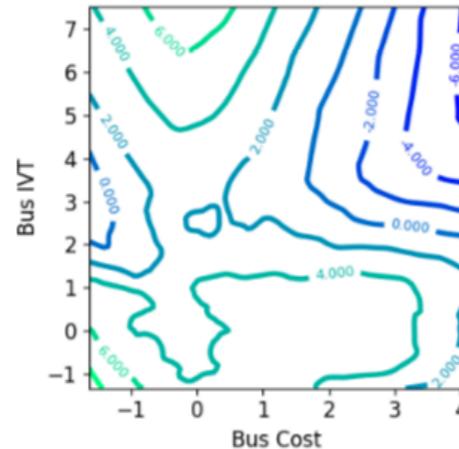
TB-ResNet



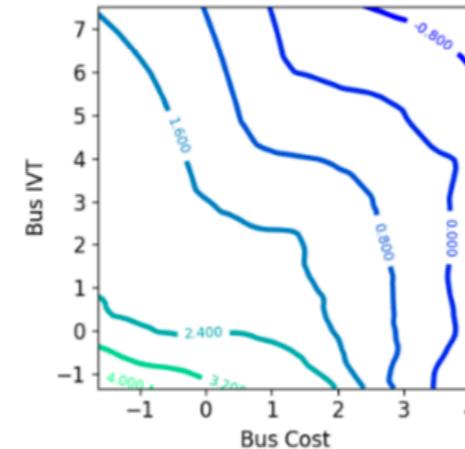
2. Interpretability of Utility Function in the CM Scenario



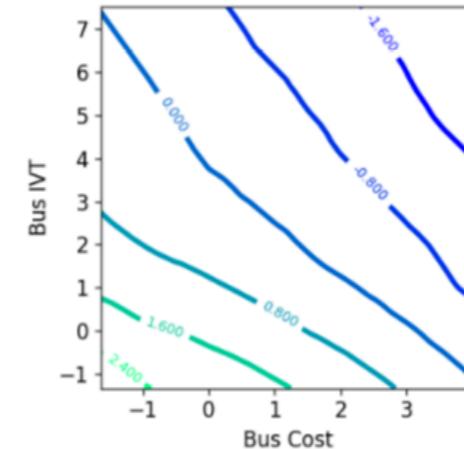
(a) DNN (55.2%)
 $1e - 10$; 56.4%)



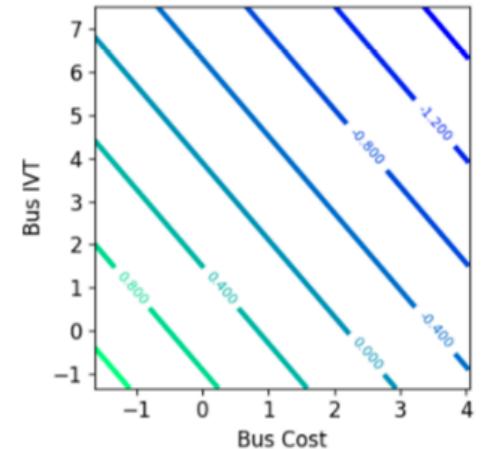
(b) CM Resnet ($\lambda = 1e - 10$; 56.4%)



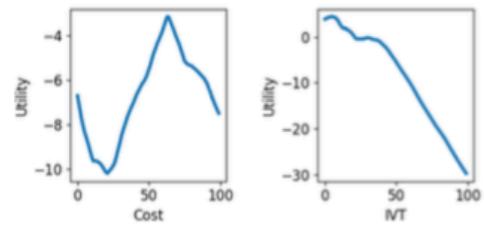
(c) CM Resnet ($\lambda = 0.005$; 57.3%)



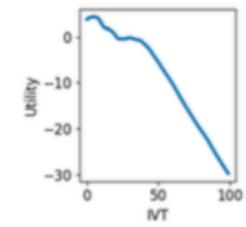
(d) CM Resnet ($\lambda = 0.01$; 56.8%)



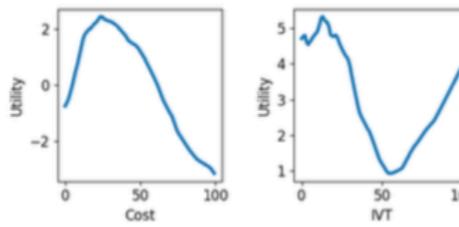
(e) CM (44.7%)



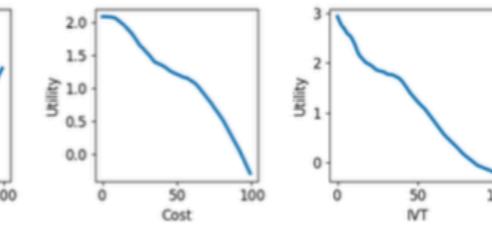
(f) x0



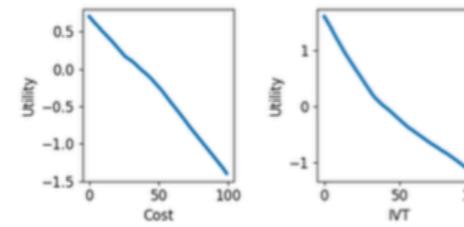
(g) x0



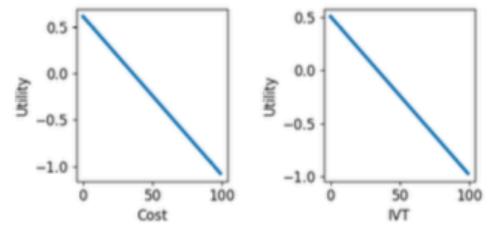
(h) x0



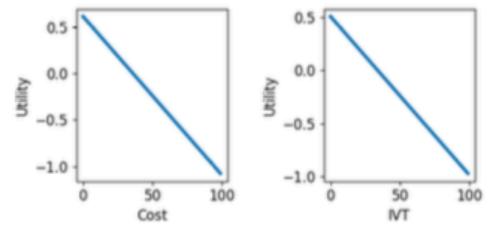
(i) x1



(j) x0

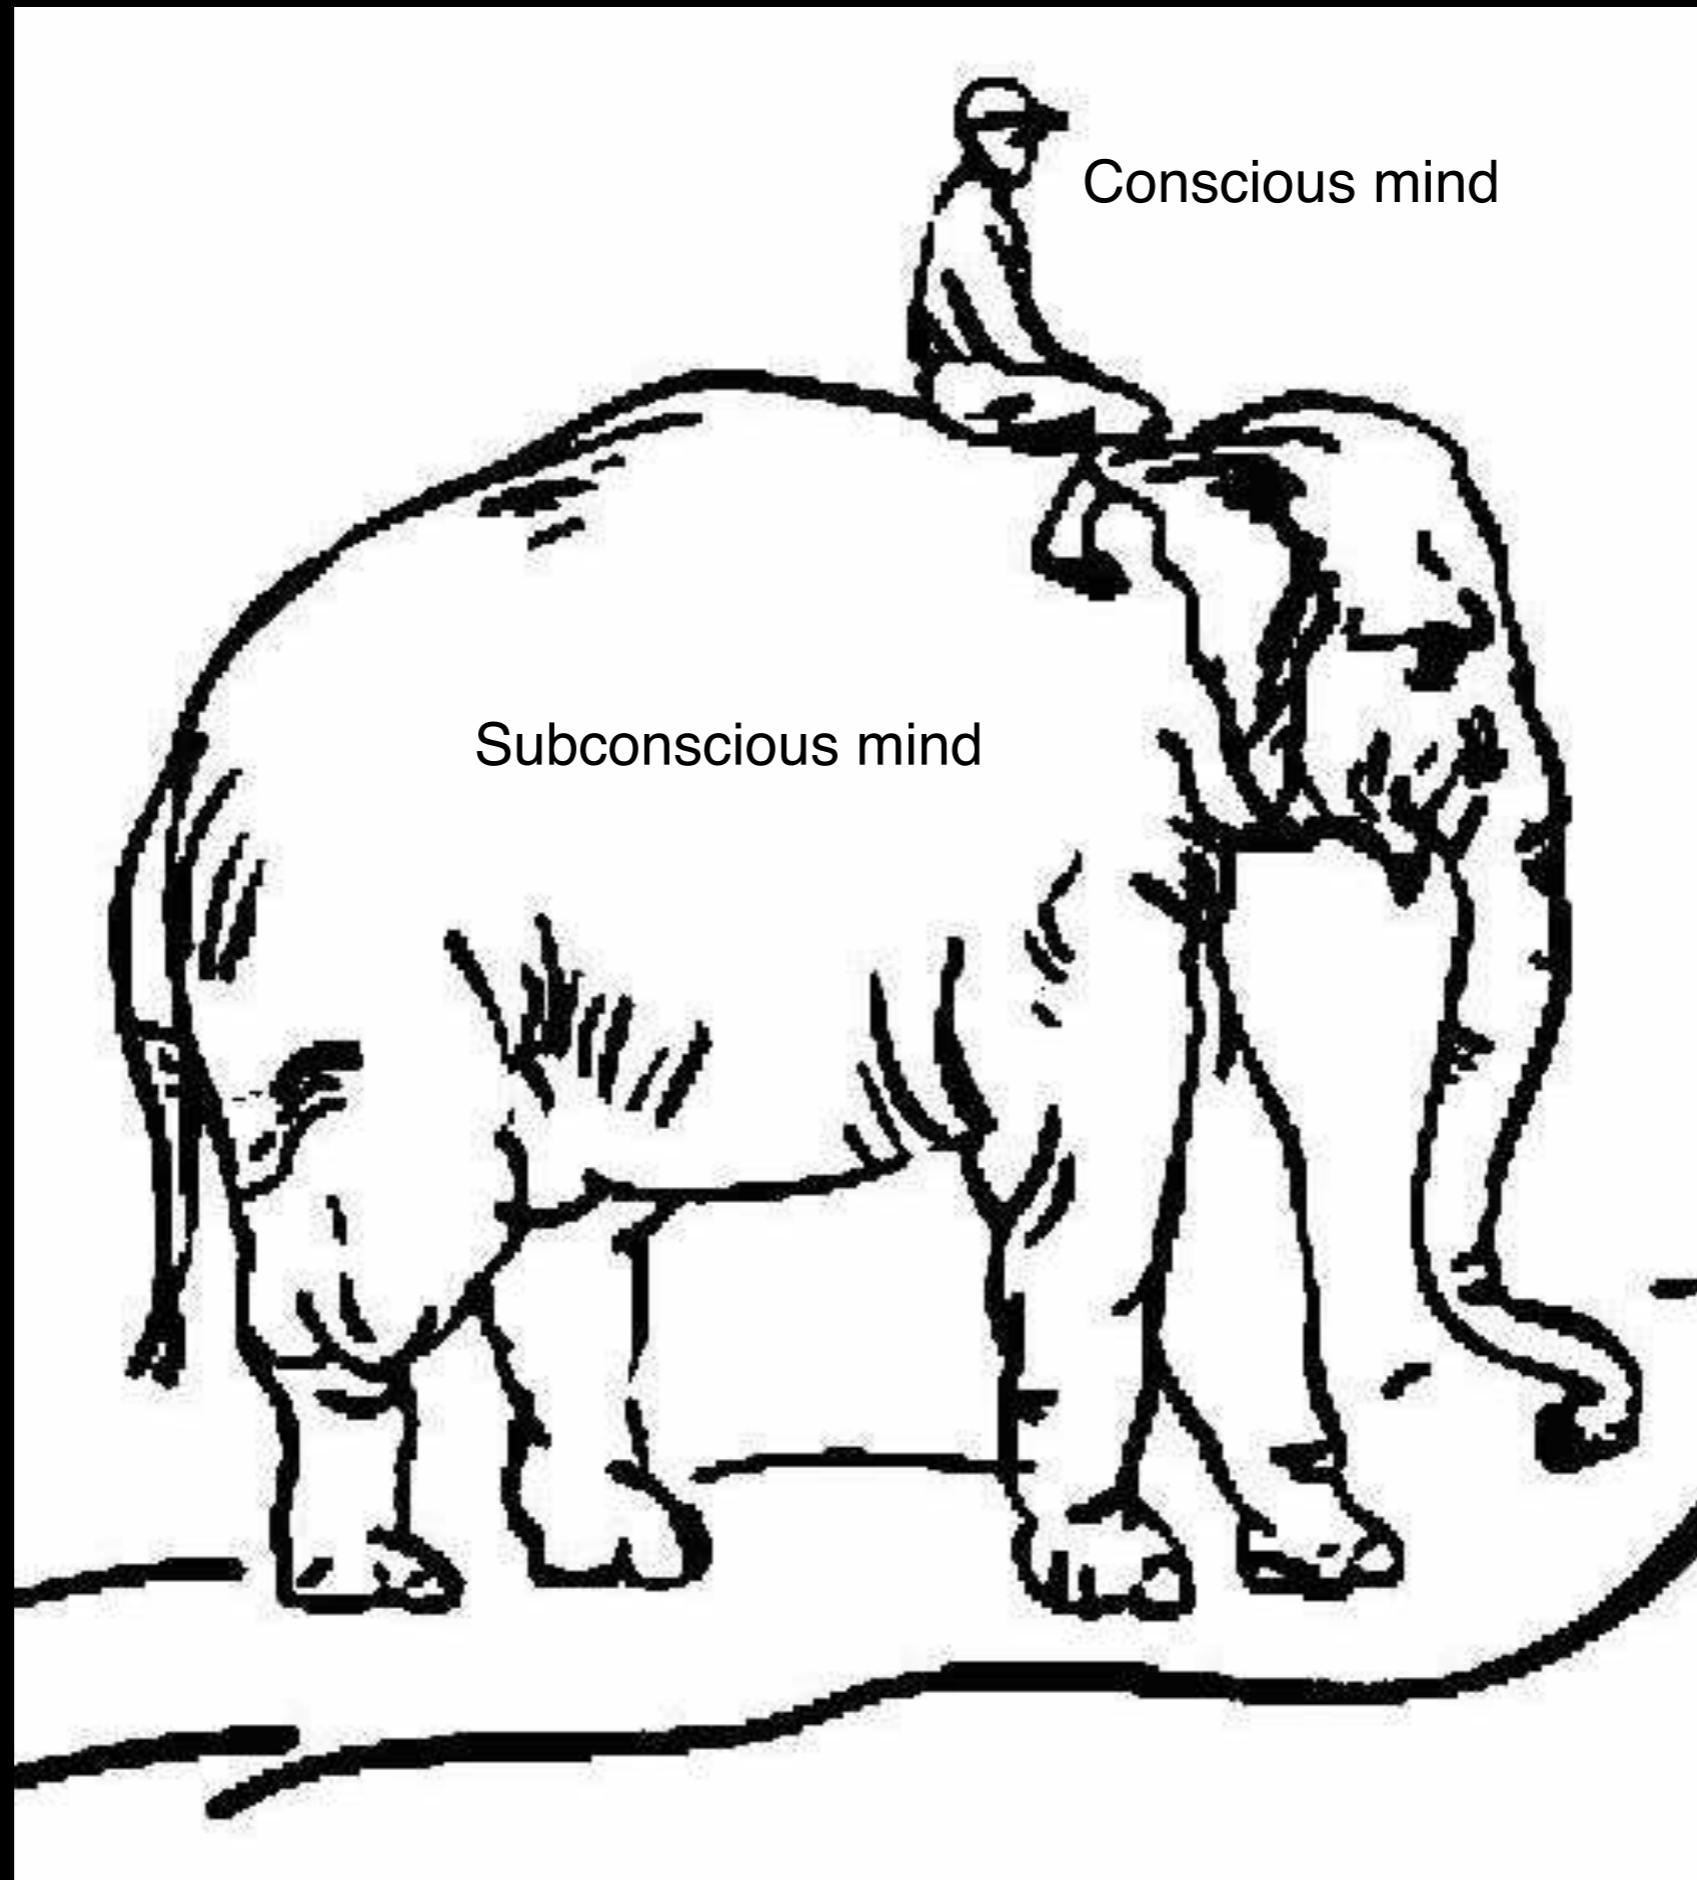


(n) x0



(o) x1

A Behavioral Perspective



Homo Econ vs. Homo Sapiens

- Rational Agent Assumption
 - perception rationality
 - preference rationality
 - process rationality
- Behavioral deviation
 - loss aversion
 - overweighing small probability
 - mental accounting
 - hyperbolic discounting
 - framing effects
 - altruism
 - price saliency

Behavioral
Science



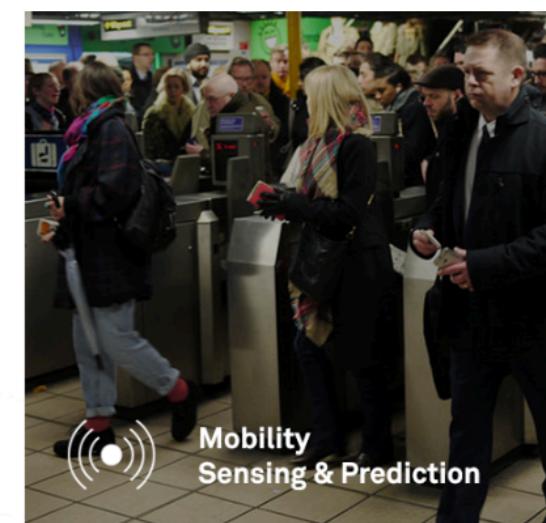
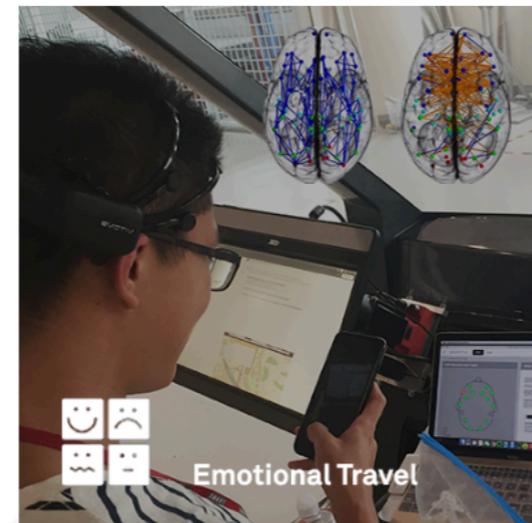
Transportation
Technology

- Emotional
- Social
- Perceptual

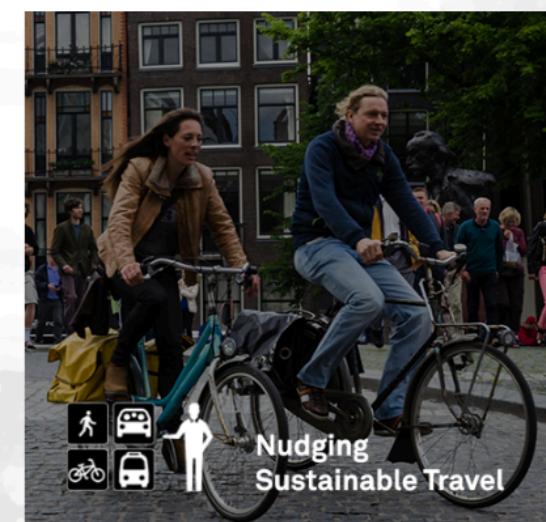
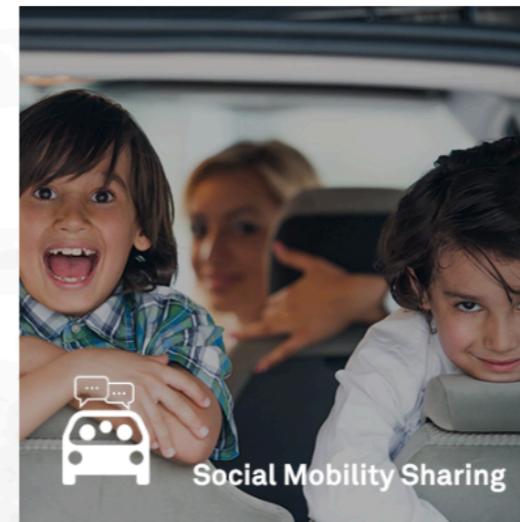
- Sharing
- Automation
- On-demand

Data Analytics and Machine Learning

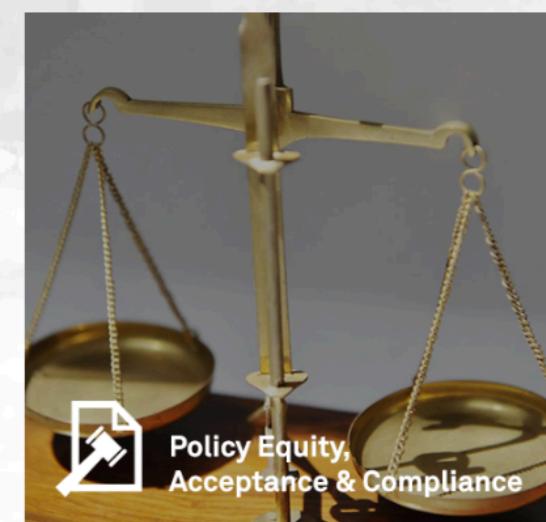
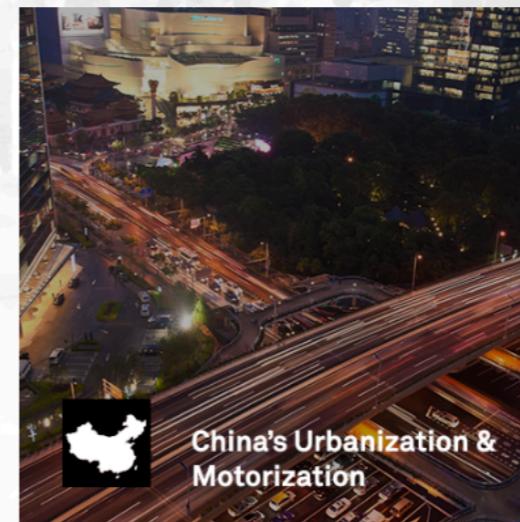
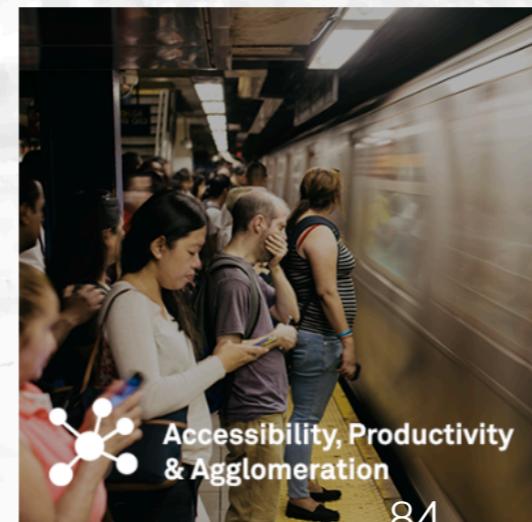
SECTION I. BEHAVIORAL FOUNDATION OF URBAN MOBILITY



SECTION II. MOBILITY SYSTEM: DESIGN WITH BEHAVIORAL PERSPECTIVE



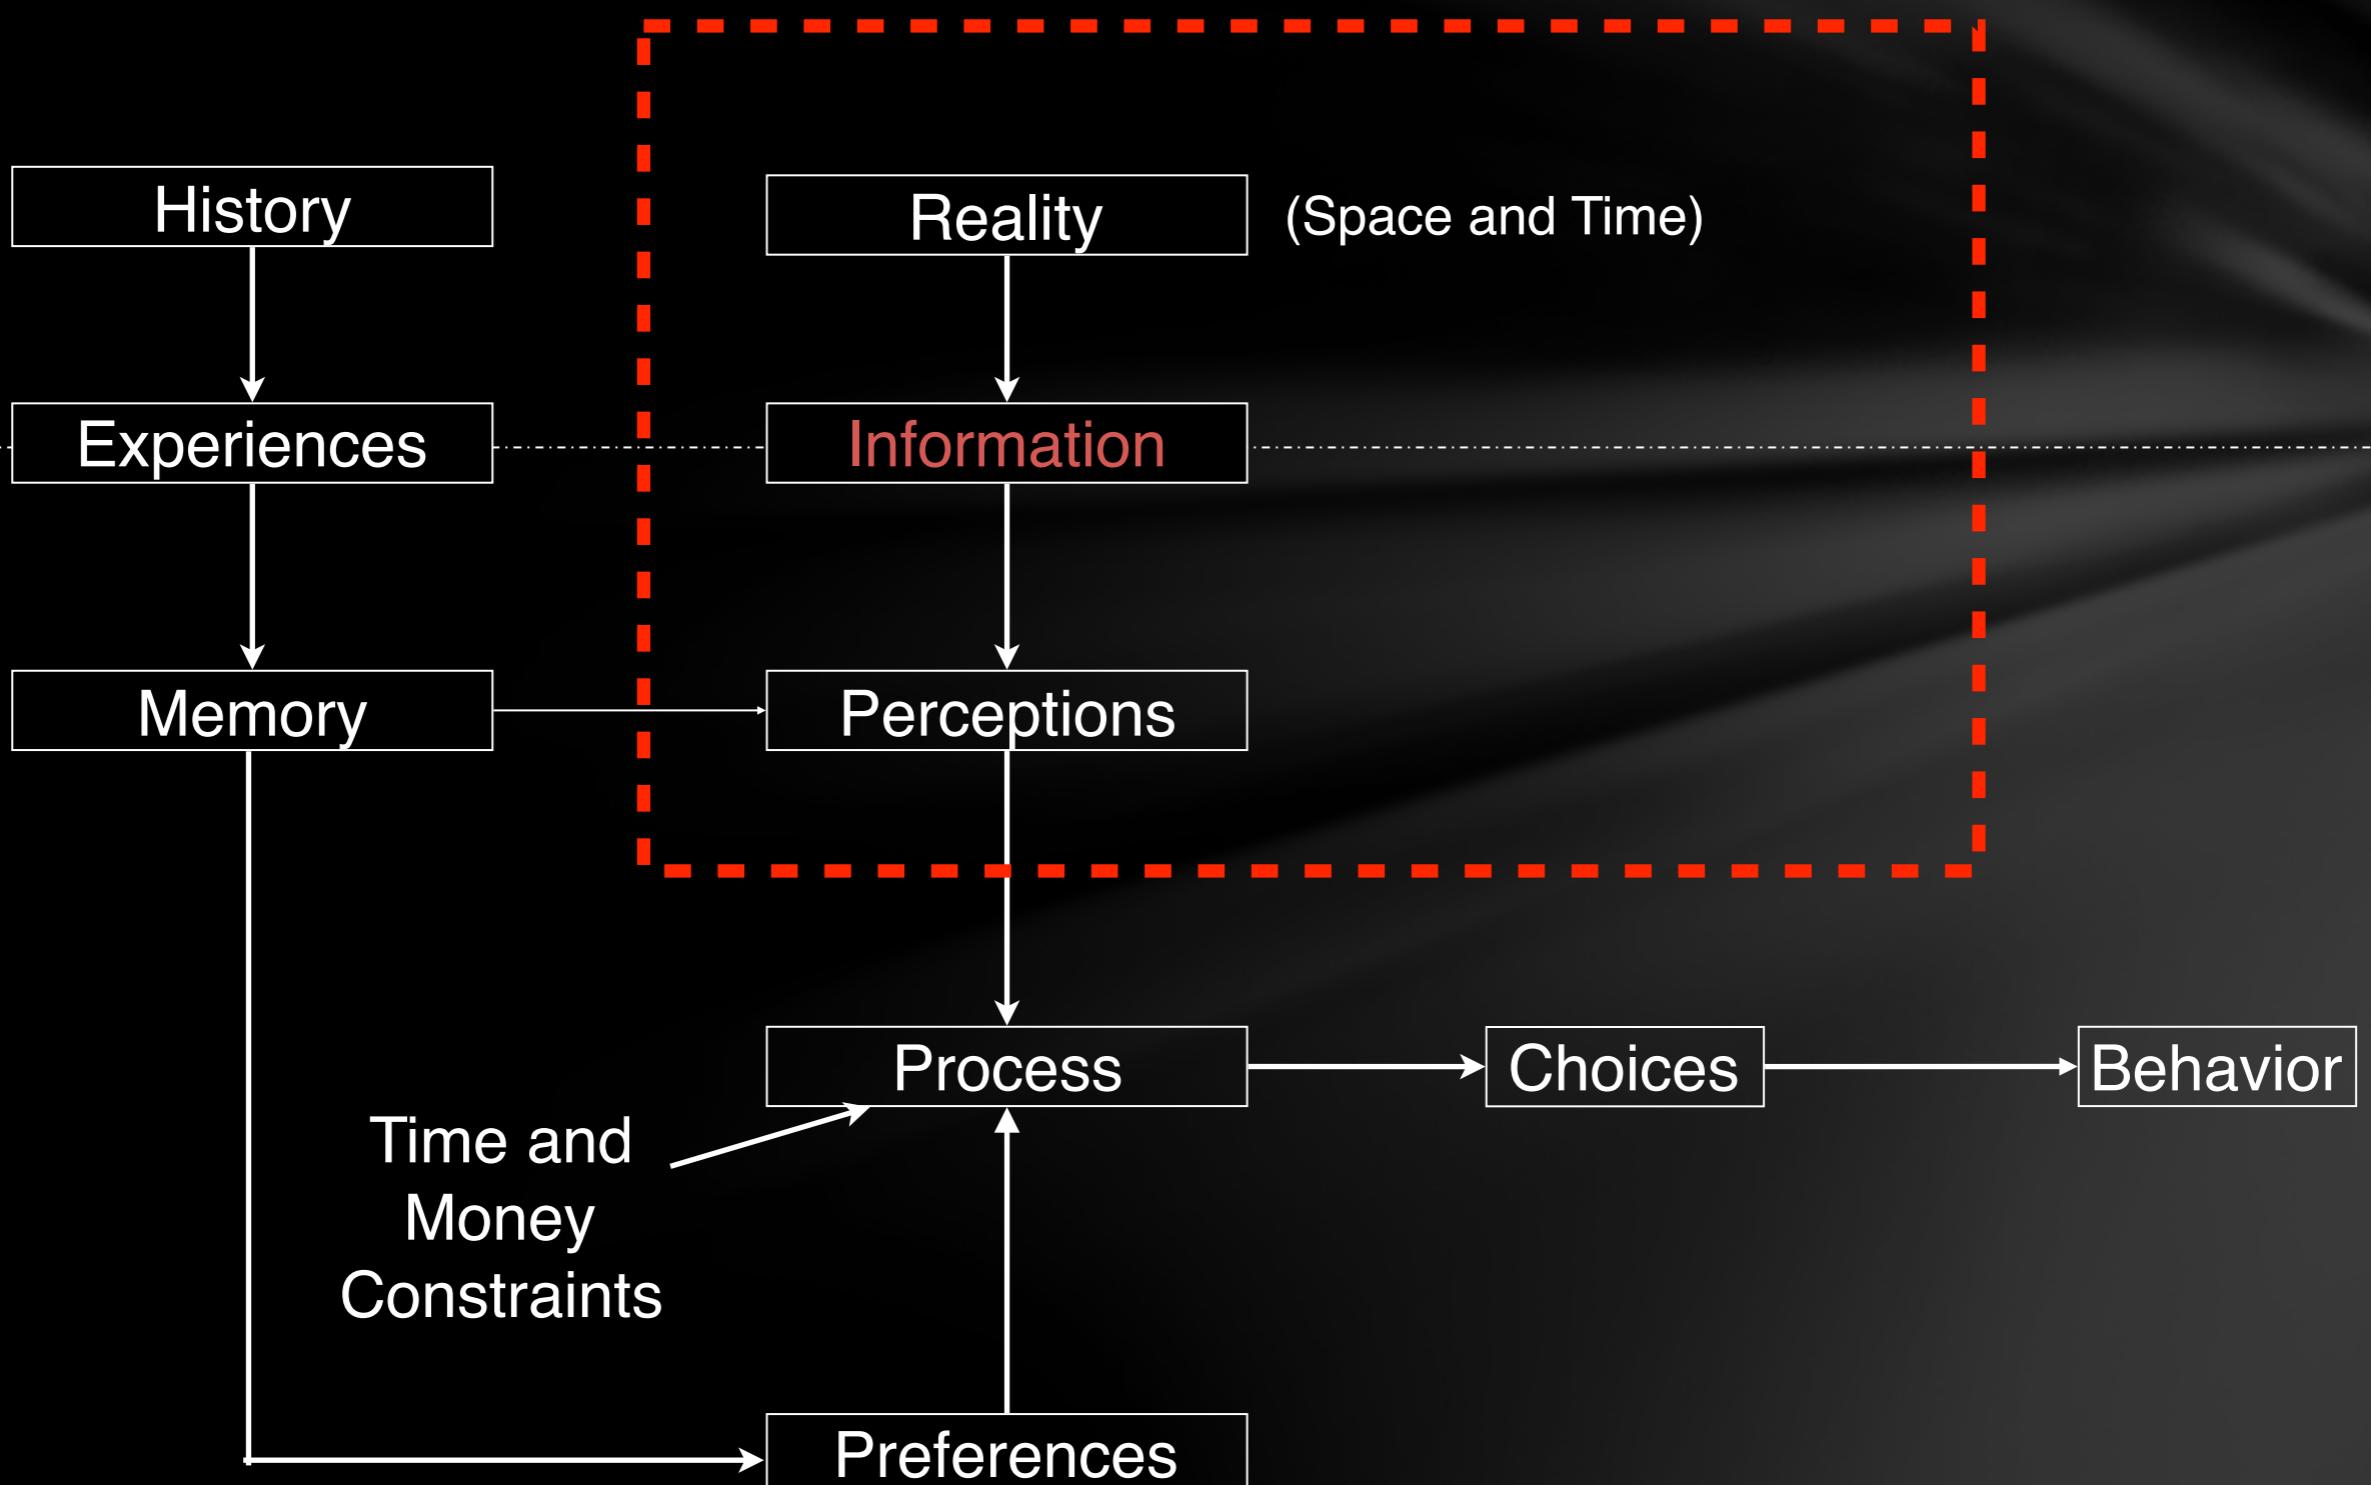
SECTION III. MOBILITY POLICY: FORMULATE, IMPLEMENT AND EVALUATE



JTL Urban Mobility Lab
mobility.mit.edu

Behavior – Theory – Policy

A Theory of Behavioral Choice



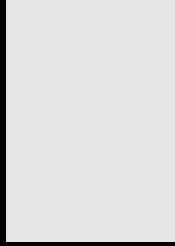
Theoretical Challenges in DNN

- Approximation
- Optimization
- Generalization

Correlation vs. Causality

Umbrella

CAUSAL HIERARCHY

| Level | Typical Activity | Typical Question | Examples |
|--|---|---|---|
|  Association $P(y x, c)$ | <p>Seeing random sampling</p> | <p>What is? How would seeing X change my belief in Y?</p> | <p>What does a symptom tell us about the disease?</p> |
|  Intervention $P(y \text{do}(x), c)$ | <p>Doing random experimentation</p> | <p>What if? What if I do X?</p> | <p>What if I take aspirin, will my headache be cured?</p> |
|  Counterfactual $P(y_x x', y')$ | <p>Imagining, Retrospecting</p> | <p>What if I had acted differently? Why?</p> | <p>Was it the aspirin that stopped my headache?</p> |

Machine Learning for Economics

SUSAN ATHEY (STANFORD GSB)

What's New About ML?

Flexible, rich, data-driven models

Increase in personalization and precision

Methods to avoid overfitting

The contrast between routine statistical analysis and data generated by machine learning can be quite stark.

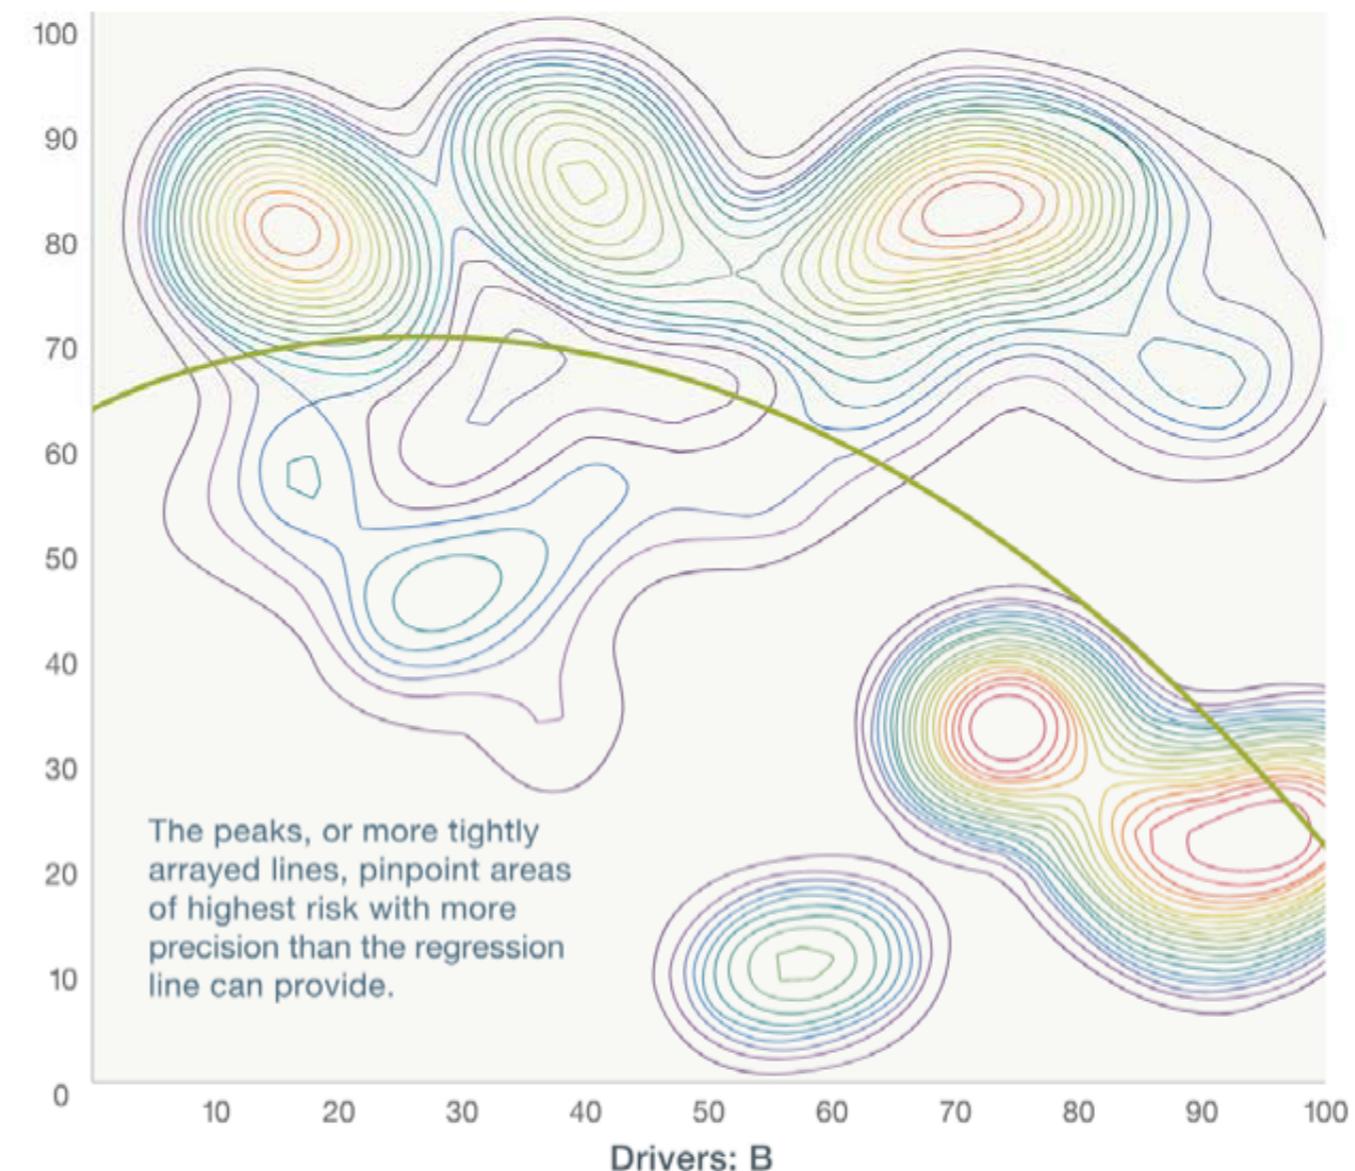
Value at risk from customer churn, telecom example

Classic regression analysis

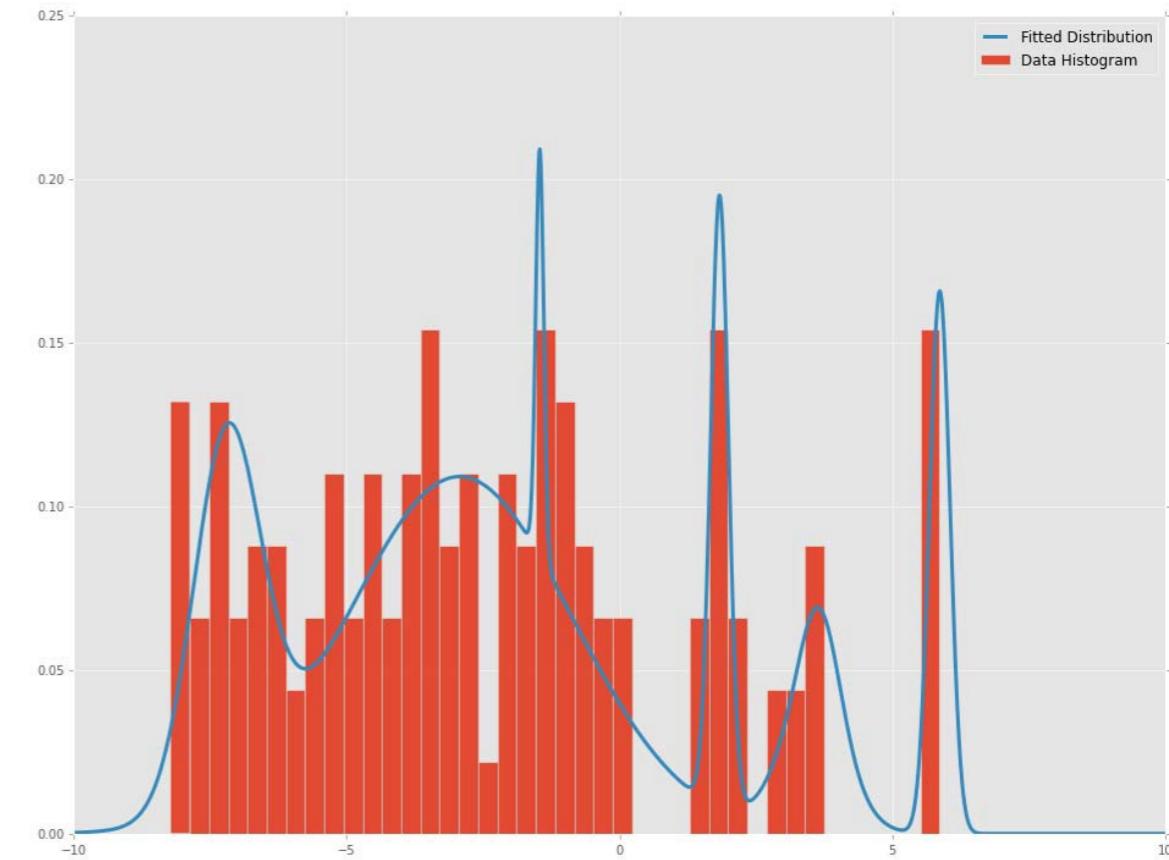
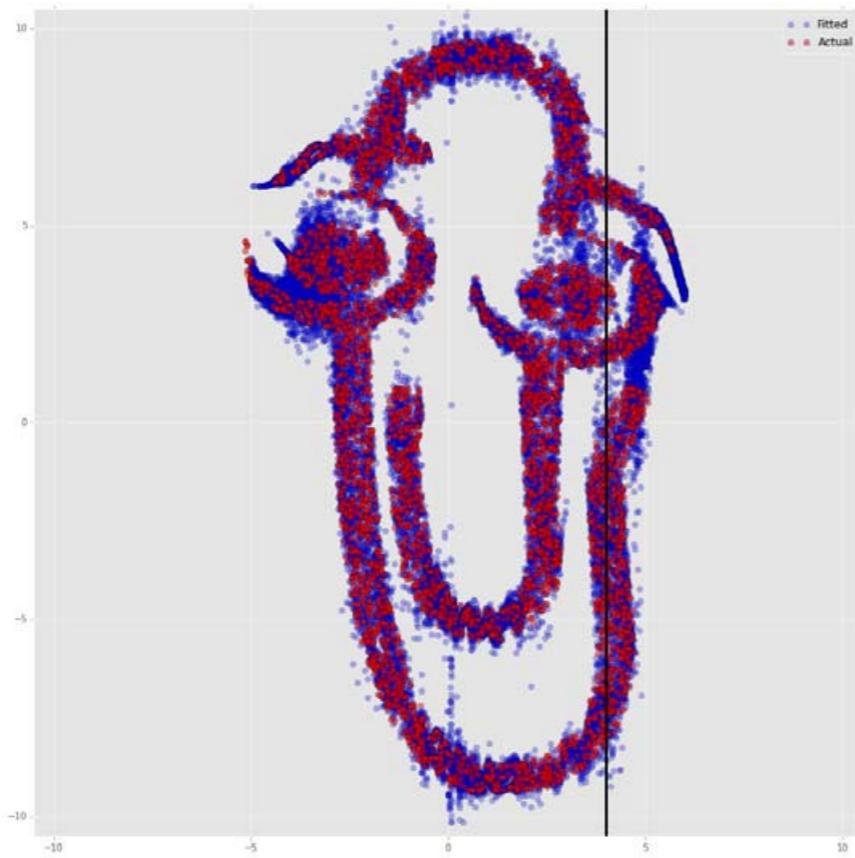


Isobar graph facilitated by machine learning: warmer colors indicate higher degrees of risk

Drivers: A



Ability to Fit Complex Shapes



Contrast with Traditional Econometrics

Economists have focused on the case with substantially more observations than covariates ($N \gg P$)

- In-sample MSE is a good approximation to out-of-sample MSE
- OLS is BLUE, and if overfitting is not a problem, then no need to incur bias
- OLS uses all the data and minimizes in-sample MSE

OLS obviously fails due to overfitting when $P \sim N$ and fails entirely when $P > N$

- ML methods generally work when $P > N$

Economists worry about estimating causal effects and identification

- Causal effects
- Counterfactual predictions
- Separating correlation from causality
- Standard errors
- Structural models incorporating behavioral assumptions

Identification problems can not be evaluated using a hold-out set

- If joint dist'n of observable same in training and test, will get the same results in both

Causal methods sacrifice goodness-of-fit to focus only on variation in data that identifies parameters of interest

What We Say v. What We Do (Econometrics)

What We Say

- Causal inference and counterfactuals
- God gave us the model
- We report estimated causal effects and appropriate standard errors
- Plus a few additional specifications for robustness

What we do

- Run OLS or IV regressions
 - Try a lot of functional forms
 - Report standard errors as if we ran only one model
 - Have research assistants run hundreds of regressions and pick a few “representative” ones
- Use complex structural models
 - Make a lot of assumptions without a great way to test them

Key Lessons for Econometrics

Many problems can be decomposed into predictive and causal parts

- Can use off-the-shelf ML for predictive parts

Data-driven model selection

- Tailored to econometric goals
- Focus on parameters of interest
- Define correct criterion for model
- Use data-driven model selection where performance can be evaluated
- While retaining ability to do inference

ML-Inspired Approaches for Robustness

Validation

- ML always has a test set
- Econometrics can consider alternatives
 - Ruiz, Athey and Blei (2017) evaluate on days with unusual prices
 - Athey, Blei, Donnelly and Ruiz (2017) evaluate change in purchases before and after price changes
 - Tech firm applications have many A/B tests and algorithm changes

Other computational approaches for structural models

- Stochastic gradient descent
- Variational Inference (Bayesian models)

See Sendhil Mullainathan et al (JEP, AER) for key lessons about prediction in economics

- See also Athey (Science, 2017)

Empirical Economics in Five Years: Susan's Predictions

Regularization/data-driven model selection will be the standard for economic models

Prediction problems better appreciated

Measurement using ML techniques an important subfield

Textual analysis standard (already many examples)

Models will explicitly distinguish causal parts and predictive parts

Reduced emphasis on sampling variation

Model robustness emphasized on equal footing with standard errors

Models with lots of latent variables

Knowledge Discovery

So How Do Computers Discover New Knowledge?

1. **Symbolists**--Fill in gaps in existing knowledge
2. **Connectionists**--Emulate the brain
3. **Evolutionists**--Simulate evolution
4. **Bayesians**--Systematically reduce uncertainty
5. **Analogizers**--Notice similarities between old and new

The Five Tribes of Machine Learning

| Tribe | Origins | Key Algorithm |
|----------------|----------------------|-------------------------|
| Symbolists | Logic, philosophy | Inverse deduction |
| Connectionists | Neuroscience | Backpropagation |
| Evolutionists | Evolutionary biology | Genetic programming |
| Bayesians | Statistics | Probabilistic inference |
| Analogizers | Psychology | Kernel machines |

SRC: Pedro Domingos ACM Webinar Nov 2015 <http://learning.acm.org/multimedia.cfm>

Representation, Evaluation and Optimization

- Representation
 - Probabilistic logic (e.g., Markov logic networks)
 - Weighted formulas → Distribution over states
- Evaluation
 - Posterior probability
 - User-defined objective function
- Optimization
 - Formula discovery: Genetic programming
 - Weighted learning: Backpropagation

SRC: Pedro Domingos ACM Webinar Nov 2015 <http://learning.acm.org/multimedia.cfm>
<http://www.ibtimes.co.uk/elon-musks-1bn-non-profit-launches-gym-train-ai-atari-games-1557362>

MIT Subjects Related to ML

- 6.036 Introduction to Machine Learning
- 6.862 Applied Machine Learning (6.862 = 6.036 + Term Project)
- 6.867 Machine Learning
- 9.520/6.860: Statistical Learning Theory and Applications
- 9.S914: Mathematical Statistics: A Non-Asymptotic Approach
- 6.883: Online Methods in Machine Learning: Theory and Applications

MIT Subjects Related to ML

- 4.S42 Machine Learning for Creative Design
- 6.S897/HST.S53: Machine Learning for Healthcare
- MAS.533 AI for Impact ~ Towards Health & Sustainability from People to Planet
- 6.268 Network Science and Models
- 11.s938/11.s196 Deep Learning for Urban Transportation

Deep Learning for Urban Transportation

Part I: Introduction and DNN Basics

Sep 10 Introduction: Deep Learning Meets Transportation

Part II: Passenger and Traffic Flow Prediction

Sep 17 Demand prediction and Deep learning basics

Sep 24 Advanced modeling techniques: ConvLSTM, Attention, individualized predictions

Oct 1 Guest Lecture Prof. Justin Dauwels

Oct 8 Advanced Applications: Generative models (VAE), graph embeddings,

Part III: DNN and Demand Analysis

Oct 22 DNN and Discrete Choice 1

Oct 29 DNN and Discrete Choice 2 and Guest Lecture by Prof. Hai Wang

Nov 5 DNN and Prospect Theory

Part IV: Reinforcement Learning and Control in Transportation

Nov 12 RL and Control Part 1

Nov 19 RL and Control Part 2

Nov 26 Guest Lecture Prof. Cathy Wu

Part V: FAT and Summary

Dec 3 Fairness, Accountability and Transparency

Dec 10 Student Presentation