

# **Deep Learning: Introducing the toolbox**

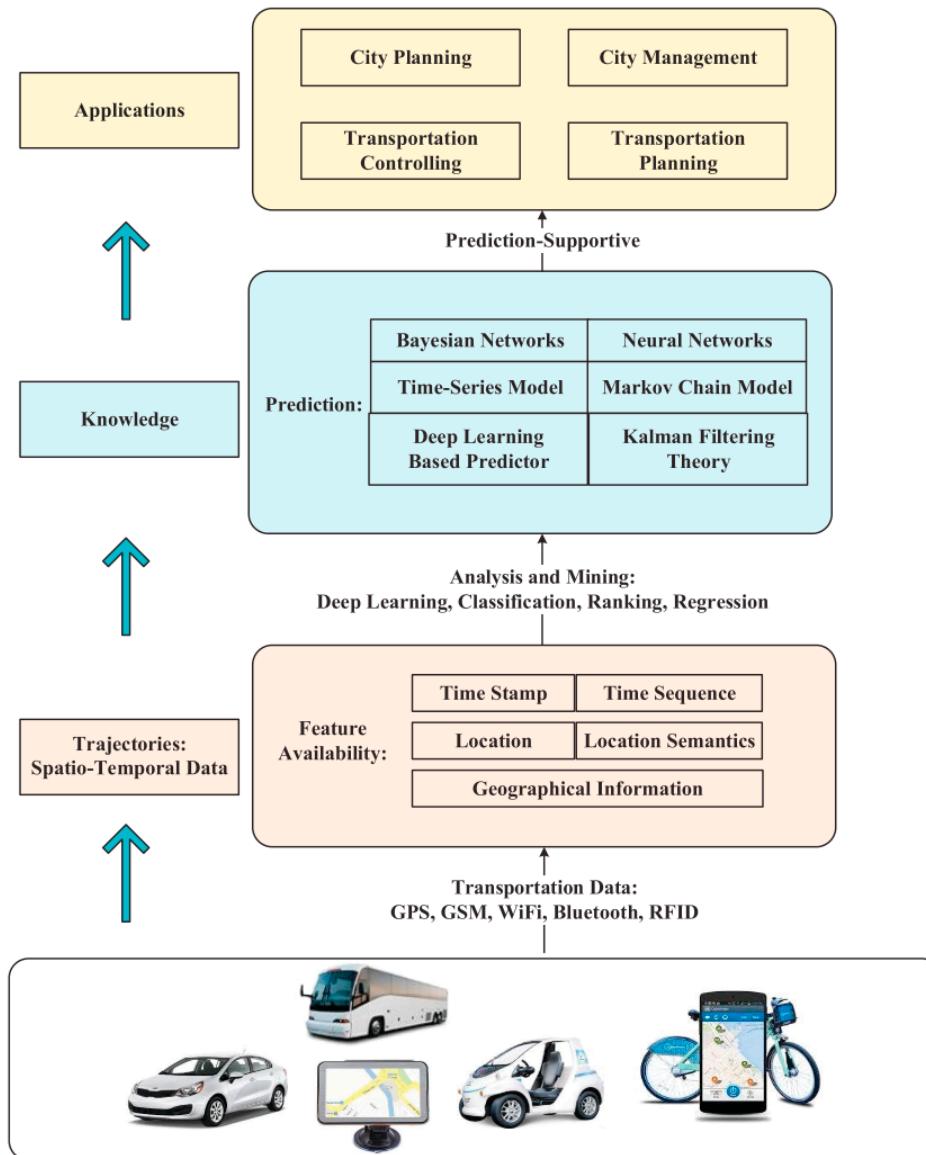
# Overview

## The Deep Learning Toolbox

- Data
- Architecture
- Activation functions
- Loss function
- Regularization
- Software frameworks



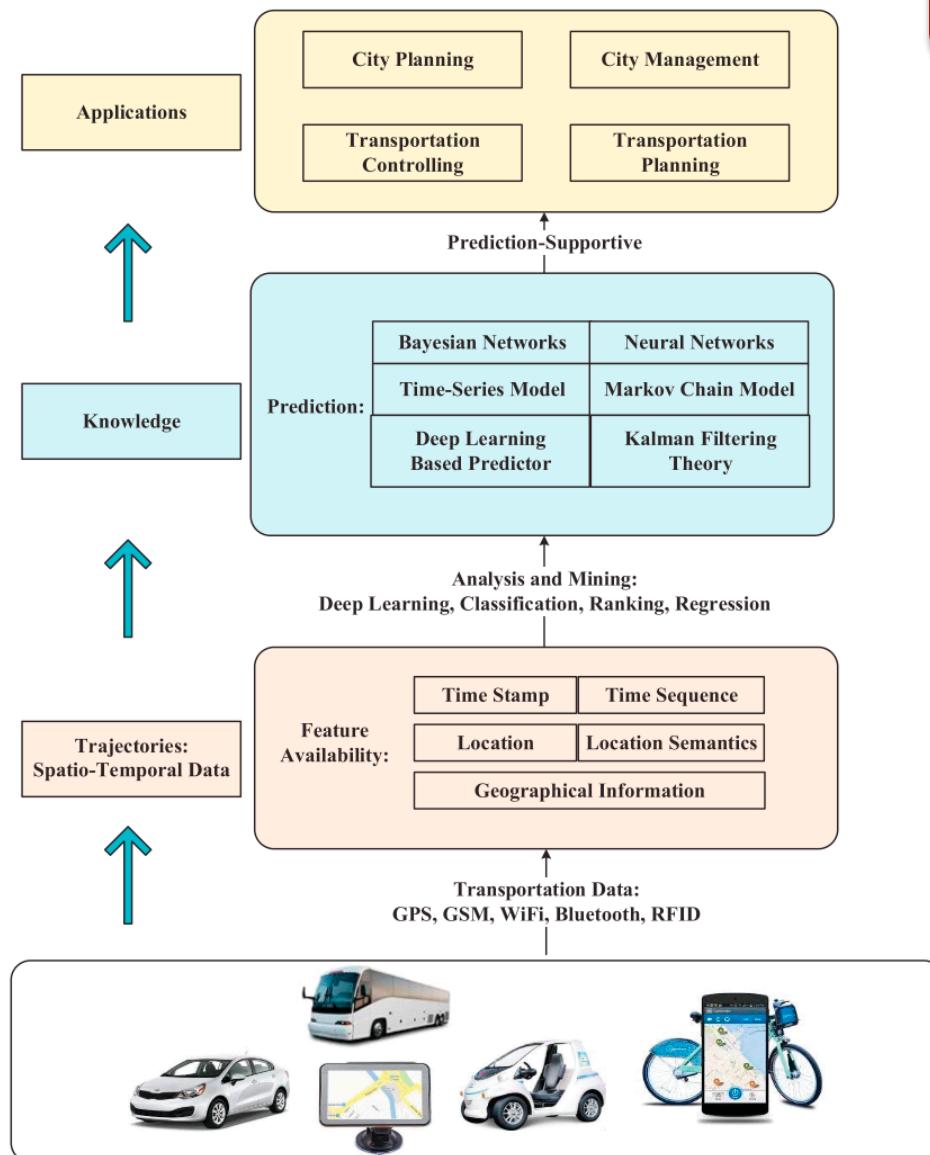
# Data



Source	Tools	Data
Smart Card	Smart Card	OD Flows, Travel Time
GPS	GPS	Vehicle Position, Vehicle Density, Vehicle Speed
Video	Video Camera	Vehicle Position, Vehicle Speed, Vehicle Density, Vehicle Classification
Road Site Sensor	Induction Loops, Road Tubes, Microwave Radar, LIDAR/Infared Acoustic, Toll Plazas	Vehicle Position, Vehicle Speed, Vehicle Density, Vehicle Classification
Floating Car Sensor	License Plate Recognition, Transponders	Travel Time, OD Flows
Wide Area Sensor	GPS, Cell phone Tracking, Airborne Sensors	Travel Time, OD Flows
Connected Vehicles	Diverse Sensors and Autonomous Vehicles (CAVs)	Coordinate, speed, acceleration, safety data,
Passive Collection	Social Mobile Data	Travel Time, OD Flows
Other Sources	Smart Grid, Smart Meters, Cellular Service, Dedicated Tests	Electric and Energy Consumption, Location, Channel Data

# Data

Most common!



Source	Tools	Data	
Smart Card	Smart Card	OD	Flows, Travel Time
GPS	GPS	Vehicle Position, Vehicle Density, Vehicle Speed	Vehicle
Video	Video Camera	Position, Vehicle Speed, Vehicle Density, Vehicle Classification	Position, Vehicle Speed, Vehicle Density, Vehicle Classification
Road Site Sensor	Induction Loops, Road Tubes, Microwave Radar, LIDAR/Infared Acoustic, Toll Plazas	Vehicle Position, Vehicle Speed, Vehicle Density, Vehicle Classification	Vehicle Position, Vehicle Speed, Vehicle Density, Vehicle Classification
Floating Car Sensor	License Plate Recognition, Transponders	Travel Time, OD Flows	Travel Time, OD Flows
Wide Area Sensor	GPS, Cell phone Tracking, Airborne Sensors	Travel Time, OD Flows	Travel Time, OD Flows
Connected Vehicles	Diverse Sensors and Autonomous Vehicles (CAVs)	Coordinate, speed, acceleration, safety data,	Coordinate, speed, acceleration, safety data,
Passive Collection	Social Mobile Data	Media, Phone	Travel Time, OD Flows
Other Sources	Smart Grid, Smart Meters, Cellular Service, Dedicated Tests	Electric and Energy Consumption, Location, Channel Data	Electric and Energy Consumption, Location, Channel Data

# Data

Smart Card Data: a sample of two weeks' data of a major subway system

New York City TLC Trip Record [[Link](#)]

NHTS 2017

Blue Bikes Comprehensive Trip Histories [[Link](#)]

Open transit data toolkit [[Link](#)]

- MTA Bus Time® Historical Data [[Link](#)]
  - Bus speeds
  - headways
- Subway Turnstile Data
  - Turnstile data tracks the number of entries and exits of passengers in the subway system

Didi Opendata [[Link](#)]

New York City MTA GTFS data [[Link](#)]

NYC open data including many transportation datasets [[Link](#)]

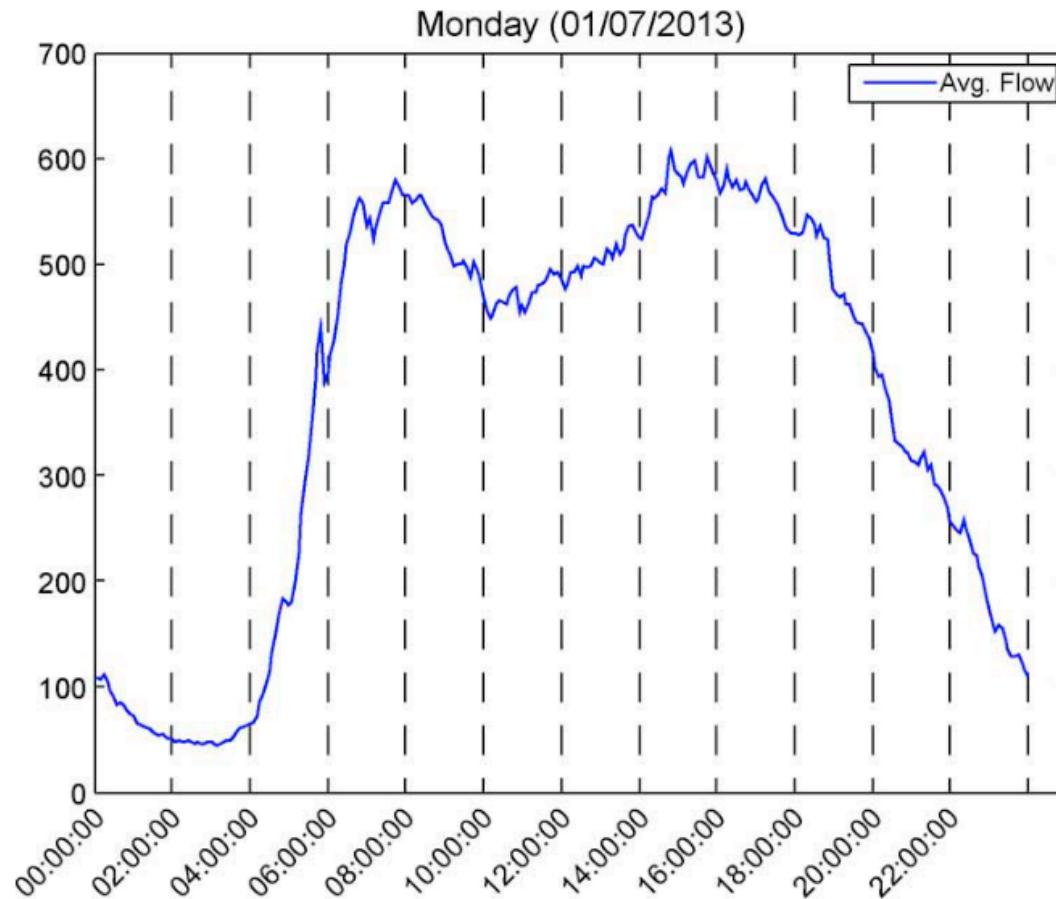
Flow provides a way of creating a simulation based on OpenStreetMap data [[Link](#)]

Google Street View [[Link](#)]

IOT data [[link](#)]

# A simple problem

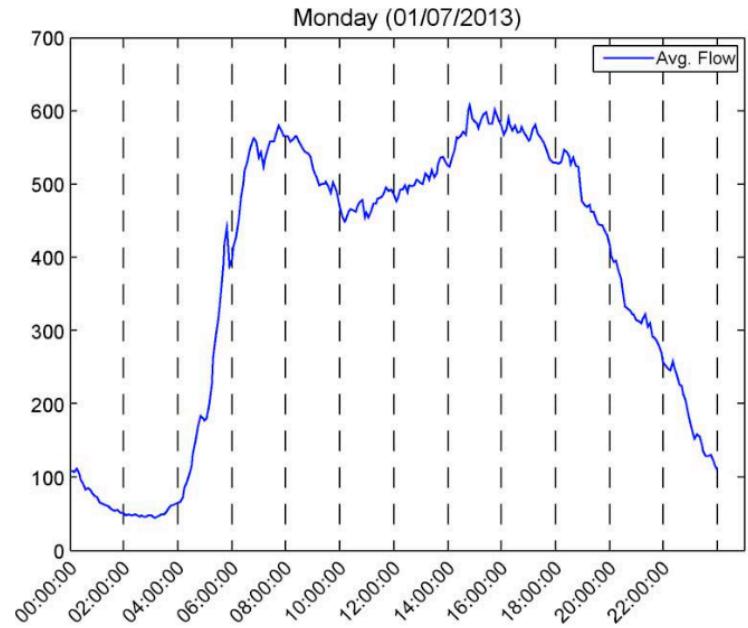
- Predict traffic flow for the next 15 minutes



# Traditional approaches

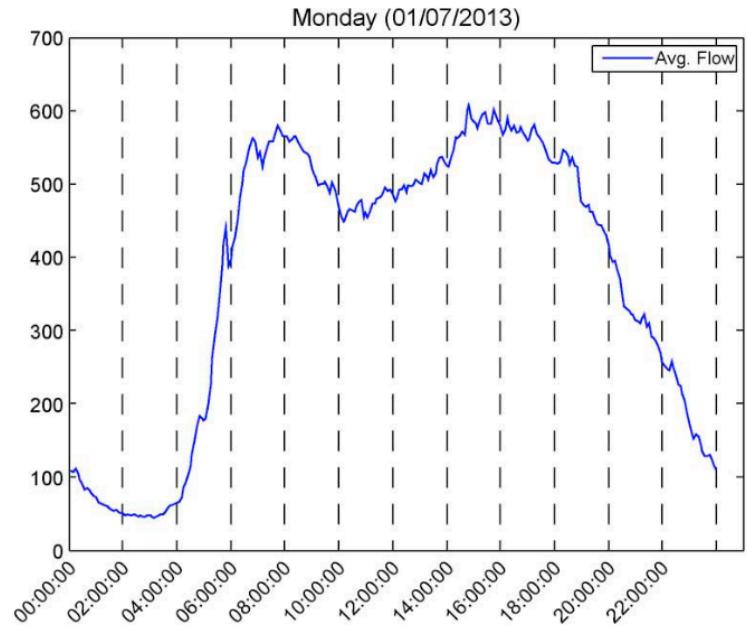
- ARIMA
  - ARIMAX
  - SARIMA

$$y_t = \sum_{j=1}^r \phi_j y_{t-j} + \sum_{j=1}^{r-1} \psi_j \epsilon_{t-j} + \epsilon_t$$



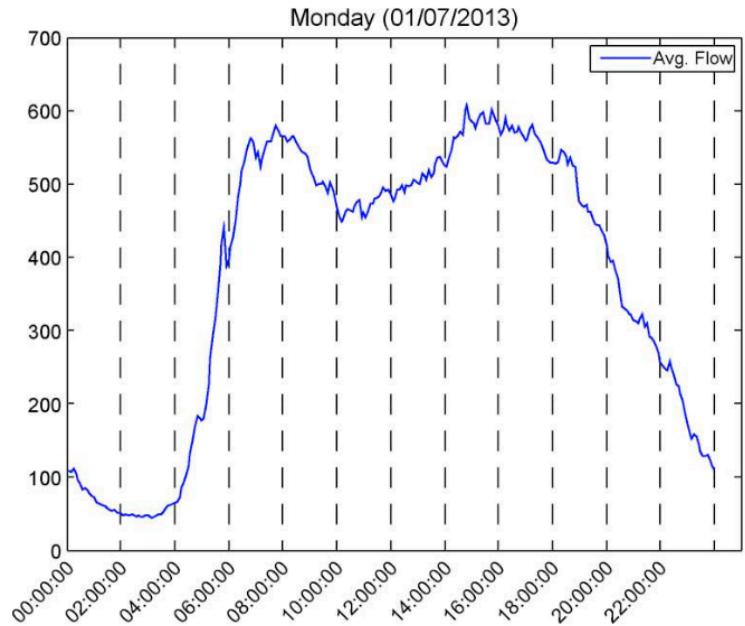
# Traditional approaches

- ARIMA
  - ARIMAX
  - SARIMA
- KNN



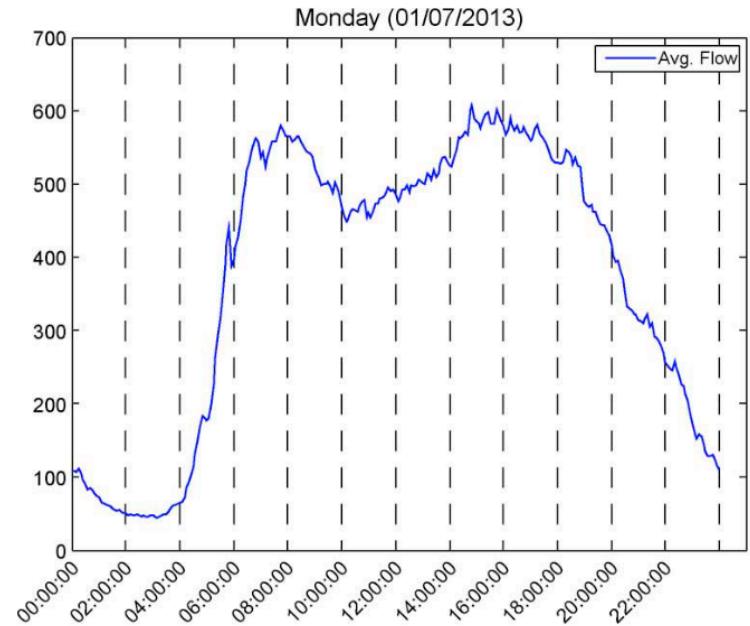
# Traditional approaches

- ARIMA
  - ARIMAX
  - SARIMA
- KNN
- ANN (1 hidden layer)



# Traditional approaches

- ARIMA
  - ARIMAX
  - SARIMA
- KNN
- ANN (1 hidden layer)
- SVM



# Our first DL model

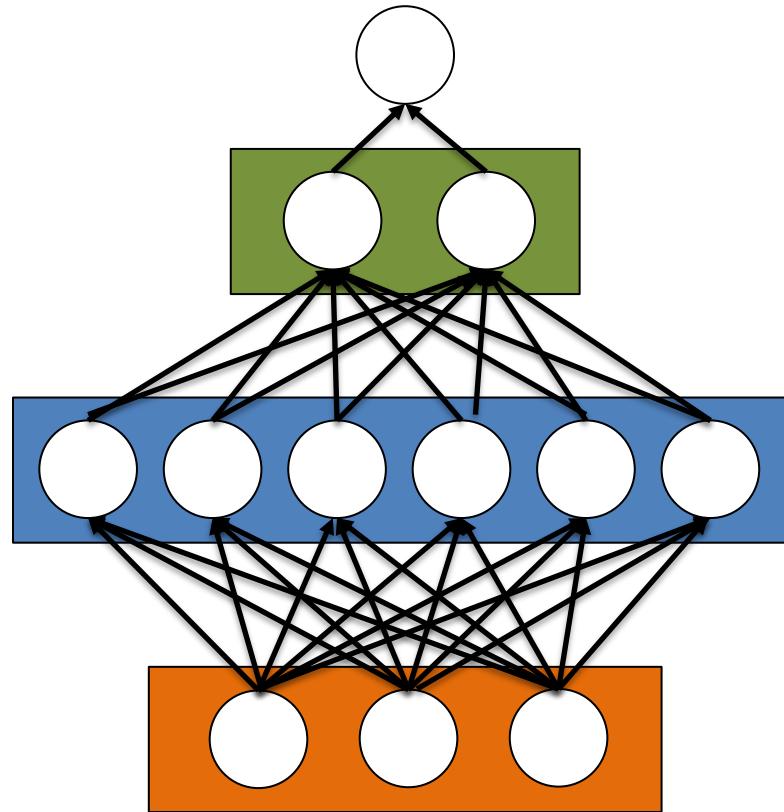
- What is “deep”, exactly?

# Our first DL model

- What is “deep”, exactly?
  - More than 1 hidden layer

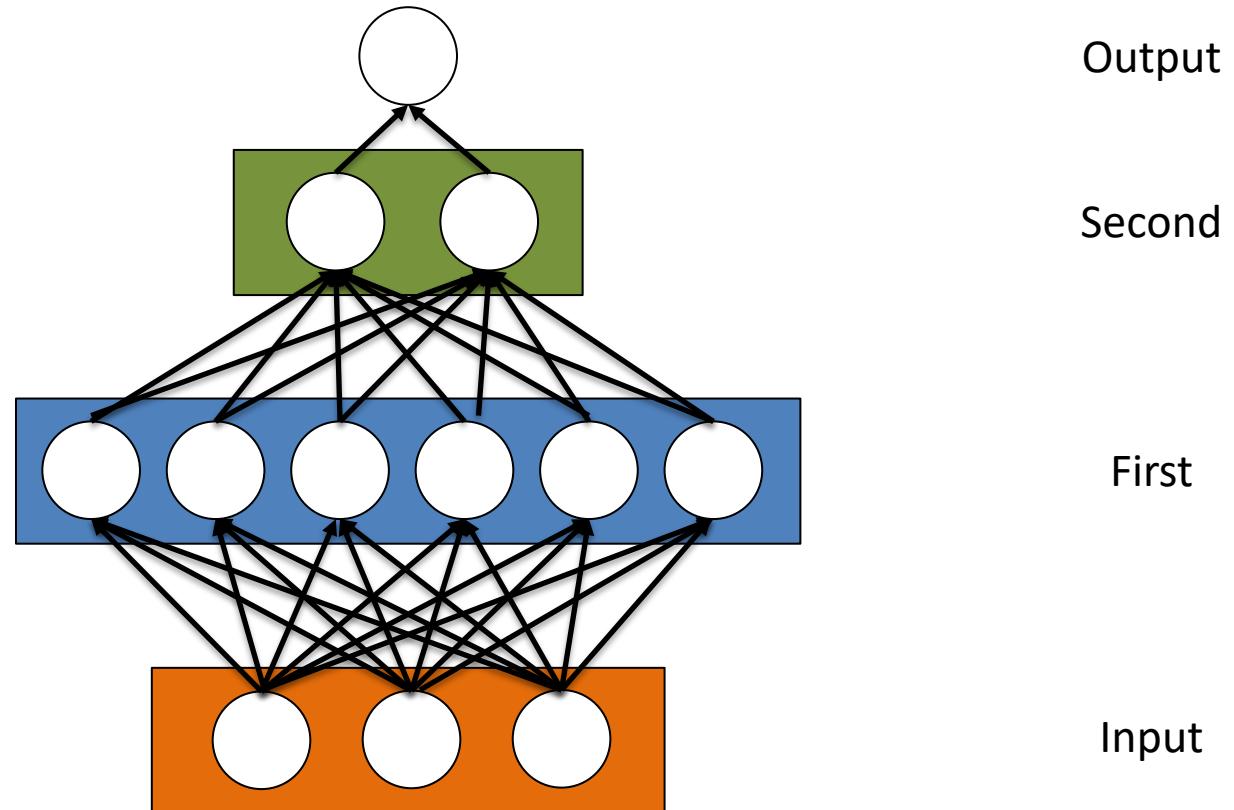
# Our first DL model

- What is “deep”, exactly?
  - More than 1 hidden layer



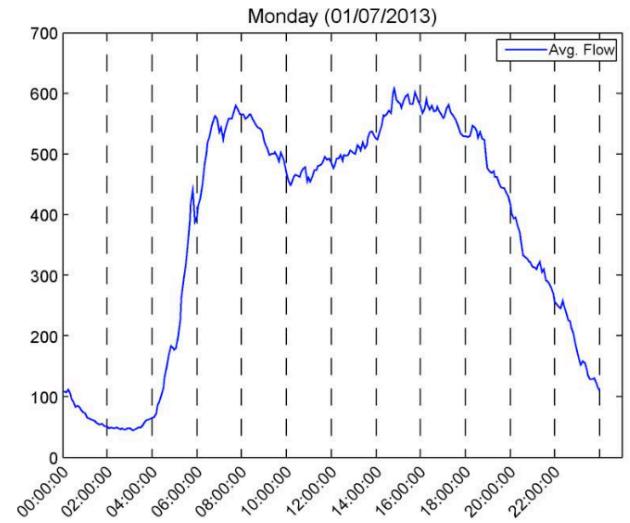
# Our first DL model

- What is “deep”, exactly?
  - More than 1 hidden layer



# Our first DL model

- How should we represent our input data?
  - Data collected every 30 s
  - Aggregated over 5 min intervals
  - $D = \{x_t, x_{t-1}, \dots, x_1\}$

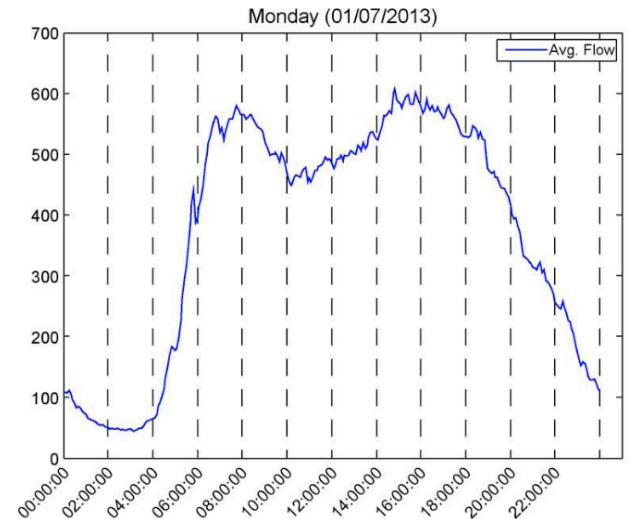


# Our first DL model

- $D = \{x_t, x_{t-1}, \dots, x_1\}$
- Let's turn it into a supervised learning problem
  - Choose a lookback horizon  $h$

$$\{x_t\} \quad \{x_{t-1}, x_{t-2}, \dots, x_{t-h}\}$$

$$\{x_{t-1}\} \quad \{x_{t-2}, x_{t-1}, \dots, x_{t-h-1}\}$$



# Our first DL model

- $D = \{x_t, x_{t-1}, \dots, x_1\}$
- Let's turn it into a supervised learning problem
  - Choose a lookback horizon  $h$

$$\{x_t\} \quad \{x_{t-1}, x_{t-2}, \dots, x_{t-h}\}$$

$$\{x_{t-1}\} \quad \{x_{t-2}, x_{t-3}, \dots, x_{t-h-1}\}$$

- Note: fixed input size

# Our first DL model

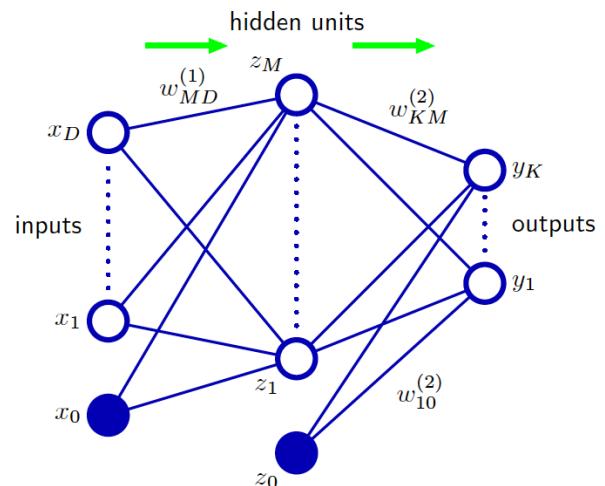
- Data
- Architecture
- Activation functions
- Losses
- Regularization
- Software frameworks

# Let's start with the simplest architecture

D-dimensional input

1 hidden layer comprising M hidden units

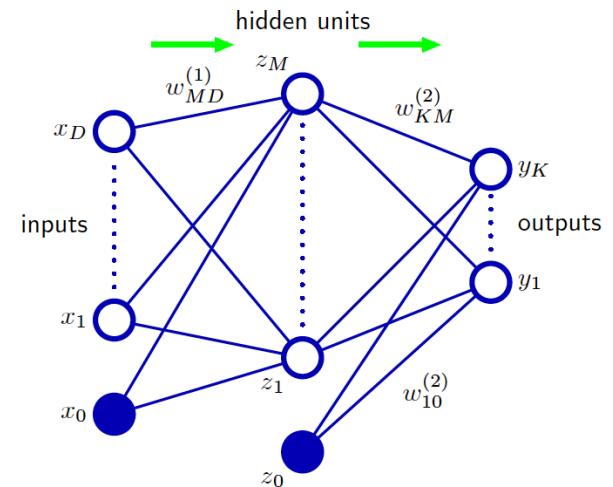
K-dimensional output



# Let's start with the simplest architecture

First, we construct M linear combinations  
of the input variables

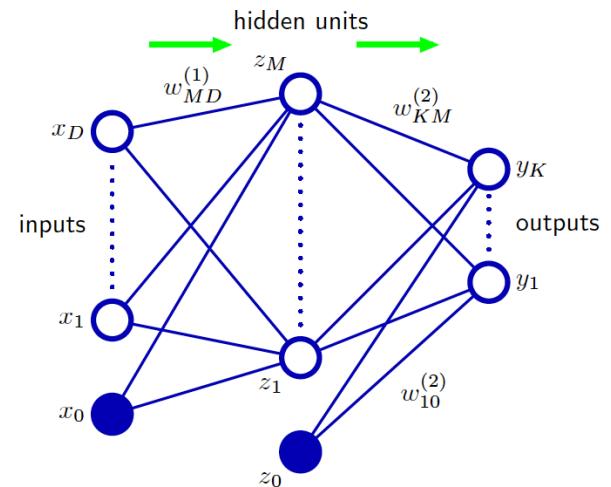
$$a_j = \sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)}$$



# Let's start with the simplest architecture

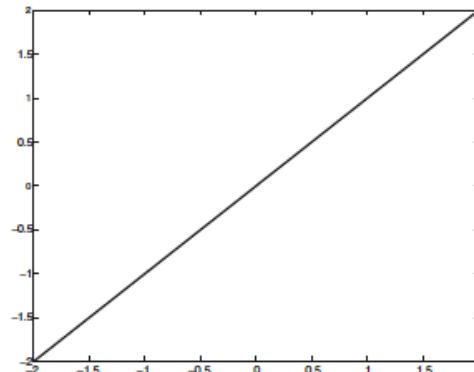
Then, we pass the output through a nonlinear activation function

$$z_j = h(a_j).$$

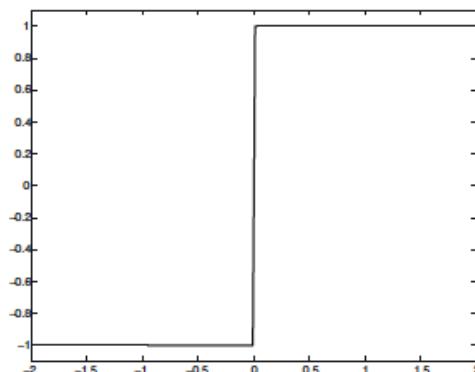


# Let's start with the simplest architecture

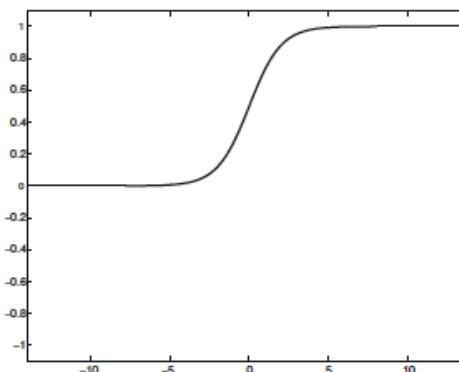
Examples of activation functions



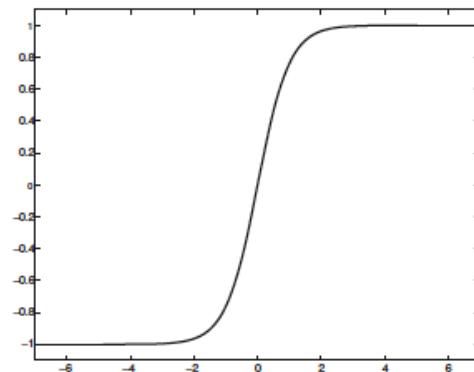
(a) Identity



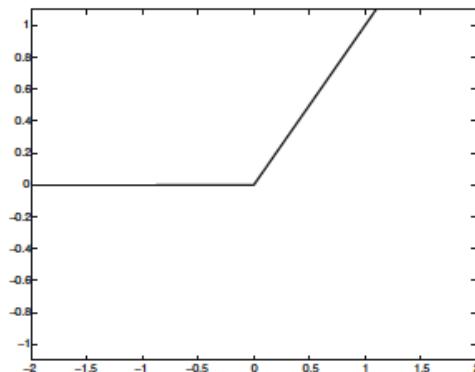
(b) Sign



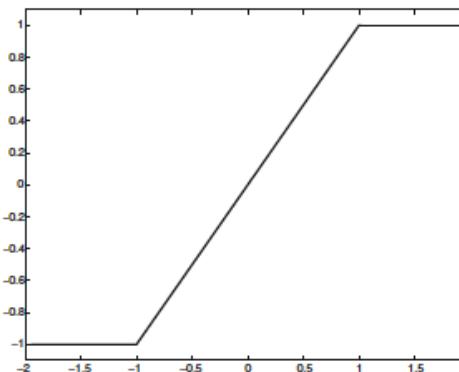
(c) Sigmoid



(d) Tanh



(e) ReLU

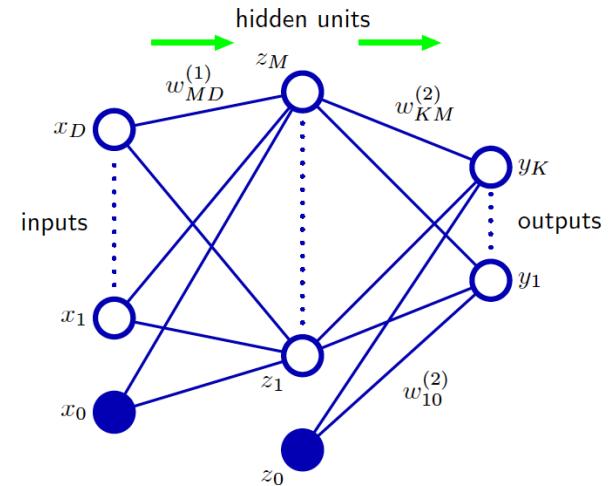


(f) Hard Tanh

# Let's start with the simplest architecture

The outputs are then again linearly combined to give output unit activations

$$a_k = \sum_{j=1}^M w_{kj}^{(2)} z_j + w_{k0}^{(2)}$$



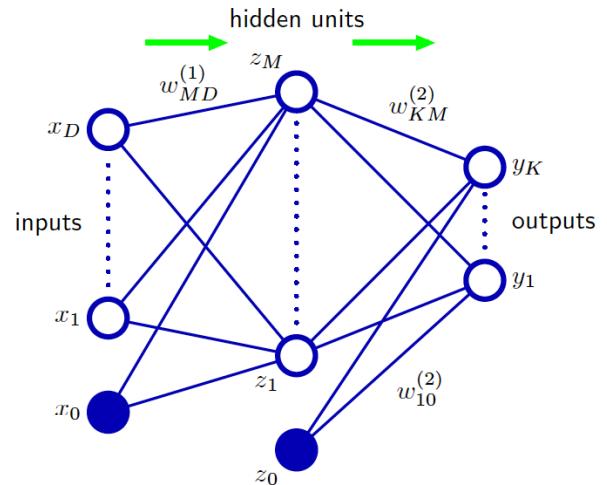
# Let's start with the simplest architecture

Finally, the output is passed through another activation function

$$y_k = \sigma(a_k)$$

For regression, it's the identity function

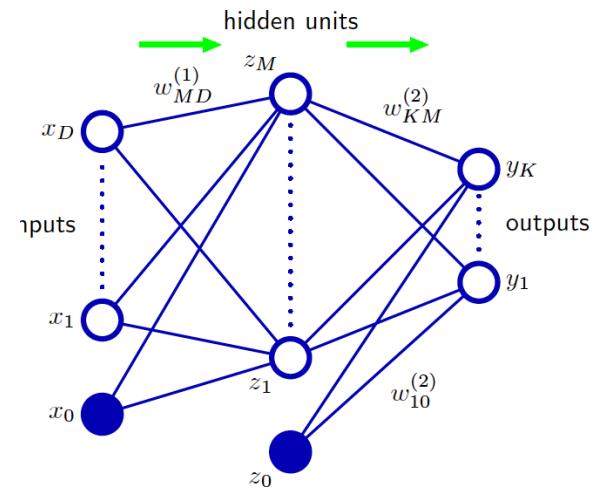
For classification, it's SoftMax  $\frac{\exp(a_k)}{\sum_j \exp(a_j)}$



# Let's start with the simplest architecture

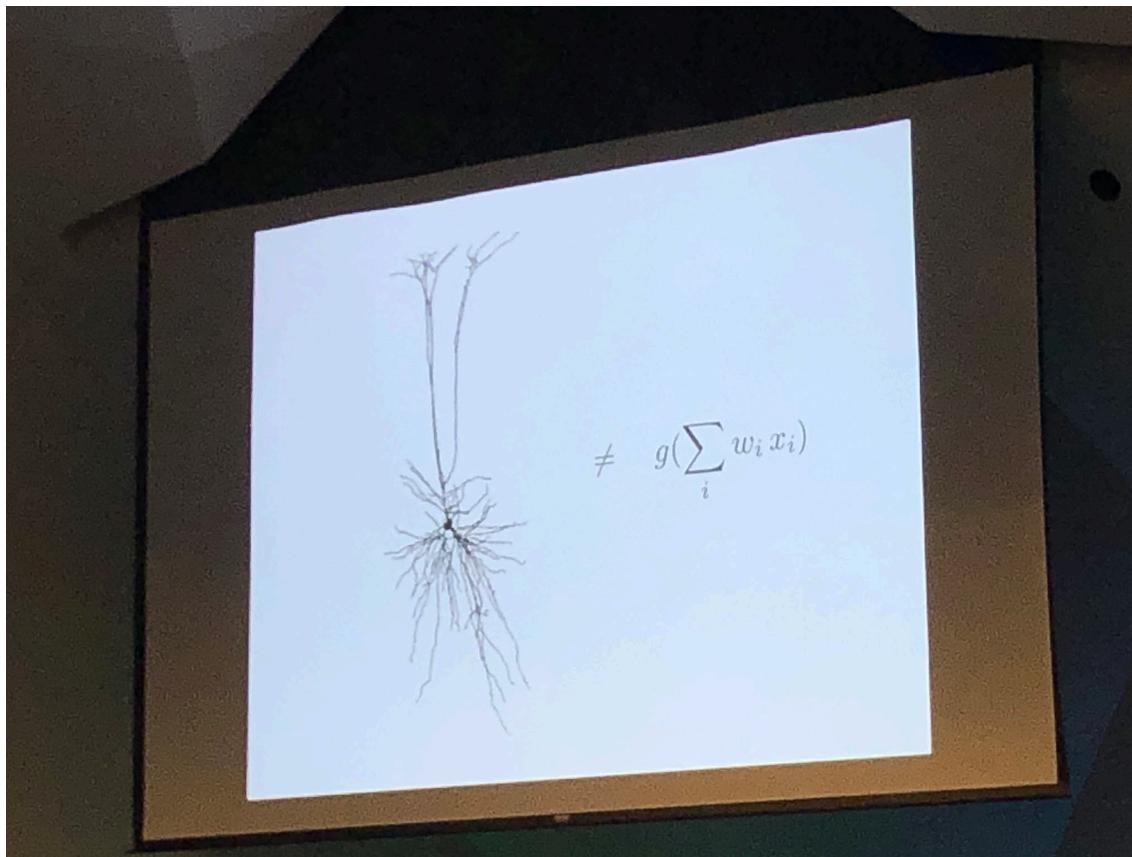
In summary

$$y_k(\mathbf{x}, \mathbf{w}) = \sigma \left( \sum_{j=1}^M w_{kj}^{(2)} h \left( \sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)} \right) + w_{k0}^{(2)} \right)$$



# Let's start with the simplest architecture

$$y_k(\mathbf{x}, \mathbf{w}) = \sigma \left( \sum_{j=1}^M w_{kj}^{(2)} h \left( \sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)} \right) + w_{k0}^{(2)} \right)$$



# Our first DL model

- ~~Data~~
- ~~Architecture~~
- ~~Activation functions~~
- Losses
- Regularization
- Software frameworks

# Our first DL model

- Loss function
  - Quantifies how unhappy you will be if your predictions are wrong

# Our first DL model

- Loss function
  - Quantifies how unhappy you will be if your predictions are wrong
  - For classification

$$L(f(x_i), y_i) = \mathbf{1}_{f(x_i) \neq y_i}$$

- For regression

$$L(f(x_i), y_i) = (f(x_i) - y_i)^2$$

# Module: tf.losses

## Classes

`class Reduction`: Types of loss reduction.

## Functions

`absolute_difference(...)`: Adds an Absolute Difference loss to the training procedure.

`add_loss(...)`: Adds a externally defined loss to the collection of losses.

`compute_weighted_loss(...)`: Computes the weighted loss.

`cosine_distance(...)`: Adds a cosine-distance loss to the training procedure. (deprecated arguments)

`get_losses(...)`: Gets the list of losses from the loss\_collection.

`get_regularization_loss(...)`: Gets the total regularization loss.

`get_regularization_losses(...)`: Gets the list of regularization losses.

`get_total_loss(...)`: Returns a tensor whose value represents the total loss.

`hinge_loss(...)`: Adds a hinge loss to the training procedure.

`huber_loss(...)`: Adds a Huber Loss term to the training procedure.

`log_loss(...)`: Adds a Log Loss term to the training procedure.

`mean_pairwise_squared_error(...)`: Adds a pairwise-errors-squared loss to the training procedure.

`mean_squared_error(...)`: Adds a Sum-of-Squares loss to the training procedure.

`sigmoid_cross_entropy(...)`: Creates a cross-entropy loss using `tf.nn.sigmoid_cross_entropy_with_logits`.

`softmax_cross_entropy(...)`: Creates a cross-entropy loss using `tf.nn.softmax_cross_entropy_with_logits_v2`.

`sparse_softmax_cross_entropy(...)`: Cross-entropy loss using  
`tf.nn.sparse_softmax_cross_entropy_with_logits`.

# Our first DL model

- Optimize!

$$\underset{\text{_____}}{\text{minimize}} \sum_{i=1}^n L(f(x_i), y_i)$$

# Our first DL model

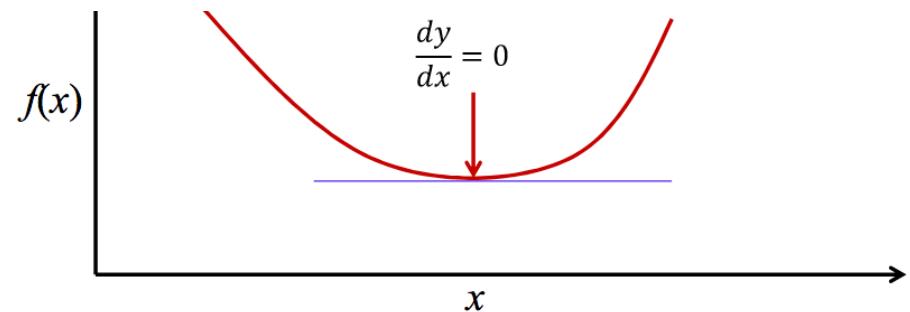
- Optimize!

$$\underset{\text{minimize}}{\sum_{i=1}^n} L(f(x_i), y_i)$$

- One-dimensional

$$\frac{df(x)}{dx} = 0$$

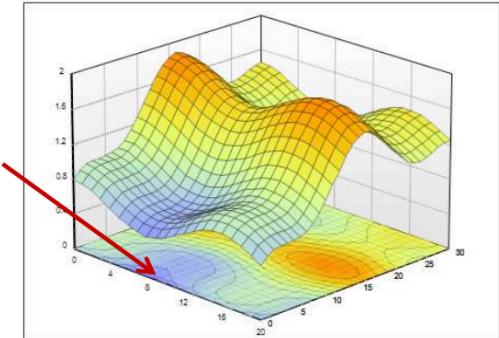
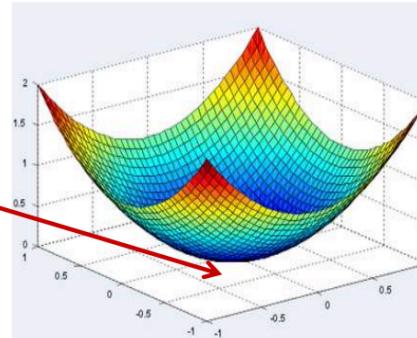
$$f''(x_{soln}) = \frac{df'(x_{soln})}{dx}$$



# Our first DL model

- Multidimensional

$$\nabla f(X) = \begin{bmatrix} \frac{\partial f(X)}{\partial x_1} & \frac{\partial f(X)}{\partial x_2} & \dots & \frac{\partial f(X)}{\partial x_n} \end{bmatrix}$$

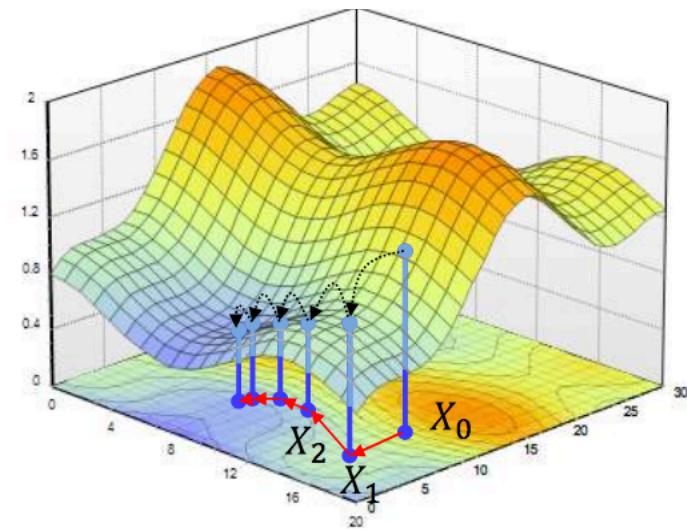
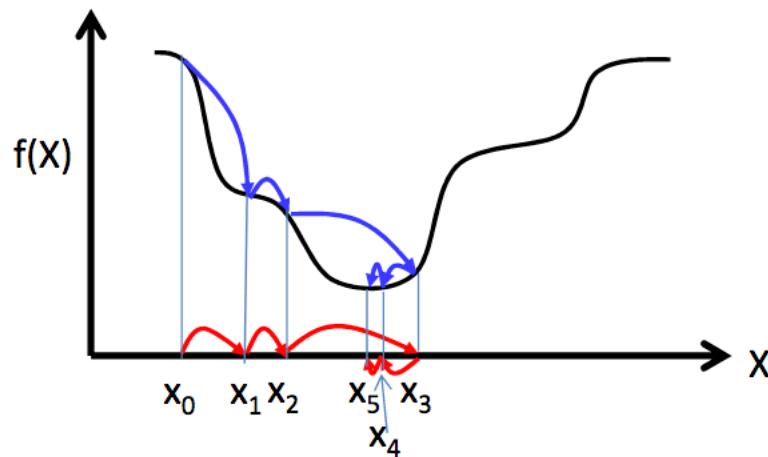


- Hessian matrix for the second derivatives

$$\nabla^2 f(x_1, \dots, x_n) := \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

# Our first DL model

- Often it is not possible to simply solve this optimization problem
  - The function may have an intractable form
- In these situations, iterative solutions are used
  - Begin with a guess for the optimal values, and then refine it iteratively until the correct value is obtained



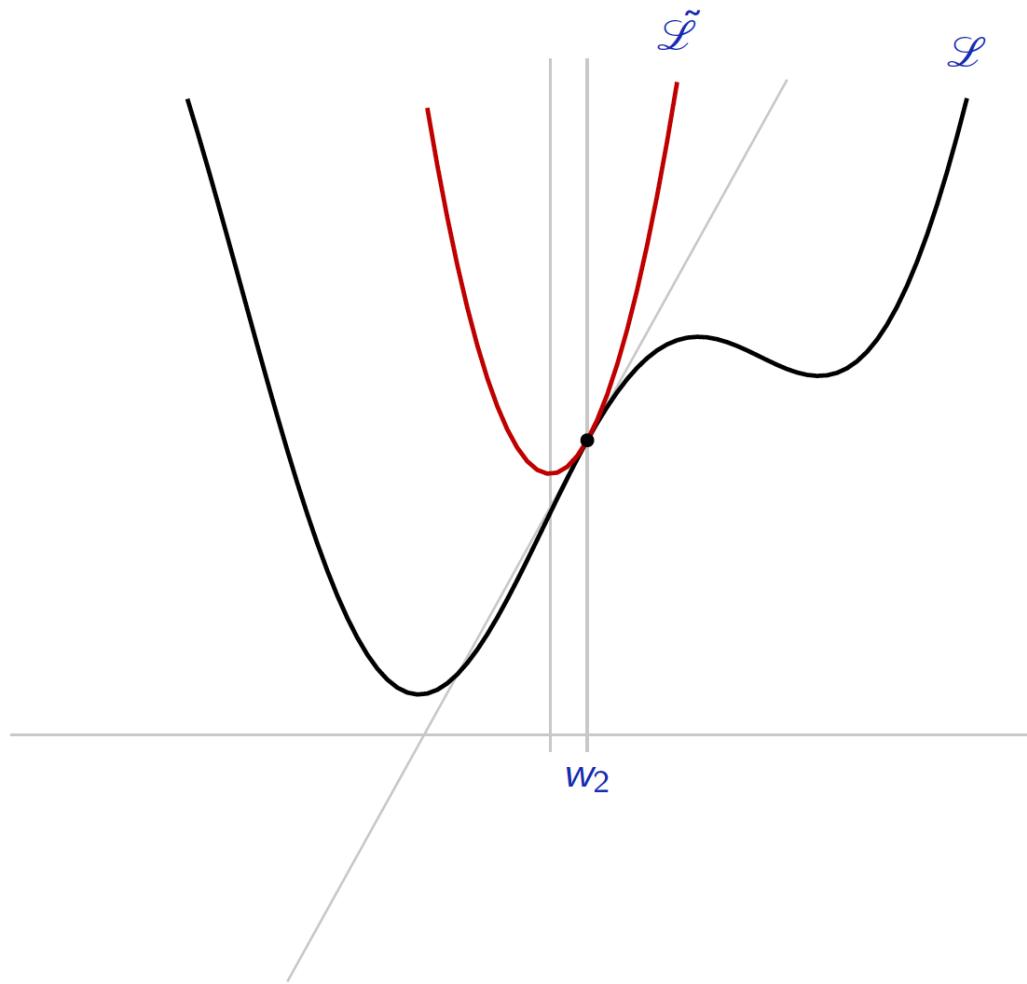
# Our first DL model

- (Vanilla) Gradient descent

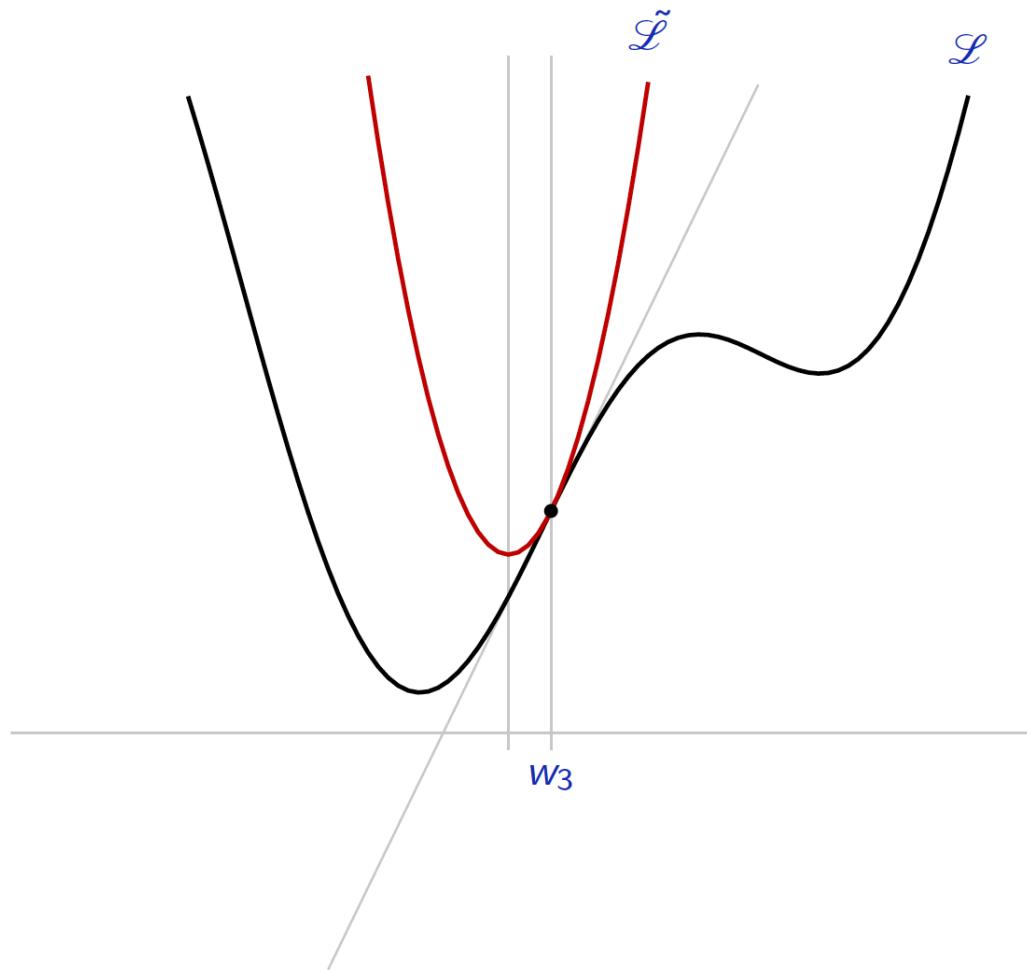
$$w^{k+1} = w^k - \eta^k \nabla f(w^k)^T$$

- Step-size (learning rate) is a hyperparameter

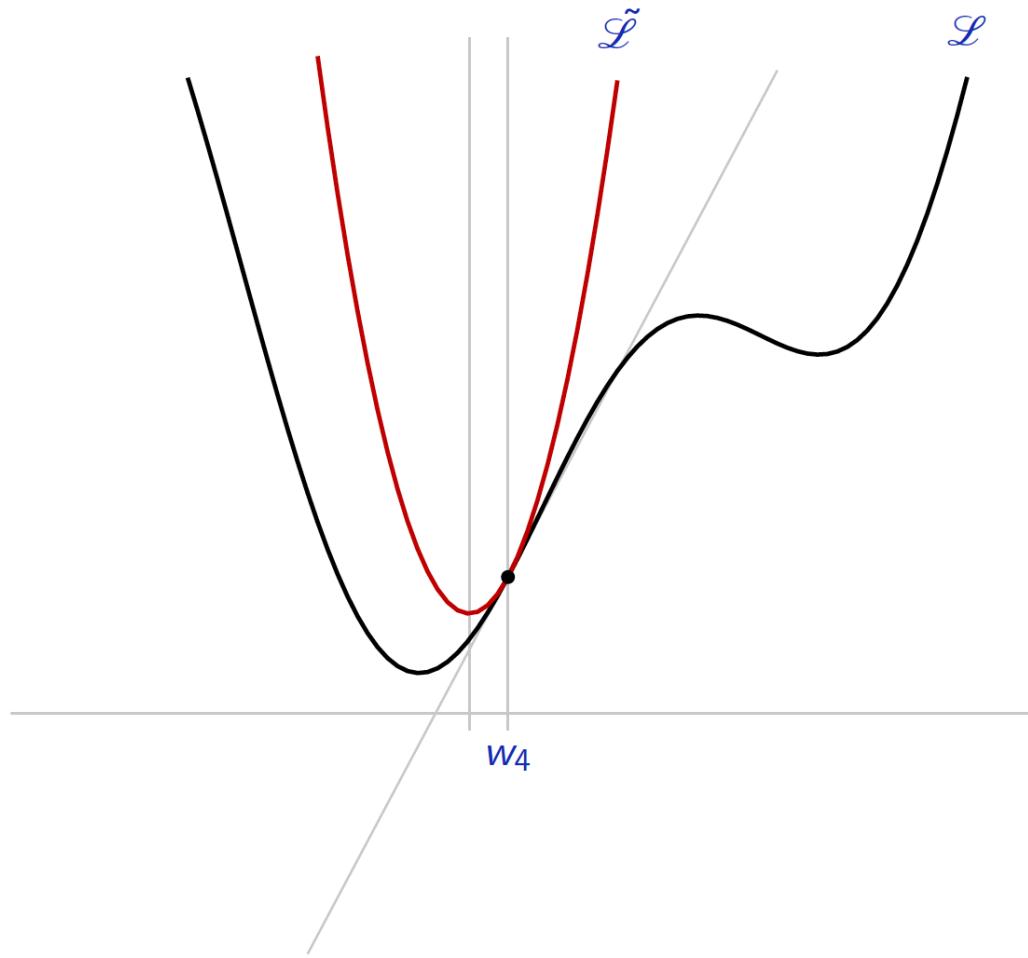
$$\eta = 0.125$$



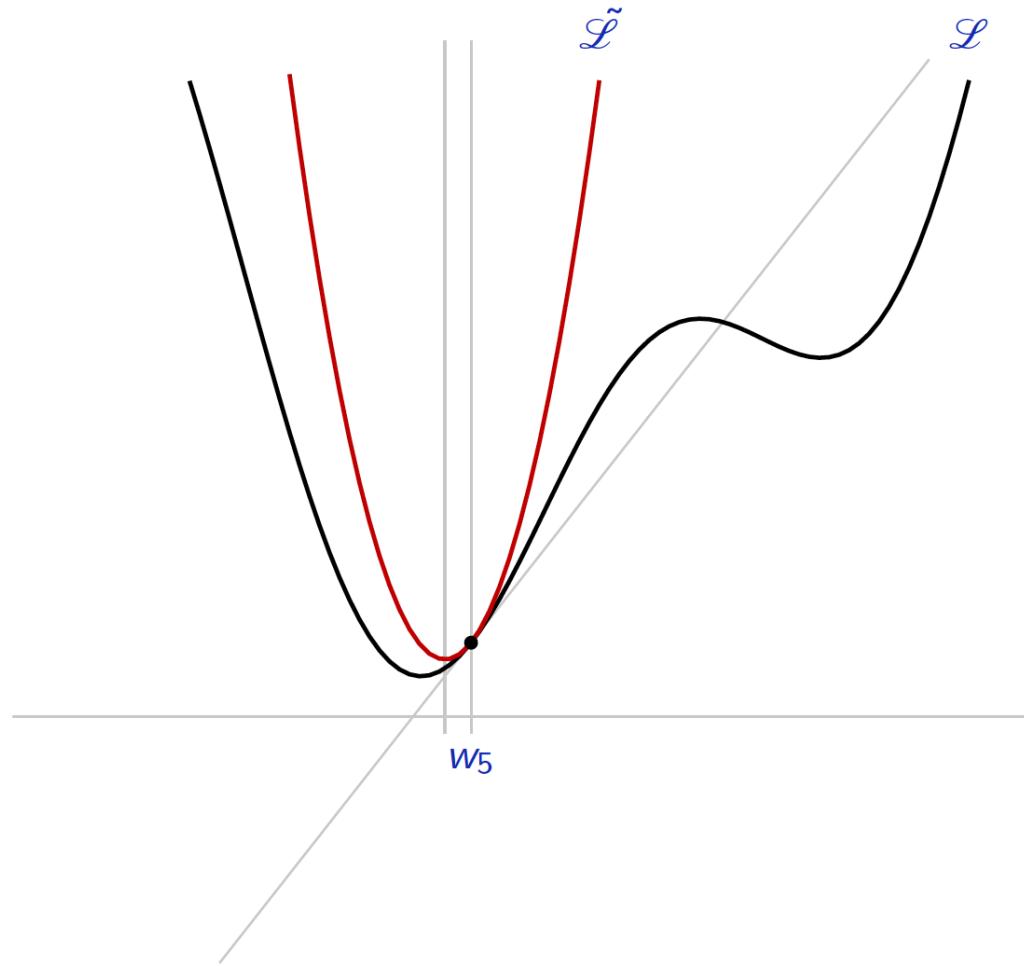
$$\eta = 0.125$$



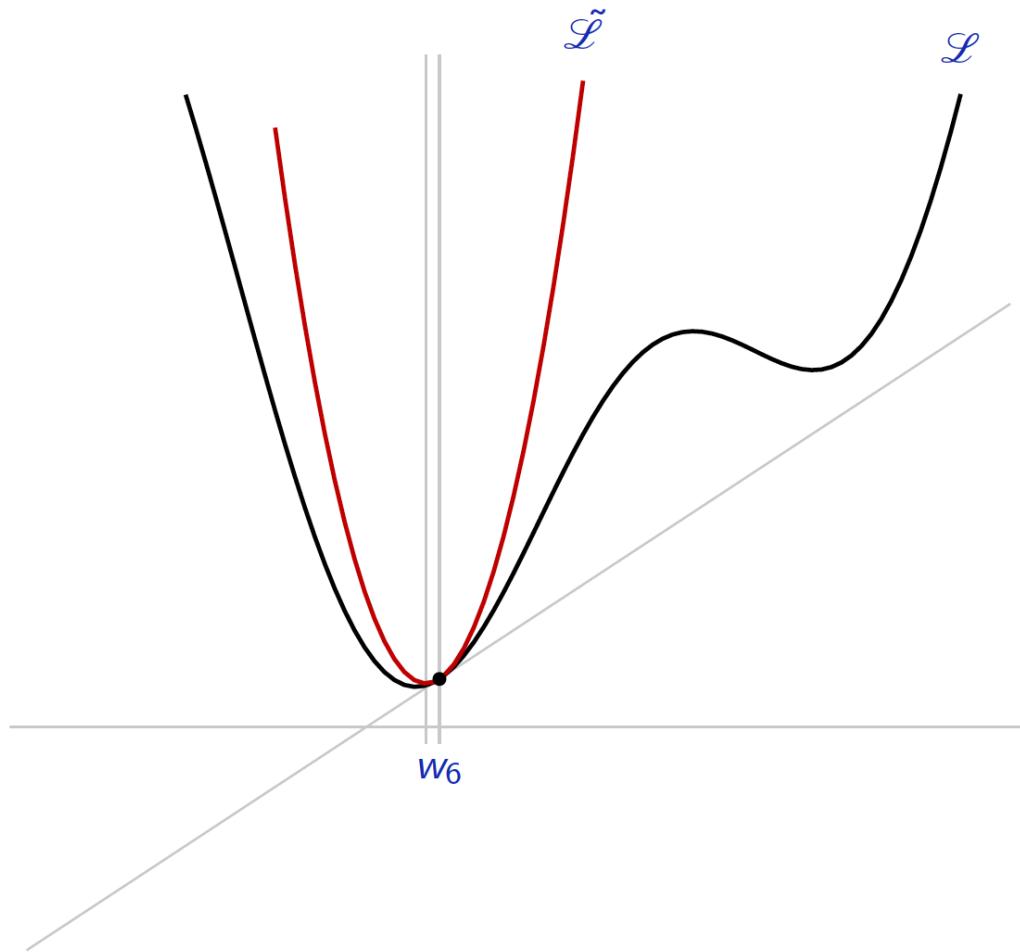
$$\eta = 0.125$$



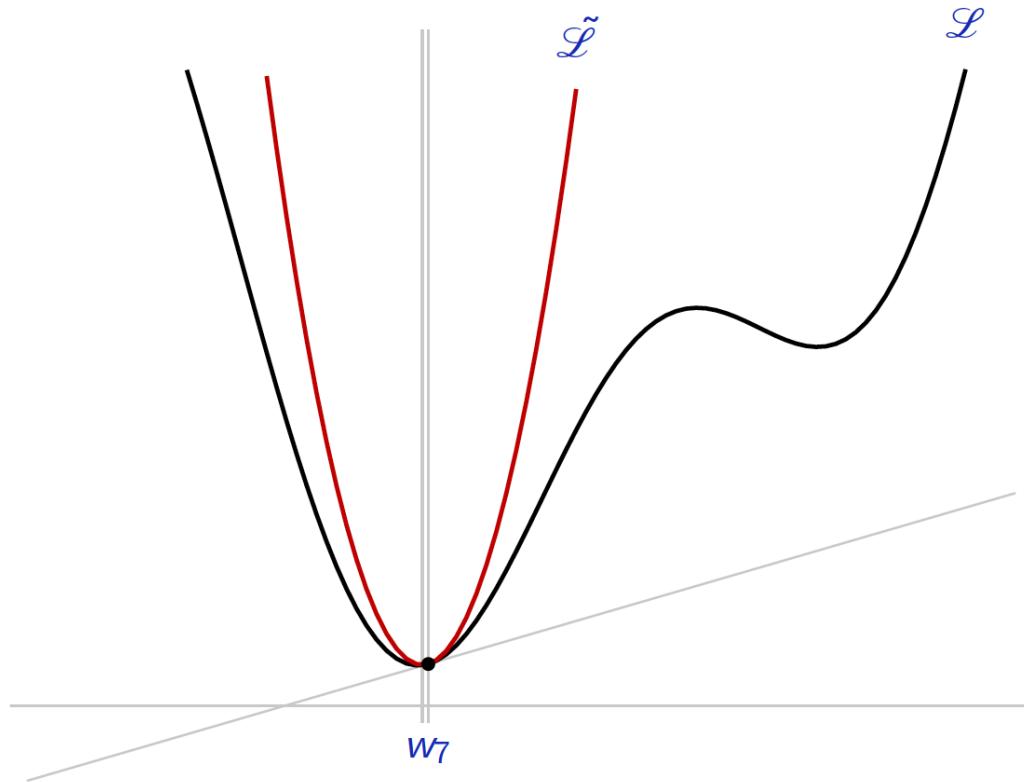
$$\eta = 0.125$$



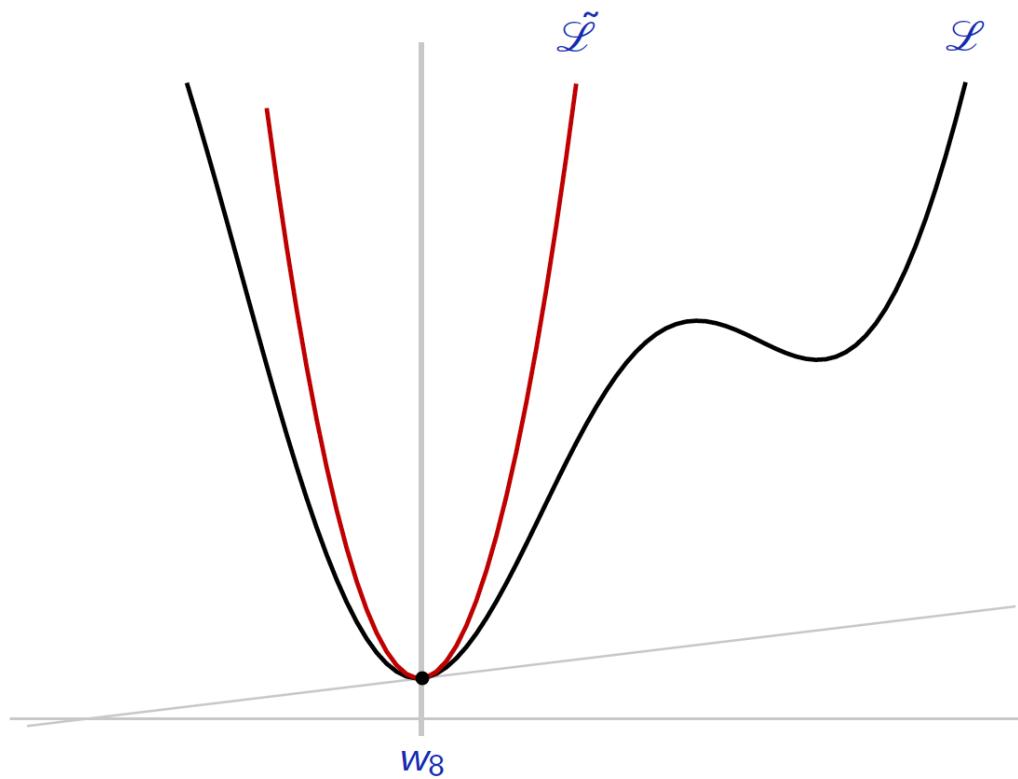
$$\eta = 0.125$$



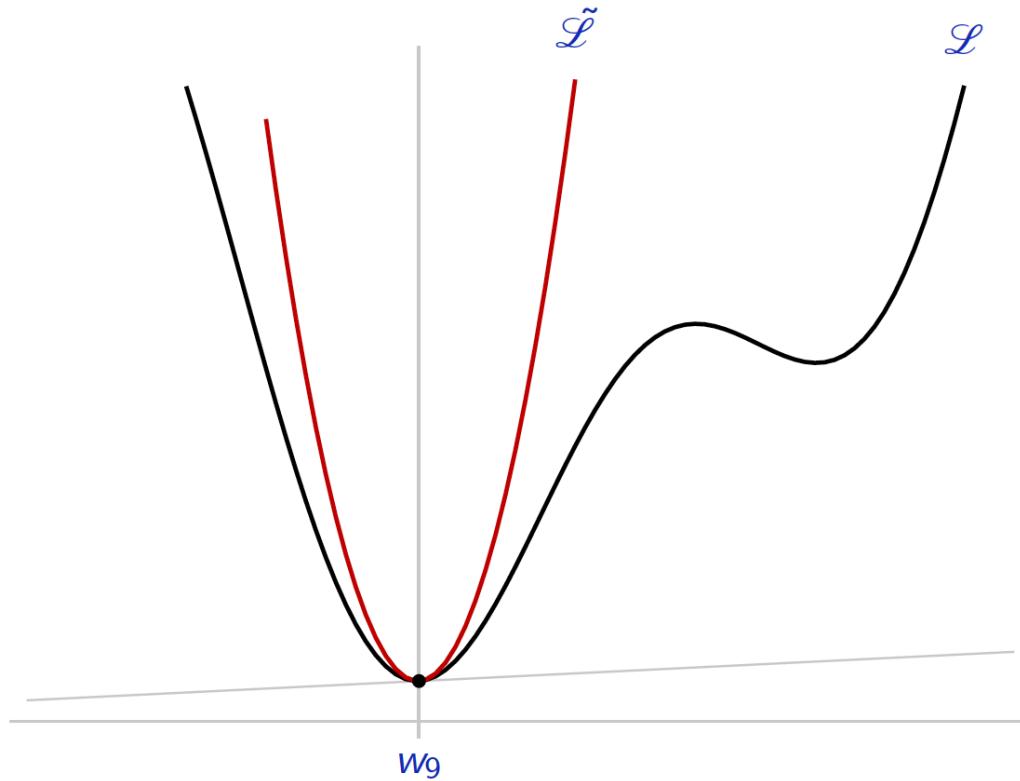
$$\eta = 0.125$$



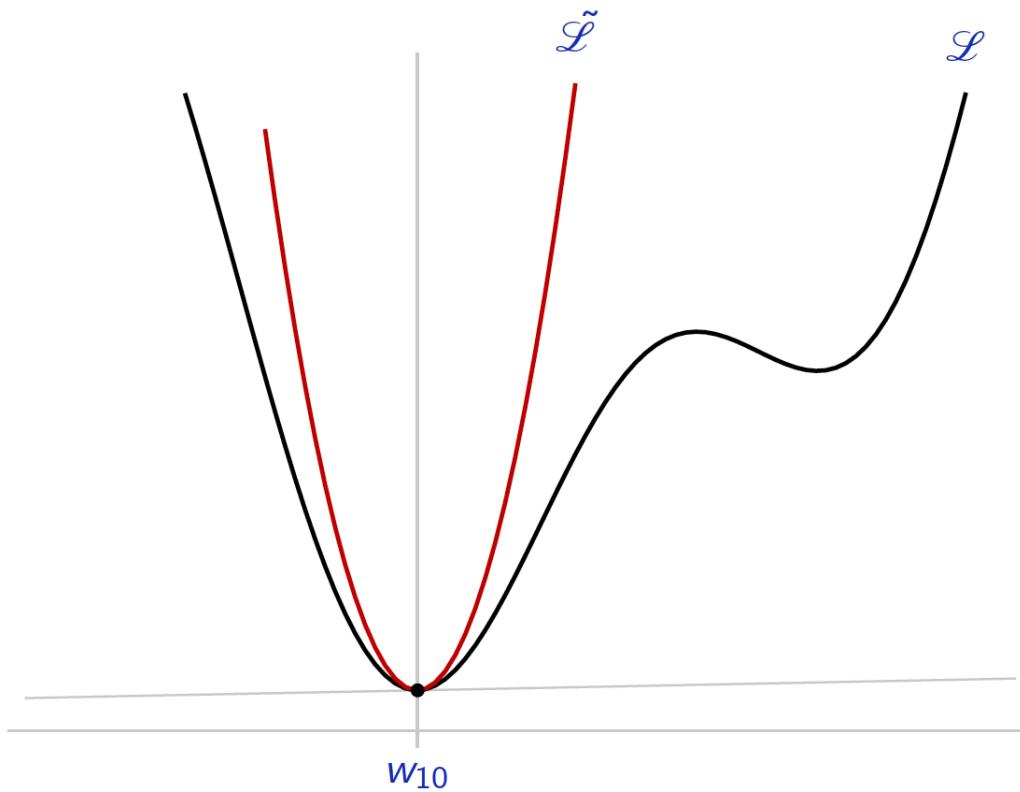
$$\eta = 0.125$$



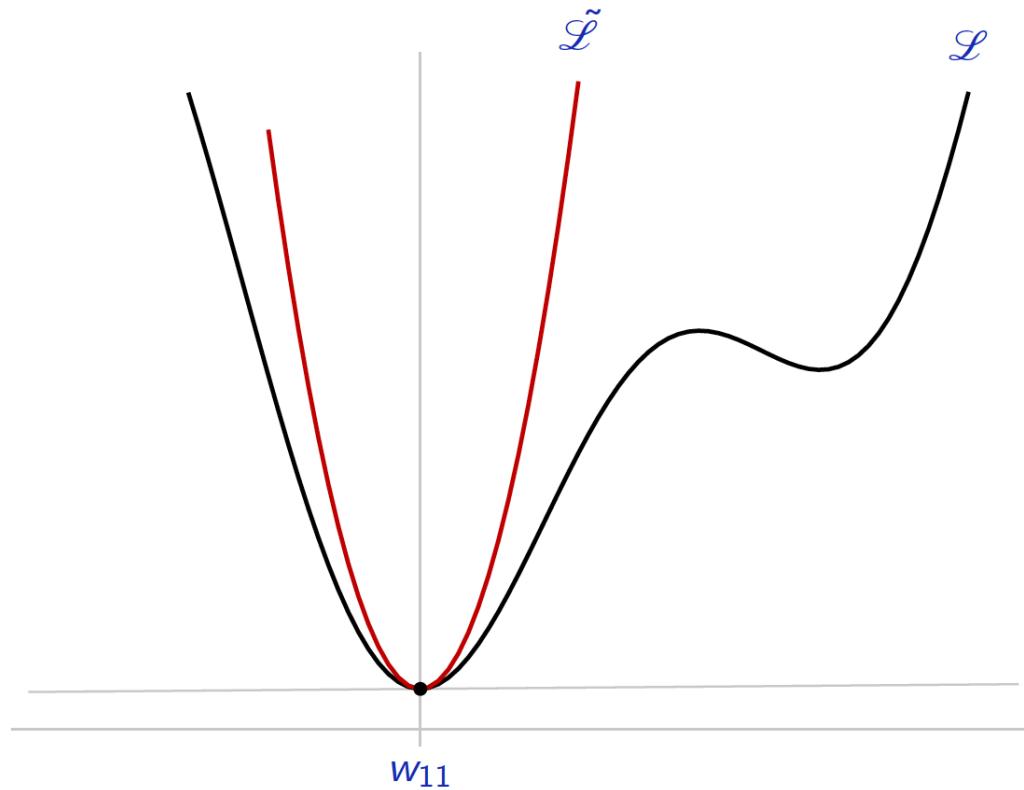
$$\eta = 0.125$$



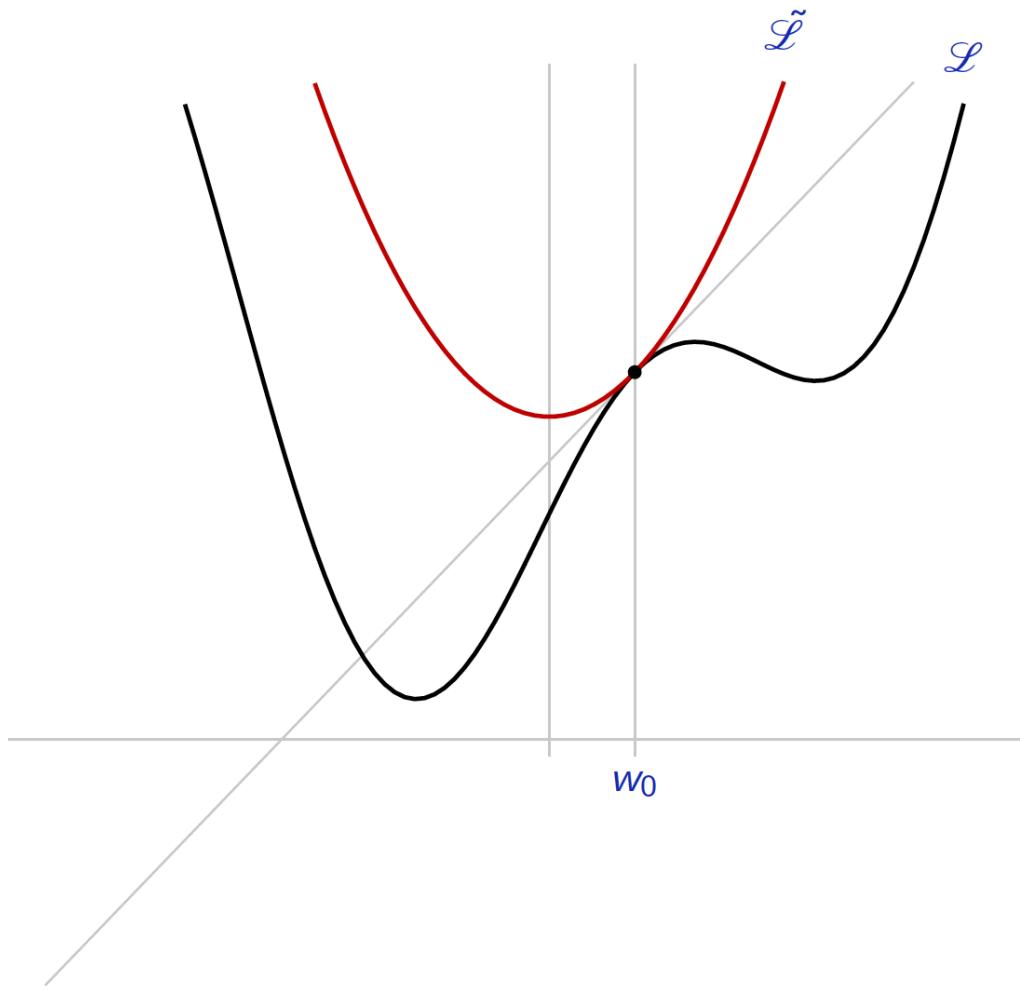
$$\eta = 0.125$$



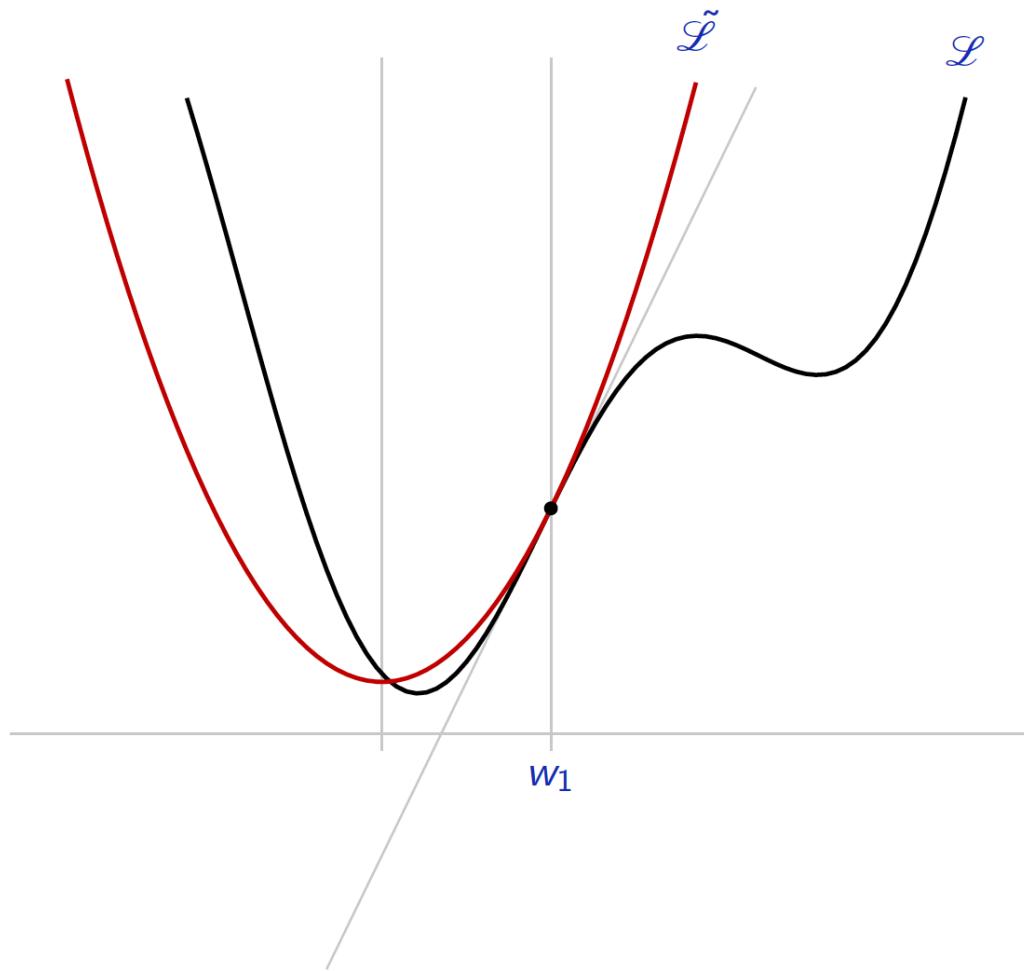
$$\eta = 0.125$$



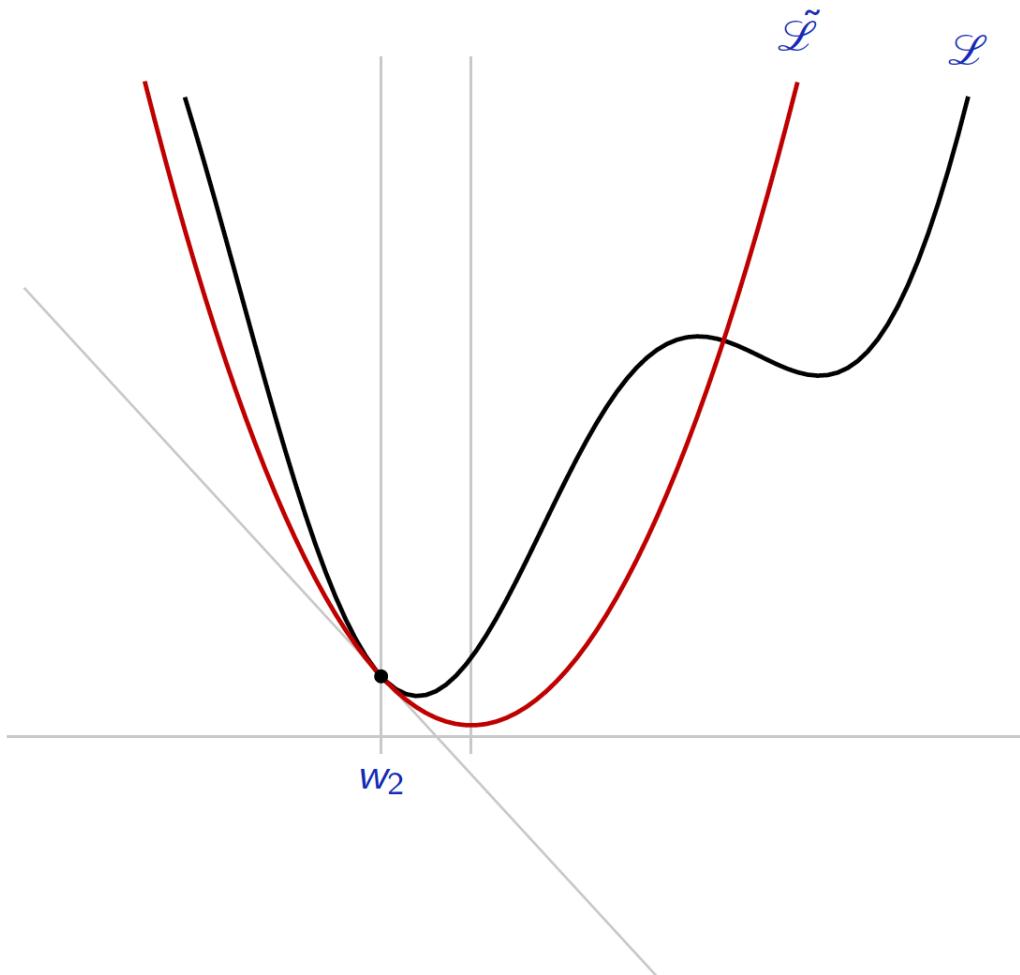
$$\eta = 0.5$$



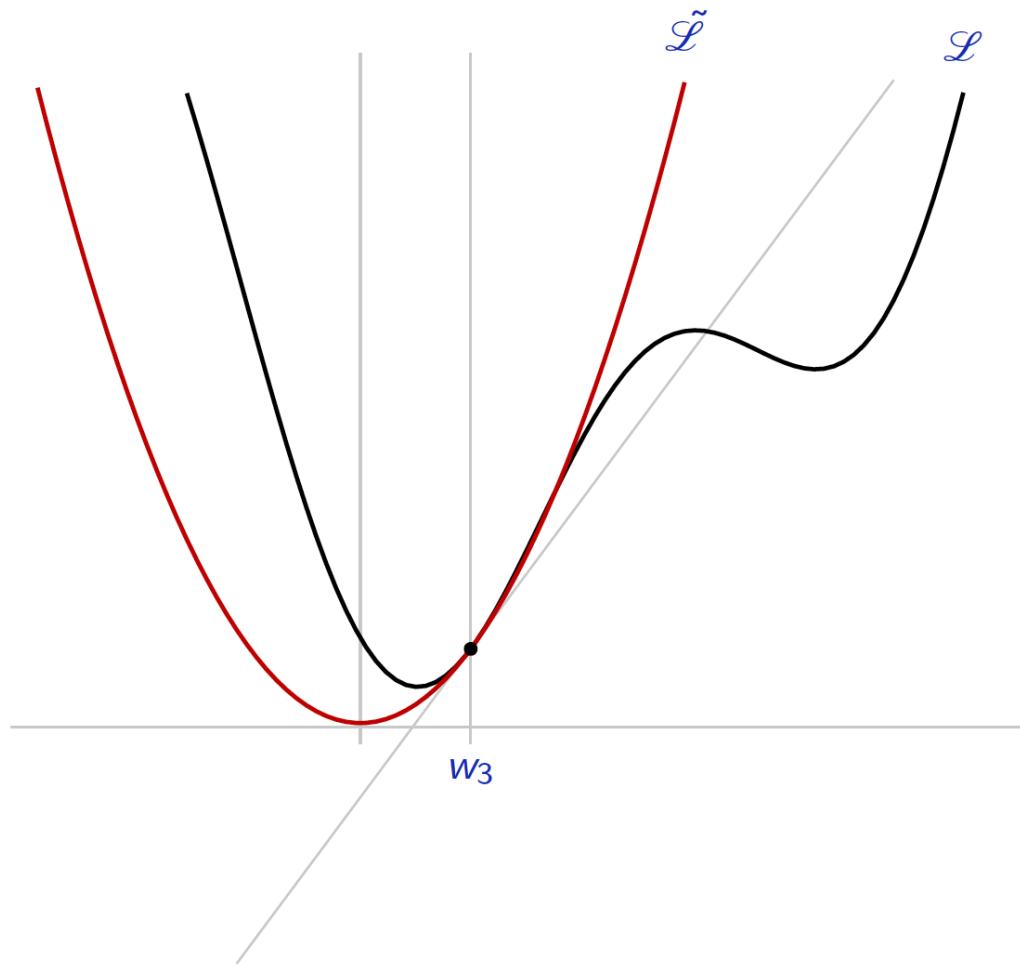
$$\eta = 0.5$$



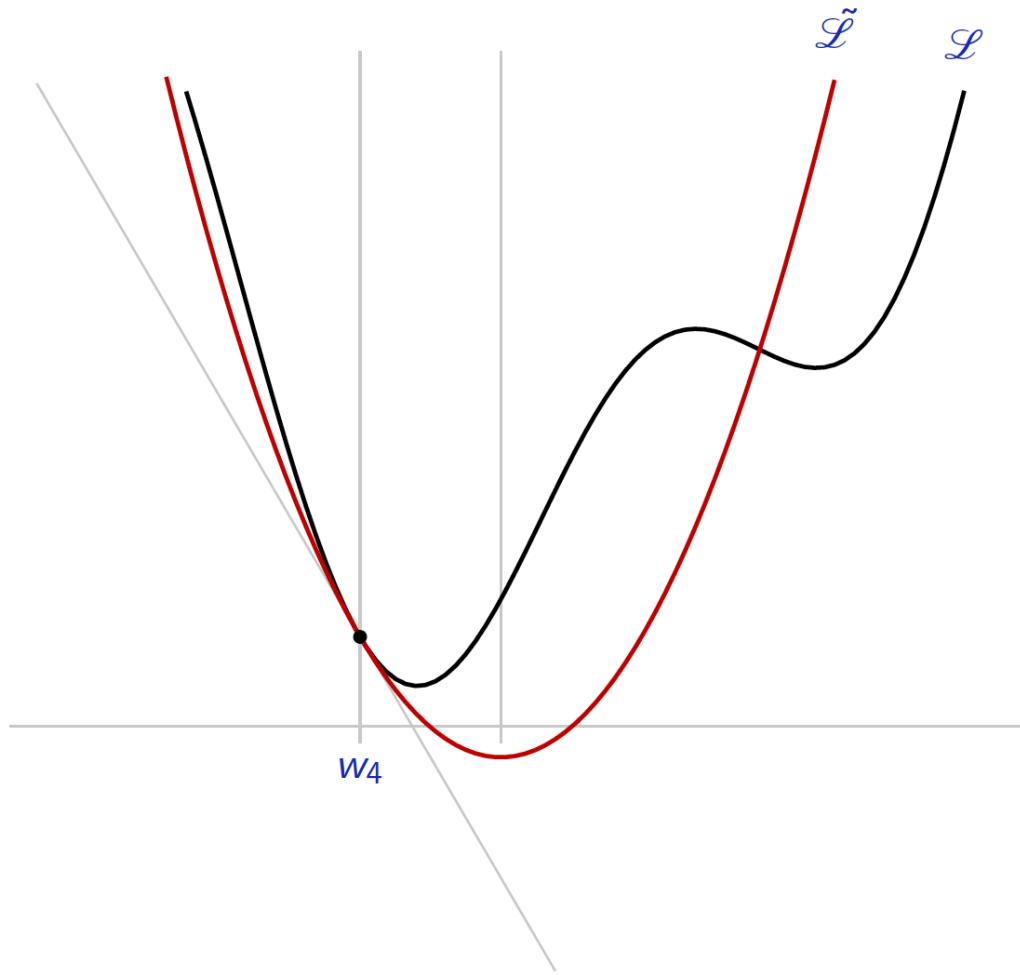
$$\eta = 0.5$$



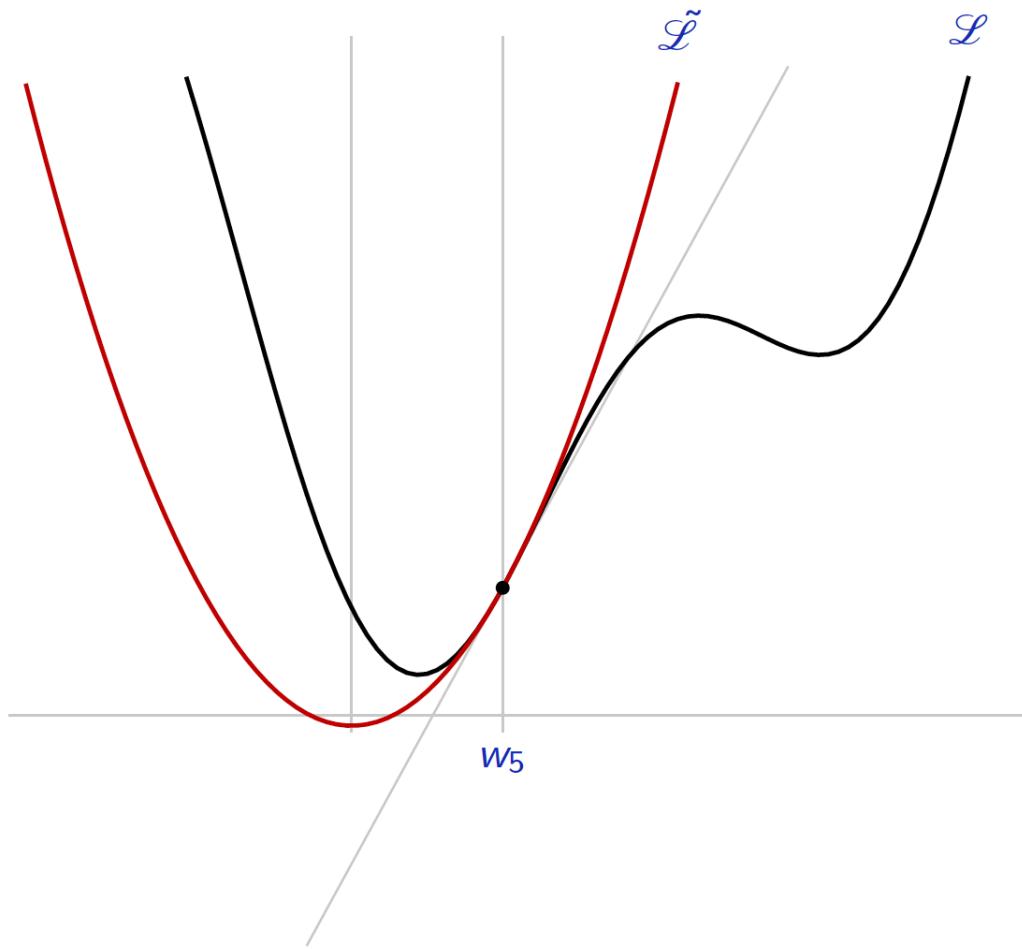
$$\eta = 0.5$$



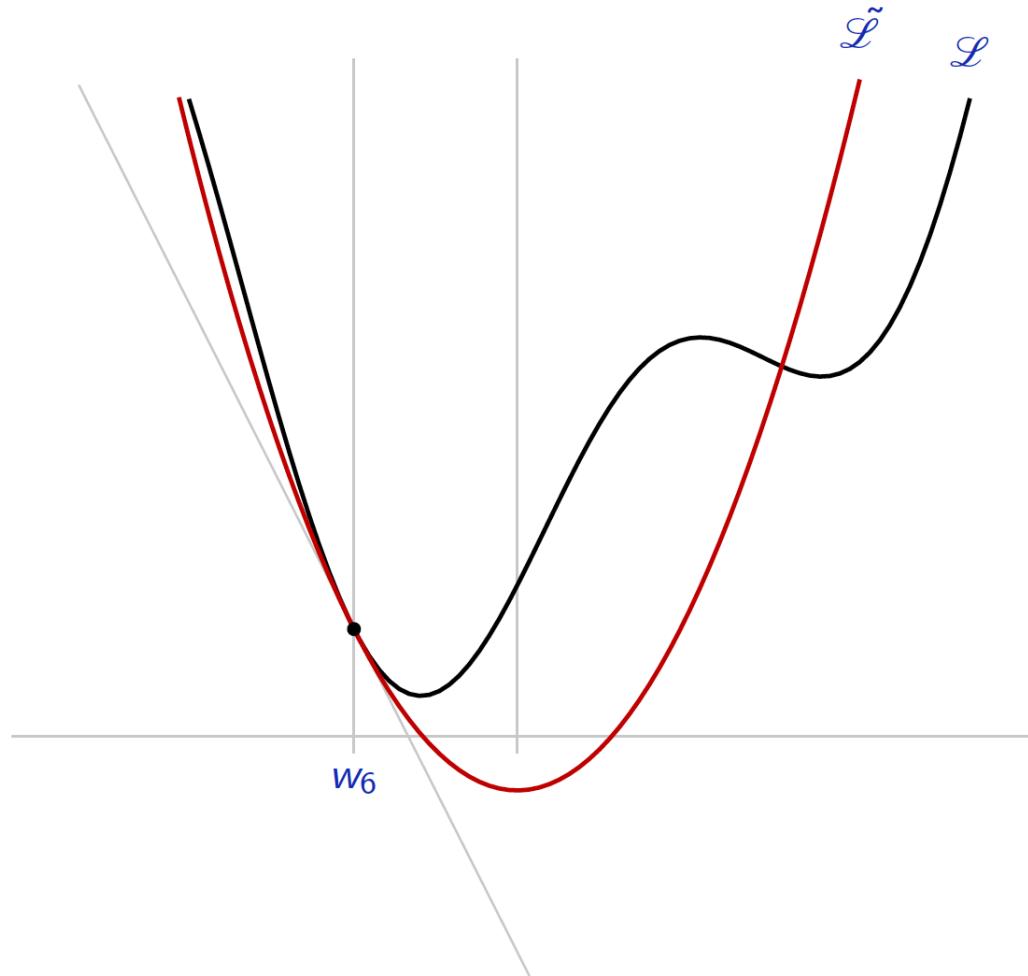
$$\eta = 0.5$$



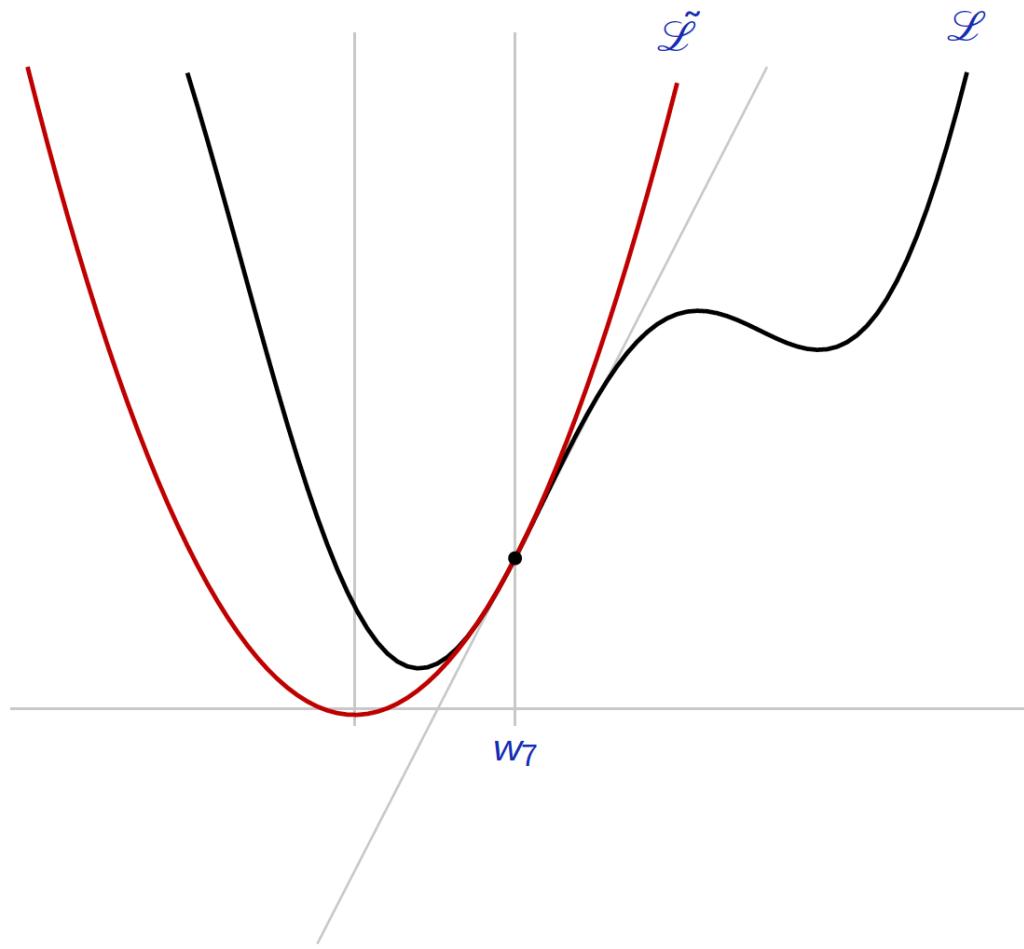
$$\eta = 0.5$$



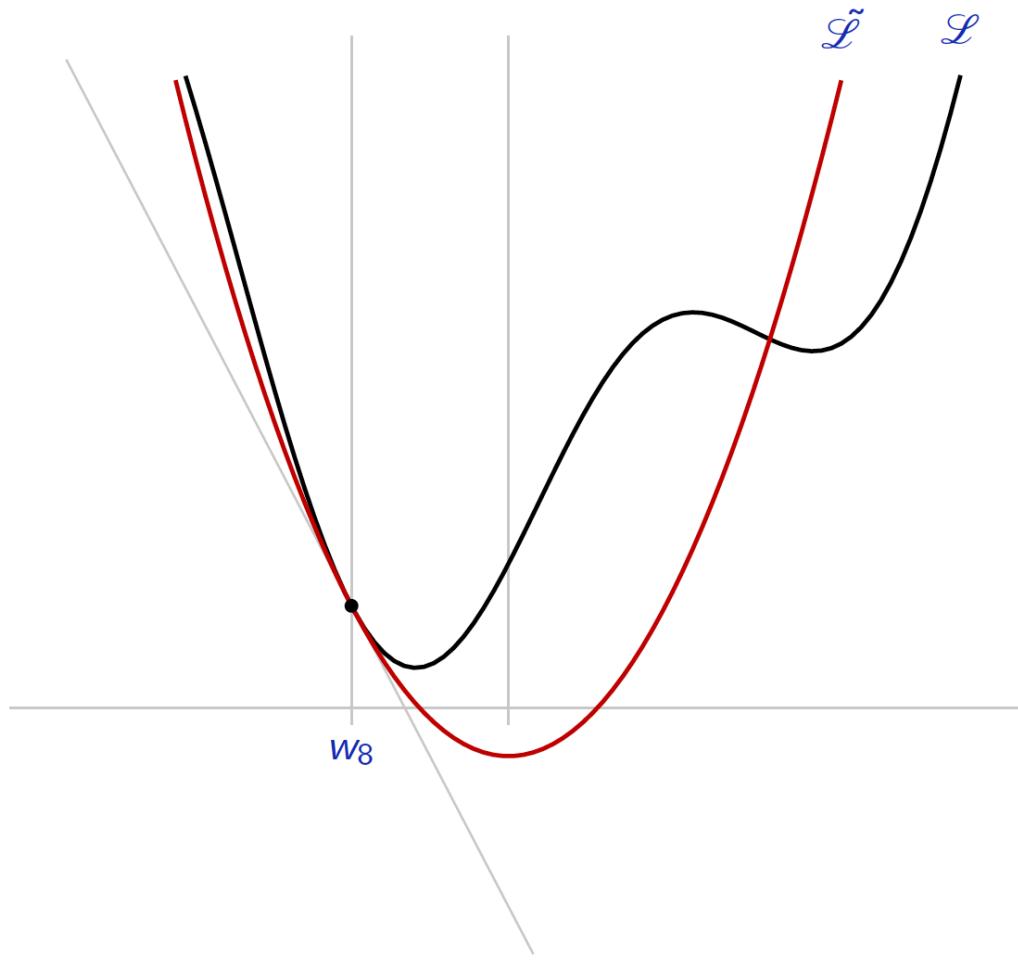
$$\eta = 0.5$$



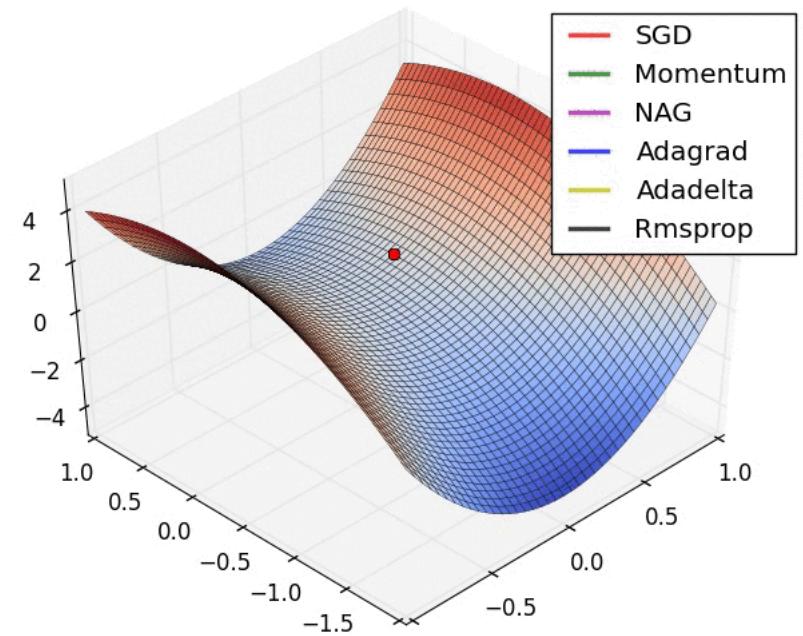
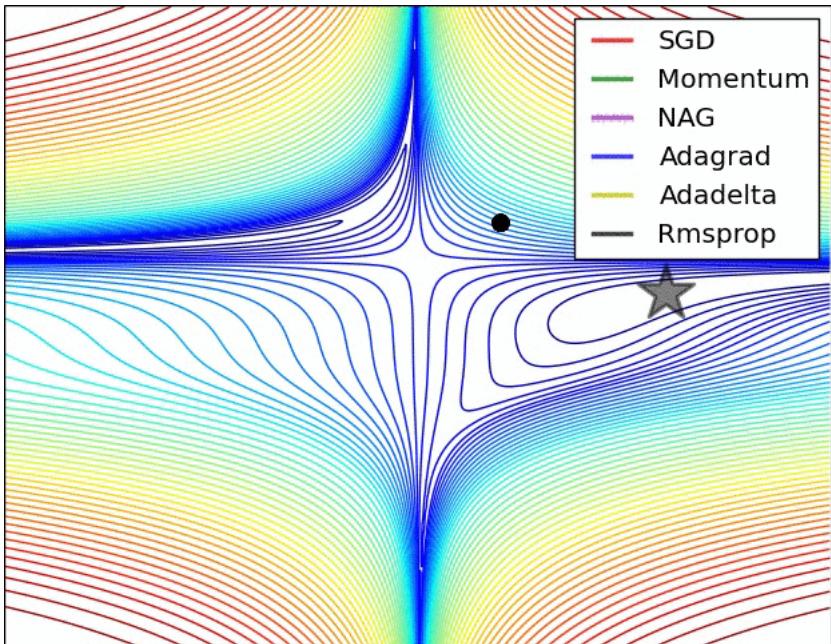
$$\eta = 0.5$$



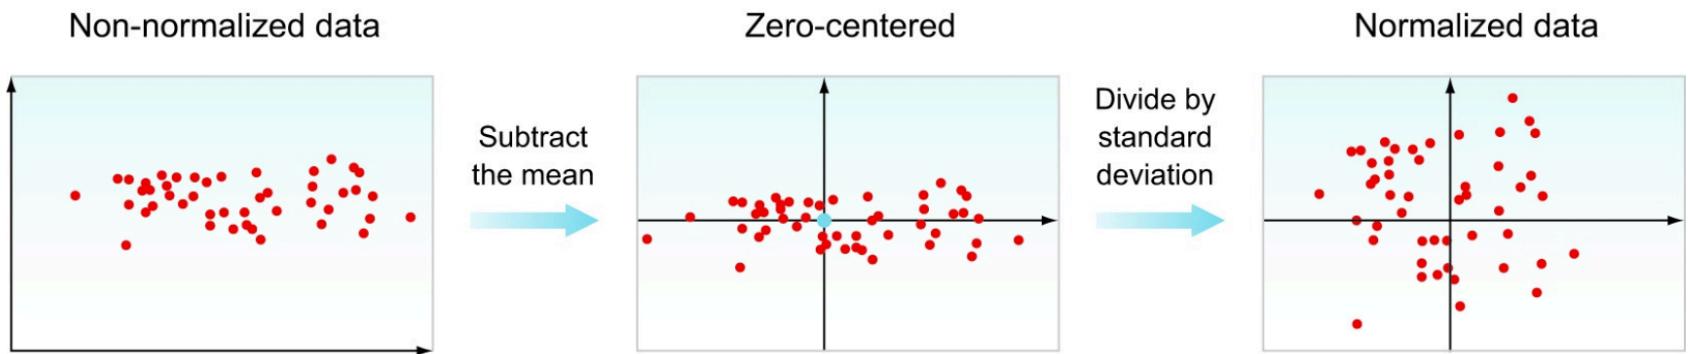
$$\eta = 0.5$$



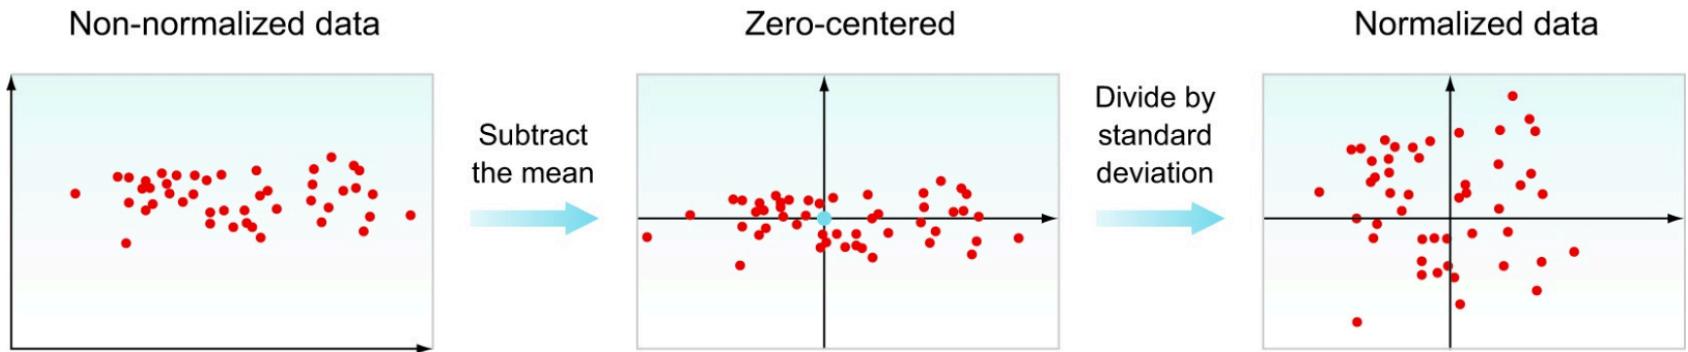
# Gradient descent variants



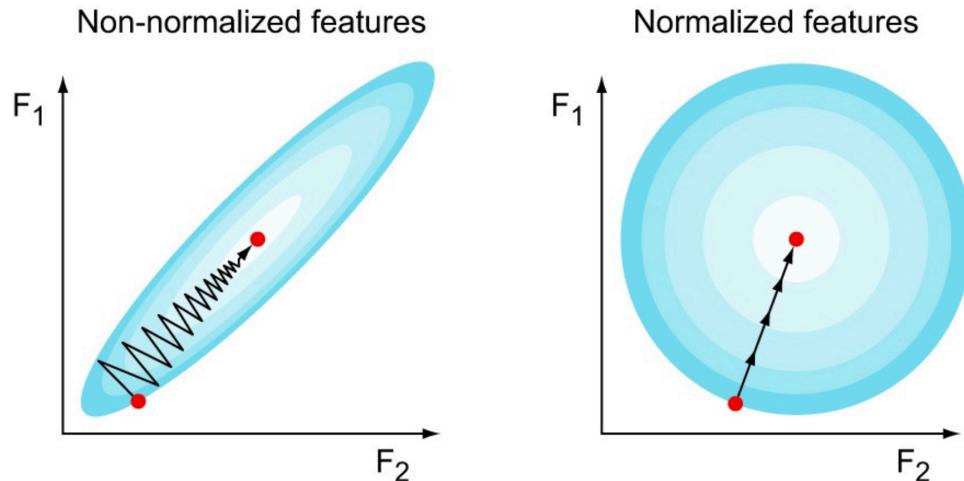
# Always normalize your data



# Always normalize your data



Gradient descent with and without feature scaling

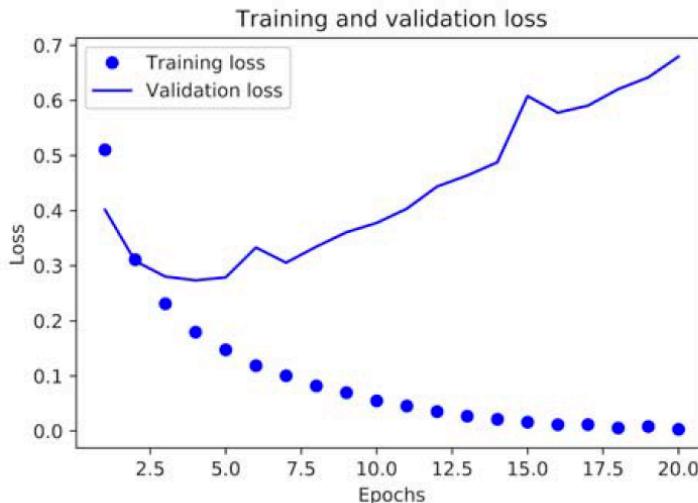


# Our first DL model

- Data  $\{x_t\}$   $\{x_{t-1}, x_{t-2}, \dots, x_{t-h}\}$   
 $\{x_{t-1}\}$   $\{x_{t-2}, x_{t-1}, \dots, x_{t-h-1}\}$
- Architecture 1 hidden layer
- Activation functions ReLU
- Loss function  $L(f(x_i), y_i) = (f(x_i) - y_i)^2$
- Regularization
- Software frameworks

# Overfitting

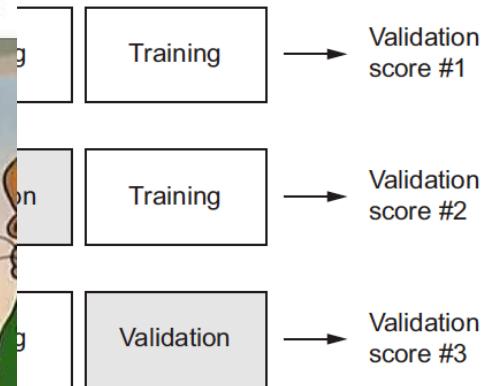
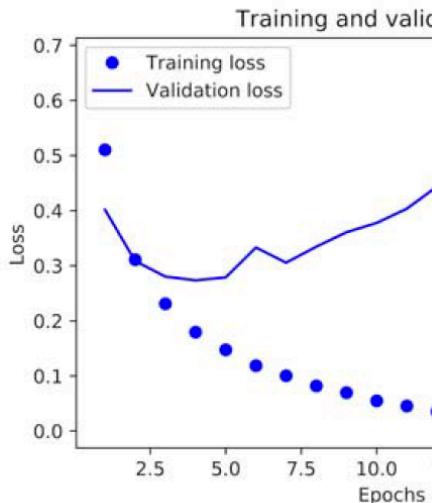
- If a model has a high capacity, it will overfit to the training data and perform poorly on unseen data



# Overfitting

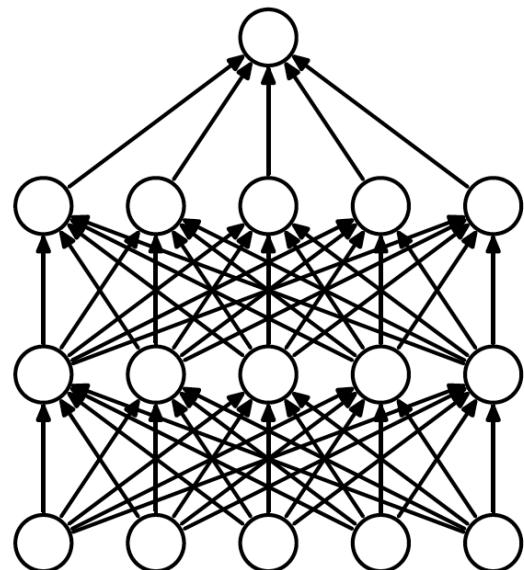
- If a model has a high capacity, it will overfit to the training data and perform poorly on unseen data

**When your network  
seems to be overfitting..**

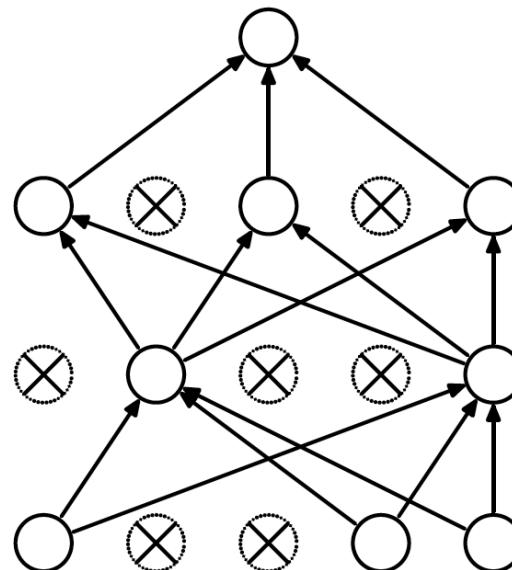


# Regularization

- Dropout



(a) Standard Neural Net

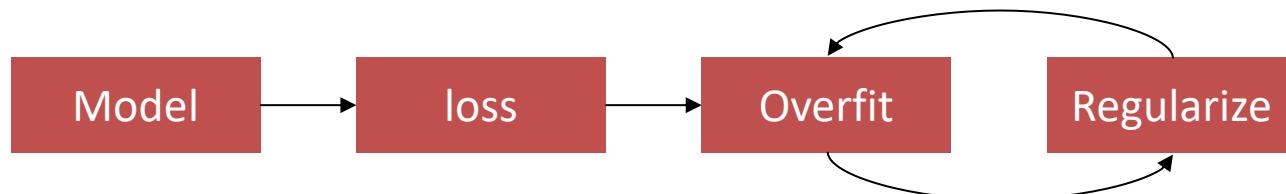
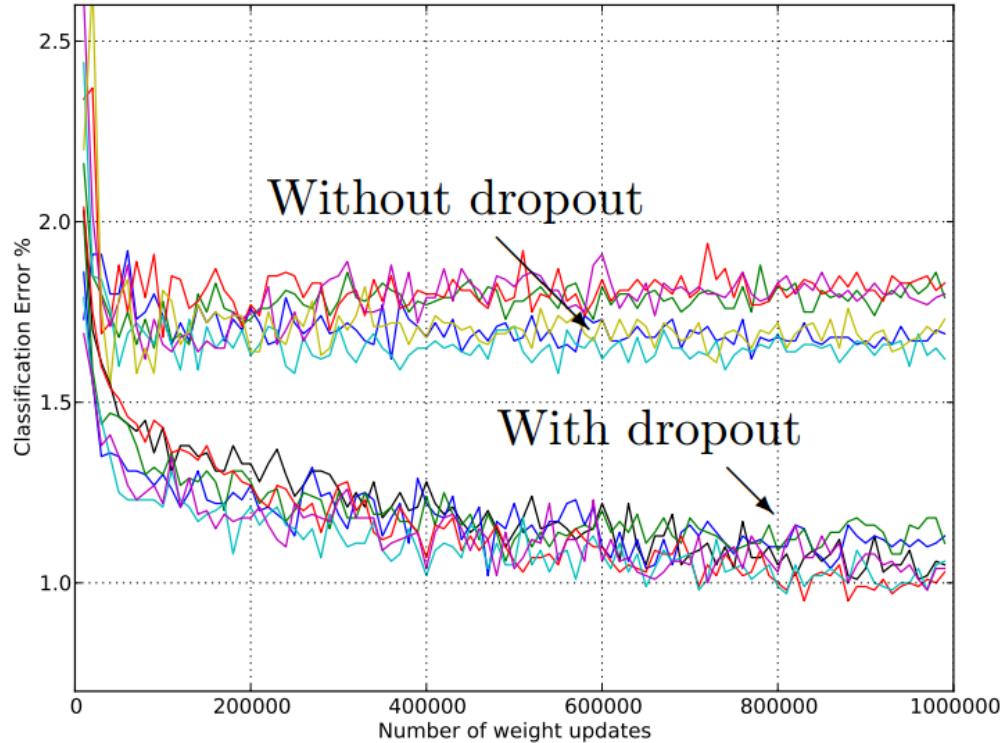


(b) After applying dropout.

<https://tinyurl.com/yxhx39tb>

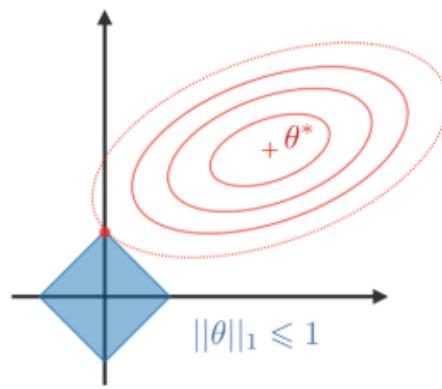
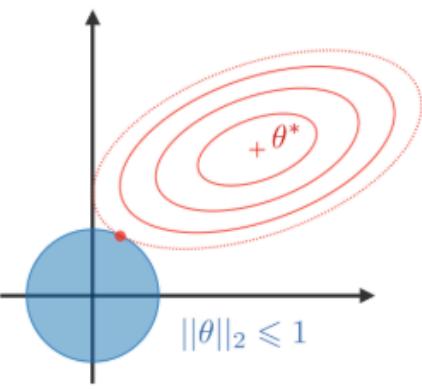
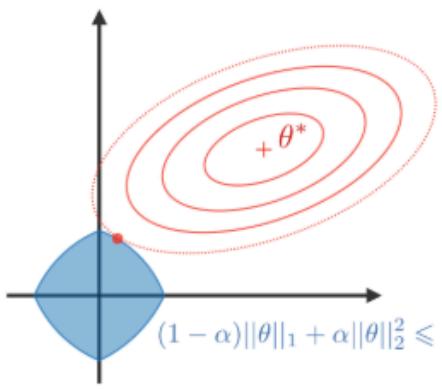
# Regularization

- Dropout



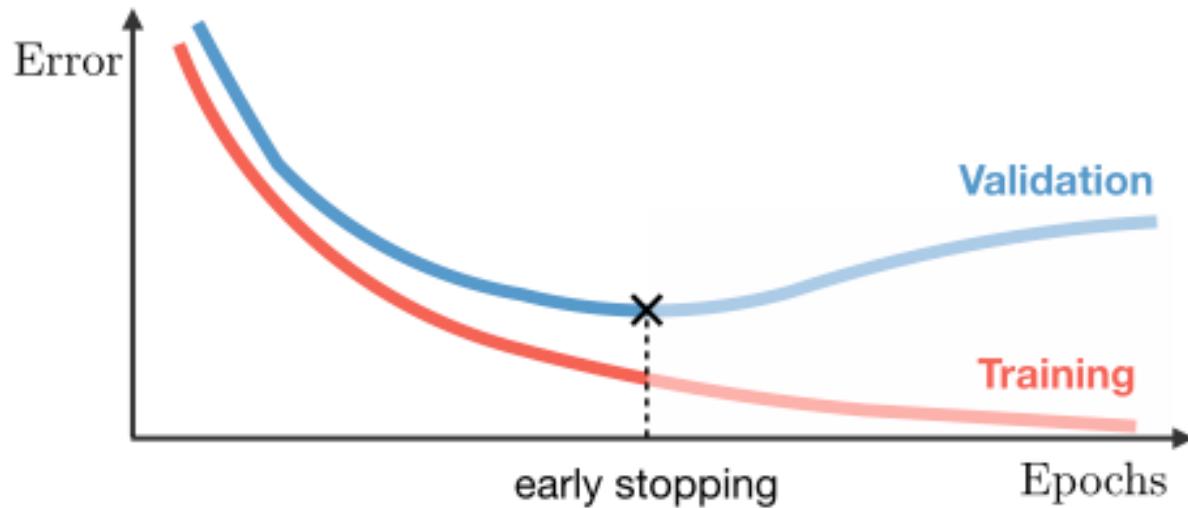
# Regularization

- Weight regularization

LASSO	Ridge	Elastic Net
<ul style="list-style-type: none"><li>• Shrinks coefficients to 0</li><li>• Good for variable selection</li></ul>	Makes coefficients smaller	Tradeoff between variable selection and small coefficients
 <p>A 2D plot showing red elliptical contours of a function. A blue diamond-shaped region represents the L1 ball, centered at the origin. A red dot labeled <math>\theta^*</math> is located inside the diamond. The inequality <math>\ \theta\ _1 \leq 1</math> is written below the x-axis.</p>	 <p>A 2D plot showing red elliptical contours of a function. A blue circular region represents the L2 ball, centered at the origin. A red dot labeled <math>\theta^*</math> is located inside the circle. The inequality <math>\ \theta\ _2 \leq 1</math> is written below the x-axis.</p>	 <p>A 2D plot showing red elliptical contours of a function. A blue diamond-shaped region represents the L1 ball, centered at the origin. A blue circular region represents the L2 ball, also centered at the origin. A red dot labeled <math>\theta^*</math> is located inside the intersection of the two shapes. The inequality <math>(1 - \alpha)\ \theta\ _1 + \alpha\ \theta\ _2 \leq 1</math> is written below the x-axis.</p>
$\dots + \lambda \ \theta\ _1$ $\lambda \in \mathbb{R}$	$\dots + \lambda \ \theta\ _2^2$ $\lambda \in \mathbb{R}$	$\dots + \lambda \left[ (1 - \alpha)\ \theta\ _1 + \alpha\ \theta\ _2^2 \right]$ $\lambda \in \mathbb{R}, \alpha \in [0, 1]$

# Regularization

- Early stopping



<https://tinyurl.com/yxhx39tb>

# Our first DL model

- Data
- Architecture
- Activation functions
- Losses
- Regularization
- Software frameworks

# Software frameworks

- Many, many frameworks



# Lab 1

Keras

TensorFlow / Theano / CNTK / ...

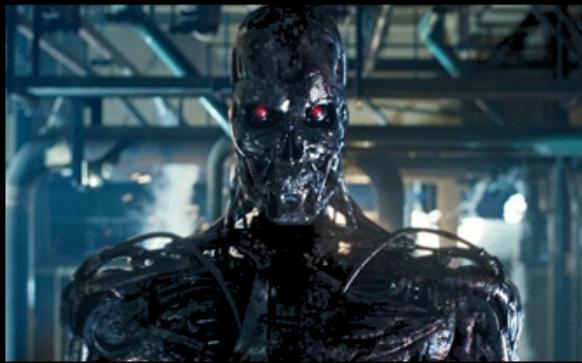
CUDA / cuDNN

BLAS, Eigen

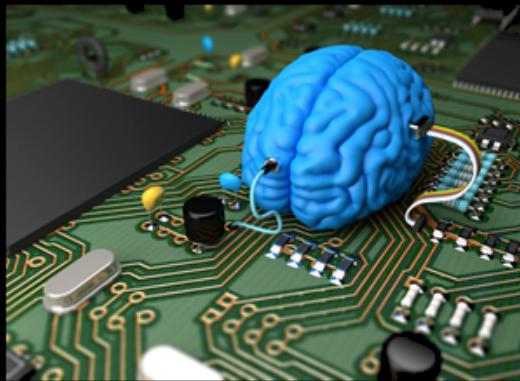
GPU

CPU

# Deep Learning



What society thinks I do



What my friends think I do



What other computer scientists think I do



What mathematicians think I do



What I think I do

```
In [1]:  
import keras  
Using TensorFlow backend.
```

What I actually do

# In the real-world ...



Contents lists available at [ScienceDirect](#)

Transportation Research Part C

journal homepage: [www.elsevier.com/locate/trc](http://www.elsevier.com/locate/trc)



Deep learning for short-term traffic flow prediction

Nicholas G. Polson <sup>a</sup>, Vadim O. Sokolov <sup>b,\*</sup>

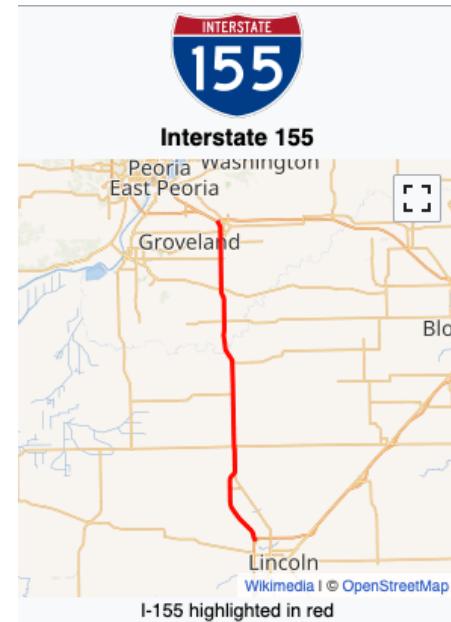
<sup>a</sup> Booth School of Business, University of Chicago, Chicago, IL 60637, USA

<sup>b</sup> Systems Engineering and Operations Research, George Mason University, Fairfax, VA 22030, USA

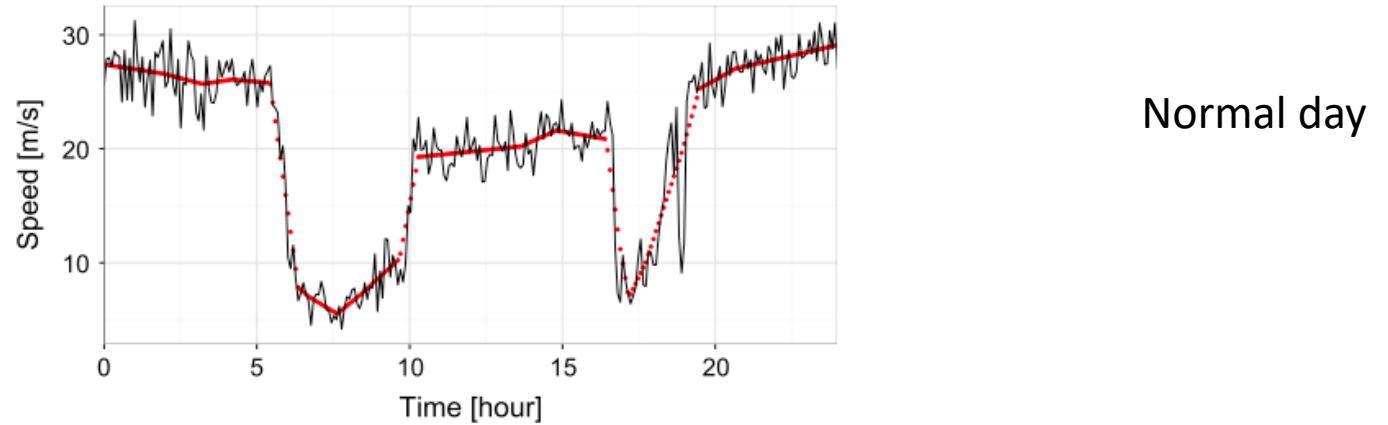


- Road sensor data from Interstate I-155

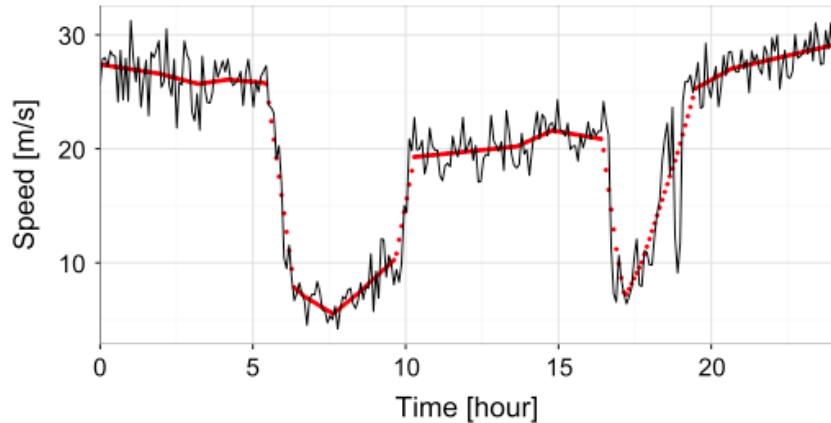
$$\begin{bmatrix} x_{1,t-40} & \dots & x_t \\ \vdots & \vdots & \vdots \\ x_{n,t-40} & \dots & x_{n,t} \end{bmatrix}$$



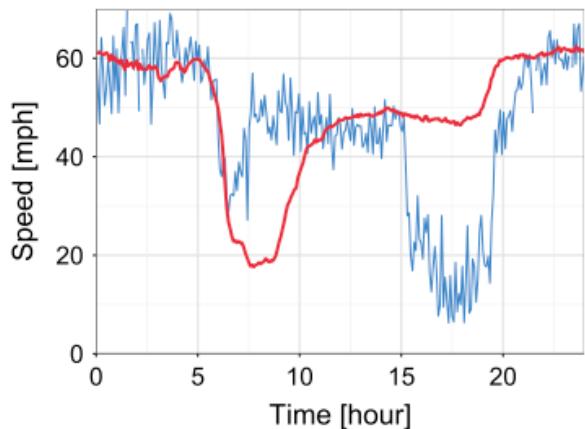
# In the real-world ...



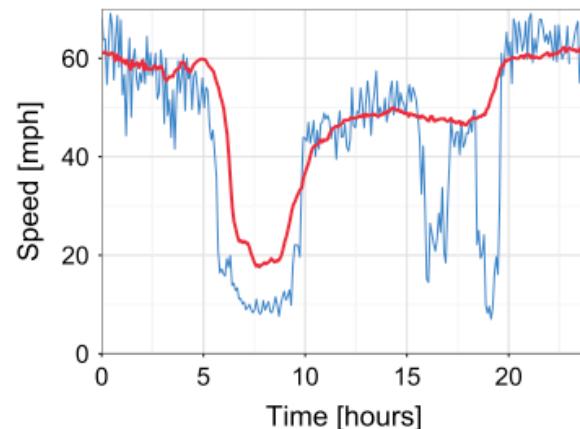
# In the real-world ...



Normal day



(a) Chicago Bears football game

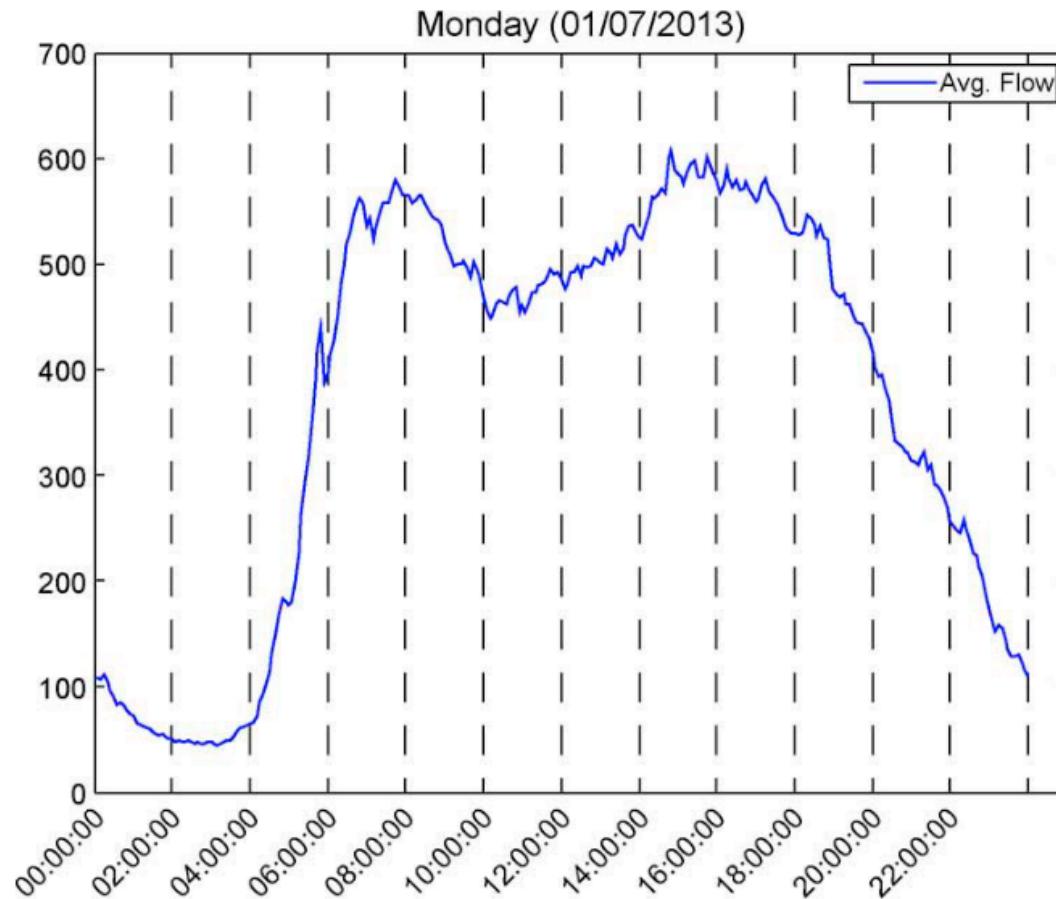


Events

(b) Snow weather

# A simple problem

- Predict traffic flow for the next 15 minutes



## A bit more complicated problem

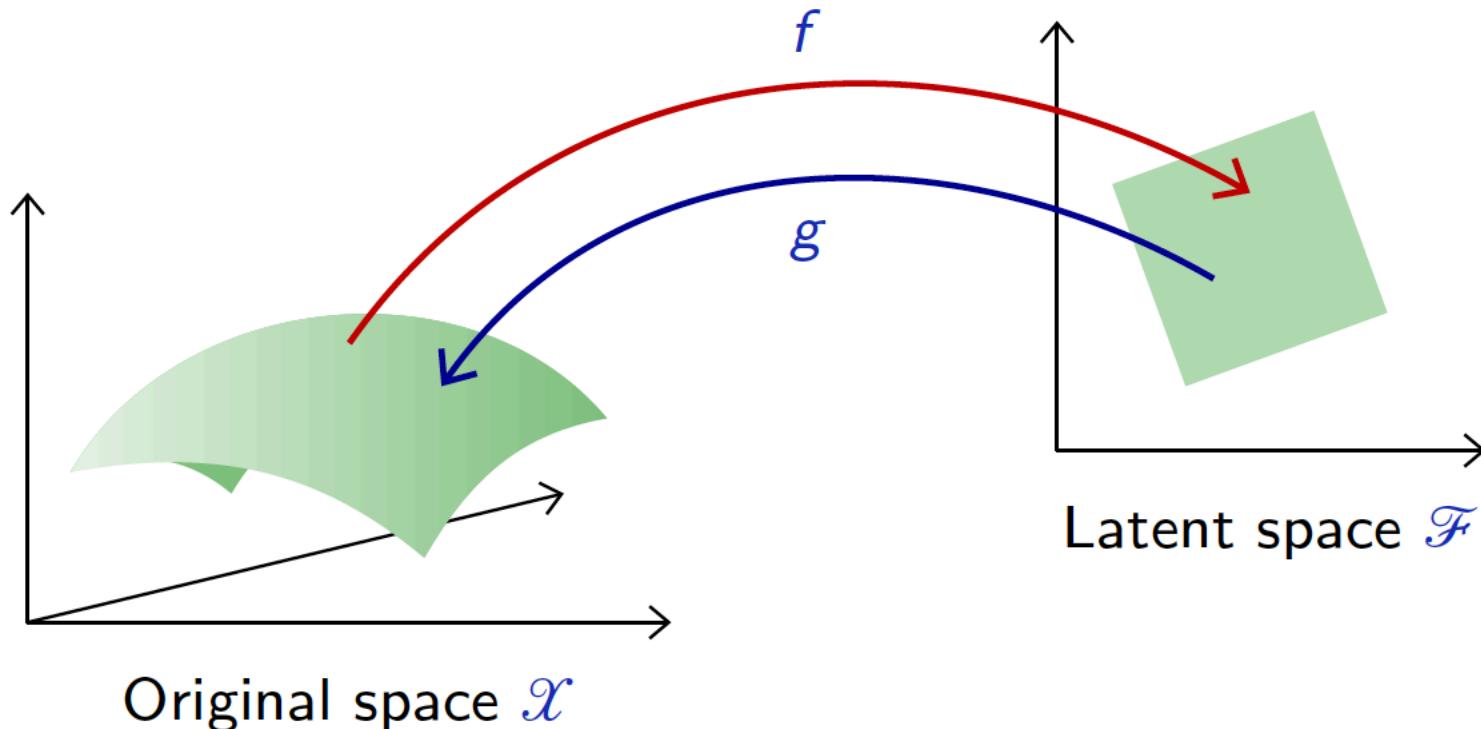
- Predict traffic flow for the next 15 minutes
- There are 15,000 detectors
- Do we train 15,000 models?

# A bit more complicated problem

- Predict traffic flow for the next 15 minutes
- There are 15,000 detectors
- Do we train 15,000 models?
- First idea:
  - There is a lot of overlap between the detectors
  - Signal/noise ratio is low
  - The useful information must be in a lower dimension than 15,000
  - **Dimensionality reduction!**

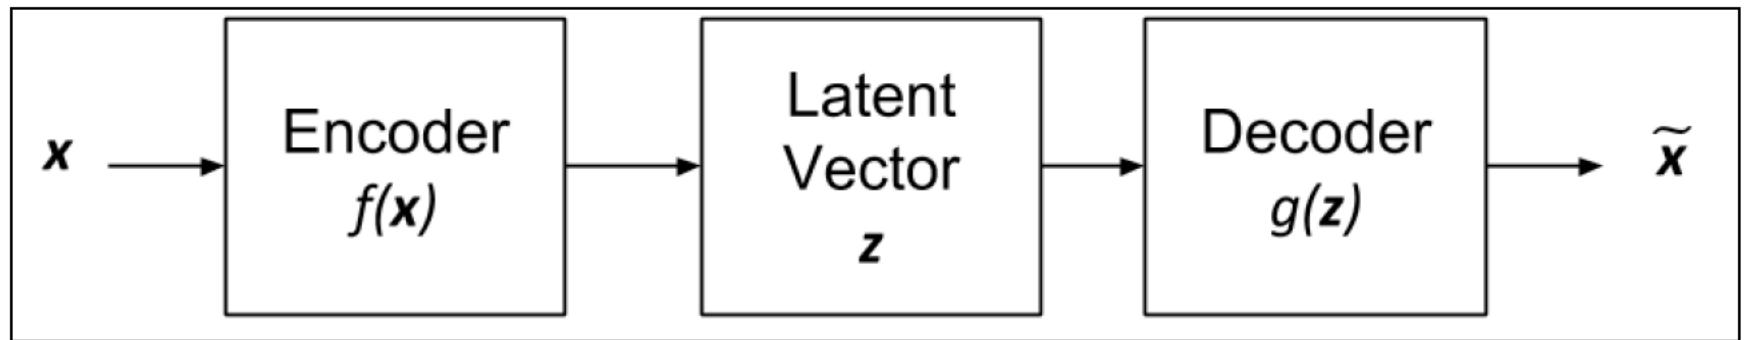
# Deep Autoencoder

- Try to learn a latent representation of the input



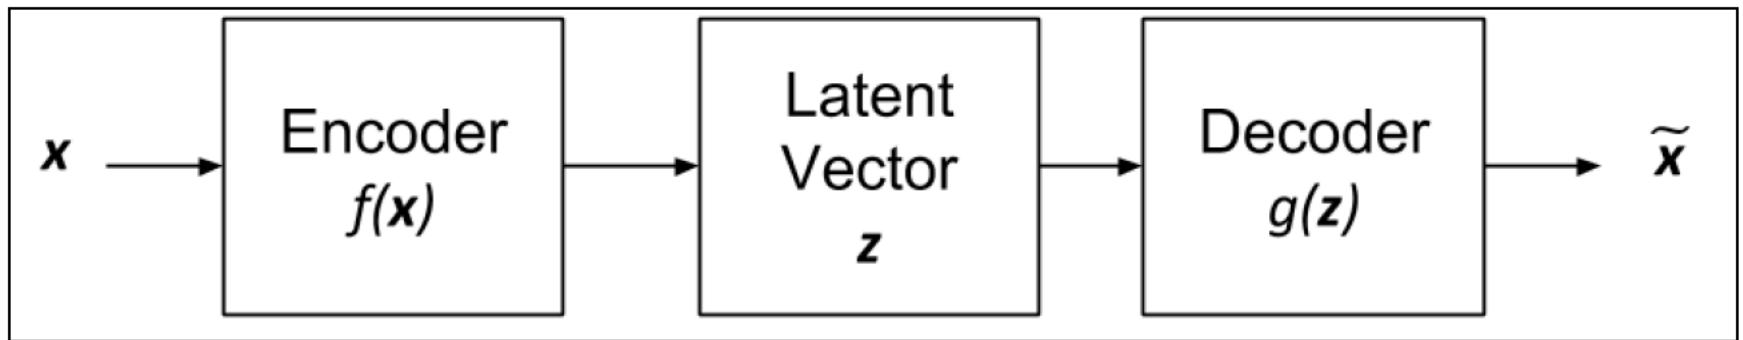
# Deep Autoencoder

- Try to learn a latent representation of the input



# Deep Autoencoder

- Try to learn a latent representation of the input

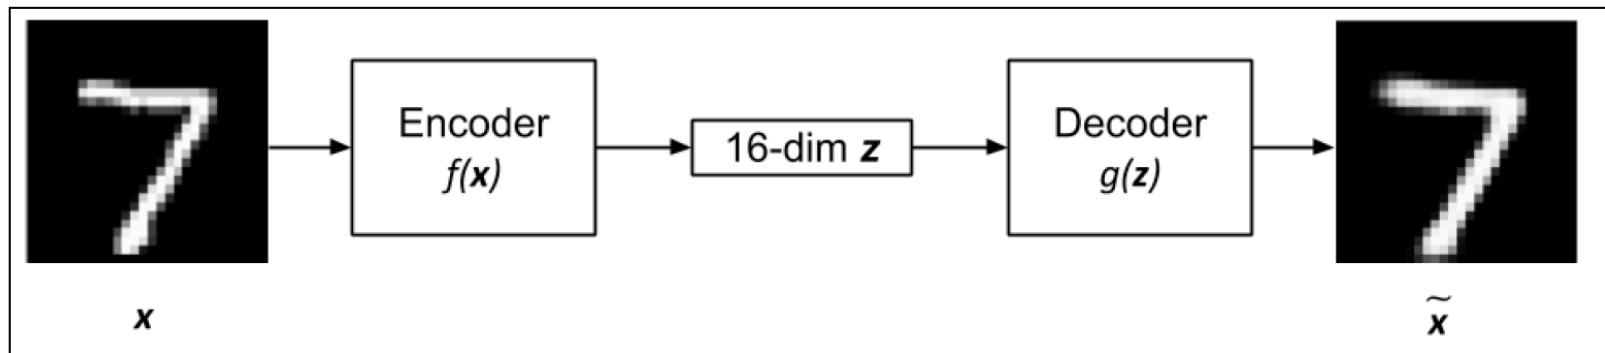
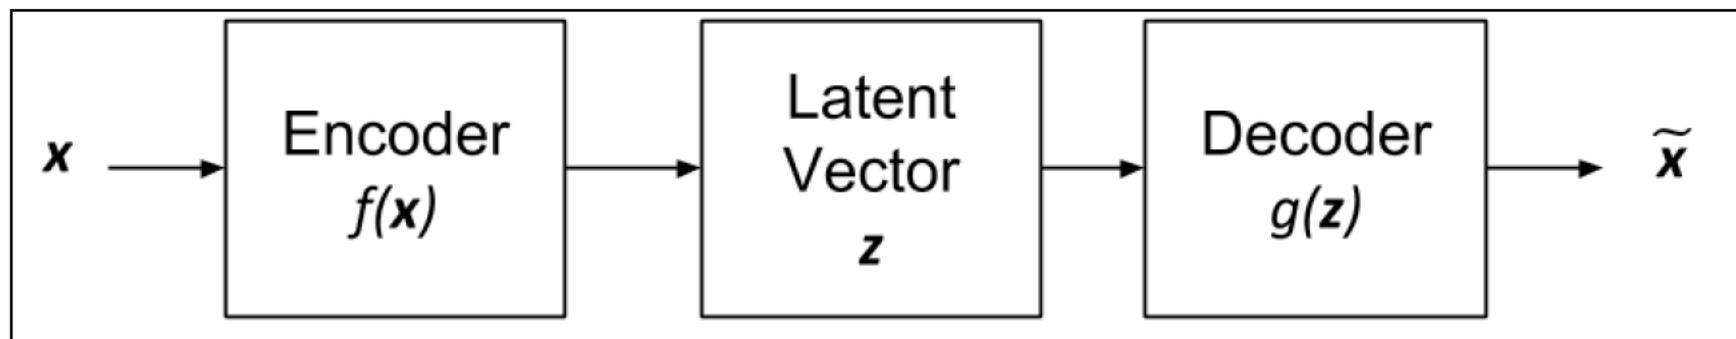


$$y(x) = f(W_1x + b)$$

$$z(x) = g(W_2y(x) + c)$$

# Deep Autoencoder

- Try to learn a latent representation of the input



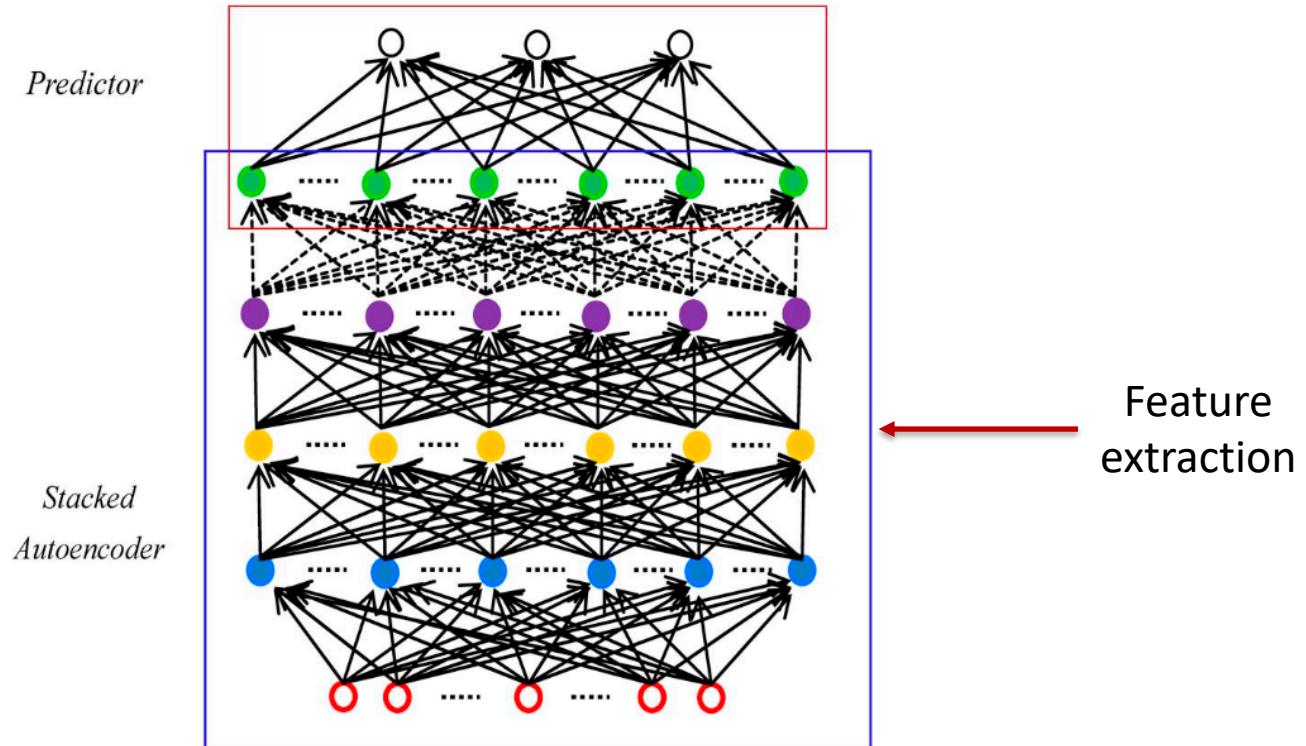
# In the real-world ...

IEEE TRANSACTIONS ON INTELLIGENT TRANSPORTATION SYSTEMS, VOL. 16, NO. 2, APRIL 2015

865

## Traffic Flow Prediction With Big Data: A Deep Learning Approach

Yisheng Lv, Yanjie Duan, Wenwen Kang, Zhengxi Li, and Fei-Yue Wang, *Fellow, IEEE*



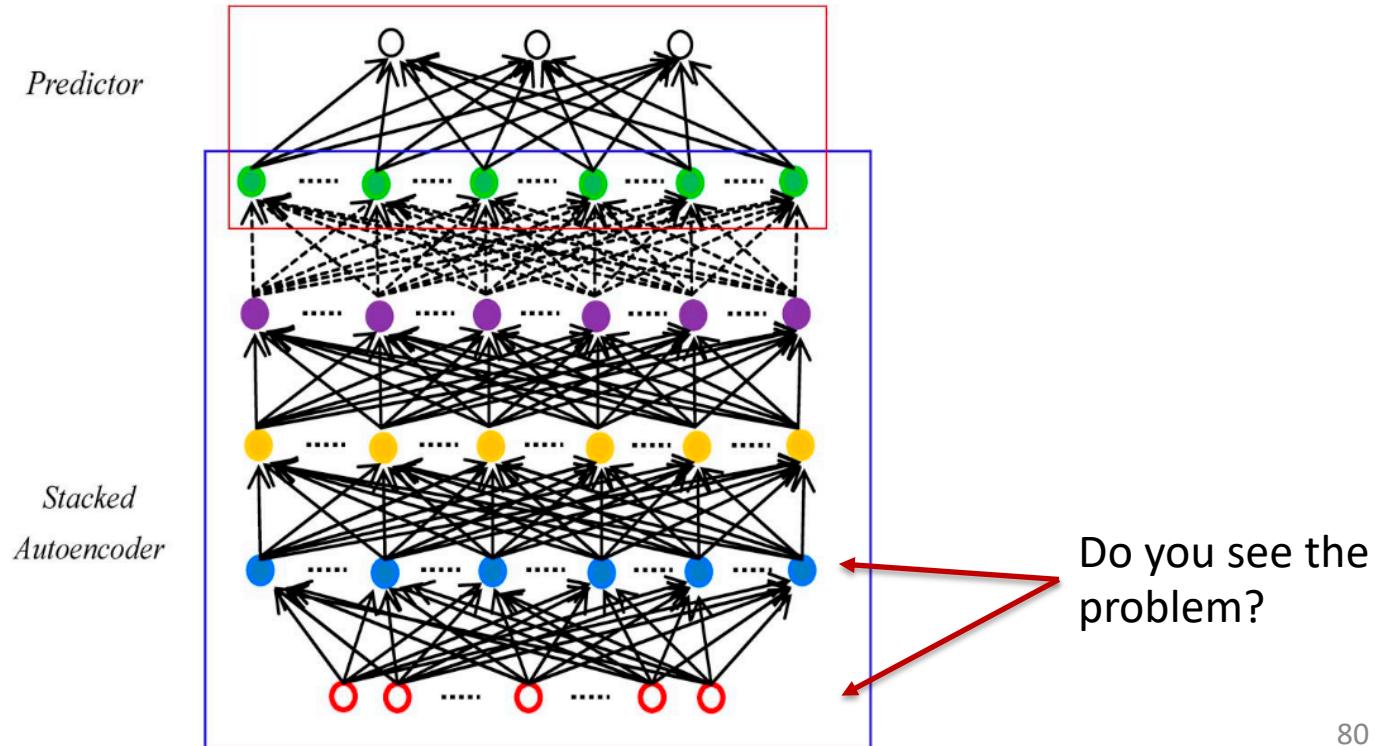
# In the real-world ...

IEEE TRANSACTIONS ON INTELLIGENT TRANSPORTATION SYSTEMS, VOL. 16, NO. 2, APRIL 2015

865

## Traffic Flow Prediction With Big Data: A Deep Learning Approach

Yisheng Lv, Yanjie Duan, Wenwen Kang, Zhengxi Li, and Fei-Yue Wang, *Fellow, IEEE*



# Lab 2

Keras

TensorFlow / Theano / CNTK / ...

CUDA / cuDNN

BLAS, Eigen

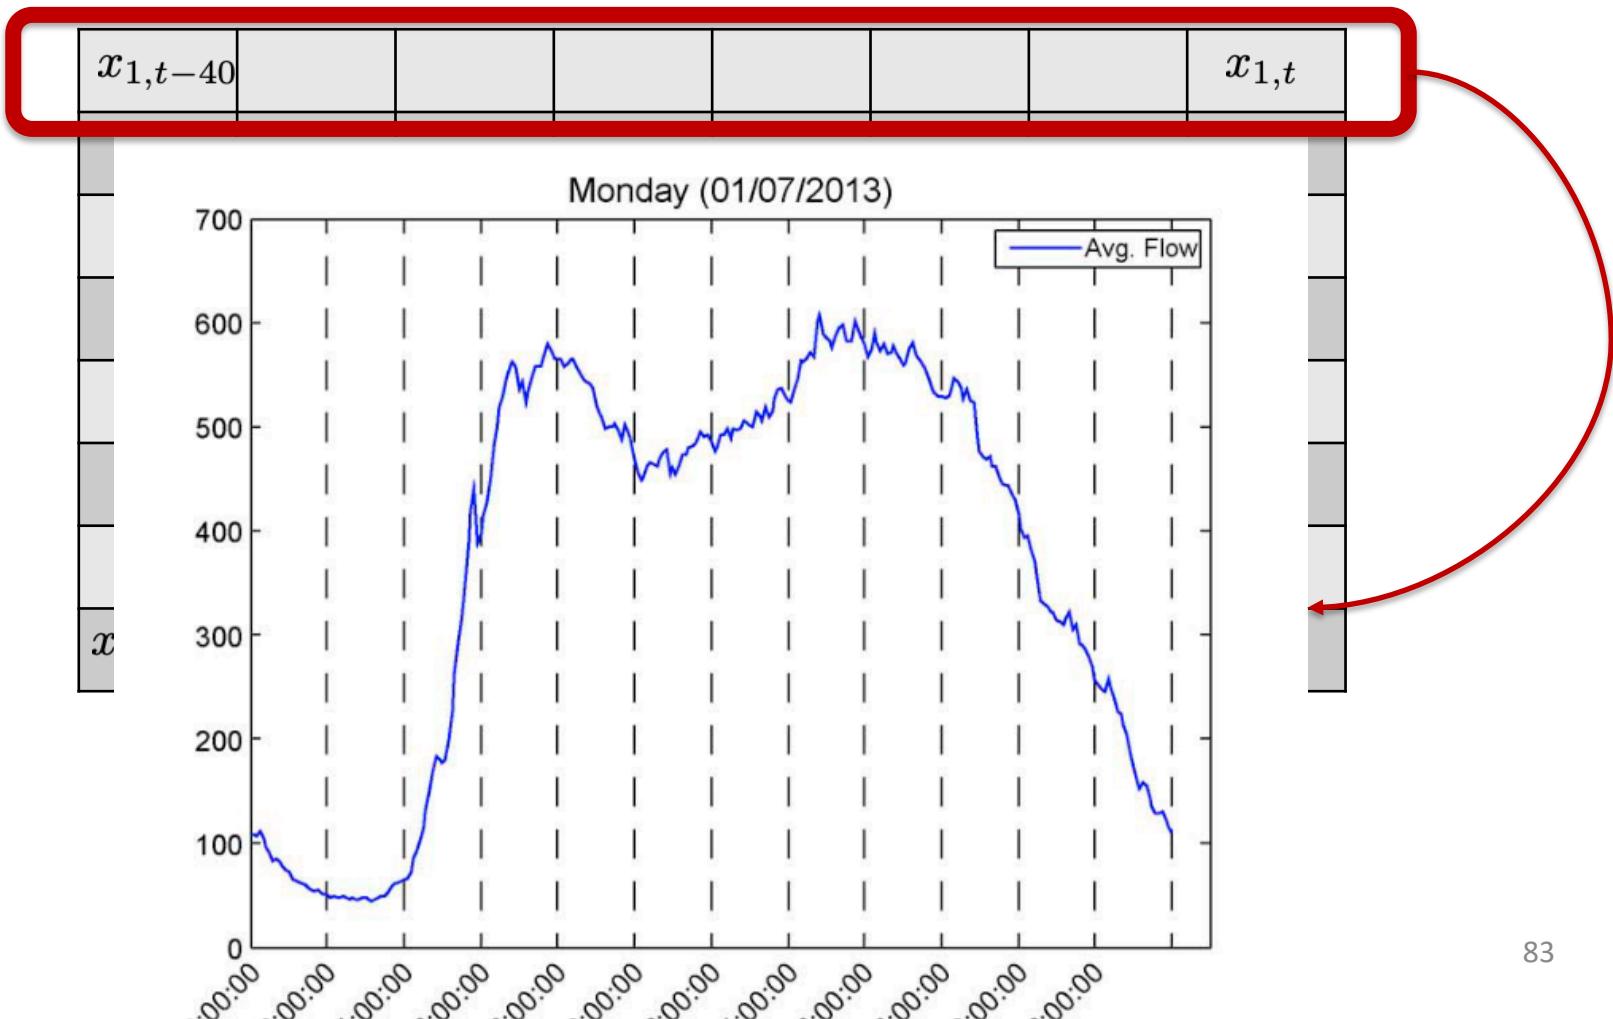
GPU

CPU

# Problems with MLP

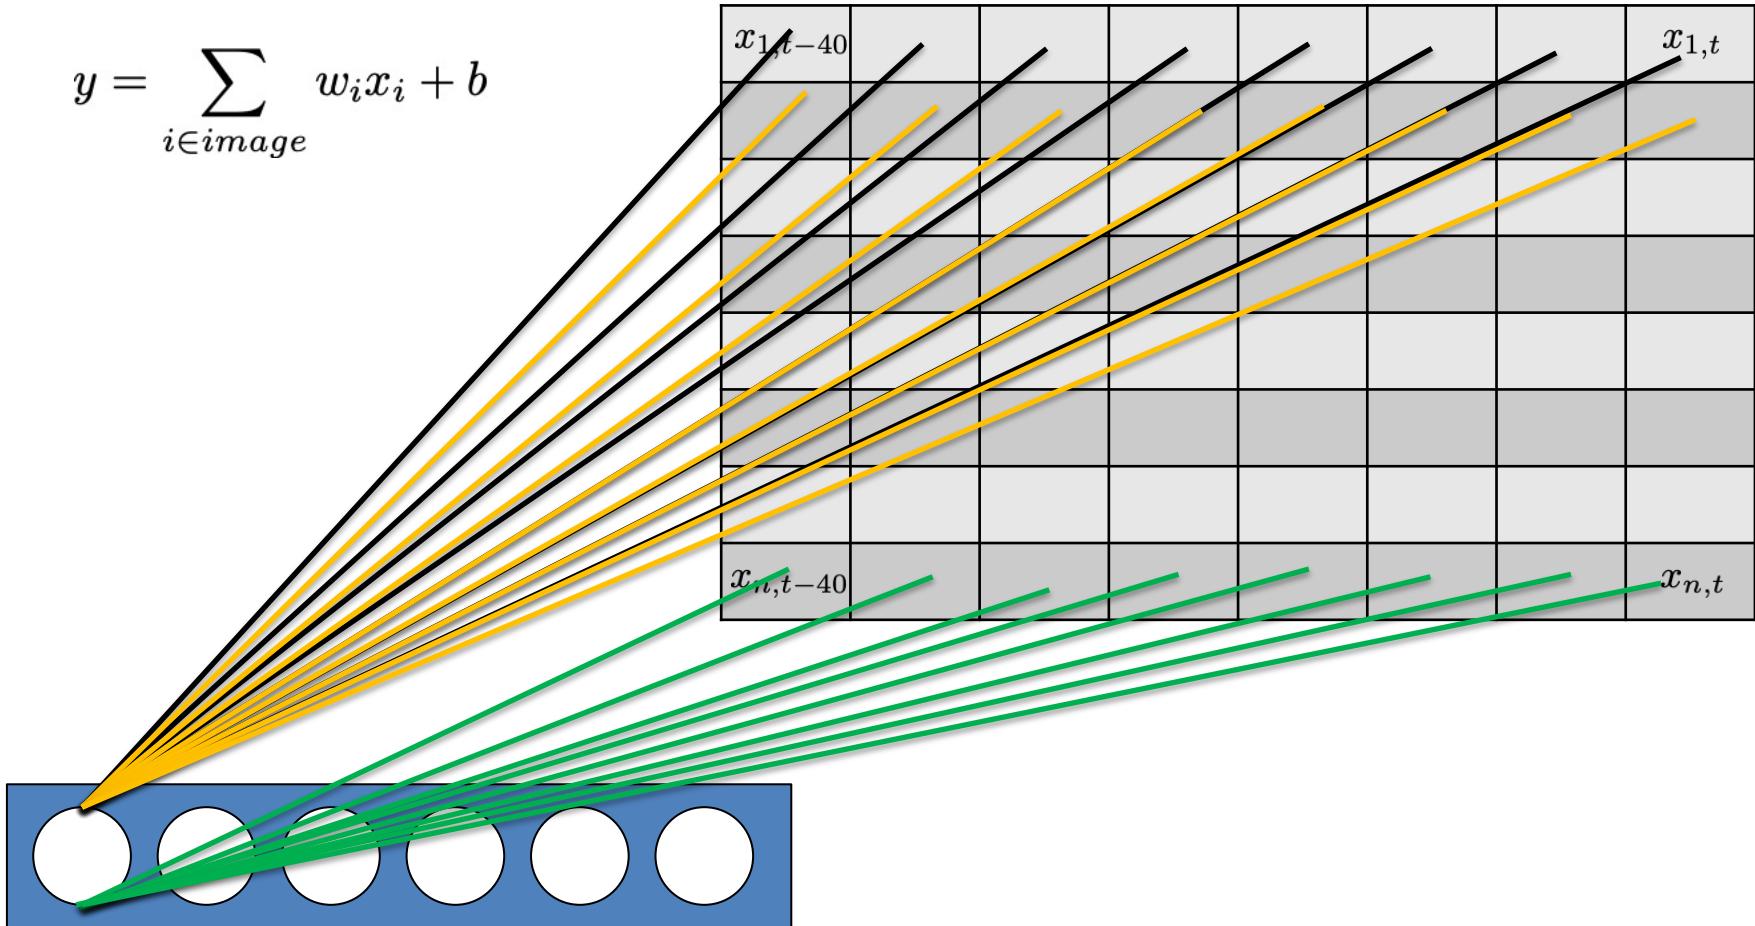
$x_{1,t-40}$								$x_{1,t}$
$x_{n,t-40}$								$x_{n,t}$

# Problems with MLP



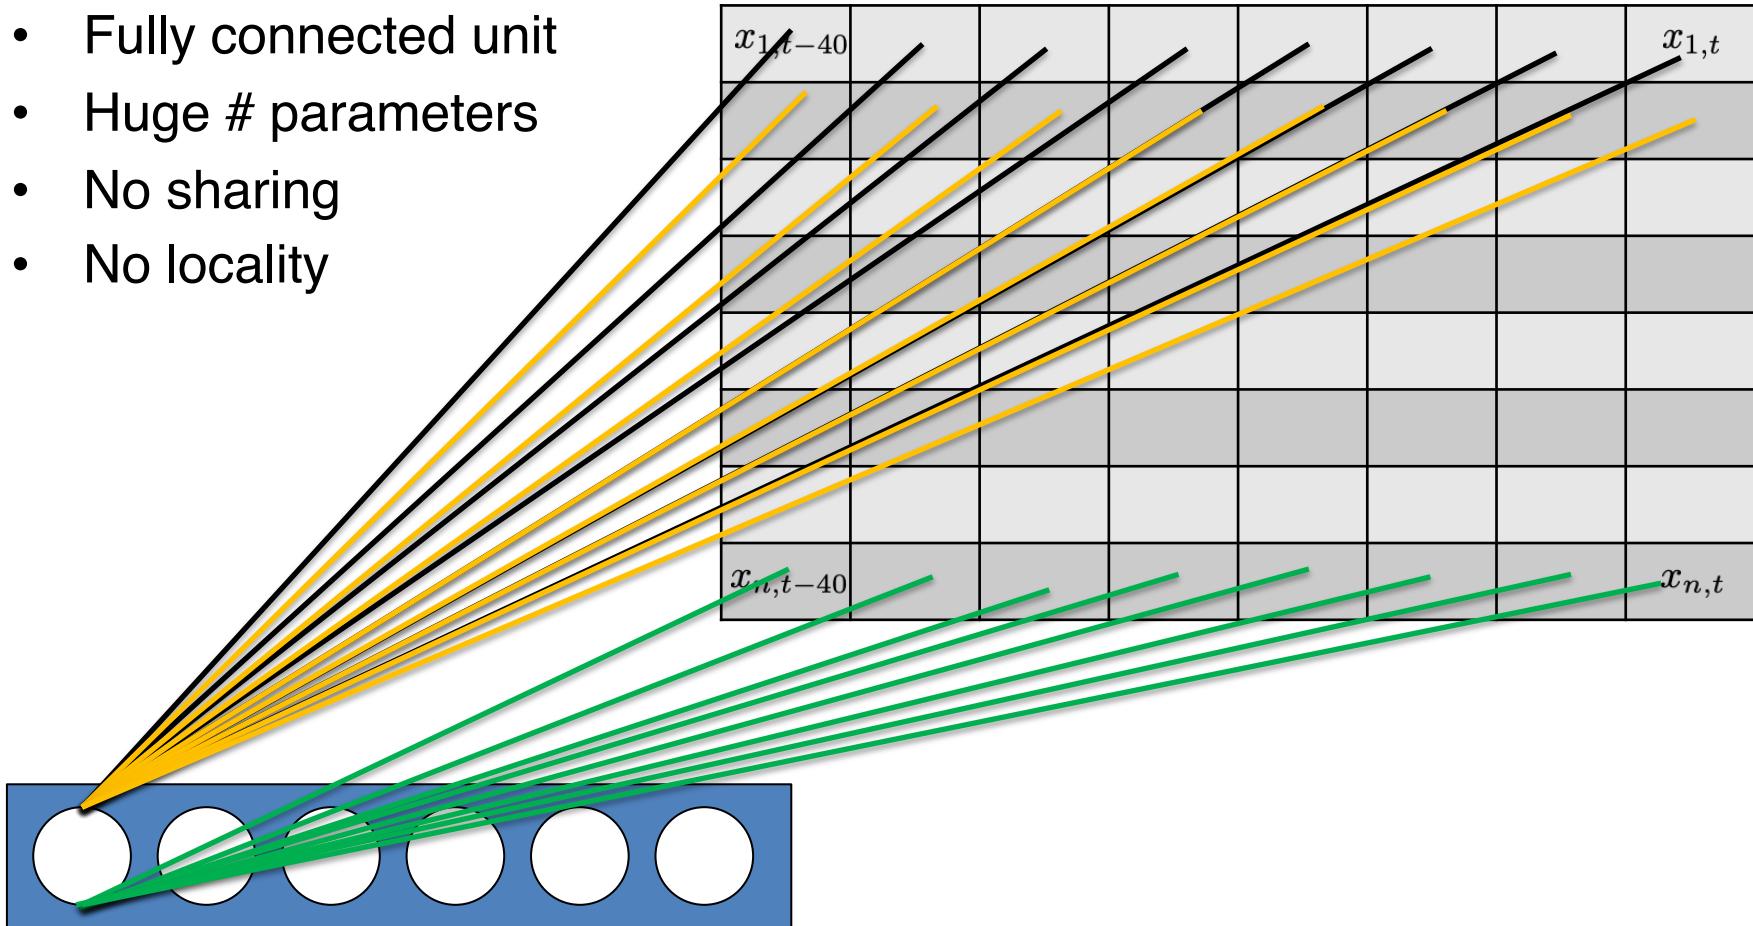
# Problems with MLP

$$y = \sum_{i \in \text{image}} w_i x_i + b$$



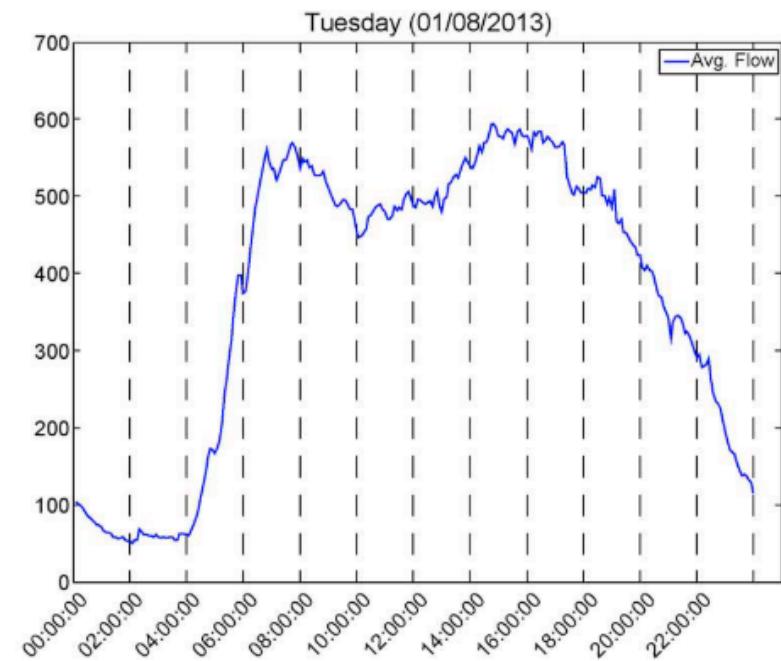
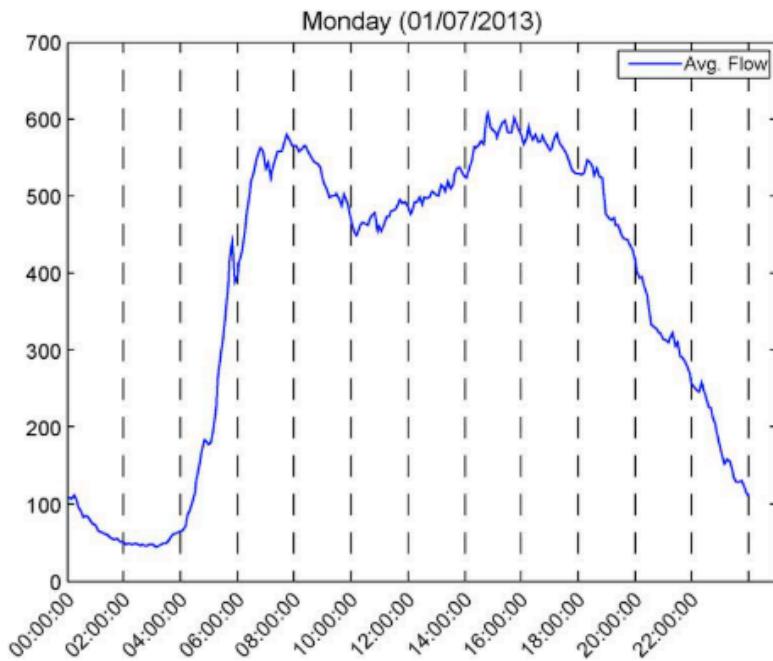
# Problems with MLP

- Fully connected unit
- Huge # parameters
- No sharing
- No locality



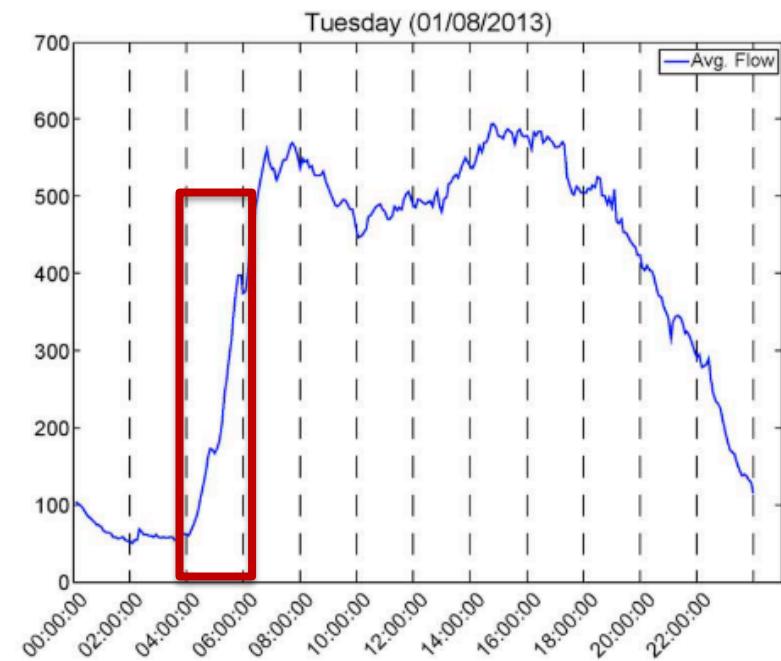
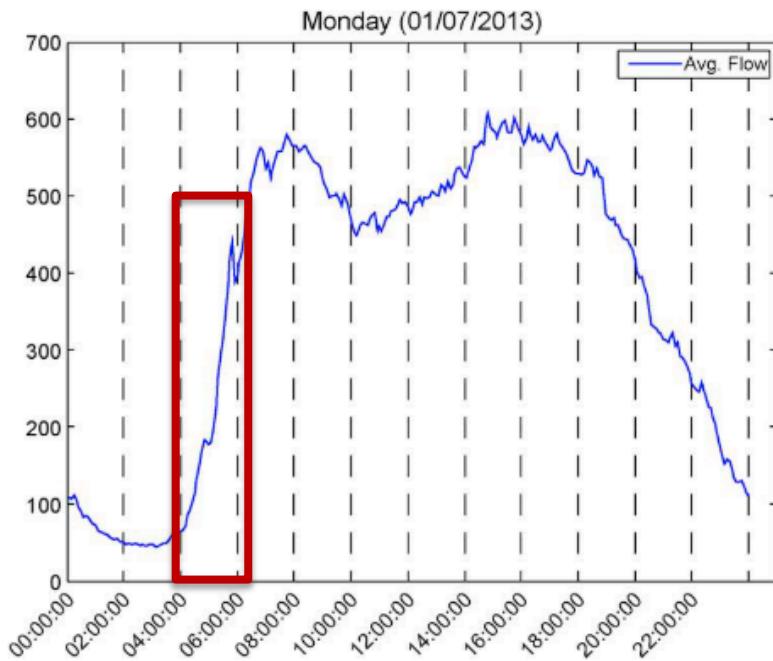
# Problems with MLP

- locality



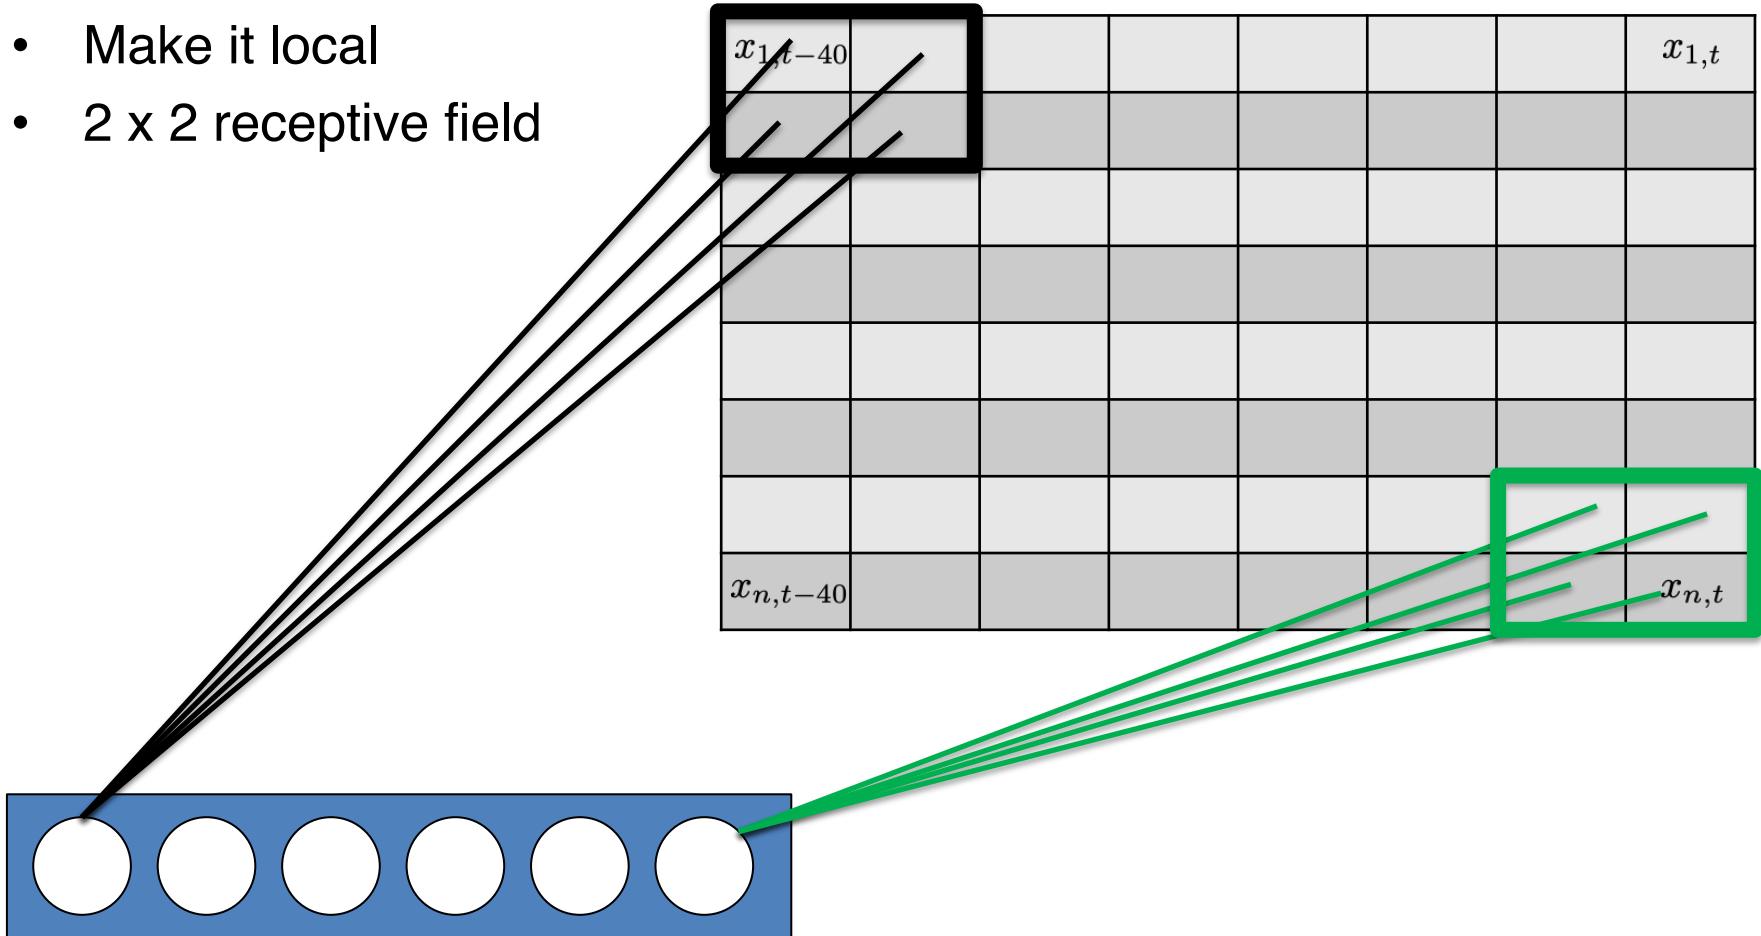
# Problems with MLP

- locality



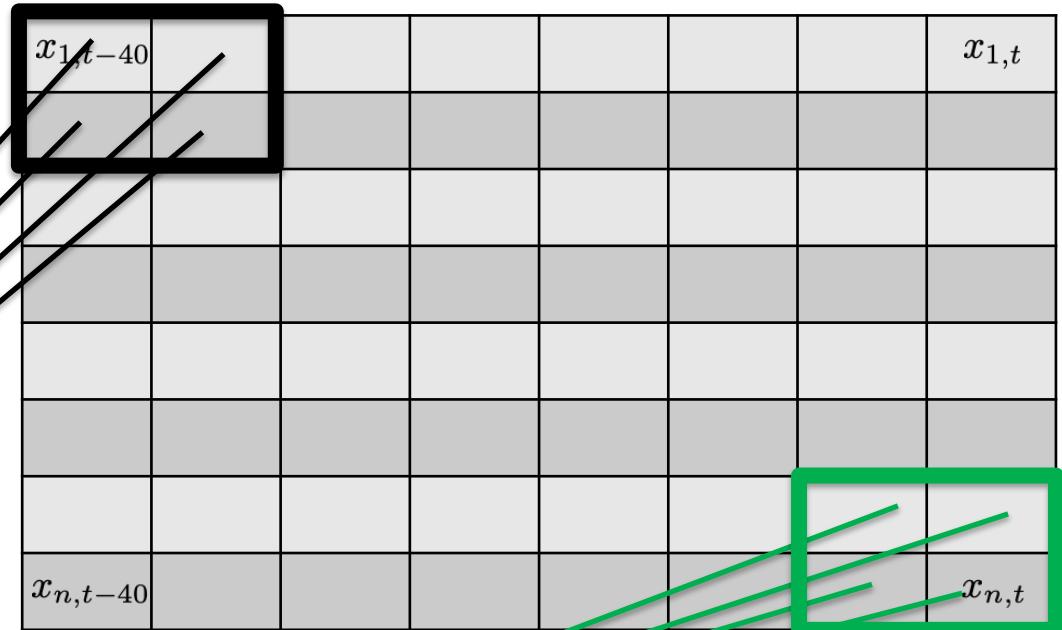
# Make it local

- Make it local
- $2 \times 2$  receptive field



# Make it local

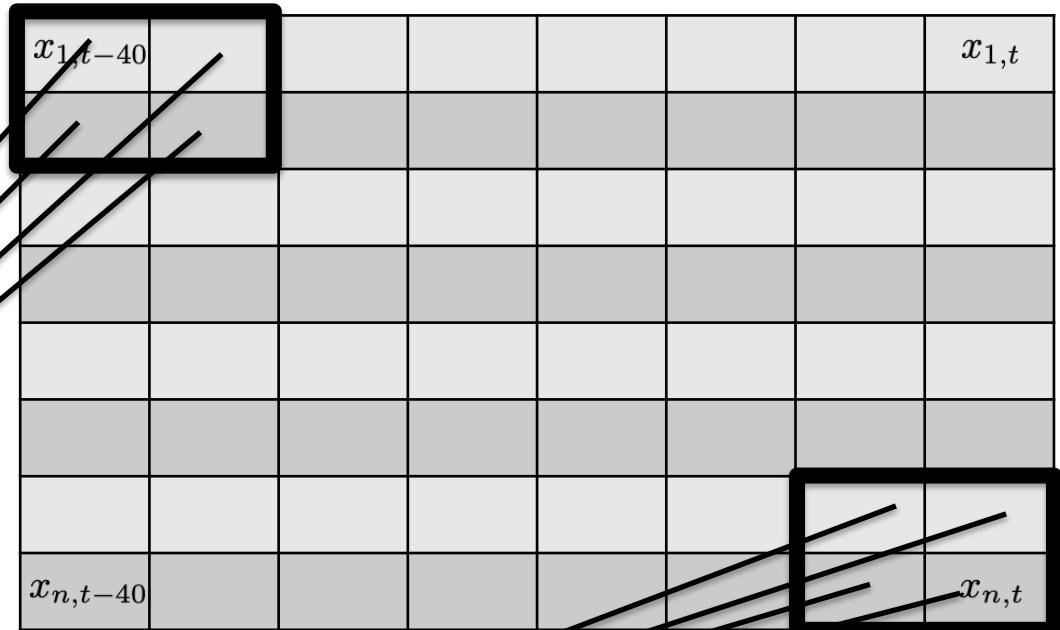
$$y = \sum_{i \in receptive\ field} w_i x_i + b$$



$$y = \sum_{i \in receptive\ field} w'_i x_i + b'$$

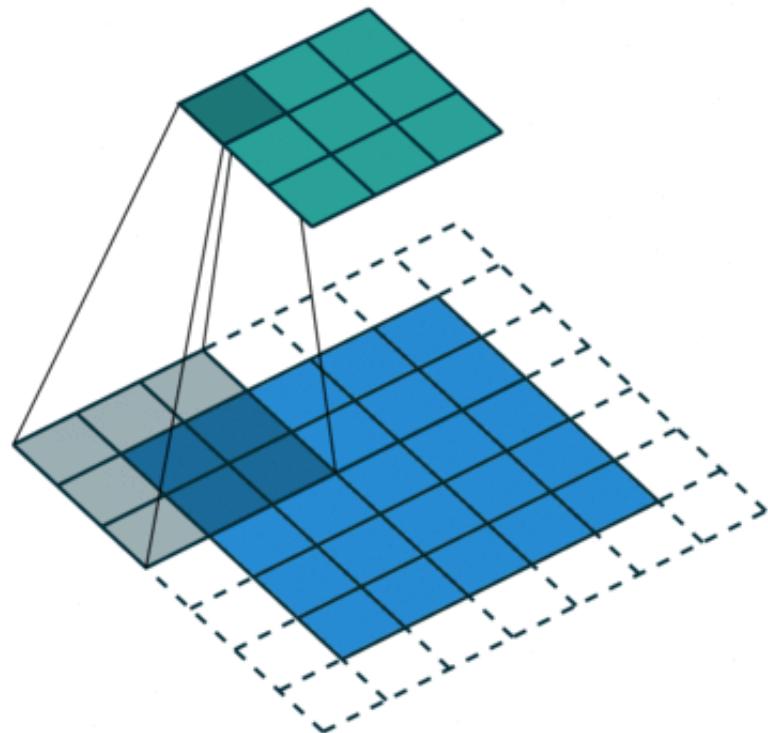
# Share weights

$$y = \sum_{i \in receptive\ field} w x_i + b$$

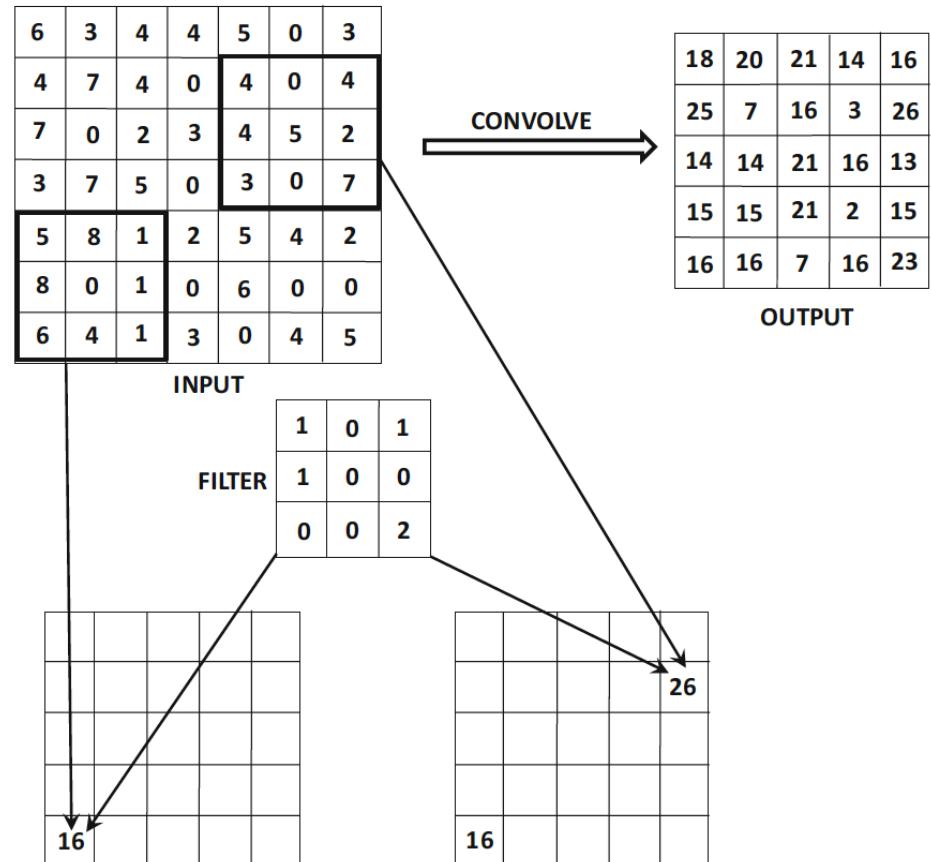
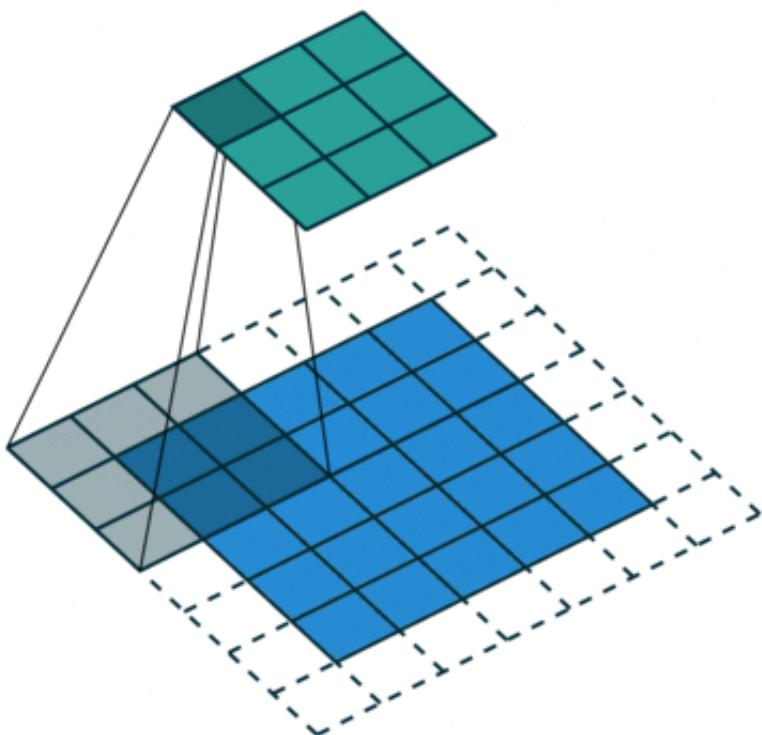


$$y = \sum_{i \in receptive\ field} w x_i + b$$

# Convolutional Neural Net



# Convolutional Neural Net



# Convolutional Neural Net

<https://tinyurl.com/y3wgholx>

- Example:

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input

1	0	1
0	1	0
1	0	1

Filter / Kernel

# Convolutional Neural Net

<https://tinyurl.com/y3wgholx>

- Example:

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input

1	0	1
0	1	0
1	0	1

Filter / Kernel

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

Input x Filter

4		

Feature Map

# Convolutional Neural Net

<https://tinyurl.com/y3wgholx>

- Example:

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input

1	0	1
0	1	0
1	0	1

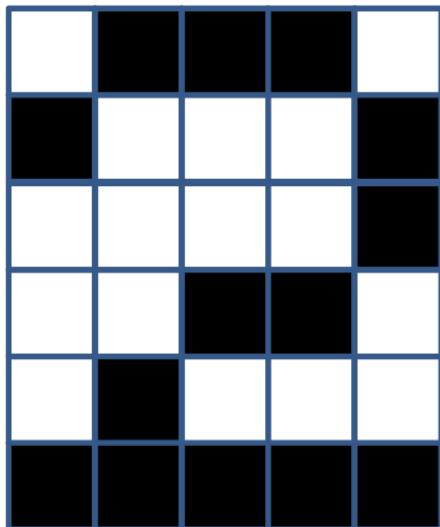
Filter / Kernel

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

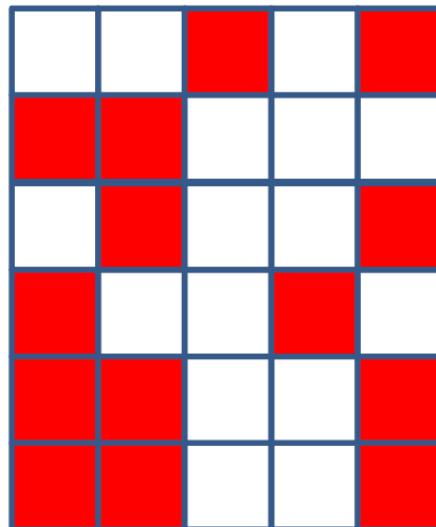
4		

# Another interpretation

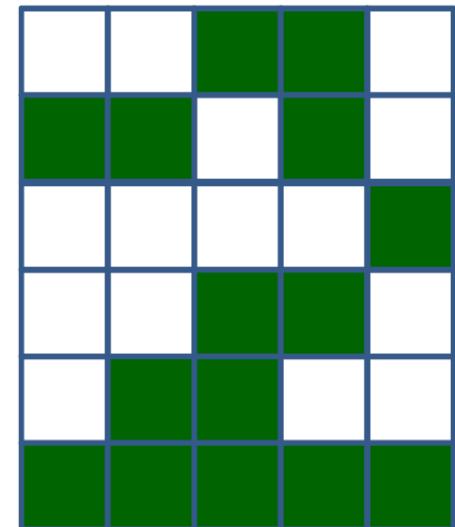
**w**



**x**



**x**



Correlation = 0.57

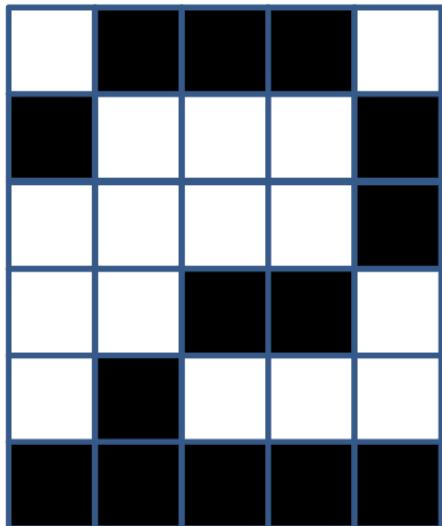
$$y = \begin{cases} 1 & \text{if } \sum_i w_i x_i \geq T \\ 0 & \text{else} \end{cases}$$

Correlation = 0.82

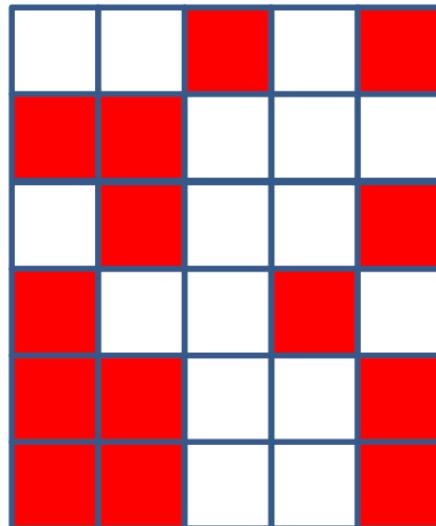


# Another interpretation

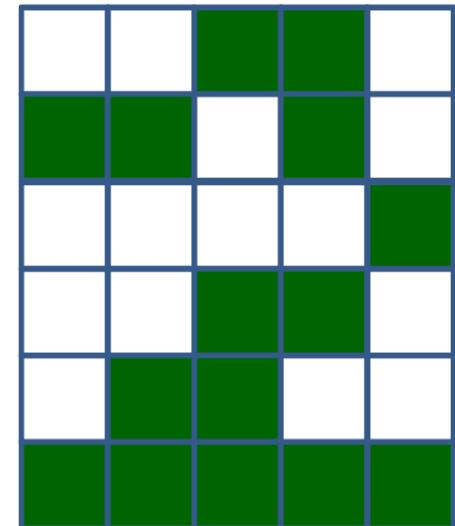
**W**



**X**



**X**



Correlation = 0.57

Correlation = 0.82

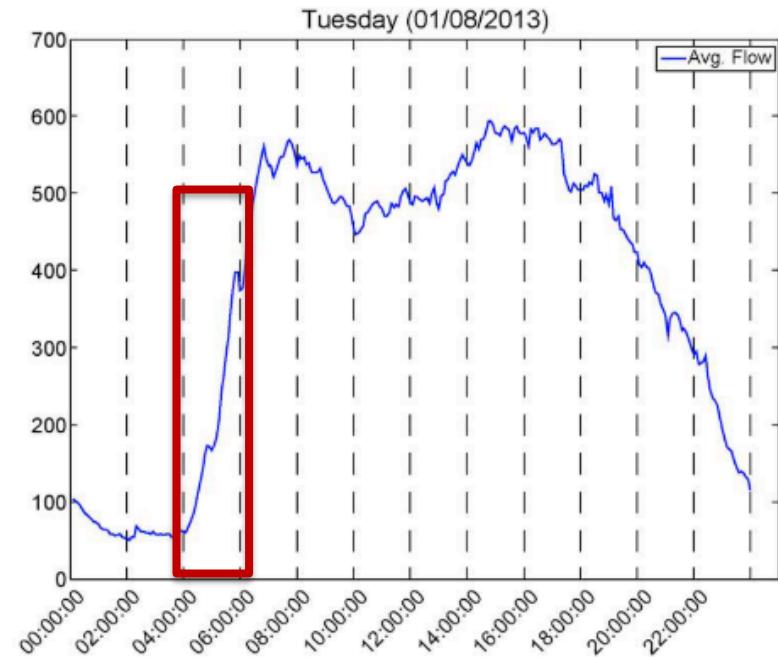
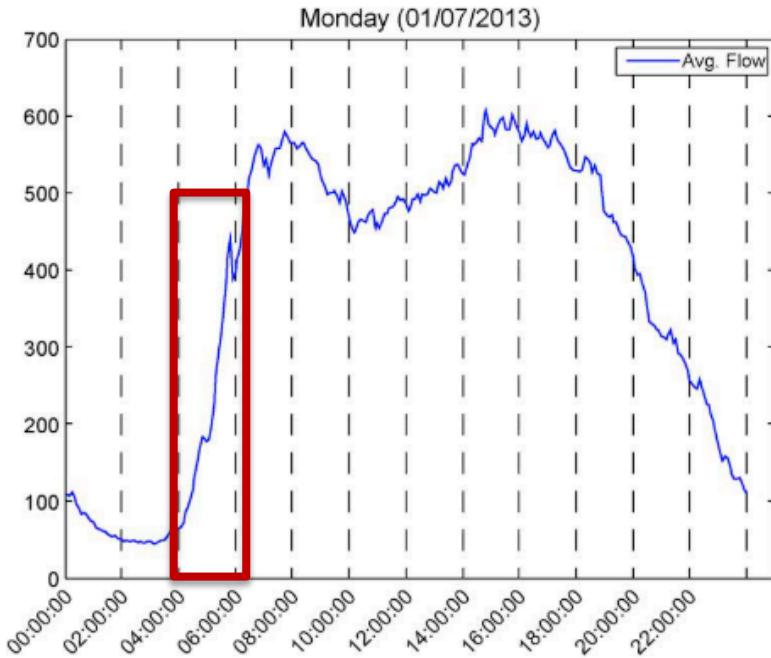


$$y = \begin{cases} 1 & \text{if } \sum_i w_i x_i \geq T \\ 0 & \text{else} \end{cases}$$

The green pattern is more correlated with the weights

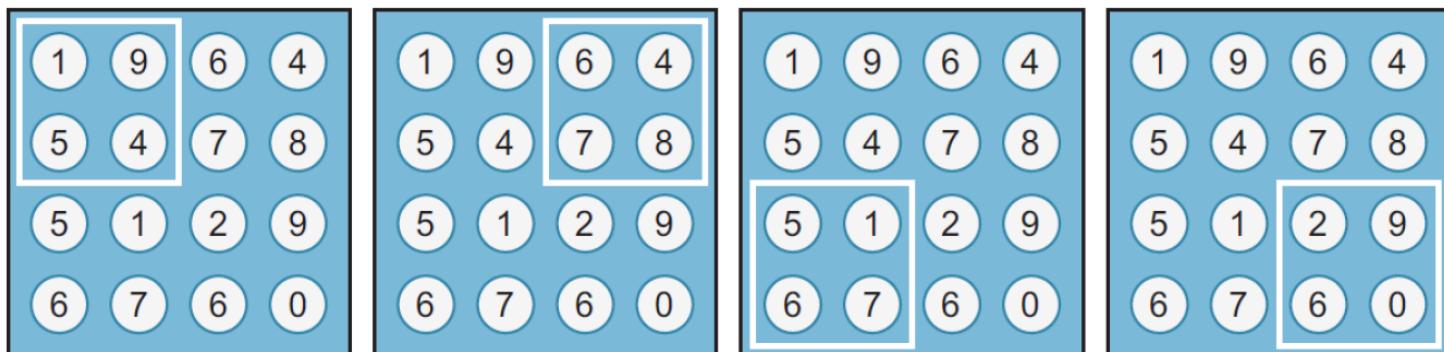
# CNN

- What it inputs differ only slightly?
  - How do we account for jitter?



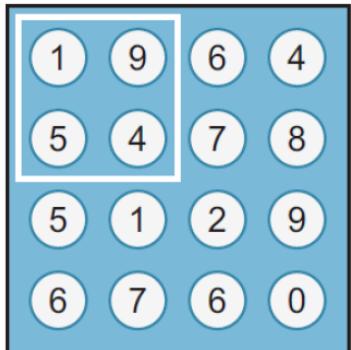
# CNN

- What its inputs differ only slightly?
  - How do we account for jitter?
- Pooling!



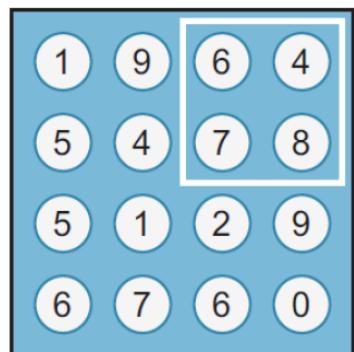
# CNN

- What its inputs differ only slightly?
  - How do we account for jitter?
- Pooling!



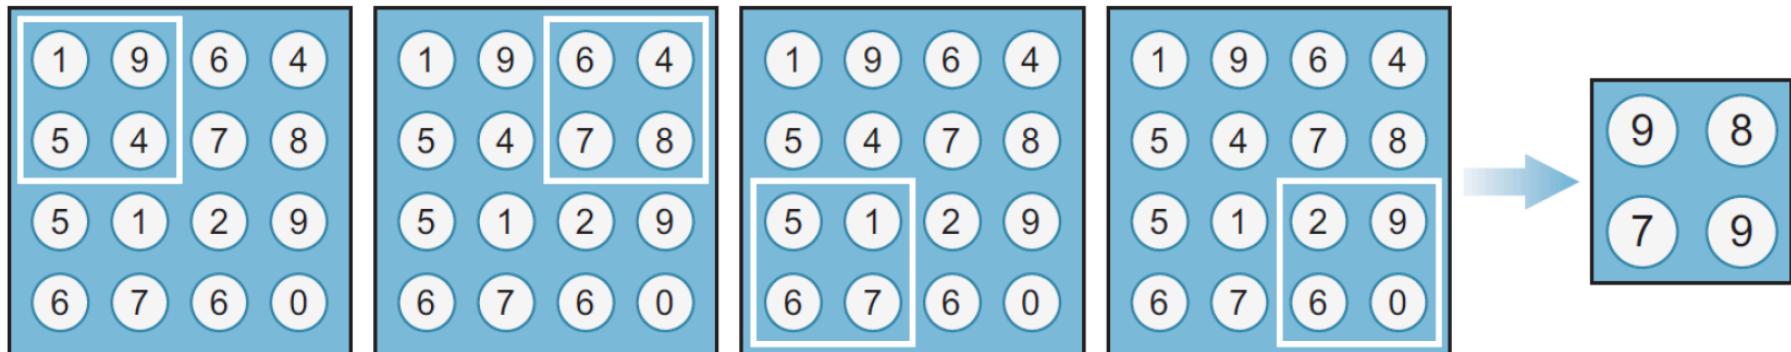
# CNN

- What its inputs differ only slightly?
  - How do we account for jitter?
- Pooling!



# CNN

- What its inputs differ only slightly?
  - How do we account for jitter?
- Pooling!



# A typical CNN architecture



Contents lists available at [ScienceDirect](#)

Transportation Research Part C

journal homepage: [www.elsevier.com/locate/trc](http://www.elsevier.com/locate/trc)



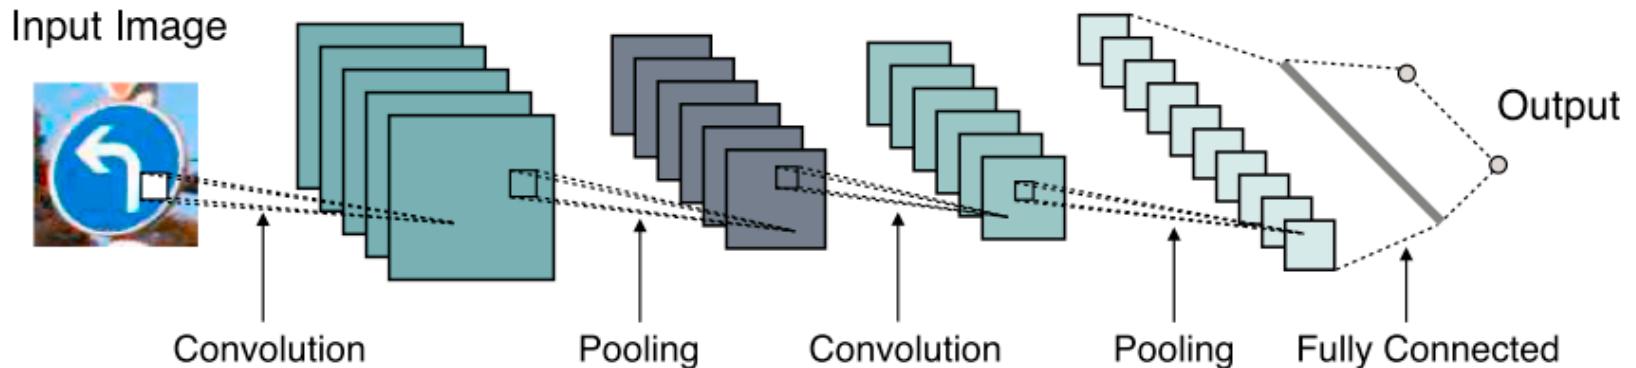
Review

## Enhancing transportation systems via deep learning: A survey

Yuan Wang<sup>a</sup>, Dongxiang Zhang<sup>b,\*</sup>, Ying Liu<sup>b</sup>, Bo Dai<sup>b</sup>, Loo Hay Lee<sup>a</sup>

<sup>a</sup> Department of Industrial Systems Engineering and Management, National University of Singapore, Singapore

<sup>b</sup> School of Computer Science and Engineering, University of Electronic Science and Technology of China, China



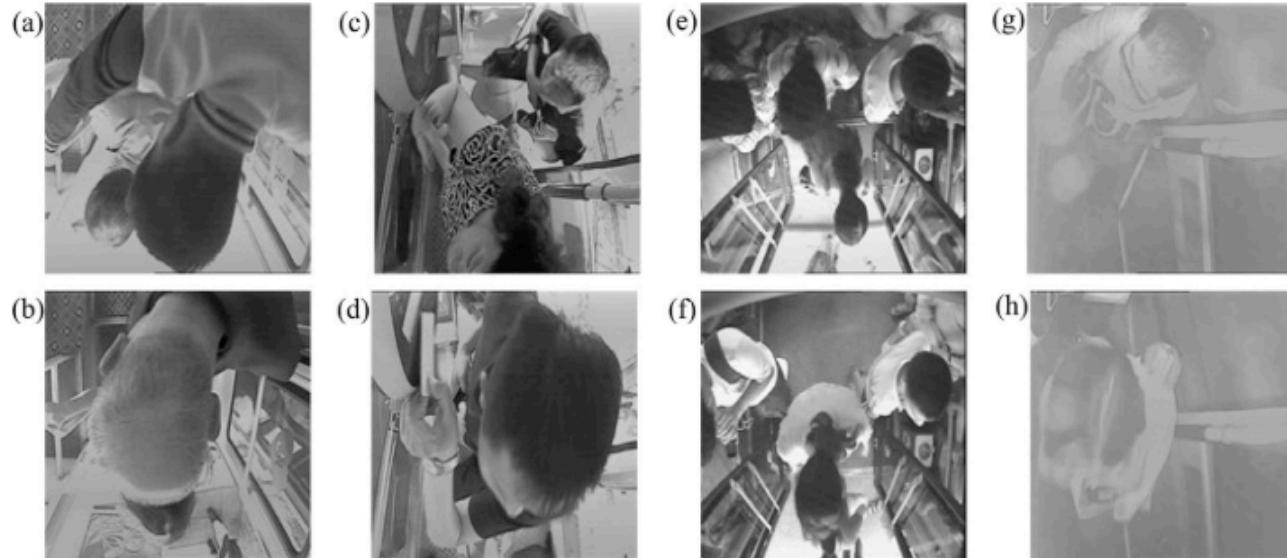
**Fig. 3.** Model architecture of convolutional neural network.

# **Applications based on “true” image data**

- Need to estimate passenger flow
- Traditionally, automatic passenger counting can be done by contact-type counters, optical sensors, and vision-based systems.
- Contact-type:
  - Time consuming, problematic when demand is high
- Optical sensors:
  - Typically undercount
- How about vision-based models?

# Applications based on “true” image data

- Camera feed from buses

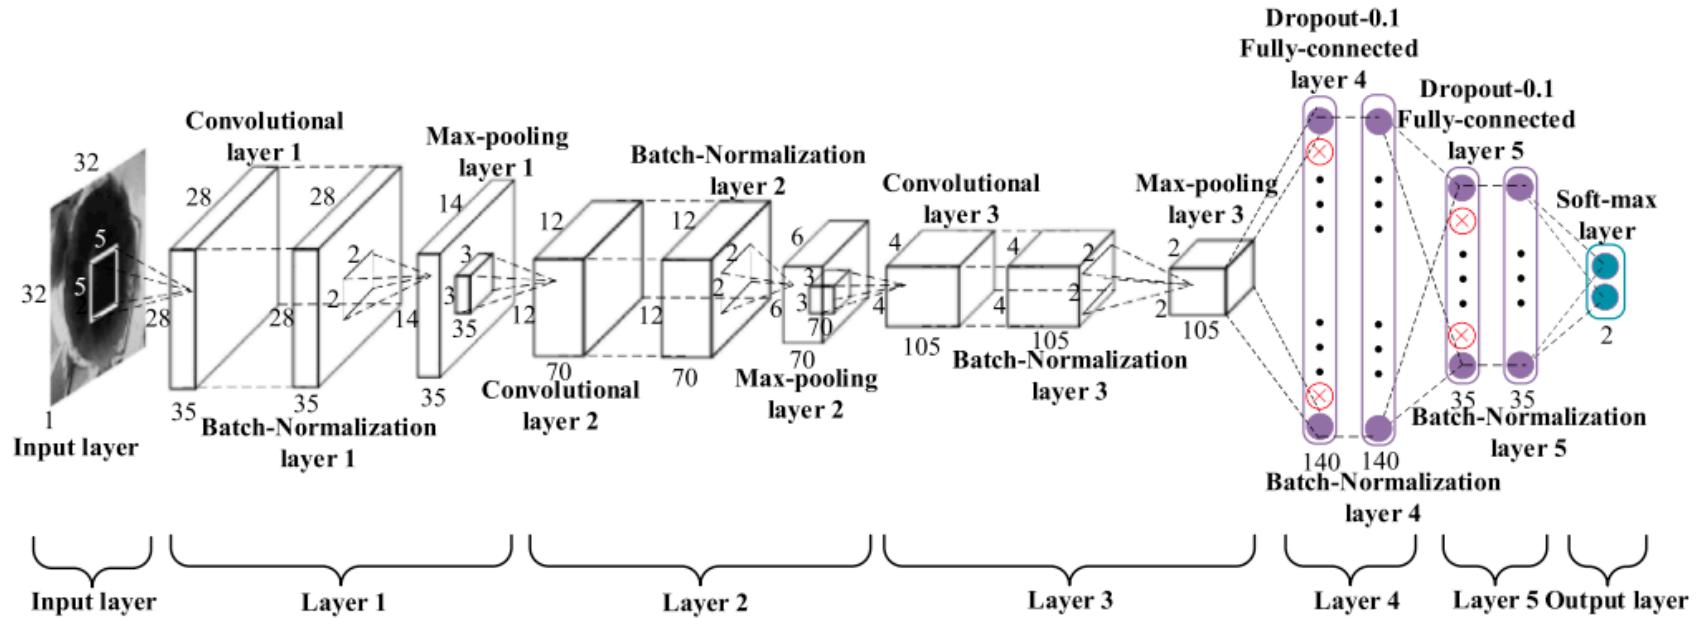


**Fig. 1.** Examples of the public transportation scenarios.

"Passenger flow estimation based on convolutional neural network in public transportation system."  
*Knowledge-Based Systems* 123 (2017)

# Applications based on “true” image data

- Camera feed from buses



"Passenger flow estimation based on convolutional neural network in public transportation system."

*Knowledge-Based Systems* 123 (2017)

# Applications based on “true” image data

- Camera feed from buses



"Passenger flow estimation based on convolutional neural network  
in public transportation system."  
*Knowledge-Based Systems* 123 (2017)

# **Let's clarify what an image is**



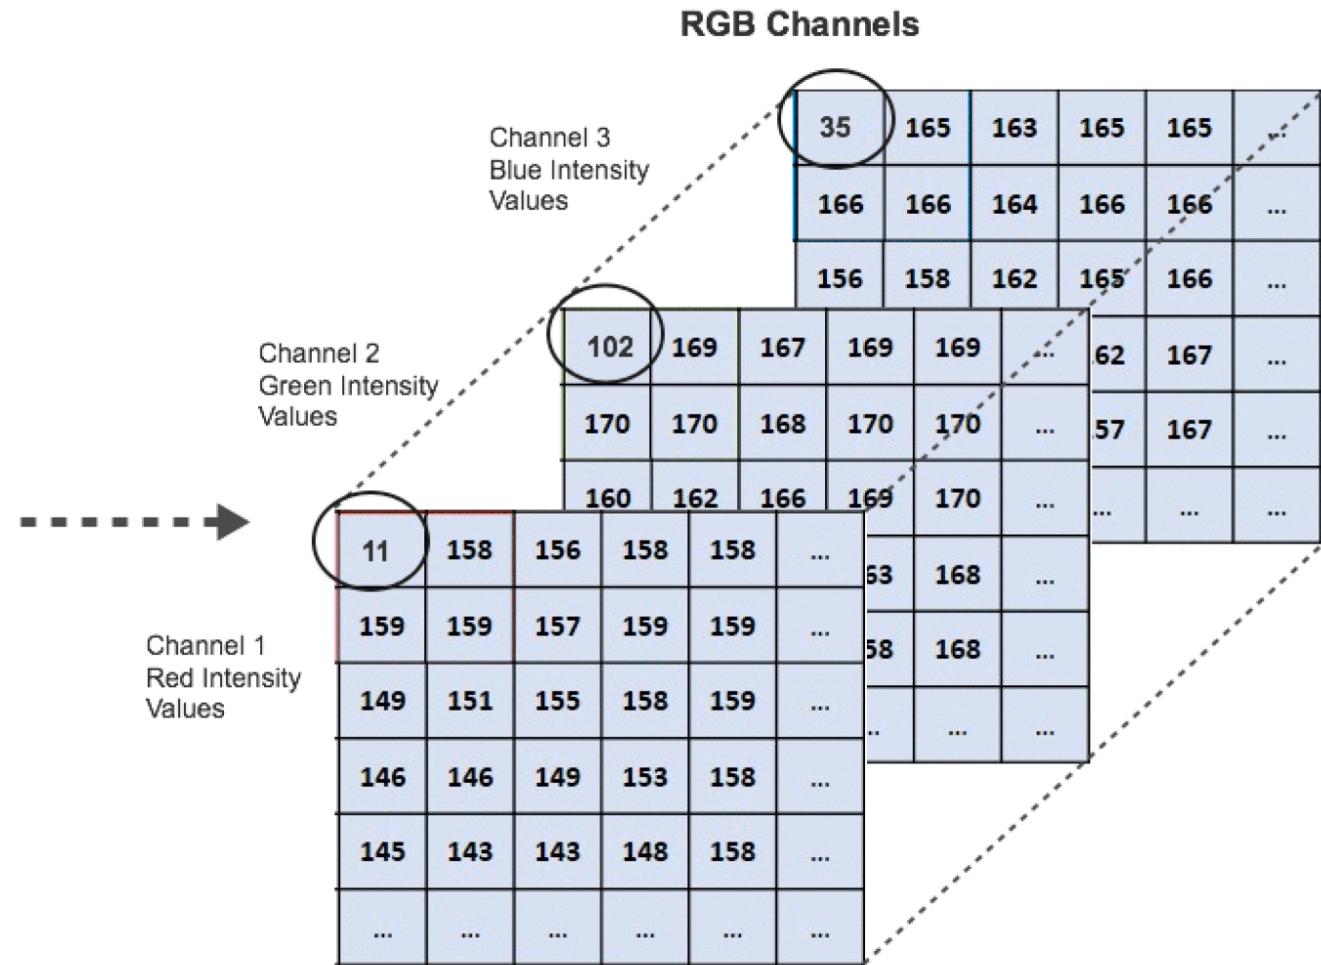
$$28 \times 28 \\ = 784 \text{ pixels}$$

# Let's clarify what an image is

$$F(0, 0) = [11, 102, 35]$$

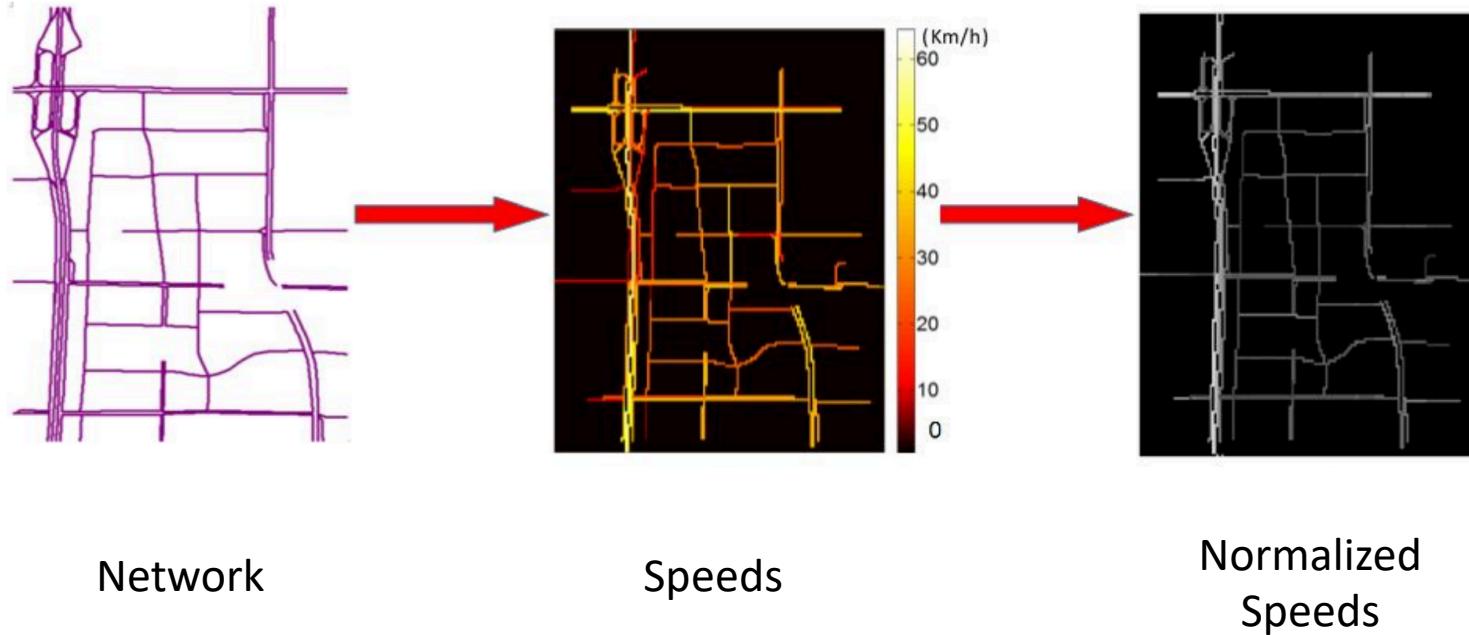


Color Image



# Other applications of CNN

- Traffic Prediction in Transportation Networks



Network

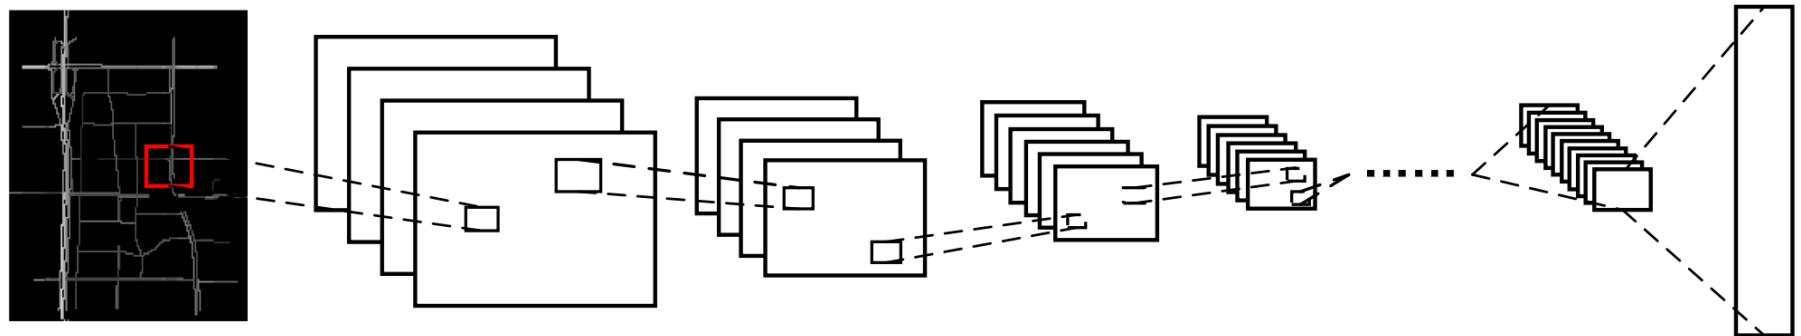
Speeds

Normalized  
Speeds

"Spatiotemporal recurrent convolutional networks for traffic prediction in transportation networks."  
*Sensors* 17, no. 7 (2017)

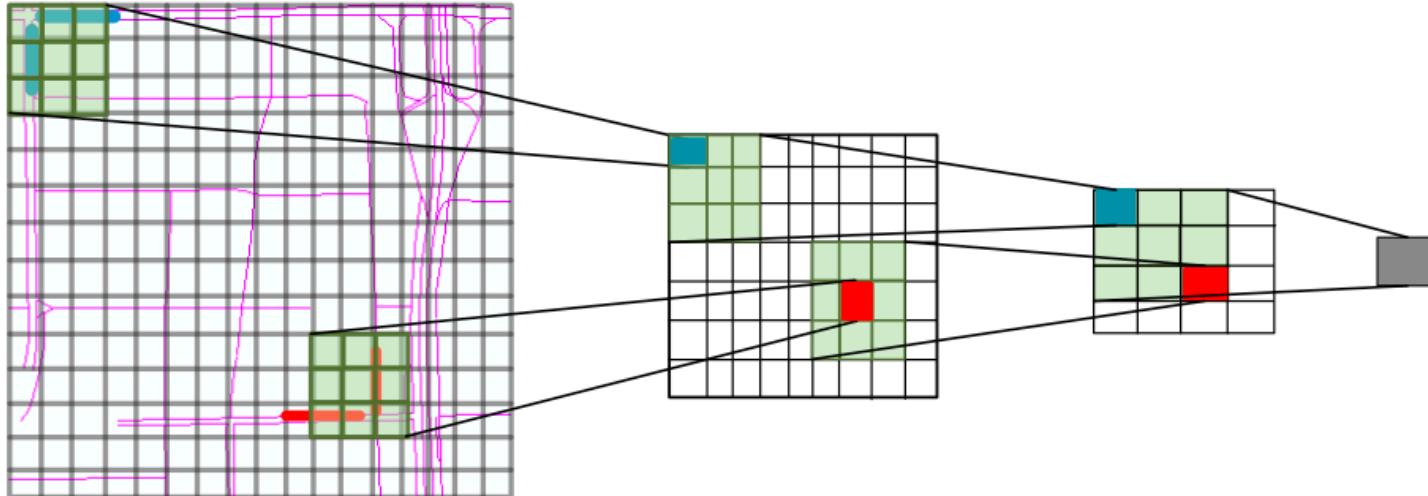
# Other applications of CNN

- Traffic Prediction in Transportation Networks



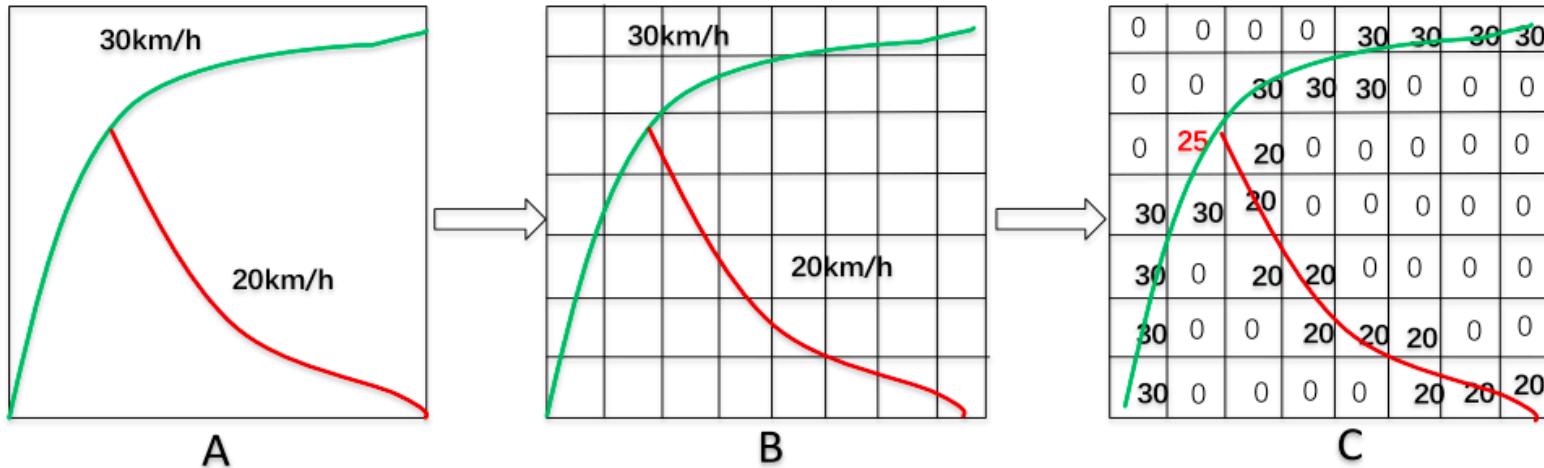
# Other applications of CNN

- Let's look a little bit deeper



# Other applications of CNN

- Let's look a little bit deeper



# Lab 3

Keras

TensorFlow / Theano / CNTK / ...

CUDA / cuDNN

BLAS, Eigen

GPU

CPU

# References

Deep Learning: Practice and Trends  
Reed, de Freitas, Vinyals

CMU Deep learning course 2019  
Bhiksha Raj